A Course in Differential Privacy

Manuel V. Loureiro

Version 2026-10-26

Table of contents

Preface 1

I. Privacy Failures 2

1. Attempts at Data Privacy 4
1.1. Pseudonyms: the taxi trips . . . 4
1.2. Rich records: the Netflix Prize . . . 5
1.3. Releasing a model: memorization . . . 6
1.4. Releasing statistics: genomic studies . . . 7
1.5. Removing identifiers: k-anonymity . . . 8
1.6. Summary . . . 9
1.7. References . . . 10

2. Reconstruction Attacks 12
2.1. Reconstruction from exact statistics . . . 12
2.2. Subset queries and blatant non-privacy . . . 13
2.3. Asking every query . . . 14
2.4. Random queries . . . 15
2.5. Reconstruction in practice . . . 20
2.6. Summary . . . 21
2.7. References . . . 22

II. The Basics of Differential Privacy 23

3. Randomized Response and the Definition 25
3.1. Randomized response . . . 25
3.2. Differential privacy . . . 27
3.3. Randomized response revisited . . . 30
3.4. Summary . . . 31
3.5. References . . . 32

4. The Laplace Mechanism 34
4.1. Sensitivity . . . 34
4.2. The Laplace distribution . . . 35
4.3. The Laplace mechanism . . . 35
4.4. Counting queries . . . 37
4.5. Histograms . . . 38
4.6. Properties of differential privacy . . . 39
4.7. Summary . . . 41

4.8. References 42

5. Approximate Differential Privacy **43**
5.1. The definition 43
5.2. The privacy loss random variable 44
5.3. How small δ must be 47
5.4. The Gaussian mechanism 49
5.5. Properties 54
5.6. Summary 56
5.7. References 57

6. Advanced Composition **58**
6.1. The theorem 58
6.2. Reduction to four outcomes 61
6.3. Composing the four-outcome pair 64
6.4. Summary 66
6.5. References 67

7. The Exponential Mechanism **68**
7.1. Selecting an object 68
7.2. Privacy 69
7.3. Utility 70
7.4. Applications 71
7.5. Summary 77
7.6. References 77

III. Query Release and Its Limits **78**

8. Private Multiplicative Weights **80**
8.1. Linear queries 80
8.2. Learning from experts 83
8.3. Multiplicative weights for queries 86
8.4. Private multiplicative weights 88
8.5. Summary 91
8.6. References 92

9. Sparse Vector **93**
9.1. Answering queries online 93
9.2. AboveThreshold 94
9.3. Sparse vector 96
9.4. Numeric sparse vector 98
9.5. Online private multiplicative weights 99
9.6. Summary 101
9.7. References 101

10. Beyond Global Sensitivity **102**
10.1. Local sensitivity 102
10.2. Propose-test-release 103

10.3. Stability: the mode and the histogram . . . 105
10.4. Privately bounding the local sensitivity . . . 108
10.5. Smooth sensitivity . . . 109
10.6. Summary . . . 110
10.7. References . . . 111

11. Packing Lower Bounds **112**
11.1. Two problems and their upper bounds . . . 112
11.2. The packing argument . . . 114
11.3. One-way marginals . . . 116
11.4. Histograms . . . 117
11.5. Summary . . . 119
11.6. References . . . 119

IV. Private Learning and Statistics **121**

12. What Differential Privacy Protects **123**
12.1. What the definition promises . . . 123
12.2. Learning a rule is not a violation . . . 124
12.3. Inferences from what people reveal . . . 126
12.4. What a violation looks like . . . 127
12.5. Summary . . . 128
12.6. References . . . 128

13. Private Empirical Risk Minimization **129**
13.1. Empirical risk minimization . . . 129
13.2. Why the minimizer is not private . . . 131
13.3. Output perturbation . . . 131
13.4. Objective perturbation . . . 134
13.5. Gradient perturbation . . . 136
13.6. Summary . . . 141
13.7. References . . . 141

14. Private Deep Learning **143**
14.1. Neural networks . . . 143
14.2. Differentially private SGD . . . 144
14.3. Privacy accounting . . . 145
14.4. DP-SGD in practice . . . 150
14.5. Private aggregation of teacher ensembles . . . 150
14.6. Summary . . . 153
14.7. References . . . 153

15. Private Mean Estimation **155**
15.1. The mean of a dataset and of a distribution . . . 155
15.2. Unbounded data . . . 156
15.3. Estimating a parameter of a distribution . . . 158
15.4. Clipping the Gaussian . . . 159
15.5. A coarse estimate from a histogram . . . 161

15.6. Shrinking confidence intervals . . . 162
15.7. Many dimensions . . . 164
15.8. Heavy-tailed data . . . 166
15.9. Summary . . . 168
15.10. References . . . 170

16. Adaptive Data Analysis **172**
16.1. Statistical queries . . . 172
16.2. Adaptivity breaks the bound . . . 173
16.3. Privacy implies generalization . . . 174
16.4. An easier transfer theorem . . . 176
16.5. Summary . . . 179
16.6. References . . . 180

V. Deployments **181**

17. Local Differential Privacy **183**
17.1. The local model . . . 183
17.2. Counting bits, central and local . . . 184
17.3. The price is unavoidable . . . 185
17.4. Repeated questions: RAPPOR . . . 189
17.5. Values no one knew: Apple . . . 191
17.6. Numbers in one bit: Microsoft . . . 192
17.7. Summary . . . 195
17.8. References . . . 195

18. The Shuffle Model and Central Deployments **197**
18.1. Encode, shuffle, analyze . . . 197
18.2. The shuffle model . . . 198
18.3. Amplification by shuffling . . . 198
18.4. Counting bits in the shuffle model . . . 199
18.5. A query service: LinkedIn . . . 203
18.6. Data releases . . . 207
18.7. Summary . . . 210
18.8. References . . . 210

Appendices **212**

A. Glossary **212**

B. Index **213**

Preface

This book is a course in differential privacy. It starts from the failures that motivate the subject, builds the definition and its basic algorithms, and goes on to answering many queries, private statistics and machine learning, lower bounds, and deployments.

Each chapter takes one topic. The text is short, and the mathematics is complete. Definitions, theorems and proofs are set apart and numbered, and every proof is given in full or, where it is long and standard, cited with what it needs. The prose between them says why each result comes next and what it means.

The reader should know discrete probability, including expectation, variance, independence and the inequalities of Markov, Chebyshev and Hoeffding, and basic linear algebra and calculus. Results from probability that the course needs beyond these are stated and proved where they are first used.

Notation. A dataset is $D = (x_1, \dots, x_n) \in \mathcal{X}^n$, a list of n records from a universe $\mathcal{X}$. Two datasets are neighbours, $D \sim D'$, when they differ in one record, which is the replace-one relation. A mechanism M maps a dataset to a random output in a set $\mathcal{Y}$. The privacy parameters are $\varepsilon \geq 0$ and $\delta \in [0, 1]$. The natural logarithm is written ln, and the logarithm to base 2 is written $\log_2$.

Part I.

Privacy Failures

Privacy is easy to promise and hard to deliver. Before the course defines differential privacy, this part shows what goes wrong without a definition. Its two chapters look at the problem from opposite sides, first through real releases that failed, then through theorems that say exactly when a release must fail.

The first chapter goes through five releases of sensitive data whose owners took care to protect them. They replaced identifiers by pseudonyms, published only a trained model or only summary statistics, or removed names. Every protection was defeated, and almost always in the same way. The attacker combined the release with information from another source, such as a public register, a review website or a voter roll. No protection chosen by intuition said what it guaranteed against such side information. The chapter ends with k-anonymity, which coarsens records until each looks like several others, and shows that it fails when two releases are combined.

The second chapter turns these anecdotes into mathematics. It asks how much an analyst can learn from counting queries about a dataset of n secret bits when every answer may be wrong by a bounded amount. If all possible queries are answered, errors must be of order n to stop the analyst from rebuilding almost every bit. If only about n random queries are answered, errors of order $\sqrt{n}$ are needed, and an efficient attack based on linear programming does the rebuilding. The same attack broke a deployed query system whose answers were too accurate.

Together the chapters set the requirements for a definition of privacy. It must hold against any side information, apply to any kind of output, and degrade gracefully as more is released. It must also allow enough error, which the reconstruction attacks show cannot be avoided. Differential privacy, defined in the next part, is built to meet these requirements.

Table 1.: The chapters of this part.

Chapter	Question it answers
Attempts at Data Privacy	Why do pseudonyms, removing identifiers, releasing models or statistics, and k-anonymity fail to protect people?
Reconstruction Attacks	How much error must answers to counting queries have before an analyst cannot rebuild the data?

1. Attempts at Data Privacy

Before defining privacy, it helps to see how hard it is to get right. This chapter goes through five releases of sensitive data whose owners believed them safe. They are a pseudonymized table of taxi trips, a table of film ratings, a language model, the summary statistics of genetic studies and hospital records with names removed. Each was undone by an attacker who combined the release with information from somewhere else.

The cases share a pattern. A protection was chosen by intuition, and nobody could say exactly what it guaranteed or against whom. The last section looks at k-anonymity, the first formal answer to these attacks, which asks every record to look like at least $k-1$ others. It fails in the same way. The next chapter makes the attacks quantitative.

1.1. Pseudonyms: the taxi trips

In 2014, a public-records request obtained from New York City's taxi regulator the record of every taxi trip and fare in the city in 2013, about 19 gigabytes of data. A trip record holds the pickup and drop-off times, their coordinates, the distance, the duration and the number of passengers. It also holds two identifiers, the number of the cab's medallion and the number of the driver's licence. These two would reveal where each driver worked and how much they earned, so the regulator did not publish them as they were. It replaced each by a string of 32 hexadecimal digits.

Replacing each identifier by a code in this way is called pseudonymization. The definition below names the one property such a code must have, that the same identifier always gets the same code, because the attack depends on exactly that.

Definition 1.1 (Pseudonymization)**.** Let $\mathcal{U}$ be a set of identifiers, such as licence numbers. A **pseudonymization** replaces each identifier $u \in \mathcal{U}$ in a table by $h(u)$, where h is a fixed injective function, so that the same person gets the same pseudonym in every row.

The pseudonyms did not hold for long. An analyst noticed that one licence pseudonym earned far more per day than any driver could. That pseudonym turned out to be the MD5 hash of the string `0`, a placeholder for a missing licence, so many drivers' trips had been pooled under it. The discovery revealed the function h, which was MD5 with no secret key. A hash function is hard to invert on inputs drawn from a huge set, but the identifier formats allowed only about 22 million strings, about 3 million licence numbers and 19 million medallion numbers. With so few possible inputs, anyone can apply the function to all of them and look each pseudonym up, as the next result says.

Proposition 1.1 (Dictionary attack)**.** *Let the set of identifiers $\mathcal{U}$ be finite, and let the pseudonymization function h be injective and computable by anyone. Then $|\mathcal{U}|$ evaluations of h recover the identifier behind every pseudonym in the table.*

Proof.

1. Compute $h(u)$ for every $u \in \mathcal{U}$ and store the pairs $(h(u), u)$ in a table sorted by $h(u)$. This takes $|\mathcal{U}|$ evaluations of h.
2. For each pseudonym v in the released table, look up v. Since h is injective, exactly one stored pair has first entry v, and its second entry is the identifier.

□

A licence number is not yet a name. The second step joined the decoded table with a public list of licensed drivers, which gives each licence number its holder's name. Every driver's trips, home area and income were then public. Joining a release with a second table that names people is the step that turns every attack in this chapter into a disclosure, so it gets its own definition.

Definition 1.2 (Linkage attack)**.** Let T be a released table and T' an auxiliary table that names people, and let Q be a set of attributes that both tables contain. A **linkage attack** joins T and T' on Q. A row of T is **re-identified** when its values on Q match exactly one row of T' and no other row of T.

Replacing MD5 by random pseudonyms would stop the dictionary attack but not linkage. A passenger who takes one ride knows the time and place of the pickup and can read the driver's licence number in the cab. That single row of the table carries the driver's pseudonym, which then points to every other trip by the same driver. Passengers are exposed too. Someone who sees a colleague hail a cab outside the office at a known time can find that trip, and its drop-off point is likely the colleague's home. The fare table even shows the tip. The attacker needs no identifier at all, because the time and place of a pickup identify the trip.

1.2. Rich records: the Netflix Prize

From 2006 to 2009, Netflix ran a contest with a prize of one million dollars for the team that most improved its film recommendations. The winning entry blended many models, among them matrix factorization. To train their models, contestants received about 100 million ratings from about 480,000 subscribers. Each rating was a tuple of a user pseudonym, a film, a score and a date. Viewing histories are sensitive, since they can hint at a person's politics or sexuality, and in the United States the Video Privacy Protection Act of 1988 restricts their disclosure.

Narayanan and Shmatikov showed in 2008 that the pseudonyms protected little. Many subscribers also post reviews on the Internet Movie Database (IMDb) under their name or a stable handle. A linkage attack in the sense of Definition 1.2, with Q the films, scores and dates, matches the two sources. The match need not be exact. They scored each candidate pair of users by how many films both rated with similar scores on nearby dates, and they weighted rarely rated films more heavily. They reported that 8 ratings, 2 of which may be wrong, with dates known only to within 14 days, single out 99% of the records in the dataset.

The lesson is that in a sparse table with many columns, almost every record is unique, so no attribute is harmless. Even a handful of a person's ratings identifies them. The release also had concrete consequences, because a class action over it was settled out of court and Netflix cancelled a planned second contest.

1.3. Releasing a model: memorization

The first two releases were whole tables. Releasing only a function of the data seems safer, such as a model trained on it. Carlini and coauthors tested this in 2019 on language models, which assign a probability to every sequence of words. Such a model predicts each word from the ones before it, using a vector of parameters θ learned from a corpus of text. To test whether a model has memorized a string, the authors measured how surprised the model is by it, counted in bits. A string the model finds likely gets a small score.

Definition 1.3 (Log-perplexity)**.** Write $\Pr_\theta$ for the probability that the model with parameters θ assigns to a sequence of words, or to the next word given the words before it. The **log-perplexity** of a sequence $x_1, \dots, x_\ell$ of ℓ words under this model is

$$P_\theta(x_1, \dots, x_\ell) = -\log_2 \Pr_\theta[x_1, \dots, x_\ell] = \sum_{i=1}^{\ell} -\log_2 \Pr_\theta\left[x_i \mid x_1, \dots, x_{i-1}\right]. \tag{1.1}$$

The second form is the chain rule of probability. Low log-perplexity means high probability. A model of English gives "the cat sat on the mat" a low score and a string of unrelated words a high one. The question is whether a rare string that appears in the training text, such as a password or an identity number, gets a much lower score than it otherwise would. If it does, the model has memorized it, and anyone who can query the model may find it.

To measure this, the authors inserted a **canary** into the training text: a sentence of a fixed format with a random secret, such as "my identity number is 314-15-9265". They then compared its score with those of all the other sentences of the same format. Exposure turns that comparison into a number, which is large when the model ranks the true canary above most of the alternatives.

Definition 1.4 (Exposure)**.** Let $\mathcal{R}$ be the finite set of candidate secrets, all sentences of the canary's format, and let $N = |\mathcal{R}|$. Sort $\mathcal{R}$ by increasing log-perplexity under the model, and let $\operatorname{rank}(r) \in \{1, \dots, N\}$ be the position of r, so that rank 1 is the most likely candidate. The **exposure** of r is

$$\operatorname{exposure}(r) = \log_2 N - \log_2 \operatorname{rank}(r). \tag{1.2}$$

Exposure ranges from 0, for the least likely candidate, to $\log_2 N$, for the most likely one. A canary at the top of the list can be read off the model by trying the candidates in order. The next result says how much exposure to expect when the model has not seen the canary, which calibrates the measurement.

Proposition 1.2 (Exposure without memorization)**.** *Suppose the canary r is drawn uniformly from $\mathcal{R}$, independently of the model, and the ranking breaks ties in a fixed way. Then for every $t \geq 0$,*

$$\Pr[\operatorname{exposure}(r) \geq t] \leq 2^{-t}, \qquad \mathbb{E}[\operatorname{exposure}(r)] < \frac{1}{\ln 2} \approx 1.443. \tag{1.3}$$

Proof.

1. The ranking depends only on the model, and r is uniform and independent of it. So $\mathrm{rank}(r)$ is uniform on $\{1, \dots, N\}$.
2. The exposure is at least t exactly when $\mathrm{rank}(r) \le N2^{-t}$. By step 1 this has probability $\lfloor N2^{-t} \rfloor / N \le 2^{-t}$.
3. By step 1, the average exposure is $\log_2 N - \frac{1}{N}\sum_{k=1}^{N} \log_2 k = \frac{1}{N \ln 2}\,(N \ln N - \ln N!)$.
4. Since ln is increasing, $\ln k$ is at least the integral of $\ln x$ over $[k-1, k]$ for $k \ge 2$. Summing over k gives

$$\ln N! \ge \int_1^N \ln x \, \mathrm{d}x = N \ln N - N + 1. \tag{1.4}$$

5. Steps 3 and 4 give $\mathbb{E}[\mathrm{exposure}(r)] \le \frac{N-1}{N \ln 2} < \frac{1}{\ln 2}$.

□

So an exposure of t bits is as surprising as an event of probability 2^{-t} under the hypothesis that the model never saw the canary. With nine-digit secrets, $N = 10^9$, and the largest possible exposure is $\log_2 10^9 \approx 29.9$.

The experiments gave both answers. A translation model trained on about 100,000 pairs of sentences was given a canary with a nine-digit identity number, and after four insertions the canary reached the largest possible exposure. Gmail's Smart Compose, a model trained on billions of word sequences, was given canaries whose secret part was a few words from a vocabulary of tens of thousands, so N was of order 10^{12}. Even after 10,000 insertions their exposure reached only about 10, far from the roughly $\log_2 10^{12} \approx 39.9$ that extraction would need. Trying all N candidates is not needed to extract a secret. A shortest-path search over the digits, in the manner of Dijkstra's algorithm with $-\log_2$ of each conditional probability as the edge length, finds the most likely candidates first. Standard defences against overfitting, such as dropout and regularization, did not reduce exposure. Training with differential privacy, the subject of the next part, was the only method the authors found that did.

1.4. Releasing statistics: genomic studies

Medical and genetic studies need sensitive data and are expected to publish their findings. Aggregate statistics seemed a safe middle ground. Suppose a study enrols only people with a stigmatized disease. Then revealing that someone took part reveals their diagnosis, and a test that decides membership from the published statistics is an attack.

Homer and coauthors showed in 2008 that such a test exists for genetic data. They took a person's genotype, the published allele frequencies of a study group and those of a reference population. From these alone, they could tell whether the person's DNA was in the group, under conditions on the number of genetic markers and the size of the group. Others found that the test makes many more false positives in practice than this analysis suggests, because its threshold depends on how well the reference population matches the study group. Even so, the National Institutes of Health in the United States withdrew open access to aggregate genomic statistics, including

allele frequencies, odds ratios and p-values, and placed them behind an approval process. The attack cost open science. A method that releases such statistics with a guarantee would let them be public again.

1.5. Removing identifiers: k-anonymity

In the mid-1990s, the Group Insurance Commission of Massachusetts let researchers request the hospital records of every state employee. Each record held the patient's name, social security number, ZIP code, date of birth, sex and diagnosis. The commission removed the names and social security numbers, and the governor, William Weld, assured the public that patients' privacy was protected.

Latanya Sweeney, then a graduate student, bought the voter roll of Cambridge, Massachusetts, for $20. It lists every registered voter's name, address, ZIP code, date of birth and sex. A linkage attack on ZIP code, date of birth and sex found the governor's hospital records, and she mailed them to him. The attack works at scale. Sweeney later estimated that 87% of the population of the United States is the only person with their combination of five-digit ZIP code, sex and date of birth.

Her answer, with Samarati, was a property of the released table. Its attributes fall into three kinds. **Identifiers**, such as names, are removed. **Quasi-identifiers**, such as ZIP code, age and sex, are harmless alone but identify people in combination. **Sensitive attributes**, such as the diagnosis, are what the release is for, so they stay. The quasi-identifiers are coarsened, for example by replacing a ZIP code by its first three digits or an age by a range, until no combination is rare. The goal is that every person hides in a crowd of at least k rows that look the same on the quasi-identifiers, and the definition states that goal.

Definition 1.5 (k-anonymity)**.** Let $k \geq 1$ be an integer. A table is **k-anonymous** with respect to a set of quasi-identifiers if every combination of quasi-identifier values that appears in the table appears in at least k rows. The rows that share a combination form an **equivalence class**.

A k-anonymous table stops the attack on the governor: a linkage on the quasi-identifiers matches every person to at least k rows, not one. But the definition constrains only the quasi-identifiers. It says nothing about the sensitive values within a class, and nothing about other releases. Two examples show what goes wrong, using the two tables below, released by two hospitals that serve the same town.

Table 1.1.: Hospital A's release, 3-anonymous with respect to ZIP code and age.

ZIP code	Age	Diagnosis
021**	20–29	influenza
021**	20–29	asthma
021**	20–29	diabetes
021**	30–39	hepatitis
021**	30–39	hepatitis
021**	30–39	hepatitis
024**	40–59	influenza

ZIP code	Age	Diagnosis
024**	40–59	migraine
024**	40–59	asthma

Table 1.2.: Hospital B's release, 3-anonymous with respect to ZIP code and age.

ZIP code	Age	Diagnosis
02***	18–24	influenza
02***	18–24	asthma
02***	18–24	fracture
02***	25–34	diabetes
02***	25–34	migraine
02***	25–34	fracture

The first example uses only hospital A's table. It shows that a class whose rows all share one sensitive value hides nothing.

Example 1.1 (Homogeneity)**.** A neighbour aged 35, living in ZIP code 02139, was a patient of hospital A. In Table 1.1, the neighbour must be in the class with ZIP code 021** and age 30–39. All three of its rows say hepatitis, so the attacker learns the diagnosis without knowing which row is the neighbour's.

The second example needs classes with varied diagnoses, so homogeneity is not the problem. Instead, it combines the two hospitals' releases.

Example 1.2 (Intersection)**.** A colleague aged 27, living in ZIP code 02139, was treated at both hospitals for the same condition. In Table 1.1 the colleague's class allows influenza, asthma or diabetes. In Table 1.2 the class with age 25–34 allows diabetes, migraine or fracture. Only diabetes is in both lists, so the two releases together reveal the diagnosis, although each is 3-anonymous.

The second example is a failure of composition. Each hospital checked its own table, and neither could know what the other released. Ganta, Kasiviswanathan and Smith showed in 2008 that such attacks succeed in practice, and that they defeat refinements of k-anonymity as well.

The five cases point to three requirements for a definition. It must hold whatever side information an attacker has, since every attack above used some. It must apply to any output, a table, a model or a statistic, and not only to tables. And it must degrade gracefully when several releases are combined. Differential privacy, defined in the chapter on randomized response, meets all three.

1.6. Summary

Table 1.3.: The chapter's results.

Result	Statement
Definition 1.1	replace each identifier u by $h(u)$ for a fixed injective h
Proposition 1.1	a public h on a set of N identifiers is inverted by N evaluations
Definition 1.2	join a release with named auxiliary data on shared attributes
Definition 1.4	$\log_2 N - \log_2 \operatorname{rank}(r)$ for a canary among N candidates
Proposition 1.2	without memorization, exposure $\geq t$ has probability at most 2^{-t} and mean below $1/\ln 2$
Definition 1.5	every quasi-identifier combination appears in at least k rows
Example 1.1	a class with one sensitive value reveals it
Example 1.2	two k-anonymous releases together can reveal a value

1.7. References

- Vijay Pandurangan, "On Taxis and Rainbows", 2014. The inversion of the hashed medallion and licence numbers.
- Arvind Narayanan and Vitaly Shmatikov, "Robust De-anonymization of Large Sparse Datasets", 2008. The linkage of the Netflix Prize data with public reviews, and the statistics of how few ratings identify a record.
- Nicholas Carlini, Chang Liu, Úlfar Erlingsson, Jernej Kos and Dawn Song, "The Secret Sharer: Evaluating and Testing Unintended Memorization in Neural Networks", 2019. Canaries, exposure, the extraction search, the experiments on a translation model and on Smart Compose, and the experiments on defences.
- Nils Homer and coauthors, "Resolving Individuals Contributing Trace Amounts of DNA to Highly Complex Mixtures Using High-Density SNP Genotyping Microarrays", 2008. The membership test on aggregate genetic statistics.
- Rosemary Braun, William Rowe, Carl Schaefer, Jinghui Zhang and Kenneth Buetow, "Needles in the Haystack: Identifying Individuals Present in Pooled Genomic Data", 2009. A study showing that the test's false-positive rate is sensitive to its assumptions about the reference population.
- Latanya Sweeney, "Simple Demographics Often Identify People Uniquely", 2000. The estimate that ZIP code, sex and date of birth identify most of the population of the United States.
- Pierangela Samarati and Latanya Sweeney, "Generalizing Data to Provide Anonymity when Disclosing Information", 1998. The definition of k-anonymity.

- Srivatsava Ranjit Ganta, Shiva Prasad Kasiviswanathan and Adam Smith, "Composition Attacks and Auxiliary Information in Data Privacy", 2008. Intersection attacks on independently anonymized releases.

2. Reconstruction Attacks

The previous chapter's attacks were clever, but each exploited a particular release. This chapter asks a sharper question. When does answering counting questions about a dataset, each answer only slightly wrong, let an attacker rebuild almost the whole dataset? It does not yet define what it means for an algorithm to be private. It defines a failure so flagrant that nobody would call it private, and it shows that a wide class of algorithms fails in this way.

The answers are quantitative. Exact statistics give the data away, as a census example shows. If a curator answers every possible counting query, it must make errors of order n, the number of people, to avoid reconstruction. If it answers only about n random queries, errors of order $\sqrt{n}$ are needed. The last section shows the second attack breaking a deployed system.

2.1. Reconstruction from exact statistics

A census collects **microdata**, one record per person with fields such as age, sex and race, and publishes tables of aggregates for each small area, such as a city block. The tables give counts, medians and means for groups of residents. Cells about fewer than three people are suppressed, but suppression alone does not protect anyone, because tables overlap. Suppose a block has four residents of some group, and three of them are women. Then the group has exactly one man, and a published median age for the group's men would be his age.

Every published number is a constraint on the unknown microdata, and enough constraints leave only one possibility. The next example counts how many age combinations remain possible for three people after one line of a census table, which gives only their median and mean ages.

Example 2.1 (One line of a census table)**.** A block's table says that it has three men, with median age 35 and mean age 47. Call the youngest man's age u_1, the middle one's u_2 and the oldest one's u_3, so that $u_1 \le u_2 \le u_3$. Ages are whole numbers between 0 and 125.

1. Without the statistics, the number of such triples is the number of multisets of size 3 from 126 values, $\binom{128}{3} = 341{,}376$.
2. The median of three numbers is the middle one, so $u_2 = 35$. Then u_1 has 36 values and u_3 has 91, which leaves $36 \times 91 = 3{,}276$ triples.
3. The mean gives $u_1 + u_2 + u_3 = 141$, so $u_3 = 106 - u_1$. Each u_1 from 0 to 35 gives a u_3 between 71 and 106, and every such u_3 lies between u_2 and 125. So exactly 36 triples remain, from $(0, 35, 106)$ to $(35, 35, 71)$.

One line of the table has cut the possibilities by a factor of almost 10,000. The other lines, about the same people grouped differently, add further constraints. A reconstruction attack writes all of them down and searches for microdata that satisfies every one. This search is an integer program,

which is hard to solve in the worst case. In practice, though, solvers for integer programs and for Boolean satisfiability handle such instances quickly. For a fictional block of seven people, Garfinkel, Abowd and Martindale found that the constraints of its full table have exactly one solution, which a solver finds in a fraction of a second.

The attack scales. The United States Census Bureau ran it on its own published tables from the 2010 census. The reconstructed records matched the confidential ones exactly for 46% of the population, and for 71% if ages could be off by one year. Linking the reconstruction with commercial databases re-identified more than 50 million people by name. In response, the Bureau adopted differential privacy for the 2020 census.

2.2. Subset queries and blatant non-privacy

To analyse such attacks, strip the setting to its core. There are n people. Each has public identifying attributes, such as a name and a postal code, and one secret bit, such as whether they have a disease. The secret bits form the vector $d = (d_1, \dots, d_n) \in \{0,1\}^n$. An analyst may ask how many of the people who satisfy a condition on the identifiers have secret bit 1. A condition such as "the name is Alice, Bob or Eve" can single out any set of people. So a query is a subset $S \subseteq [n] = \{1, \dots, n\}$, which we identify with its indicator vector $S \in \{0,1\}^n$. Its entry i is 1 when person i is in the set and 0 otherwise.

A curator holds the secret bits and answers such queries. The definition assumes only one thing about it, a bound E on how far each answer may be from the truth.

Definition 2.1 (Subset queries and a curator with error E)**.** Let $d \in \{0,1\}^n$ be the vector of secret bits. The true answer to a **subset query** $S \in \{0,1\}^n$ is the number of people in S whose bit is 1,

$$a(S) = \sum_{i \in S} d_i = \langle S, d \rangle, \tag{2.1}$$

where $\langle S, d \rangle = \sum_i S_i d_i$ is the inner product of the indicator vector S with d.

A **curator with error** E answers each query S with a response $r(S)$, a number within E of the truth, $|r(S) - a(S)| \leq E$. Within that bound, the curator may choose its answers in any way, depending on d and on all the queries.

The error is not assumed to be random. The curator may pick the most misleading answer within the allowed range, so every attack below works against random noise too. Some error is needed, since for each person i the exact answer to the query $S = \{i\}$ is that person's secret bit.

To measure how much an attacker learns, count the bits it gets wrong. The attacker outputs a guess $c \in \{0,1\}^n$ of the secret vector d. The **Hamming distance** $\mathrm{dist}(c, d) = |\{i : c_i \neq d_i\}|$ is the number of positions where the guess is wrong.

The next definition names the clearest possible failure of privacy. Some attacker gets all but a vanishing fraction of the bits right, for every secret vector, except with a vanishing probability.

Definition 2.2 (Blatant non-privacy)**.** Consider a family of curators, one for each number of people n. The family is **blatantly non-private** if three things exist. The first is an attacker that interacts with the curator and outputs a guess $c \in \{0,1\}^n$. The second is a bound $g(n)$ on the number of wrong bits, with $g(n)/n \to 0$. The third is a failure probability $\beta(n)$, with $\beta(n) \to 0$. They must satisfy, for every n and every secret vector $d \in \{0,1\}^n$,

$$\Pr\left[\text{dist}(c,d) \leq g(n)\right] \geq 1 - \beta(n). \tag{2.2}$$

An attacker that gets all but a vanishing fraction of the bits right, written $o(n)$ wrong bits, has learned the secret of almost everyone. Whatever privacy means, this is a violation of it.

2.3. Asking every query

The first attack ignores cost. It asks all 2^n subset queries and then outputs any candidate that agrees with every answer.

Theorem 2.1 (Reconstruction from all subset queries)**.** *Let $E \geq 0$, and let a curator with error E answer all 2^n subset queries, giving the answer $r(S)$ to the query S. An attacker that outputs any $c \in \{0,1\}^n$ satisfying*

$$|\langle S, c\rangle - r(S)| \leq E \quad \textit{for every } S \subseteq [n] \tag{2.3}$$

finds such a c, and every such c differs from the secret vector d in at most $4E$ positions, $\text{dist}(c,d) \leq 4E$.

Proof.

1. The true vector d satisfies Equation 2.3 by Definition 2.1, so a candidate exists.
2. Let $I_0 = \{i : d_i = 0\}$. The attacker asked the query I_0, so Equation 2.3 gives $\langle I_0, c\rangle \leq r(I_0) + E$. The curator's bound gives $r(I_0) \leq \langle I_0, d\rangle + E = E$. Hence $\langle I_0, c\rangle \leq 2E$. This sum counts the positions of I_0 where $c_i = 1 \neq d_i$, so c and d differ in at most $2E$ positions of I_0.
3. Let $I_1 = \{i : d_i = 1\}$. In the same way, $\langle I_1, c\rangle \geq r(I_1) - E \geq \langle I_1, d\rangle - 2E = |I_1| - 2E$. The number of positions of I_1 where $c_i = 0$ is $|I_1| - \langle I_1, c\rangle \leq 2E$.
4. The sets I_0 and I_1 partition $[n]$, so $\text{dist}(c,d) \leq 4E$.

□

The attacker does not know I_0 and I_1, which is why it asks every query. The bound is already strong for large errors, and with errors that are a vanishing fraction of n it meets Definition 2.2.

Corollary 2.1. *If $E = n/401$, the attacker of Theorem 2.1 recovers more than 99% of the bits. If $E = E(n)$ with $E(n)/n \to 0$, the curator is blatantly non-private, with $g(n) = 4E(n)$ and $\beta(n) = 0$.*

Proof. For the first claim, $4n/401 < n/100$, since $4 \times 100 < 401$. The second claim is Theorem 2.1 with Definition 2.2, because the attacker never fails. □

Errors of order n are the most a counting query can bear, since every answer lies between 0 and n. The next result shows that this is also what it takes to resist the attack. With error E, no attacker can guarantee fewer than about $E/2$ wrong bits, however many queries it asks.

Proposition 2.1 (Error E hides $E/2$ bits)**.** *Let $E \geq 1$ and $k = \min(\lfloor E \rfloor, n)$. For every attacker, which may ask any queries in any adaptive order, there are a vector $d \in \{0,1\}^n$ and a curator with error E for which* $\Pr[\text{dist}(c, d) \geq k/2] \geq 1/2$.

Proof.

1. Let d be all zeros, and let d' have ones in its first k positions and zeros elsewhere, so $\text{dist}(d, d') = k$.
2. On the data d', let the curator answer every query exactly, $r(S) = \langle S, d' \rangle$. On the data d, let it give the same answers. Their error is $|\langle S, d' - d \rangle| \leq k \leq E$, so both are curators with error E.
3. The answers to every query are the same on d and on d'. So the attacker's output c has the same distribution in the two cases.
4. For every c, the triangle inequality gives $\text{dist}(c, d) + \text{dist}(c, d') \geq k$, so at least one of the two distances is at least $k/2$. Hence $\Pr[\text{dist}(c, d) \geq k/2] + \Pr[\text{dist}(c, d') \geq k/2] \geq 1$, and one of the two probabilities is at least $1/2$.

□

So for the attacker who asks everything, the threshold between blatant non-privacy and safety is error of order n, and Theorem 2.1 is tight up to a factor of 8. But an attack with 2^n queries is not a practical threat once n is more than a few dozen.

2.4. Random queries

The second attack asks about n queries, each a uniformly random subset of the people. It works against much smaller errors than the first, of order $\sqrt{n}$, which is still large compared with the error a useful statistic can tolerate.

The idea is to fix a wrong candidate c and show that a random query is likely to expose it. If c differs from d in k positions, the answers $\langle S, c \rangle$ and $\langle S, d \rangle$ differ by a sum of about $k/2$ random signs, which is typically of size $\sqrt{k}$. A curator with error E can hide a difference of at most $2E$. So when $\sqrt{k}$ is much larger than E, a random query rules c out with constant probability. To make "typically" precise, we need a lower bound on the size of a random sum, the opposite of a concentration inequality.

Lemma 2.1 (A random subset sees a fixed vector)**.** *Let $w \in \mathbb{R}^n$, with Euclidean length $\|w\|_2 = (\sum_i w_i^2)^{1/2}$. Let S be a uniformly random subset of $[n]$, so that each i is in S independently with probability $1/2$. Then the sum $\langle S, w\rangle$ of the entries of w in S satisfies*

$$\Pr\left[\,|\langle S, w\rangle| \geq \tfrac{1}{4}\|w\|_2\,\right] \geq \frac{3}{32}. \tag{2.4}$$

Proof.

1. If $w = 0$ the event always holds, so let $w \neq 0$. Write $S_i = (1 + \xi_i)/2$, where the signs $\xi_i \in \{-1, +1\}$ are independent and uniform. Then $\langle S, w\rangle = \frac{1}{2}(a + Z)$, with the constant $a = \sum_i w_i$ and the random sum $Z = \sum_i \xi_i w_i$.
2. *Moments of Z.* Expand the powers and use independence. A product of the ξ_i has mean 0 unless each index appears an even number of times. So $\mathbb{E}[Z^2] = \sum_i w_i^2 = \|w\|_2^2$, and

$$\mathbb{E}[Z^4] = \sum_i w_i^4 + 3\sum_{i\neq j} w_i^2 w_j^2 = 3\|w\|_2^4 - 2\sum_i w_i^4 \leq 3\|w\|_2^4. \tag{2.5}$$

3. *Paley–Zygmund inequality.* Let $Y \geq 0$ have $\mathbb{E}[Y] > 0$, and let $s \in [0, 1]$. Split $\mathbb{E}[Y]$ according to whether $Y > s\,\mathbb{E}[Y]$, and apply the Cauchy–Schwarz inequality to the second part:

$$\mathbb{E}[Y] \leq s\,\mathbb{E}[Y] + \mathbb{E}\left[Y\mathbf{1}[Y > s\,\mathbb{E}[Y]]\right] \leq s\,\mathbb{E}[Y] + \sqrt{\mathbb{E}[Y^2]\,\Pr[Y > s\,\mathbb{E}[Y]]}. \tag{2.6}$$

Rearranging gives $\Pr[Y > s\,\mathbb{E}[Y]] \geq (1 - s)^2\,\mathbb{E}[Y]^2/\mathbb{E}[Y^2]$.

4. Apply step 3 to $Y = Z^2$ with $s = 1/4$, using step 2. The event $|Z| > \|w\|_2/2$ has probability at least $(3/4)^2/3 = 3/16$.
5. *Removing the constant a.* For every z, $|a+z|+|a-z| \geq 2|z|$, so $|a+Z| \geq |Z|$ or $|a-Z| \geq |Z|$. Flipping all signs maps Z to $-Z$ and preserves the distribution and the event of step 4. Hence the events $\{|Z| > \|w\|_2/2,\ |a + Z| \geq |Z|\}$ and $\{|Z| > \|w\|_2/2,\ |a - Z| \geq |Z|\}$ have the same probability. Together they cover the event of step 4, so each has probability at least $3/32$.
6. On the first of these events, $|\langle S, w\rangle| = \frac{1}{2}|a + Z| \geq \frac{1}{2}|Z| > \frac{1}{4}\|w\|_2$.

□

With the lemma, the union bound over all candidates gives the attack. The candidates are the 2^n vectors in $\{0, 1\}^n$, so about n queries suffice to rule out all the wrong ones at once.

Theorem 2.2 (Reconstruction from random queries)**.** *Let $E \geq 0$ and $\beta \in (0, 1)$, and let*

$$m \geq \frac{32}{3}\left(n \ln 2 + \ln \frac{1}{\beta}\right). \tag{2.7}$$

The attacker draws m independent, uniformly random subsets $S_1, \ldots, S_m$ of $[n]$ before seeing any answer. A curator with error E answers them with $r_1, \ldots, r_m$. The attacker outputs any

$c \in \{0,1\}^n$ with $|\langle S_j, c\rangle - r_j| \le E$ for all j. Then such a c exists and, with probability at least $1-\beta$ over the queries, every such c has

$$\operatorname{dist}(c,d) \le 64E^2. \tag{2.8}$$

Proof.

1. The true vector d satisfies all m constraints by Definition 2.1, so a candidate exists.
2. *One far candidate.* Fix c with $k = \operatorname{dist}(c,d) > 64E^2$, and let $w = c - d$. Its entries are in $\{-1,0,1\}$ and exactly k are nonzero, so $\|w\|_2 = \sqrt{k} > 8E$. By Lemma 2.1, each query has $|\langle S_j, w\rangle| \ge \sqrt{k}/4 > 2E$ with probability at least $3/32$. The queries are independent, so the probability that $|\langle S_j, w\rangle| \le 2E$ for every j is at most $(1-3/32)^m \le e^{-3m/32}$.
3. *All far candidates.* There are at most 2^n candidates. By the union bound, the event B that some candidate with $\operatorname{dist}(c,d) > 64E^2$ has $|\langle S_j, c-d\rangle| \le 2E$ for every j has probability at most $2^n e^{-3m/32}$, which is at most β by the choice of m.
4. *Consistent candidates are close.* Suppose B does not occur. If c is consistent with the answers, the triangle inequality gives $|\langle S_j, c-d\rangle| \le |\langle S_j, c\rangle - r_j| + |r_j - \langle S_j, d\rangle| \le 2E$ for every j. So c is not far, which is Equation 2.8. The event B depends only on the queries, so this holds whatever answers the curator chose.

□

The number of queries is about $7.4n$: for $n = 1{,}000$ and $\beta = 0.01$ the theorem asks for $m = 7{,}443$. Setting $\beta = 2^{-n}$ gives the asymptotic statement.

Corollary 2.2. *If $E = E(n)$ with $E(n)/\sqrt{n} \to 0$, the curator is blatantly non-private against $15n$ random queries, with $g(n) = 64E(n)^2$ and $\beta(n) = 2^{-n}$. If $E = \alpha\sqrt{n}$, the same attacker gets at most $64\alpha^2 n$ bits wrong with probability at least $1 - 2^{-n}$.*

Proof. With $\beta = 2^{-n}$, the bound on m in Theorem 2.2 is $\frac{32}{3} \cdot 2n \ln 2 \approx 14.79\, n \le 15n$. Then $g(n) = 64E(n)^2 = 64n\,(E(n)/\sqrt{n})^2$, and $g(n)/n \to 0$. □

2.4.1. An efficient attack

The attacker of Theorem 2.2 asks few queries but searches 2^n candidates. Relaxing the search to fractional vectors makes it a linear program, so the attack is called linear programming (LP) decoding. It looks for a vector of fractions $x \in [0,1]^n$ with $|\langle S_j, x\rangle - r_j| \le E$ for all j. A linear program with n variables and $2m + 2n$ inequalities is solvable in time polynomial in n and m. The attacker then rounds x to the nearest bit vector. The analysis is the same as before, with one extra step, because a fractional candidate is not one of 2^n points. A fine grid replaces the set of candidates.

Theorem 2.3 (Reconstruction by linear programming)**.** *Let $E > 0$, $\beta \in (0,1)$, $K = \lceil n/(2E) \rceil$, and*

$$m \geq \frac{32}{3}\left(n \ln(2K+1) + \ln\frac{1}{\beta}\right). \tag{2.9}$$

The attacker draws m random subsets and receives answers from a curator with error E, as in Theorem 2.2. It finds any $x \in [0,1]^n$ with $|\langle S_j, x\rangle - r_j| \leq E$ for all j, and it outputs $c_i = 1$ if $x_i \geq 1/2$ and $c_i = 0$ otherwise. With probability at least $1-\beta$ over the queries,

$$\operatorname{dist}(c,d) \leq 676E^2. \tag{2.10}$$

Proof.

1. The vector d is feasible, so the linear program has a solution x. Let $z = x - d \in [-1,1]^n$. As in step 4 of Theorem 2.2, $|\langle S_j, z\rangle| \leq 2E$ for every j.
2. *Rounding errors.* If $c_i \neq d_i$, then either $d_i = 0$ and $x_i \geq 1/2$, or $d_i = 1$ and $x_i < 1/2$. In both cases $|z_i| \geq 1/2$, so $(2z_i)^2 \geq 1$. Summing over i gives $\operatorname{dist}(c,d) \leq 4\|z\|_2^2$.
3. *The grid.* Let $h = 1/K \leq 2E/n$, and let G be the set of vectors whose entries are multiples of h in $[-1,1]$. It has $(2K+1)^n$ points. Round each entry of z to the nearest multiple of h to get $z' \in G$, with $|z_i - z'_i| \leq h/2 \leq E/n$. Then $|\langle S_j, z - z'\rangle| \leq \sum_i |z_i - z'_i| \leq E$, so $|\langle S_j, z'\rangle| \leq 3E$ for every j. Also $\|z - z'\|_2 \leq \sqrt{n} \cdot E/n \leq E$.
4. *Far grid points are ruled out.* Fix $g \in G$ with $\|g\|_2 > 12E$. By Lemma 2.1, each query has $|\langle S_j, g\rangle| \geq \|g\|_2/4 > 3E$ with probability at least $3/32$. As in Theorem 2.2, the probability that $|\langle S_j, g\rangle| \leq 3E$ for every j is at most $e^{-3m/32}$. A union bound over the $(2K+1)^n$ grid points bounds the probability that this happens for some far grid point by $(2K+1)^n e^{-3m/32} \leq \beta$.
5. Outside that event, step 3 shows that z' is not far, so $\|z'\|_2 \leq 12E$ and $\|z\|_2 \leq 13E$. Step 2 gives $\operatorname{dist}(c,d) \leq 4 \cdot 169E^2 = 676E^2$.

□

For $1 \leq E \leq n$, the factor $\ln(2K+1)$ is at most $\ln(n+3)$, so the attack runs in polynomial time with $O(n \log n)$ queries, and it gets at most $676E^2$ bits wrong. Error $o(\sqrt{n})$ is again blatantly non-private. The constants 64 and 676 are far from the truth. Figure 2.1 shows the attack on simulated data, where it recovers almost every bit at errors that the theorem cannot handle.

A sharper analysis removes the logarithm. It bounds the smallest singular value of the random query matrix, which shows directly that every fractional z with small answers on all m queries is short. Then $O(n)$ queries suffice for the linear program too (Dwork and Roth, 2014, Section 8.1). A linear program that minimizes the total error instead of bounding each error even tolerates a constant fraction of answers with arbitrary error (Dwork, McSherry and Talwar, 2007). That proof needs bounds on random matrices that this book does not develop. The simpler, inefficient version of the tolerance to arbitrary answers follows from the argument above.

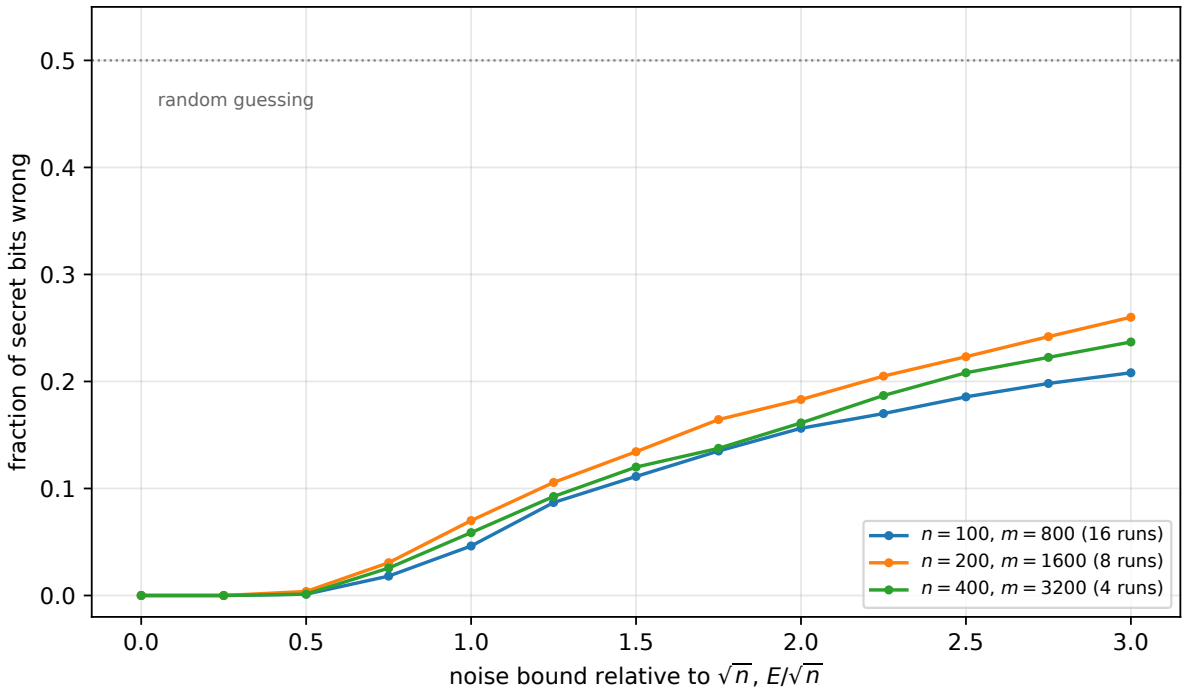


Figure 2.1.: The linear-programming attack on simulated data. For each n, the attacker asks $m = 8n$ random subset queries. The curator adds to each true answer an error drawn uniformly from $[-E, E]$. The attacker finds a point of the feasible set of Theorem 2.3, up to a violation of at most 0.01, by projected gradient descent and rounds it. Each point averages $1600/n$ runs, and within a run all values of E share the data, the queries and the uniform draws. The curves for different n nearly coincide, so the error the attack tolerates grows like $\sqrt{n}$. For $n = 400$, a single run gets none of the bits wrong at $E = \sqrt{n}/2$ and 17% of them at $E = 2\sqrt{n}$.

Proposition 2.2 (Tolerating some arbitrary answers)**.** *Let $E \geq 0$, let $\beta \in (0,1)$ be a failure probability, let $\rho \in [0, 3/128)$ be the fraction of answers that may be arbitrary, and let*

$$m \geq \frac{2048}{9}\left(n \ln 2 + \ln \frac{1}{\beta}\right). \tag{2.11}$$

The attacker draws m random subsets as in Theorem 2.2. The curator answers at least $(1-\rho)m$ of them with error at most E and the others arbitrarily, choosing which after seeing the queries. The attacker outputs any $c \in \{0,1\}^n$ with $|\langle S_j, c\rangle - r_j| \leq E$ for at least $(1-\rho)m$ of the queries. With probability at least $1-\beta$, every such c has $\operatorname{dist}(c,d) \leq 64E^2$.

Proof.

1. The true vector d meets the condition, so a candidate exists.
2. Fix c with $\operatorname{dist}(c,d) > 64E^2$ and let $V_j = \mathbf{1}[|\langle S_j, c-d\rangle| > 2E]$ indicate that query j exposes c. As in step 2 of Theorem 2.2, the V_j are independent with mean at least $3/32$. Hoeffding's inequality gives $\Pr\left[\sum_j V_j < 3m/64\right] \leq \exp(-2m(3/64)^2) = e^{-9m/2048}$.
3. A union bound over at most 2^n candidates shows that, with probability at least $1 - 2^n e^{-9m/2048} \geq 1-\beta$, every far candidate has $\sum_j V_j \geq 3m/64$.
4. Suppose this holds, and let c be an output. The queries where c is consistent and those answered with error at most E each number at least $(1-\rho)m$, so at least $(1-2\rho)m$ queries are of both kinds. On each of them, the triangle inequality gives $|\langle S_j, c-d\rangle| \leq 2E$, so $V_j = 0$. Hence $\sum_j V_j \leq 2\rho m < 3m/64$, and c is not far.

□

These attacks are nearly tight. The chapters on approximate differential privacy and advanced composition show that a curator can answer about n subset queries with error of order $\sqrt{n}$ and still meet the definition of differential privacy. The bound holds up to logarithmic factors and factors that depend on the privacy parameters. With m queries, error of order $\sqrt{m}$ suffices. So the line between blatant non-privacy and a formal guarantee runs at error $\sqrt{n}$ for n queries.

2.5. Reconstruction in practice

The random-query attack looks theoretical, since it asks questions about arbitrary sets of people. A deployed system showed otherwise. The company Aircloak built Diffix, a query interface that let analysts ask an unlimited number of counting queries on sensitive data. It promised privacy with less error than differential privacy would add. From December 2017 to May 2018, Aircloak ran a bounty program with prizes of up to $5,000 for a successful attack.

Diffix answered counting queries, which include subset queries, but with defences against the attacks above. Each condition in a query added its own layer of Gaussian noise with standard deviation 1, so the error of an answer had standard deviation equal to the square root of the number of conditions. The interface rejected queries that combined conditions with OR, and it suppressed small counts and altered extreme values. A random subset of n people needs about

$n/2$ conditions to describe person by person, and those conditions must be joined by OR. Both defences block that route.

Cohen and Nissim reversed the order of the attack. Instead of drawing a random set and then describing it, they drew a random description from a family of single conditions, each of which selects a set that looks random. Their conditions test a digit of an arithmetic expression in each person's numeric identifier, raised to a power. For example, take a multiplier 3 and an exponent $1/2$, and keep a person when the second decimal digit of $(3 \cdot \text{identifier})^{1/2}$ is even. The person with identifier 7 is kept, because $\sqrt{21} = 4.5825\ldots$ has second decimal digit 8. One of the queries they sent kept a row when

$$\lfloor 100\,(2 \cdot \text{identifier})^{0.7} + 0.5 \rfloor = \lfloor 100\,(2 \cdot \text{identifier})^{0.7} \rfloor\,, \tag{2.12}$$

which holds exactly when the third decimal digit of $(2 \cdot \text{identifier})^{0.7}$ is at most 4. For identifier 2000 that number is $332.232\ldots$, which is kept, and for 2001 it is $332.348\ldots$, which is not. Of the 1,001 identifiers from 2000 to 3000, the condition keeps 501.

Each such query has a fixed, small number of conditions, so its error does not grow with the number of rows, and Diffix added noise of standard deviation 4 to these queries. Cohen and Nissim targeted four ranges of identifiers that held between 73 and 142 rows, and they asked 3,500 queries about each. Gaussian noise is not bounded, so they used the linear program that minimizes the total error, which tolerates a few answers with large error in the spirit of Proposition 2.2. It recovered the secret attribute of every row in all four ranges. Aircloak then restricted the query language to block the attack.

The theorems of this chapter do not apply to these queries as stated, because a fixed formula is not a random subset. Their proofs use the randomness of the queries only through Lemma 2.1, by which each query exposes a fixed wrong candidate with constant probability. The digit tests behave that way on real identifiers, which is an empirical fact and not a theorem. That was enough to break the system.

2.6. Summary

Table 2.1.: The chapter's results.

Result	Statement
Example 2.1	a median and a mean cut 341,376 possible ages to 36
Definition 2.2	an attacker recovers all but $o(n)$ secret bits with probability tending to 1
Theorem 2.1	all 2^n queries with error E reveal all but $4E$ bits
Proposition 2.1	error E can hide $E/2$ bits from any attacker
Lemma 2.1	a random subset has $\lvert\langle S, w\rangle\rvert \geq \lVert w\rVert_2/4$ with probability at least $3/32$

Result	Statement
Theorem 2.2	$\frac{32}{3}(n \ln 2 + \ln \frac{1}{\beta})$ random queries with error E reveal all but $64E^2$ bits
Theorem 2.3	a linear program does the same in polynomial time with $O(n \log n)$ queries, all but $676E^2$ bits
Proposition 2.2	a fraction below 3/128 of arbitrary answers does not stop the attack

2.7. References

- Simson Garfinkel, John M. Abowd and Christian Martindale, "Understanding Database Reconstruction Attacks on Public Data", 2019. Reconstruction of census blocks from published tables with a satisfiability solver.
- John M. Abowd, "Staring Down the Database Reconstruction Theorem", 2019. The Census Bureau's reconstruction and re-identification experiment on the 2010 census.
- Irit Dinur and Kobbi Nissim, "Revealing Information while Preserving Privacy", 2003. Blatant non-privacy, the attack with all subset queries and the efficient attack with random queries.
- Cynthia Dwork, Frank McSherry and Kunal Talwar, "The Price of Privacy and the Limits of LP Decoding", 2007. The linear-programming attack with $O(n)$ queries, robust to a constant fraction of arbitrary answers.
- Cynthia Dwork and Aaron Roth, *The Algorithmic Foundations of Differential Privacy*, 2014. Section 8.1 presents both reconstruction attacks.
- Raymond Paley and Antoni Zygmund, "A Note on Analytic Functions in the Unit Circle", 1932. The inequality used in Lemma 2.1.
- Aloni Cohen and Kobbi Nissim, "Linear Program Reconstruction in Practice", 2020. The attack on Diffix with pseudo-random digit-test queries, and simulations of the attack.

Part II.

The Basics of Differential Privacy

The first part showed that releasing statistics about people can expose them, even when the statistics look harmless. This part gives the answer that the rest of the book builds on. It states a definition of privacy, differential privacy (DP), and develops the first algorithms that satisfy it and the rules for combining them.

The definition comes first, from the oldest private algorithm. In randomized response, each person flips their sensitive bit with some probability before reporting it, so every report is deniable. Differential privacy turns that deniability into a bound. Changing one person's data may change the probability of any outcome by at most a factor e^{ε}, where the privacy parameter ε is small.

The next two chapters give the basic tools for releasing numbers. The Laplace mechanism adds noise scaled to how much one person can change a statistic, and on a dataset of n people it estimates a fraction with error of order $1/(\varepsilon n)$ instead of the $1/(\varepsilon\sqrt{n})$ of randomized response. The same chapter proves that privacy survives post-processing, degrades gracefully for groups and adds up under composition. The third chapter relaxes the definition to approximate DP, which allows a small probability δ of failure. Its examples show that δ must be much smaller than $1/n$, and in exchange the relaxation admits the Gaussian mechanism, which is more accurate in high dimension.

The last two chapters go further. Advanced composition shows that the privacy cost of k analyses grows like $\sqrt{k}$ rather than k, at the price of a small extra δ. The exponential mechanism handles tasks whose answer is an object rather than a number, such as a price or a classifier, by sampling objects with probability that grows exponentially with their quality.

Every chapter of the later parts uses these results. Each chapter here states them with exact constants and proves them in full.

Table 2.2.: The chapters of this part and the questions they answer.

Chapter	Question it answers
Randomized response and the definition	How can one collect a sensitive bit from each person, and what does it mean for an algorithm to be private?
The Laplace mechanism	How much noise makes a numeric statistic private, and how do private algorithms combine?
Approximate differential privacy	What does a small additive slack δ allow, how small must it be, and what does the Gaussian mechanism gain?
Advanced composition	Why does the privacy cost of k analyses grow only like $\sqrt{k}$?
The exponential mechanism	How can one privately choose the best object from a set?

3. Randomized Response and the Definition

A group of people each hold one sensitive bit, and an analyst wants the fraction of ones without learning anyone's bit. This chapter solves that problem with a biased coin, the oldest private algorithm, and then turns the idea behind it, plausible deniability, into a definition. The definition is differential privacy (DP), and the rest of the book is built on it.

3.1. Randomized response

Write n for the number of people. Person number i holds a bit $x_i \in \{0, 1\}$, for instance whether they cheated on an exam, and the analyst wants the fraction of ones,

$$p = \frac{1}{n} \sum_{i=1}^{n} x_i. \tag{3.1}$$

Each person sends the analyst one message Y_i, which may depend on x_i and on coins that only person i sees. Two extreme protocols frame the problem. If $Y_i = x_i$ with probability 1, the analyst computes p exactly but sees every bit, so there is no privacy. If $Y_i = x_i$ with probability $1/2$ and $1 - x_i$ otherwise, then Y_i is a fair coin independent of x_i. The analyst learns nothing about anyone, and nothing about p either, since the average of the reports is distributed as $\frac{1}{n}\mathrm{Bin}(n, \frac{1}{2})$ whatever the data. Randomized response interpolates between the two.

Definition 3.1 (Randomized response)**.** For a bias $\kappa \in [0, 1/2]$, which sets the probability $\frac{1}{2} + \kappa$ of reporting the truth, person i reports

$$Y_i = \begin{cases} x_i & \text{with probability } \frac{1}{2} + \kappa, \\ 1 - x_i & \text{with probability } \frac{1}{2} - \kappa, \end{cases} \tag{3.2}$$

with fresh, independent coins for each person.

The value $\kappa = 1/2$ is the honest protocol and $\kappa = 0$ is the fair coin. In between, say at $\kappa = 1/4$, a person who reports $Y_i = 1$ can claim that their bit was 0 and the coin flipped it, which happens with probability $1/4$. That plausible deniability is the privacy, and it grows as κ shrinks.

The reports are biased towards the truth, so their average is not p. The bias is known, though, and the analyst can undo it. The average of one report is

$$\mathbb{E}[Y_i] = \left(\tfrac{1}{2} + \kappa\right) x_i + \left(\tfrac{1}{2} - \kappa\right)(1 - x_i) = 2\kappa x_i + \tfrac{1}{2} - \kappa, \tag{3.3}$$

which is linear in x_i. Solving for x_i gives an unbiased estimate of each bit, and averaging those estimates gives the estimator of p.

Definition 3.2 (Debiased estimator)**.** For $\kappa > 0$, the analyst outputs

$$\tilde{p} = \frac{1}{n} \sum_{i=1}^{n} \frac{Y_i - \frac{1}{2} + \kappa}{2\kappa}. \tag{3.4}$$

The next theorem checks that the estimator is right on average and measures how far it typically strays, through its variance.

Theorem 3.1. *For every $\kappa \in (0, 1/2]$ and every input $D \in \{0,1\}^n$, the estimator is unbiased, $\mathbb{E}[\tilde{p}] = p$, and its variance is*

$$\mathrm{Var}(\tilde{p}) = \frac{1 - 4\kappa^2}{16\kappa^2 n} \leq \frac{1}{16\kappa^2 n}. \tag{3.5}$$

Proof.

1. Each term $Z_i = (Y_i - \frac{1}{2} + \kappa)/(2\kappa)$ has mean x_i, by the formula for $\mathbb{E}[Y_i]$ above. Averaging over i and using linearity of expectation gives $\mathbb{E}[\tilde{p}] = p$.
2. The constants $\frac{1}{2} - \kappa$ and 2κ do not change the variance except through the scaling, so $\mathrm{Var}(Z_i) = \mathrm{Var}(Y_i)/(4\kappa^2)$.
3. The report Y_i is a Bernoulli variable that equals 1 with probability $\frac{1}{2} + \kappa$ or $\frac{1}{2} - \kappa$, depending on x_i. In both cases its variance is $(\frac{1}{2} + \kappa)(\frac{1}{2} - \kappa) = \frac{1}{4} - \kappa^2$.
4. The reports are independent because each person uses their own coins. So the variance of the sum is the sum of the variances, and

$$\mathrm{Var}(\tilde{p}) = \frac{1}{n^2} \sum_{i=1}^{n} \frac{\mathrm{Var}(Y_i)}{4\kappa^2} = \frac{\frac{1}{4} - \kappa^2}{4\kappa^2 n} = \frac{1 - 4\kappa^2}{16\kappa^2 n}. \tag{3.6}$$

□

Chebyshev's inequality turns the variance into an error bar: with probability at least 3/4, the error is at most two standard deviations, $|\tilde{p} - p| \leq 1/(2\kappa\sqrt{n})$. For a failure probability β that is small, Hoeffding's inequality does better, because each Z_i takes only two values.

Theorem 3.2. *For every $\kappa \in (0, 1/2]$ and $\beta \in (0, 1)$, with probability at least $1 - \beta$ over the coins,*

$$|\tilde{p} - p| \leq \frac{1}{2\kappa} \sqrt{\frac{\ln(2/\beta)}{2n}}. \tag{3.7}$$

Proof.

1. The term Z_i takes the values $(\frac{1}{2} + \kappa)/(2\kappa)$ and $-(\frac{1}{2} - \kappa)/(2\kappa)$, so it lies in an interval of length $1/(2\kappa)$.

2. The Z_i are independent with mean x_i. Hoeffding's inequality for the average of n independent variables, each in an interval of length ℓ, gives

$$\Pr\left[\,|\tilde{p}-p|\geq t\,\right]\leq 2\exp\!\left(-\frac{2nt^2}{\ell^2}\right)=2\exp(-8\kappa^2nt^2)\,. \tag{3.8}$$

3. Setting the right-hand side equal to β and solving for t gives Equation 3.7.

□

So the error is of order $1/(\kappa\sqrt{n})$, and it goes to 0 as n grows. Read the other way, an error of at most α with probability $1-\beta$ needs

$$n\geq\frac{\ln(2/\beta)}{8\kappa^2\alpha^2} \tag{3.9}$$

people. A smaller κ hides each bit better and needs more people for the same accuracy. To say how much better a smaller κ hides a bit, we need a definition.

3.2. Differential privacy

In privacy, the first step is to fix the setting. Here it is the **central model**. Each of n people sends their data point to a **trusted curator**, whom they trust with it and trust no one else. The curator runs an algorithm M on the dataset $D=(x_1,\dots,x_n)\in\mathcal{X}^n$ and publishes the output. Everything the curator publishes crosses a trust barrier, and anyone on the far side may be an adversary. The guarantee is a property of the algorithm, which in this field is called a **mechanism**. It says that no single person's data has a large effect on the distribution of the output. To make "one person's data" precise, compare two datasets that differ only in one person's record.

Definition 3.3 (Neighbouring datasets)**.** Two datasets $D, D'\in\mathcal{X}^n$ are **neighbours**, written $D\sim D'$, if they differ in exactly one record. This is the **replace-one** relation.

The definition asks that the distribution of the output barely change between neighbours. On every set of outputs, the probability may grow by at most a factor e^{ε}, where the privacy parameter ε measures how much change is allowed.

Definition 3.4 (Differential privacy)**.** Let $\varepsilon\geq 0$. A randomized algorithm $M\colon\mathcal{X}^n\to\mathcal{Y}$, which maps a dataset to an output in a set $\mathcal{Y}$, is **ε-differentially private** (ε-DP) if for all neighbours $D\sim D'$ and all events $S\subseteq\mathcal{Y}$,

$$\Pr[M(D)\in S]\leq e^{\varepsilon}\,\Pr[M(D')\in S]. \tag{3.10}$$

The probability is over the coins of M only. The datasets are fixed.

This is sometimes called **pure** differential privacy, to set it apart from a relaxation in a later chapter. Four features of it matter later.

First, the definition is quantitative. At $\varepsilon = 0$ the output distribution is the same on every dataset, which is perfect privacy and allows no accuracy. Any algorithm that is not randomized is ∞-DP, which is no guarantee at all. Values of ε between about 0.1 and 5 are typical, smaller is more private, and a claim far outside that range deserves suspicion. Because the bound is e^{ε}, privacy weakens quickly once ε is above 1.

Second, it is a worst-case guarantee. It must hold for every pair of neighbours, however unlikely the data, and no distribution over datasets is assumed.

Third, it is symmetric. Swapping D and D' in Equation 3.10 gives the lower bound $\Pr[M(D) \in S] \geq e^{-\varepsilon} \Pr[M(D') \in S]$.

Fourth, it bounds a ratio, not a difference. For small ε the factor e^{ε} is close to $1 + \varepsilon$. The exponential form is chosen because factors multiply when guarantees are chained, $e^{\varepsilon_1} e^{\varepsilon_2} = e^{\varepsilon_1 + \varepsilon_2}$, which the next chapter uses for group privacy. An additive notion of closeness, such as a small total variation distance, sounds natural but gives no meaningful guarantee.

This notion of neighbours replaces one record, and differential privacy under it is also called **bounded** differential privacy. The **add/remove** relation, also called **unbounded** differential privacy, instead adds or removes one record. A replacement is a removal followed by an addition, so an algorithm that is ε-DP under add/remove neighbours is 2ε-DP under replace-one neighbours, by the chaining argument of the next chapter's group privacy. The converse can fail, because under add/remove neighbours the size n of the dataset is itself private, so an algorithm may not use it as a public number. The book uses replacement, because it is more convenient mathematically. In practice, two algorithms can be compared only under the same notion.

For discrete outputs it suffices to check single outcomes, because the probability of an event is a sum of the probabilities of its outcomes.

Lemma 3.1. *Let $\mathcal{Y}$ be countable. If $\Pr[M(D) = y] \leq e^{\varepsilon} \Pr[M(D') = y]$ for all neighbours $D \sim D'$ and all outputs $y \in \mathcal{Y}$, then M is ε-DP.*

Proof. For an event S, sum the pointwise inequality over the outputs $y \in S$:

$$\Pr[M(D) \in S] = \sum_{y \in S} \Pr[M(D) = y] \leq e^{\varepsilon} \sum_{y \in S} \Pr[M(D') = y] = e^{\varepsilon} \Pr[M(D') \in S]. \tag{3.11}$$

□

A first consequence of the definition is that privacy forces randomness.

Proposition 3.1. *For $\varepsilon < \infty$, a deterministic ε-DP algorithm is constant.*

Proof.

1. Let M be deterministic and $D \sim D'$. Take the event $S = \{M(D)\}$. The definition gives $1 = \Pr[M(D) \in S] \leq e^{\varepsilon} \Pr[M(D') \in S]$, so $\Pr[M(D') = M(D)] > 0$. Since $M(D')$ is a single value, $M(D') = M(D)$.

2. Any two datasets in $\mathcal{X}^n$ are joined by a path of at most n neighbours, changing one record at a time. By step 1, M takes the same value at every step of the path, so it is constant.

□

3.2.1. Hypothesis testing

The definition has an operational reading, from the adversary's side. An adversary knows that the dataset is one of two neighbours, D or D', and sees the output of M. It must decide between the **null hypothesis** and the **alternative**,

$$H_0 \colon \text{the dataset is } D, \qquad H_1 \colon \text{the dataset is } D'. \tag{3.12}$$

A test is a function $\phi \colon \mathcal{Y} \to [0,1]$, the probability of answering H_1 given the output, so randomized tests are allowed. Statistics writes the two error rates α and β, but this chapter already uses those letters for accuracy. The test's **false positive rate** FPR is the chance of answering H_1 when H_0 holds, and its **false negative rate** FNR is the chance of answering H_0 when H_1 holds:

$$\mathrm{FPR} = \mathbb{E}[\phi(M(D))], \qquad \mathrm{FNR} = 1 - \mathbb{E}[\phi(M(D'))]. \tag{3.13}$$

An adversary wants both error rates small. The next theorem says that differential privacy is exactly a limit on how small they can be at once, for every test.

Theorem 3.3. *A mechanism M is ε-DP if and only if, for every pair of neighbours and every test, the false positive rate* FPR *and the false negative rate* FNR *satisfy*

$$\mathrm{FPR} + e^{\varepsilon}\mathrm{FNR} \geq 1 \qquad \textit{and} \qquad e^{\varepsilon}\mathrm{FPR} + \mathrm{FNR} \geq 1. \tag{3.14}$$

Proof.

1. *DP extends from events to tests.* Let M be ε-DP and let ϕ take values in $[0,1]$. Write its average as an integral of tail probabilities, $\mathbb{E}[\phi(M(D))] = \int_0^1 \Pr[\phi(M(D)) > t]\,\mathrm{d}t$. Each tail is the probability of the event $S_t = \{y : \phi(y) > t\}$, which Equation 3.10 bounds by $e^{\varepsilon}\Pr[M(D') \in S_t]$. Integrating over t gives $\mathbb{E}[\phi(M(D))] \leq e^{\varepsilon}\,\mathbb{E}[\phi(M(D'))]$, and the same holds with D and D' swapped.
2. *DP implies the inequalities.* Apply step 1 to ϕ, from D' to D: $1 - \mathrm{FNR} = \mathbb{E}[\phi(M(D'))] \leq e^{\varepsilon}\,\mathbb{E}[\phi(M(D))] = e^{\varepsilon}\mathrm{FPR}$. This is the second inequality. Apply step 1 to $1-\phi$, from D to D': $1 - \mathrm{FPR} = \mathbb{E}[1-\phi(M(D))] \leq e^{\varepsilon}\,\mathbb{E}[1-\phi(M(D'))] = e^{\varepsilon}\mathrm{FNR}$. This is the first.
3. *The inequalities imply DP.* Take an event S and the test $\phi(y) = \mathbf{1}[y \in S]$, which answers H_1 exactly when the output lands in S. Its error rates are $\mathrm{FPR} = \Pr[M(D) \in S]$ and $\mathrm{FNR} = 1 - \Pr[M(D') \in S]$. The second inequality says $e^{\varepsilon}\Pr[M(D) \in S] + 1 - \Pr[M(D') \in S] \geq 1$, that is, $\Pr[M(D') \in S] \leq e^{\varepsilon}\Pr[M(D) \in S]$. The pair of neighbours was arbitrary, so the same holds with the roles of D and D' swapped, which is Equation 3.10.

□

At $\varepsilon = 0$ the two inequalities force $\mathrm{FPR} + \mathrm{FNR} \geq 1$, which is what a test that ignores the output achieves, such as guessing at random ($\mathrm{FPR} = \mathrm{FNR} = 1/2$) or always answering H_0 ($\mathrm{FPR} = 0$, $\mathrm{FNR} = 1$). As ε grows, the region of achievable error pairs widens, and the adversary can beat blind guessing by more.

3.2.2. What the definition promises

In words, the probability of any event is about the same whether a given person's data is in the dataset or is replaced by someone else's. So any harm that can come from the output is not much more likely because the person took part. A linkage attack that succeeds with their record would succeed about as often without it, with any auxiliary data that exists now or is released later. Reconstruction attacks are prevented too, up to the number of queries that a later chapter quantifies. At $\varepsilon = 0.1$, the chance of any bad event grows by a factor of at most $e^{0.1} \approx 1.105$. The guarantee holds against an adversary with unlimited computation and unlimited side information, so it is information theoretic, unlike most of cryptography.

The definition does not prevent inference from aggregate patterns. Suppose a study run with differential privacy finds that smoking causes cancer, and an insurer raises the premium of a smoker who took part. The study's conclusion would have been the same without that person's record, so the harm comes from a fact about the population, not from their participation. Learning such facts is the point of statistics, and the definition is designed to allow it. For the same reason, a task whose goal is to single out an individual, such as tracing one person's contacts, is at odds with the definition. An aggregate version of the same task, such as warning everyone who attended an event at which many people later tested positive, fits it.

3.3. Randomized response revisited

Randomized response now has a privacy guarantee. The theorem treats the whole vector of reports as the output. The estimator Equation 3.4 is a function of that vector, so it inherits the guarantee by post-processing, a property proved in the next chapter.

Theorem 3.4. *Randomized response with bias $\kappa \in [0, 1/2)$, releasing $(Y_1, \ldots, Y_n)$, is ε-DP for*

$$\varepsilon = \ln \frac{\frac{1}{2} + \kappa}{\frac{1}{2} - \kappa}. \tag{3.15}$$

Equivalently, to obtain ε-DP, report the truth with probability $\frac{1}{2} + \kappa = e^{\varepsilon}/(1 + e^{\varepsilon})$.

Proof.

1. Fix an output $y \in \{0, 1\}^n$. The coins are independent across people, so $\Pr[M(D) = y] = \prod_{i=1}^{n} \Pr[Y_i = y_i]$, where each factor depends only on x_i.
2. Let D and D' differ only in record j, and write Y_j' for the report of person j on D'. All factors with $i \neq j$ are equal for the two datasets and cancel in the ratio:

$$\frac{\Pr[M(D) = y]}{\Pr[M(D') = y]} = \frac{\Pr[Y_j = y_j]}{\Pr[Y'_j = y_j]}. \tag{3.16}$$

3. Each probability on the right is $\frac{1}{2} + \kappa$ or $\frac{1}{2} - \kappa$, so the ratio is at most $(\frac{1}{2} + \kappa)/(\frac{1}{2} - \kappa) = e^{\varepsilon}$.
4. Lemma 3.1 turns this pointwise bound into the definition.
5. Solving $e^{\varepsilon} = (\frac{1}{2} + \kappa)/(\frac{1}{2} - \kappa)$ for $\frac{1}{2} + \kappa$ gives $e^{\varepsilon}/(1 + e^{\varepsilon})$.

□

For small κ the privacy parameter is $\varepsilon = 4\kappa + O(\kappa^3)$, so a bias of order ε gives ε-DP. Exactly, the bias that gives ε-DP is

$$2\kappa = \frac{e^{\varepsilon} - 1}{e^{\varepsilon} + 1} = \tanh \frac{\varepsilon}{2}, \tag{3.17}$$

so Theorem 3.2 and Theorem 3.1 become, with probability at least $1 - \beta$,

$$|\tilde{p} - p| \le \coth\Big(\frac{\varepsilon}{2}\Big) \sqrt{\frac{\ln(2/\beta)}{2n}}, \qquad \mathrm{Var}(\tilde{p}) = \frac{e^{\varepsilon}}{(e^{\varepsilon} - 1)^2\, n}. \tag{3.18}$$

Here p is again the fraction of ones and $\tilde{p}$ its estimate. The factor $\coth(\varepsilon/2)$ lies between $2/\varepsilon$ and $2/\varepsilon + \varepsilon/6$ for every $\varepsilon > 0$, so the error is $O(1/(\varepsilon\sqrt{n}))$ for $\varepsilon \le 1$. Figure 3.1 compares this bound with simulation.

A worked case shows the sizes involved at a typical privacy level.

Example 3.1. With $\varepsilon = 1$, each person tells the truth with probability $e/(1 + e) = 0.7311$. With $n = 10{,}000$ and $\beta = 0.05$, the bound Equation 3.18 is $\coth(0.5)\sqrt{\ln 40/20{,}000} = 2.1640 \times 0.01358 = 0.0294$. The standard deviation from Equation 3.18 is $\sqrt{e/((e-1)^2 \cdot 10{,}000)} = 0.0096$.

Randomized response gives more than the definition asks for. Each person randomizes their own bit before sending it, so they need not trust the curator, and no one ever sees a true bit. This is **local differential privacy**, the subject of a later chapter, and it implies the central guarantee proved here. The price is accuracy. The $\sqrt{n}$ in the error is not necessary under central privacy: the next chapter's Laplace mechanism estimates p with error $O(1/(\varepsilon n))$.

3.4. Summary

Table 3.1.: The chapter's results.

Result	Statement
Theorem 3.2	randomized response with bias κ has error $\frac{1}{2\kappa}\sqrt{\ln(2/\beta)/(2n)}$

Result	Statement
Definition 3.4	$\Pr[M(D) \in S] \le e^{\varepsilon} \Pr[M(D') \in S]$ for all neighbours and events
Lemma 3.1	for discrete outputs, single outcomes suffice
Proposition 3.1	a deterministic private algorithm is constant
Theorem 3.3	ε-DP holds exactly when every test has $\mathrm{FPR} + e^{\varepsilon}\mathrm{FNR} \ge 1$ and $e^{\varepsilon}\mathrm{FPR} + \mathrm{FNR} \ge 1$
Theorem 3.4	truth with probability $e^{\varepsilon}/(1+e^{\varepsilon})$ gives ε-DP
Equation 3.18	error $\coth(\varepsilon/2)\sqrt{\ln(2/\beta)/(2n)} = O(1/(\varepsilon\sqrt{n}))$

3.5. References

- Stanley L. Warner, "Randomized Response: A Survey Technique for Eliminating Evasive Answer Bias", 1965. The original protocol, designed for surveys on sensitive questions.
- Cynthia Dwork, Frank McSherry, Kobbi Nissim and Adam Smith, "Calibrating Noise to Sensitivity in Private Data Analysis", 2006. The definition of differential privacy.
- Larry Wasserman and Shuheng Zhou, "A Statistical Framework for Differential Privacy", 2010. The hypothesis-testing reading of the definition.
- Peter Kairouz, Sewoong Oh and Pramod Viswanath, "The Composition Theorem for Differential Privacy", 2015, and Jinshuo Dong, Aaron Roth and Weijie Su, "Gaussian Differential Privacy", 2019. Two later uses of the hypothesis-testing reading, for optimal composition and for a definition built on testing.
- Cynthia Dwork and Aaron Roth, *The Algorithmic Foundations of Differential Privacy*, 2014. The standard reference for the definition and its interpretation, including the smoking example.
- Salil Vadhan, "The Complexity of Differential Privacy", 2017. Section 1.6 explains why total variation distance gives no meaningful guarantee.
- Thomas Steinke and Jonathan Ullman, "The Pitfalls of Average-Case Differential Privacy", 2020, a post on the DifferentialPrivacy.org blog. Why the worst case over datasets matters.

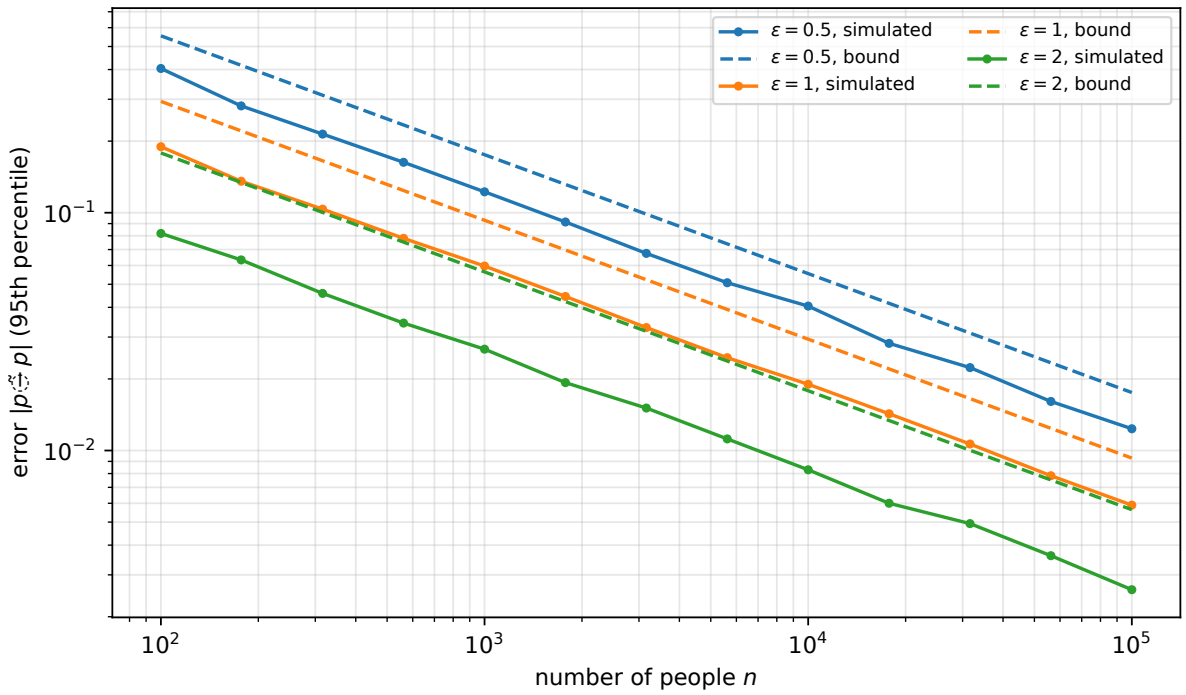


Figure 3.1.: Error of randomized response. For each ε, the solid line is the 95th percentile of $|\tilde{p} - p|$ over 2,000 simulated runs with $p = 0.3$, and the dashed line is the bound of Equation 3.18 with $\beta = 0.05$. Both fall as $1/\sqrt{n}$. The bound is 1.4 to 2.2 times the simulated error, and it is loosest for large ε, because Hoeffding's inequality uses only the range of each term and not its smaller variance.

4. The Laplace Mechanism

Randomized response privatizes one bit per person, and it extends to a value from a few categories, but no further. Most questions an analyst asks are numbers computed from the whole dataset, such as a count, an average or a histogram. This chapter answers them with the Laplace mechanism, which computes the true answer and adds noise scaled to how much one person can change it. On the fraction of ones that randomized response estimates, its error is of order $1/(\varepsilon n)$ rather than $1/(\varepsilon\sqrt{n})$.

The second half of the chapter proves three properties of differential privacy (DP) that make it possible to build private algorithms from private parts. Post-processing says that a private output stays private whatever is done with it. Group privacy bounds what is revealed about several people at once, and basic composition bounds the privacy cost of running several private algorithms on the same data.

4.1. Sensitivity

Differential privacy asks that the output hide the contribution of any one person. So the noise must be at least as large as the change that one person can cause in the answer, and the sensitivity measures the largest such change. Recall from the randomized response chapter that two datasets are neighbours, written $D \sim D'$, when they differ in exactly one record. This is the replace-one relation.

Definition 4.1 ($\ell 1$ sensitivity)**.** Let $f\colon \mathcal{X}^n \to \mathbb{R}^k$ map a dataset of n records to a vector of d numbers. The ℓ_1 **sensitivity** of f is the largest change in f, measured by the sum of the absolute changes of its coordinates, between neighbours $D \sim D'$:

$$\Delta_1 f = \sup_{D \sim D'} \|f(D) - f(D')\|_1, \tag{4.1}$$

where $\|v\|_1 = \sum_{j=1}^{d} |v_j|$ for a vector $v \in \mathbb{R}^d$.

The choice of the ℓ_1 norm is what makes the privacy proof of Theorem 4.1 work. Other mechanisms use other norms, such as the ℓ_2 norm for the Gaussian mechanism of the next chapter. For $d = 1$ all these norms are the absolute value of the difference. For $d > 1$ the ℓ_1 and ℓ_2 sensitivities can differ by a factor of up to $\sqrt{d}$, as the next chapter proves.

The running example of this chapter is the statistic that randomized response estimates, the fraction of ones. Its sensitivity is small because one person moves the fraction only a little.

Example 4.1. Let $\mathcal{X} = \{0, 1\}$ and let $f(D) = \frac{1}{n}\sum_{i=1}^{n} x_i$ be the fraction of ones. Changing one bit changes the sum by at most 1, so $\Delta_1 f = 1/n$.

4.2. The Laplace distribution

The noise must hide a shift of the true answer by up to the sensitivity. A density that falls exponentially in the distance from zero does this, because shifting it by a fixed amount changes its value by at most a fixed factor. The Laplace distribution is the simplest such density.

Definition 4.2 (Laplace distribution)**.** The **Laplace distribution** $\mathrm{Lap}(b)$ with scale $b > 0$ has density

$$\frac{1}{2b}\exp\left(-\frac{|x|}{b}\right), \qquad x \in \mathbb{R}. \tag{4.2}$$

It is a two-sided exponential distribution. The exponential distribution has density proportional to $e^{-x/b}$ on $[0, \infty)$, and the Laplace distribution reflects it to the negative half-line. Its tails decay like $e^{-|x|/b}$, more slowly than the $e^{-x^2/(2\sigma^2)}$ of a Gaussian with standard deviation σ, but still exponentially. The accuracy of the mechanism rests on three facts about this distribution, its mean, its variance and the exact size of its tails.

Lemma 4.1. *Let $Z \sim \mathrm{Lap}(b)$. Then $\mathbb{E}[Z] = 0$, $\mathrm{Var}(Z) = 2b^2$, and for every $t \geq 0$,*

$$\Pr[|Z| \geq tb] = e^{-t}. \tag{4.3}$$

Proof.

1. The density is symmetric about 0, so the mean is 0.
2. The absolute value $|Z|$ has density $\frac{1}{b}e^{-u/b}$ on $[0, \infty)$, the exponential distribution with mean b. Its second moment is $\int_0^\infty u^2 \frac{1}{b} e^{-u/b}\,\mathrm{d}u = 2b^2$, by integrating by parts twice. Since the mean of Z is 0, $\mathrm{Var}(Z) = \mathbb{E}[Z^2] = 2b^2$.
3. The tail of $|Z|$ is $\int_{tb}^\infty \frac{1}{b} e^{-u/b}\,\mathrm{d}u = e^{-t}$.

□

4.3. The Laplace mechanism

The mechanism adds independent Laplace noise to each coordinate, with scale equal to the sensitivity divided by ε.

Definition 4.3 (Laplace mechanism)**.** For $f\colon \mathcal{X}^n \to \mathbb{R}^d$ with $0 < \Delta_1 f < \infty$ and for $\varepsilon > 0$, the Laplace mechanism outputs

$$M(D) = f(D) + (Z_1, \ldots, Z_d), \tag{4.4}$$

where $Z_1, \ldots, Z_d$ are independent with distribution $\mathrm{Lap}(\Delta_1 f/\varepsilon)$.

Before proving that the mechanism is private, consider its accuracy on the running example. The mechanism is unbiased, and its error follows from Lemma 4.1.

Example 4.2. For $f(D) = p$, the fraction of ones, Example 4.1 gives $\Delta_1 f = 1/n$. The mechanism outputs $\tilde{p} = p + Z$ with $Z \sim \mathrm{Lap}(1/(\varepsilon n))$. So $\mathbb{E}[\tilde{p}] = p$, and $\mathrm{Var}(\tilde{p}) = 2/(\varepsilon n)^2$. Chebyshev's inequality already gives an error of order $1/(\varepsilon n)$: with probability at least $3/4$, the error is at most two standard deviations, $2\sqrt{2}/(\varepsilon n)$. The exact tail does better for small failure probabilities β. By Equation 4.3 with $t = \ln(1/\beta)$, with probability at least $1 - \beta$,

$$|\tilde{p} - p| \leq \frac{\ln(1/\beta)}{\varepsilon n}. \tag{4.5}$$

With $\varepsilon = 1$, $n = 10{,}000$ and $\beta = 0.05$, this is $\ln(20)/10{,}000 = 3.0 \times 10^{-4}$. Randomized response with the same ε, n and β has the bound 0.0294, which is 98 times larger.

The gap grows with n, because randomized response has error of order $1/(\varepsilon\sqrt{n})$ and the Laplace mechanism of order $1/(\varepsilon n)$. The Laplace mechanism needs a trusted curator who sees the true data, while randomized response does not. That is the price of the smaller error. It remains to prove that the mechanism is private.

Theorem 4.1 (Privacy of the Laplace mechanism)**.** *For every $f\colon \mathcal{X}^n \to \mathbb{R}^d$ with $0 < \Delta_1 f < \infty$ and every $\varepsilon > 0$, the Laplace mechanism is ε-DP.*

Proof.

1. Fix neighbours $D \sim D'$, and write p and q for the densities of $M(D)$ and $M(D')$ on $\mathbb{R}^d$. In this proof p is a density, not the fraction of ones. Let $b = \Delta_1 f/\varepsilon$. The coordinates are independent, so at a point $y \in \mathbb{R}^d$,

$$\frac{p(y)}{q(y)} = \prod_{j=1}^{d} \frac{\exp(-|f(D)_j - y_j|/b)}{\exp(-|f(D')_j - y_j|/b)} = \exp\left(\frac{\varepsilon}{\Delta_1 f} \sum_{j=1}^{d} \left(|f(D')_j - y_j| - |f(D)_j - y_j| \right) \right). \tag{4.6}$$

2. By the triangle inequality, $|f(D')_j - y_j| - |f(D)_j - y_j| \leq |f(D')_j - f(D)_j|$ for each j. Summing over j gives at most $\|f(D) - f(D')\|_1$, which is at most $\Delta_1 f$ by Definition 4.1. So $p(y) \leq e^{\varepsilon} q(y)$ for every y.
3. For an event $S \subseteq \mathbb{R}^d$, integrate the pointwise bound over S: $\Pr[M(D) \in S] = \int_S p \leq e^{\varepsilon} \int_S q = e^{\varepsilon} \Pr[M(D') \in S]$.

□

The proof shows why the Laplace distribution fits the definition. Its density falls exponentially in the distance from the true value, and the definition bounds a ratio by an exponential, so moving the centre by the sensitivity changes the density by a factor of at most e^{ε}.

Figure 4.1 shows the two densities for a one-dimensional statistic. They have the same shape, shifted by the sensitivity, and the ratio of their heights never exceeds e^{ε}.

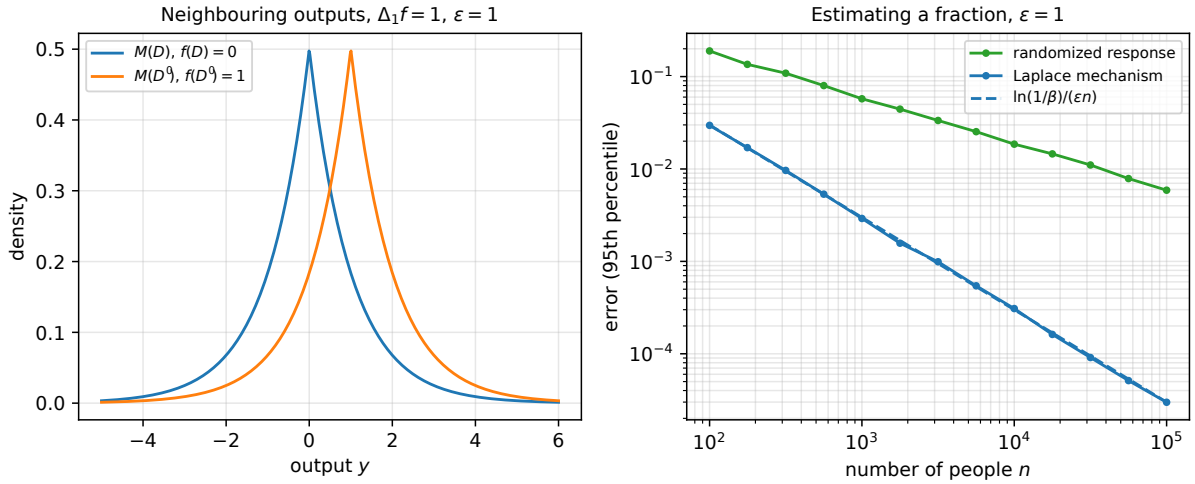


Figure 4.1.: Left, the densities of the Laplace mechanism on two neighbours whose values differ by the sensitivity $\Delta_1 f = 1$, at $\varepsilon = 1$: their ratio stays between e^{-1} and e. Right, the error in estimating a fraction of ones, as the 95th percentile over 2,000 simulated runs with $p = 0.3$. Randomized response falls as $1/\sqrt{n}$ and the Laplace mechanism as $1/n$. The dashed line is the exact 95th percentile $\ln(20)/(\varepsilon n)$ from Equation 4.3.

4.4. Counting queries

A **counting query** asks how many people in the dataset have some property. It is given by a predicate $\phi\colon \mathcal{X} \to \{0,1\}$ and has the value $f(D) = \sum_{i=1}^{n} \phi(x_i)$. Changing one record changes the count by at most 1, so $\Delta_1 f = 1$. The Laplace mechanism adds $\mathrm{Lap}(1/\varepsilon)$, whose error is at most $\ln(1/\beta)/\varepsilon$ with probability $1-\beta$, whatever the size of the dataset. On a large dataset, an error of a few units is small next to a count of thousands.

Many counting queries $f = (f_1, \dots, f_k)$, with predicates $\phi_1, \dots, \phi_k$, are answered together on the same data. For example, ϕ_1 might ask whether a person smokes and ϕ_2 whether their income is above some threshold. One person can then change all of the answers at once.

Proposition 4.1. *The vector of k counting queries has ℓ_1 sensitivity at most k, with equality when some record satisfies every predicate and another satisfies none. For every $\beta \in (0,1)$, with probability at least $1-\beta$, the Laplace mechanism answers all k queries with error at most*

$$\frac{k \ln(k/\beta)}{\varepsilon}. \tag{4.7}$$

Proof.

1. Let D and D' differ in one record, x in D and x' in D'. In each query the other terms cancel, so $\|f(D) - f(D')\|_1 = \sum_{j=1}^{k} |\phi_j(x) - \phi_j(x')| \le k$. The sum equals k when $\phi_j(x) = 1$ and $\phi_j(x') = 0$ for all j.
2. The noise in each answer is $\mathrm{Lap}(k/\varepsilon)$. By Equation 4.3 with $t = \ln(k/\beta)$, each noise is at least $k \ln(k/\beta)/\varepsilon$ in absolute value with probability β/k.

3. The union bound over the k answers gives a failure probability of at most β.

□

This analysis needs all k queries to be fixed in advance. The composition theorem of Section 4.6 gives the same privacy when each query is chosen after seeing the earlier answers.

The error of order k/ε per query should be compared with the reconstruction attacks of the reconstruction chapter. There, about n counting queries answered with error $o(\sqrt{n})$ let an attacker recover almost the whole dataset. Here, n counting queries are answered privately with error of order n/ε, up to a logarithm. Between $\sqrt{n}$ and n there is a wide gap. Either better attacks exist or less noise suffices, and the second is true. Advanced composition, under a relaxation of the definition, lowers the error to about $\sqrt{k}/\varepsilon$ per query, up to logarithmic factors.

4.5. Histograms

When the queries have structure, a sharper sensitivity bound is possible. In a **histogram**, each record falls into exactly one of k bins, such as an age in whole years or an eye colour, and the query counts the records in each bin: $f_j(D) = |\{i : x_i \text{ is in bin } j\}|$.

Proposition 4.2. *A histogram with $k \geq 2$ bins has ℓ_1 sensitivity 2. For every $\beta \in (0, 1)$, with probability at least $1 - \beta$, the Laplace mechanism, which adds $\mathrm{Lap}(2/\varepsilon)$ to each bin, answers every bin with error at most*

$$\frac{2\ln(k/\beta)}{\varepsilon}. \tag{4.8}$$

Proof.

1. Let D and D' differ in one record, which lies in bin a in D and in bin a' in D'. If $a = a'$ the histograms are equal. Otherwise count a drops by 1 and count a' rises by 1, so $f(D) - f(D') = e_a - e_{a'}$, where e_j is the j-th unit vector, and $\|e_a - e_{a'}\|_1 = 2$.
2. By Equation 4.3 with $t = \ln(k/\beta)$, each of the k noises is at least $2\ln(k/\beta)/\varepsilon$ in absolute value with probability β/k. The union bound over the bins finishes the proof.

□

The sensitivity is 2 because a replacement moves a record from one bin to another. Under neighbours that add or remove a record, only one count changes, and the sensitivity is 1. A single bin has error at most $2\ln(1/\beta)/\varepsilon$ with probability $1 - \beta$, and the proposition pays only $\ln k$ more to control all bins at once. So the noise per bin does not depend on the number of bins, and the error over all bins grows only like $\ln k$, which makes private histograms over very large domains practical. For $k = 100$ bins, $\varepsilon = 1$ and $\beta = 0.05$, the bound is $2\ln(2000) = 15.20$, and in a simulation of 20,000 runs the 95th percentile of the largest error is 15.17. The same 100 bins treated as arbitrary counting queries would get the bound $100\ln(2000) = 760.1$ from Proposition 4.1.

4.6. Properties of differential privacy

Three properties let one reason about a private algorithm piece by piece.

4.6.1. Post-processing

Once an output is private, no computation on it can undo the privacy, as long as the computation does not look at the data again. The randomized response chapter already used this, when it turned the private reports into an estimate.

Theorem 4.2 (Post-processing)**.** *Let $M\colon \mathcal{X}^n \to \mathcal{Y}$ be ε-DP, and let $F\colon \mathcal{Y} \to \mathcal{Z}$ be a randomized map whose coins are independent of those of M. Then $F \circ M$ is ε-DP.*

Proof.

1. A randomized map is a random choice of a deterministic one. Write $F(y) = g_R(y)$, where R is the coins of F and each $g_r\colon \mathcal{Y} \to \mathcal{Z}$ is deterministic.
2. Fix neighbours $D \sim D'$ and an event $S \subseteq \mathcal{Z}$. For fixed coins, $g_R(M(D)) \in S$ exactly when $M(D) \in g_R^{-1}(S)$, so

$$\Pr[F(M(D)) \in S] = \mathbb{E}_R\Big[\Pr\left[M(D) \in g_R^{-1}(S)\right]\Big]. \tag{4.9}$$

3. Inside the expectation, apply the definition of ε-DP to the event $g_R^{-1}(S)$, then undo step 2 for D':

$$\mathbb{E}_R\Big[\Pr\left[M(D) \in g_R^{-1}(S)\right]\Big] \le e^{\varepsilon}\,\mathbb{E}_R\Big[\Pr\left[M(D') \in g_R^{-1}(S)\right]\Big] = e^{\varepsilon}\Pr[F(M(D')) \in S]. \tag{4.10}$$

□

4.6.2. Group privacy

The definition protects one record. Datasets that differ in several records, such as the records of one family, are still protected, with a guarantee that degrades in proportion to their number.

Theorem 4.3 (Group privacy)**.** *Let $M\colon \mathcal{X}^n \to \mathcal{Y}$ be ε-DP, and let $D, D' \in \mathcal{X}^n$ differ in k records. For every event $S \subseteq \mathcal{Y}$,*

$$\Pr[M(D) \in S] \le e^{k\varepsilon}\Pr[M(D') \in S]. \tag{4.11}$$

Proof.

1. The proof is a **hybrid argument**. Change the k differing records one at a time, from their values in D to their values in D'. This gives datasets $D = D^{(0)}, D^{(1)}, \dots, D^{(k)} = D'$ in which each consecutive pair are neighbours.

2. Apply the definition to each consecutive pair, from the first to the last:

$$\Pr[M(D^{(0)}) \in S] \le e^{\varepsilon} \Pr[M(D^{(1)}) \in S] \le e^{2\varepsilon} \Pr[M(D^{(2)}) \in S] \le \dots \le e^{k\varepsilon} \Pr[M(D^{(k)}) \in S]. \tag{4.12}$$

□

So a group of k people has the protection of $k\varepsilon$-DP, which is meaningful only while $k\varepsilon$ is small. The theorem also relates the two notions of neighbours in the randomized response chapter. A replacement is a removal followed by an addition, so a mechanism that is ε-DP when neighbours add or remove a record is 2ε-DP when neighbours replace one.

4.6.3. Basic composition

An analyst rarely runs one algorithm. In an **adaptive composition**, M_1 runs on the dataset, then M_2 is chosen in the light of the output y_1 and runs on the same dataset, and so on up to M_k. Formally, for each i and each sequence of earlier outputs $y_{<i} = (y_1, \dots, y_{i-1})$, the mechanism $M_i(\cdot\,; y_{<i})$ is a mechanism on $\mathcal{X}^n$ with fresh coins, and the composition releases $(y_1, \dots, y_k)$. Basic composition says that the privacy parameters simply add up, even though each mechanism may depend on the earlier outputs.

Theorem 4.4 (Basic composition)**.** *If each $M_i(\cdot\,; y_{<i})$ is ε_i-DP for every $y_{<i}$, then the adaptive composition of $M_1, \dots, M_k$ is $(\sum_{i=1}^{k} \varepsilon_i)$-DP. In particular, k mechanisms that are each ε-DP compose to a $k\varepsilon$-DP mechanism.*

The proof is simplest for countable outputs, where the probability of a sequence of outputs factors into conditional probabilities.

Proof. Countable outputs.

1. Fix neighbours $D \sim D'$ and outputs $y = (y_1, \dots, y_k)$. Each M_i uses fresh coins, so the chain rule gives $\Pr[M(D) = y] = \prod_{i=1}^{k} \Pr[M_i(D; y_{<i}) = y_i]$, and likewise for D'.
2. Each factor for D is at most e^{ε_i} times the matching factor for D', by the privacy of $M_i(\cdot\,; y_{<i})$ applied to the event $\{y_i\}$. Multiplying the k bounds gives

$$\Pr[M(D) = y] = \prod_{i=1}^{k} \Pr[M_i(D; y_{<i}) = y_i] \le \prod_{i=1}^{k} e^{\varepsilon_i} \Pr[M_i(D'; y_{<i}) = y_i] = \exp\Big(\sum_{i=1}^{k} \varepsilon_i\Big) \Pr[M(D') = y]. \tag{4.13}$$

3. The pointwise lemma of the randomized response chapter turns this bound on single outcomes into the definition.

□

For outputs that are not countable, such as the real numbers that the Laplace mechanism releases, the same idea works with averages in place of products. It uses a fact proved in the randomized response chapter, as the first step of its hypothesis-testing theorem. If M is ε-DP and a function h takes values in $[0, 1]$, then for all neighbours

$$\mathbb{E}[h(M(D))] \le e^{\varepsilon}\, \mathbb{E}[h(M(D'))]. \tag{4.14}$$

Proof. General outputs.

1. The composition of the first $k-1$ mechanisms is itself a mechanism, so by induction on k it suffices to treat $k = 2$.
2. Fix neighbours $D \sim D'$ and an event S of pairs. For an output y_1, write $S_{y_1} = \{y_2 : (y_1, y_2) \in S\}$ and $g_D(y_1) = \Pr[M_2(D; y_1) \in S_{y_1}]$, and define $g_{D'}$ in the same way. Conditioning on the first output gives $\Pr[M(D) \in S] = \mathbb{E}_{y_1 \sim M_1(D)}[g_D(y_1)]$, and likewise for D'.
3. By the privacy of $M_2(\cdot\,; y_1)$, $g_D(y_1) \le e^{\varepsilon_2} g_{D'}(y_1)$ for every y_1.
4. The function $g_{D'}$ takes values in $[0, 1]$, so the fact above, applied to M_1, gives

$$\Pr[M(D) \in S] \le e^{\varepsilon_2}\, \mathbb{E}_{y_1 \sim M_1(D)}[g_{D'}(y_1)] \le e^{\varepsilon_1 + \varepsilon_2}\, \mathbb{E}_{y_1 \sim M_1(D')}[g_{D'}(y_1)] = e^{\varepsilon_1 + \varepsilon_2} \Pr[M(D') \in S]. \tag{4.15}$$

□

Basic composition recovers the analysis of Proposition 4.1. Answering each counting query with $\mathrm{Lap}(k/\varepsilon)$ noise is (ε/k)-DP, and k of them compose to ε-DP, now even when the queries are chosen adaptively. The cost of k analyses grows linearly in k. It can be made to grow like $\sqrt{k}$, but only under the relaxation of the definition that the next chapter introduces.

4.7. Summary

Table 4.1.: The chapter's results.

Result	Statement
Definition 4.1	$\Delta_1 f$ is the largest ℓ_1 change in f between neighbours
Lemma 4.1	$\mathrm{Lap}(b)$ has variance $2b^2$ and $\Pr[\lvert Z\rvert \ge tb] = e^{-t}$
Theorem 4.1	adding $\mathrm{Lap}(\Delta_1 f/\varepsilon)$ to each coordinate is ε-DP
Example 4.2	a fraction of ones with error $\ln(1/\beta)/(\varepsilon n)$
Proposition 4.1	k counting queries: sensitivity k, error $k \ln(k/\beta)/\varepsilon$
Proposition 4.2	a histogram: sensitivity 2, error $2\ln(k/\beta)/\varepsilon$ over all bins
Theorem 4.2	post-processing keeps ε-DP

Result	Statement
Theorem 4.3	datasets differing in k records: factor $e^{k\varepsilon}$
Theorem 4.4	adaptive composition adds the ε_i

4.8. References

- Cynthia Dwork, Frank McSherry, Kobbi Nissim and Adam Smith, "Calibrating Noise to Sensitivity in Private Data Analysis", 2006. The Laplace mechanism and the notion of sensitivity.
- Irit Dinur and Kobbi Nissim, "Revealing Information while Preserving Privacy", 2003. The reconstruction attacks that the error of the Laplace mechanism is compared with.
- Cynthia Dwork and Aaron Roth, *The Algorithmic Foundations of Differential Privacy*, 2014. Section 3.3 covers the Laplace mechanism, counting queries and histograms, and Section 2.3 and Section 3.5 cover post-processing, group privacy and composition.
- Cynthia Dwork, Guy Rothblum and Salil Vadhan, "Boosting and Differential Privacy", 2010. Advanced composition, which improves the linear growth of basic composition.

5. Approximate Differential Privacy

Pure differential privacy (DP) bounds the ratio of two output probabilities by e^{ε} on every event. This chapter adds a small additive slack δ to that bound. The relaxed definition, (ε, δ)-DP, lets an algorithm fail with probability about δ, and in exchange it admits the Gaussian mechanism, which adds less noise than the Laplace mechanism to statistics with many coordinates. It is also the setting in which the next chapter's advanced composition holds.

The chapter defines the relaxation and reads it through the privacy loss random variable. Two examples then show how badly a mechanism may behave with probability δ, which forces δ to be much smaller than $1/n$. The Gaussian mechanism comes next, with a full privacy proof, and the chapter ends by checking that post-processing, group privacy and composition survive the relaxation.

5.1. The definition

An additive bound alone is the wrong notion of privacy. Suppose we asked only that, for all neighbours and all events,

$$\Pr[M(D) \in S] \leq \Pr[M(D') \in S] + \delta. \tag{5.1}$$

This allows no useful accuracy when δ is small, and it allows harmful algorithms when δ is large (Vadhan, 2017, Section 1.6). The second failure appears in Section 5.3. Combined with the multiplicative bound of pure DP, however, an additive term gives the right definition. Recall that two datasets in $\mathcal{X}^n$ are neighbours, written $D \sim D'$, when they differ in exactly one record (the replace-one relation).

Definition 5.1 (Approximate differential privacy)**.** Let $\varepsilon \geq 0$ and $\delta \in [0, 1]$. A randomized algorithm $M \colon \mathcal{X}^n \to \mathcal{Y}$ is (ε, δ)**-differentially private** if for all neighbours $D \sim D'$ and all events $S \subseteq \mathcal{Y}$,

$$\Pr[M(D) \in S] \leq e^{\varepsilon} \Pr[M(D') \in S] + \delta. \tag{5.2}$$

At $\delta = 0$ this is the ε-DP of the randomized response chapter, which from now on is called **pure DP**. A definition with $\delta > 0$ is called **approximate DP**. The rest of the chapter is about what the extra δ permits.

5.2. The privacy loss random variable

Definition 5.1 bounds probabilities of events, but it is easier to understand one outcome at a time. Fix neighbours $D \sim D'$. An observer who sees an output y and wonders whether the dataset was D or D' compares how likely y is under each. Write P and Q for the distributions of $M(D)$ and $M(D')$, and $p(y)$ and $q(y)$ for the likelihood of the outcome y under each. When $\mathcal{Y}$ is countable, these are the probabilities $p(y) = \Pr[M(D) = y]$ and $q(y) = \Pr[M(D') = y]$. When both outputs have ordinary densities, p and q are those densities. (In general, p and q are densities with respect to the measure $P + Q$, which always exist.)

The privacy loss is the logarithm of the ratio $p(y)/q(y)$. The ratio needs a convention in one case: an outcome that can occur under D but never under D', so that $p(y)$ is positive while $q(y)$ is zero. Seeing such an outcome rules D' out entirely, and its loss is set to $+\infty$.

Definition 5.2 (Privacy loss)**.** Let $p(y)$ and $q(y)$ be the probabilities, or densities, of an outcome y under $M(D)$ and under $M(D')$. The **privacy loss** of an outcome y with $p(y) > 0$ is

$$L_{D,D'}(y) = \ln \frac{p(y)}{q(y)}, \tag{5.3}$$

where $L_{D,D'}(y) = +\infty$ if $q(y) = 0$. Once the pair is fixed we write it $L(y)$. The **privacy loss random variable** is $L = L(Y)$ with $Y = M(D)$: the privacy loss of the outcome the mechanism actually produces on D.

The outcome $M(D)$ has $p(M(D)) > 0$ with probability 1, so L is always defined. A large value of L means that the observed outcome is much more likely under D than under D', so it is evidence that the dataset is D. A negative value is evidence for D', and an infinite value means that the outcome is impossible under D'.

The Laplace mechanism of the previous chapter gives a first example. It applies to a function $f\colon \mathcal{X}^n \to \mathbb{R}$ whose value changes by at most $\Delta_1 f$ between neighbours. The mechanism releases $f(D) + Z$, where the noise Z has the Laplace density $\frac{1}{2b}e^{-|z|/b}$ with scale $b = \Delta_1 f/\varepsilon$, and it is ε-DP.

Example 5.1. Take neighbours with $f(D) = a$ and $f(D') = a - \Delta_1 f$, the largest gap allowed. The densities give

$$L(y) = \frac{\varepsilon}{\Delta_1 f}(|y - a + \Delta_1 f| - |y - a|), \tag{5.4}$$

which equals ε for $y \geq a$, equals $-\varepsilon$ for $y \leq a - \Delta_1 f$, and is linear in between.

1. The event $y \geq a$ has probability $\frac{1}{2}$ under P, because the noise is symmetric. So L has an atom of mass $\frac{1}{2}$ at ε.
2. The event $y \leq a - \Delta_1 f$ needs noise below $-\Delta_1 f = -\varepsilon b$, which has probability $\frac{1}{2}e^{-\varepsilon}$. So L has an atom of mass $\frac{1}{2}e^{-\varepsilon}$ at $-\varepsilon$.

3. In between, write $y = a - s$ with $s \in (0, \Delta_1 f)$, so that $L = \varepsilon - 2\varepsilon s/(\Delta_1 f)$. The variable s has density $\frac{\varepsilon}{2\Delta_1 f} e^{-\varepsilon s/(\Delta_1 f)}$, and the change of variables gives L the density $u \mapsto \frac{1}{4} e^{(u-\varepsilon)/2}$ on $(-\varepsilon, \varepsilon)$, of total mass $\frac{1}{2}(1 - e^{-\varepsilon})$.

The three masses add up to 1. The privacy loss never leaves $[-\varepsilon, \varepsilon]$, as Figure 5.1 shows on the left.

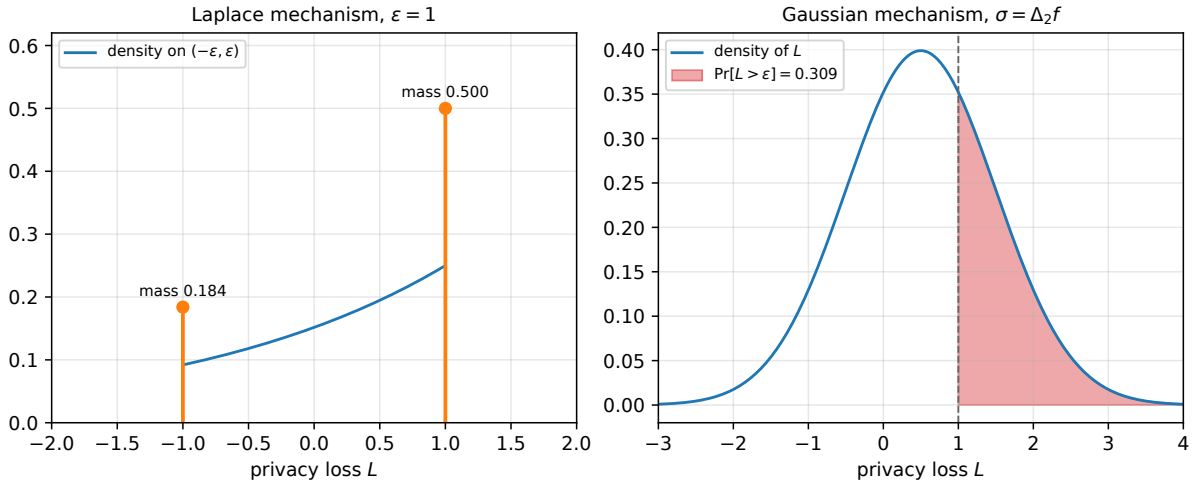


Figure 5.1.: Privacy loss for two mechanisms on neighbours whose values differ by the full sensitivity. Left, the Laplace mechanism with $\varepsilon = 1$, from Example 5.1: atoms of mass 0.5 at ε and 0.184 at $-\varepsilon$, and a density in between. Right, the Gaussian mechanism with noise standard deviation $\sigma = \Delta_2 f$, from Lemma 5.4: the loss is Gaussian with mean $1/2$ and variance 1, so it exceeds $\varepsilon = 1$ with probability 0.309.

The example is an instance of a general fact. Pure DP says that the privacy loss is bounded by ε with probability 1, and approximate DP says, roughly, that it is bounded by ε except with probability δ. Both follow from one identity, which computes the largest violation of the multiplicative bound over all events. That largest violation is the smallest δ that works for the pair. The identity writes it as an average over the privacy loss. It uses the positive part $(z)_+ = \max(z, 0)$, which keeps z when it is positive and replaces it by 0 otherwise. When the loss is infinite, the term $e^{\varepsilon - L}$ is read as 0.

Lemma 5.1. *For neighbours $D \sim D'$ and $\varepsilon \geq 0$,*

$$\sup_{S \subseteq \mathcal{Y}} \Big(\Pr[M(D) \in S] - e^{\varepsilon} \Pr[M(D') \in S] \Big) = \mathbb{E}\Big[\big(1 - e^{\varepsilon - L}\big)_+ \Big], \tag{5.5}$$

where L is the privacy loss random variable of Definition 5.2 for this pair, $(z)_+ = \max(z, 0)$ is the positive part of z, and $e^{\varepsilon - L}$ is read as 0 when $L = +\infty$. So M is (ε, δ)-DP if and only if the right-hand side is at most δ for every ordered pair of neighbours.

Proof.

1. For an event S, write the difference as an integral, $\Pr[M(D) \in S] - e^{\varepsilon} \Pr[M(D') \in S] = \int_S (p - e^{\varepsilon} q)$. The integrand is positive exactly on $B = \{y : p(y) > e^{\varepsilon} q(y)\}$, so the integral is largest for $S = B$, where it equals $\int (p - e^{\varepsilon} q)_+$.
2. Where $p(y) = 0$, the integrand $(p - e^{\varepsilon} q)_+$ is 0. Where $p(y) > 0$, it equals $p(y)(1 - e^{\varepsilon} q(y)/p(y))_+ = p(y)(1 - e^{\varepsilon - L(y)})_+$, which is also right when $q(y) = 0$ and $L(y) = \infty$.
3. Integrating $p(y)$ times a function of y is taking the expectation under P, so $\int (p - e^{\varepsilon} q)_+ = \mathbb{E}[(1 - e^{\varepsilon - L})_+]$.
4. Definition 5.1 asks that the left-hand side of Equation 5.5 be at most δ for every ordered pair of neighbours, which proves the last sentence.

□

The right-hand side of Equation 5.5 is called the **hockey-stick divergence** of P from Q, after the shape of the function $z \mapsto (1 - e^{\varepsilon - z})_+$. Its integrand is 0 when $L \le \varepsilon$ and at most 1 otherwise. So the lemma ties (ε, δ)-DP to how often L exceeds ε, and the next theorem makes that tie precise in both directions.

Theorem 5.1. *Let $M : \mathcal{X}^n \to \mathcal{Y}$ be a randomized algorithm, and let L be its privacy loss random variable for an ordered pair of neighbours $D \sim D'$.*

1. *M is ε-DP if and only if $L \le \varepsilon$ with probability 1 for every ordered pair of neighbours. In that case $|L| \le \varepsilon$ with probability 1.*
2. *If $\Pr[L > \varepsilon] \le \delta$ for every ordered pair of neighbours, then M is (ε, δ)-DP.*
3. *If M is (ε, δ)-DP, then for every $t > 0$ and every ordered pair of neighbours,*

$$\Pr[L > \varepsilon + t] \le \frac{\delta}{1 - e^{-t}}. \tag{5.6}$$

Proof.

1. With $\delta = 0$, Lemma 5.1 says that M is ε-DP exactly when $\mathbb{E}[(1 - e^{\varepsilon - L})_+] = 0$ for every pair. The integrand is nonnegative, so this holds exactly when it is 0 with probability 1, that is, when $L \le \varepsilon$ with probability 1. For the lower bound, apply the same fact to the pair (D', D). It gives $q \le e^{\varepsilon} p$ except on a set of Q-probability 0, and that set also has P-probability 0, because pure DP gives $\Pr[M(D) \in S] \le e^{\varepsilon} \Pr[M(D') \in S]$. So $L \ge -\varepsilon$ with probability 1 under P.
2. The integrand $(1 - e^{\varepsilon - L})_+$ is 0 when $L \le \varepsilon$ and at most 1 when $L > \varepsilon$. So its expectation is at most $\Pr[L > \varepsilon] \le \delta$, and Lemma 5.1 gives (ε, δ)-DP.
3. On the event $L > \varepsilon + t$, the integrand is at least $1 - e^{-t}$. So $(1 - e^{-t}) \Pr[L > \varepsilon + t] \le \mathbb{E}[(1 - e^{\varepsilon - L})_+] \le \delta$, by Lemma 5.1.

□

The converse of part 2 fails with the same parameters, which is why part 3 pays in both ε and δ. For example, the Gaussian mechanism on the right of Figure 5.1 has privacy loss above 1 with probability 0.309, yet Proposition 5.3 below shows that it is $(1, 0.127)$-DP. Taking $t = \varepsilon$, an (ε, δ)-DP mechanism has privacy loss above 2ε with probability at most $\delta/(1 - e^{-\varepsilon})$. Up to these

constants, then, (ε, δ)-DP says that the privacy loss is at most ε except with probability about δ. The definition says nothing about what happens on that rare event.

5.3. How small δ must be

On the event of probability δ, the privacy loss may be infinite. The first example shows this in its crudest form.

Example 5.2. Let $\perp$ be a fixed symbol that carries no information about the data. The mechanism $M(D)$ outputs $\perp$ with probability $1-\delta$ and outputs the whole dataset D with probability δ. For neighbours $D \neq D'$ and an event S,

$$\Pr[M(D) \in S] \leq (1-\delta)\,\mathbf{1}[\perp \in S] + \delta \leq \Pr[M(D') \in S] + \delta, \tag{5.7}$$

so M is $(0, \delta)$-DP. With probability $1-\delta$ the privacy loss is $\ln \frac{1-\delta}{1-\delta} = 0$, and with probability δ it is infinite, because the output D is impossible under D'.

This mechanism publishes everyone's data with probability δ, so δ must be small. The next example shows how small. It publishes each record on its own, and that makes δ comparable to $1/n$.

Definition 5.3 (Name and shame)**.** For $\delta \in [0,1]$, the mechanism NS_δ draws independent coins $c_1, \ldots, c_n$, each equal to 1 with probability δ, and outputs the set of pairs $\{(i, x_i) : c_i = 1\}$.

Each record, which might be a person's medical file, is published in the clear with probability δ. Even so, the mechanism meets the definition with $\varepsilon = 0$ and the same δ, because a person's record affects the output only when that person's coin comes up 1, which happens with probability δ.

Proposition 5.1. *For every $\delta \in [0,1]$, NS_δ is $(0, \delta)$-DP.*

Proof.

1. Let D and D' differ only in record i, and run the mechanism on both datasets with the same coins. Let E be the event $c_i = 1$, so $\Pr[E] = \delta$.
2. On the complement $\bar{E}$, record i is not published, and the output is a function of the coins and of the records $j \neq i$, which the two datasets share. So on $\bar{E}$ the two outputs are equal.
3. For an event S, split on E and use step 2:

$$\Pr[\mathrm{NS}_\delta(D) \in S] = \Pr[\mathrm{NS}_\delta(D) \in S,\ \bar{E}] + \Pr[\mathrm{NS}_\delta(D) \in S,\ E] \leq \Pr[\mathrm{NS}_\delta(D') \in S,\ \bar{E}] + \Pr[E]. \tag{5.8}$$

4. The first term is at most $\Pr[\mathrm{NS}_\delta(D') \in S]$ and the second is δ, which is Equation 5.2 with $\varepsilon = 0$.

□

So a mechanism that plainly harms people meets $(0, \delta)$-DP. How much harm it does depends on the chance that it publishes at least one record, which the next proposition computes.

Proposition 5.2. *The probability that* NS_δ *publishes at least one record is* $1 - (1-\delta)^n$*, and*

$$n\delta - \frac{(n\delta)^2}{2} \le 1 - (1-\delta)^n \le n\delta. \tag{5.9}$$

Proof.

1. The coins are independent, so no record is published with probability $(1-\delta)^n$, and the complement gives the formula.
2. The upper bound is the union bound over the n records, each published with probability δ.
3. For the lower bound, $1 - \delta \le e^{-\delta}$ gives $(1-\delta)^n \le e^{-n\delta}$. The inequality $e^{-x} \le 1 - x + x^2/2$ for $x \ge 0$ then gives $1 - (1-\delta)^n \ge n\delta - (n\delta)^2/2$.

□

So the chance that someone is named is about $n\delta$ when $n\delta$ is small. At $\delta = 1/n$, step 3 of the proof gives a chance of at least $1 - e^{-1} \approx 0.632$ that some record is published in the clear. A guarantee that allows this is not privacy, so δ must be well below $1/n$. Even $\delta = n^{-1.1}$ allows a chance of up to $n\delta = n^{-0.1}$, which is 0.25 for a million records. Common choices are $\delta \le 1/n^2$, or a δ that is **cryptographically small**, smaller than every inverse polynomial in n.

In both examples the privacy loss is 0 or infinite, so the rare event is a catastrophe. Without more information about a mechanism, one has to assume this worst case. Many mechanisms degrade gracefully instead. Their privacy loss exceeds ε with probability about δ but rarely exceeds it by much, and they satisfy a family of guarantees at once, such as $(1, 10^{-3})$-DP and $(2, 10^{-6})$-DP. The Gaussian mechanism of the next section is the main example, as Example 5.5 shows. A single pair (ε, δ) is a crude threshold on the distribution of the privacy loss. Rényi differential privacy and concentrated differential privacy describe the whole tail of that distribution, and they appear in later chapters.

The definition must be stated for events, not single outcomes. For pure DP the two are equivalent on countable outputs, by the pointwise lemma of the randomized response chapter. For approximate DP the pointwise version is much weaker.

Example 5.3. Let $m = \lceil 1/\delta \rceil$ for some $\delta \in (0, 1)$. The mechanism $M(D) = (D, U)$ outputs the whole dataset together with a number U drawn uniformly from $\{1, \dots, m\}$. Every single outcome has probability at most $1/m \le \delta$, so for all neighbours and all outcomes y,

$$\Pr[M(D) = y] \le e^{\varepsilon} \Pr[M(D') = y] + \delta. \tag{5.10}$$

Yet the event $S = \{(D, u) : 1 \le u \le m\}$ has probability 1 under D and 0 under D', so M is not (ε, δ_0)-DP for any ε and any $\delta_0 < 1$.

The privacy loss of this mechanism is infinite with probability 1. So Theorem 5.1, which looks at one outcome at a time through the privacy loss, catches the failure that the pointwise inequality misses.

5.4. The Gaussian mechanism

The Laplace mechanism calibrates its noise to the ℓ_1 sensitivity. The Gaussian mechanism adds Gaussian noise instead, calibrated to the ℓ_2 sensitivity, which can be much smaller.

Definition 5.4 ($\ell 2$ sensitivity)**.** The ℓ_2 **sensitivity** of a function $f\colon \mathcal{X}^n \to \mathbb{R}^d$ is

$$\Delta_2 f = \sup_{D \sim D'} \|f(D) - f(D')\|_2. \tag{5.11}$$

The ℓ_1 sensitivity $\Delta_1 f$ of the previous chapter is the same supremum with the ℓ_1 norm. The two norms, and so the two sensitivities, are within a factor $\sqrt{d}$ of each other.

Lemma 5.2. *For every $v \in \mathbb{R}^d$, $\|v\|_2 \le \|v\|_1 \le \sqrt{d}\,\|v\|_2$. Hence $\Delta_2 f \le \Delta_1 f \le \sqrt{d}\,\Delta_2 f$.*

Proof.

1. Expanding the square, $\|v\|_1^2 = \sum_i v_i^2 + \sum_{i \neq j} |v_i||v_j| \ge \|v\|_2^2$.
2. The Cauchy–Schwarz inequality applied to $(|v_1|, \dots, |v_d|)$ and the all-ones vector gives $\|v\|_1 = \sum_i |v_i| \cdot 1 \le \|v\|_2 \sqrt{d}$.
3. Applying both inequalities to $f(D) - f(D')$ and taking the supremum over neighbours gives the bounds on the sensitivities.

□

The Gaussian distribution $\mathcal{N}(\mu, \sigma^2)$ with mean μ and variance σ^2 has density

$$\frac{1}{\sqrt{2\pi\sigma^2}} \exp\!\left(-\frac{(x-\mu)^2}{2\sigma^2}\right). \tag{5.12}$$

Its tails fall like $e^{-x^2/(2\sigma^2)}$, faster than the $e^{-|x|/b}$ of the Laplace distribution. We write $\mathcal{N}(0, \sigma^2 I_d)$ for the vector of d independent $\mathcal{N}(0, \sigma^2)$ coordinates. The Gaussian mechanism adds such a vector to the true answer.

Definition 5.5 (Gaussian mechanism)**.** For $f\colon \mathcal{X}^n \to \mathbb{R}^d$ and a noise level $\sigma > 0$, the Gaussian mechanism outputs

$$M(D) = f(D) + Z, \qquad Z \sim \mathcal{N}(0, \sigma^2 I_d). \tag{5.13}$$

The theorem below says how large the noise level σ must be for a given ε and δ.

Theorem 5.2 (Privacy of the Gaussian mechanism)**.** *Let f have $0 < \Delta_2 f < \infty$, and let $\varepsilon > 0$ and $\delta \in (0,1)$. The Gaussian mechanism with*

$$\sigma = \frac{\Delta_2 f}{\varepsilon}\sqrt{2\ln(1/\delta) + \varepsilon} \tag{5.14}$$

is (ε, δ)-DP.

The proof is in Section 5.4.1. Asymptotically the noise level is $\sigma = \frac{\Delta_2 f}{\varepsilon}\sqrt{2\ln(1/\delta)}\,(1 + o(1))$ as $\delta \to 0$. Compared with the Laplace mechanism, whose noise has scale $\Delta_1 f/\varepsilon$ in each coordinate, there are two differences. The Gaussian mechanism pays a factor $\sqrt{2\ln(1/\delta)}$ and gives only approximate DP, and it measures sensitivity in ℓ_2, which can save a factor $\sqrt{d}$. The next example shows that the saving can be the larger effect.

Example 5.4. Each of n people answers d yes-or-no questions, so the dataset is $D \in \{0,1\}^{n \times d}$ with rows $x_i \in \{0,1\}^d$. The analyst wants the vector of averages $f(D) = \frac{1}{n}\sum_i x_i$, with small error in ℓ_2 norm.

1. Replacing row x by row x' changes f by $(x - x')/n$. The largest change is from the all-ones row to the all-zeros row, so $\Delta_1 f = d/n$ and $\Delta_2 f = \sqrt{d}/n$.
2. The Laplace mechanism adds $\mathrm{Lap}(d/(\varepsilon n))$ to each coordinate, and each coordinate has variance $2d^2/(\varepsilon n)^2$. The root mean squared ℓ_2 error is

$$\sqrt{\mathbb{E}\|Z\|_2^2} = \sqrt{d \cdot \frac{2d^2}{\varepsilon^2 n^2}} = \frac{\sqrt{2}\, d^{3/2}}{\varepsilon n}. \tag{5.15}$$

3. The Gaussian mechanism with Equation 5.14 adds $\mathcal{N}(0, \sigma^2)$ to each coordinate with $\sigma = \frac{\sqrt{d}}{\varepsilon n}\sqrt{2\ln(1/\delta) + \varepsilon}$, so its root mean squared error is

$$\sqrt{d}\,\sigma = \frac{d\sqrt{2\ln(1/\delta) + \varepsilon}}{\varepsilon n}. \tag{5.16}$$

4. The ratio of the two errors is $\sqrt{2d/(2\ln(1/\delta) + \varepsilon)}$. So the Gaussian mechanism is more accurate as soon as $d > \ln(1/\delta) + \varepsilon/2$, and its advantage grows like $\sqrt{d}$.

With $d = 100$ questions, $n = 10{,}000$ people, $\varepsilon = 1$ and $\delta = 10^{-6}$, the Laplace error is $\sqrt{2} \cdot 1000/10{,}000 = 0.1414$ and the Gaussian error is $100\sqrt{27.63 + 1}/10{,}000 = 0.0535$, smaller by a factor 2.64. A simulation of 20,000 runs of each mechanism gives 0.1413 and 0.0535.

Which mechanism to use is a trade-off. The Laplace mechanism gives the stronger guarantee of pure DP, and the Gaussian mechanism gives a weaker guarantee with less noise in high dimension.

5.4.1. Proof of privacy

The proof has two steps. First, the privacy loss of the Gaussian mechanism is itself Gaussian. Second, a Gaussian tail bound shows that it exceeds ε with probability at most δ, and part 2 of Theorem 5.1 concludes. The lemma below first collects four facts about Gaussians. The first two give the distribution of the privacy loss, and the last two bound its tail.

Lemma 5.3. *Let $W \sim \mathcal{N}(0,1)$.*

1. *For constants μ and $a \neq 0$, $\mu + aW \sim \mathcal{N}(\mu, a^2)$.*
2. *If $V_1 \sim \mathcal{N}(0, a_1^2)$ and $V_2 \sim \mathcal{N}(0, a_2^2)$ are independent, then $V_1 + V_2 \sim \mathcal{N}(0, a_1^2 + a_2^2)$.*
3. *For every $u \geq 0$, $\Pr[W > u] \leq e^{-u^2/2}$.*
4. *For every $u > 0$, $\Pr[W > u] \leq e^{-u^2/2}/(u\sqrt{2\pi})$.*

Proof.

1. Completing the square in the density gives the moment generating function $\mathbb{E}[e^{\lambda W}] = e^{\lambda^2/2}$. A distribution is determined by its moment generating function when that function is finite near 0, as it is for all the variables here.
2. For part 1, $\mathbb{E}[e^{\lambda(\mu + aW)}] = e^{\lambda\mu + \lambda^2 a^2/2}$, which is the moment generating function of $\mathcal{N}(\mu, a^2)$.
3. For part 2, independence makes the moment generating function of the sum the product $e^{\lambda^2 a_1^2/2} e^{\lambda^2 a_2^2/2} = e^{\lambda^2 (a_1^2 + a_2^2)/2}$.
4. For part 3, Markov's inequality applied to e^{uW} gives $\Pr[W > u] \leq e^{-u^2}\,\mathbb{E}[e^{uW}] = e^{-u^2/2}$.
5. For part 4, write $\varphi(z) = e^{-z^2/2}/\sqrt{2\pi}$ for the density, whose derivative is $-z\varphi(z)$. On $z \geq u$ the ratio z/u is at least 1, so $\Pr[W > u] \leq \int_u^\infty \frac{z}{u}\varphi(z)\,\mathrm{d}z = \frac{\varphi(u)}{u}$.

□

The first step of the plan is to compute the privacy loss. It turns out to be Gaussian, with a mean and a variance set by the distance between $f(D)$ and $f(D')$.

Lemma 5.4. *Let $D \sim D'$, let M be the Gaussian mechanism with noise level σ, and write $v = f(D) - f(D')$. The privacy loss random variable of $M(D)$ against $M(D')$ is distributed as*

$$L \sim \mathcal{N}(\rho,\ 2\rho)\,, \qquad \rho = \frac{\|v\|_2^2}{2\sigma^2}. \tag{5.17}$$

Proof.

1. The output is $f(D) + z$ with $z \sim \mathcal{N}(0, \sigma^2 I_d)$. Its density under D is proportional to $\exp(-\|z\|_2^2/(2\sigma^2))$. Under D', the same point is at distance $z + v$ from the mean $f(D')$, so its density is proportional to $\exp(-\|z + v\|_2^2/(2\sigma^2))$, with the same constant.
2. The constants cancel in the ratio, and the logarithm gives

$$L = \frac{\|z + v\|_2^2 - \|z\|_2^2}{2\sigma^2} = \frac{1}{2\sigma^2}\sum_{j=1}^{d}\left((z_j + v_j)^2 - z_j^2\right) = \frac{1}{2\sigma^2}\sum_{j=1}^{d}\left(2z_j v_j + v_j^2\right). \tag{5.18}$$

3. The constant part is $\sum_j v_j^2/(2\sigma^2) = \rho$, the claimed mean.
4. The random part is $\frac{1}{\sigma^2}\sum_j v_j z_j$, and the terms with $v_j = 0$ vanish. By part 1 of Lemma 5.3 each remaining term $v_j z_j$ is $\mathcal{N}(0, \sigma^2 v_j^2)$, and the terms are independent. By part 2 their sum is $\mathcal{N}(0, \sigma^2\|v\|_2^2)$. Dividing by σ^2 gives $\mathcal{N}(0, \|v\|_2^2/\sigma^2) = \mathcal{N}(0, 2\rho)$.

□

The mean and the variance of the privacy loss are tied, with variance twice the mean. For a single coordinate, $d = 1$, the proof is just the identity $L = vz/\sigma^2 + v^2/(2\sigma^2)$. A Gaussian privacy loss is special to this mechanism. The Laplace mechanism, by Example 5.1, has a loss with atoms.

Proof. Proof of Theorem 5.2.

1. By part 2 of Theorem 5.1, it suffices to show $\Pr[L > \varepsilon] \le \delta$ for every ordered pair of neighbours. If $v = 0$ then $L = 0$, so assume $v \neq 0$.
2. By Lemma 5.4 and part 1 of Lemma 5.3, $L = \rho + \sqrt{2\rho}\,W$ with $W \sim \mathcal{N}(0,1)$. So

$$\Pr[L > \varepsilon] = \Pr\left[W > g(\rho)\right], \qquad g(\rho) = \frac{\varepsilon - \rho}{\sqrt{2\rho}} = \frac{\varepsilon}{\sqrt{2\rho}} - \sqrt{\frac{\rho}{2}}. \tag{5.19}$$

3. The function g decreases in ρ, and $\rho \le (\Delta_2 f)^2/(2\sigma^2)$ by the definition of $\Delta_2 f$. Write $t = \sqrt{2\ln(1/\delta) + \varepsilon}$, so that Equation 5.14 reads $\sigma = t\,\Delta_2 f/\varepsilon$ and the largest value of ρ is $\varepsilon^2/(2t^2)$. Then

$$g(\rho) \ge g\left(\frac{\varepsilon^2}{2t^2}\right) = t - \frac{\varepsilon}{2t} =: u. \tag{5.20}$$

4. The value u is positive, because $t^2 \ge \varepsilon > \varepsilon/2$. Its square is $u^2 = t^2 - \varepsilon + \varepsilon^2/(4t^2) \ge t^2 - \varepsilon = 2\ln(1/\delta)$.
5. By part 3 of Lemma 5.3, $\Pr[L > \varepsilon] \le \Pr[W > u] \le e^{-u^2/2} \le e^{-\ln(1/\delta)} = \delta$.

□

Only the upper tail of L matters in this proof, because part 2 of Theorem 5.1 bounds $\Pr[L > \varepsilon]$, not $\Pr[|L| > \varepsilon]$. For $\varepsilon \le 1$ the noise level Equation 5.14 is at most $\frac{\Delta_2 f}{\varepsilon}\sqrt{2\ln(1.65/\delta)}$. The sharper tail bound of part 4 of Lemma 5.3 lowers the constant 1.65 to 1.25, which is the form of the Gaussian mechanism that later chapters use.

Corollary 5.1 (Gaussian mechanism for $\varepsilon \le 1$)**.** *Let f have $0 < \Delta_2 f < \infty$, and let $\varepsilon \in (0, 1]$ and $\delta \in (0, 1)$. The Gaussian mechanism with*

$$\sigma = \frac{\Delta_2 f}{\varepsilon}\sqrt{2\ln(1.25/\delta)} \tag{5.21}$$

is (ε, δ)-DP.

Proof.

1. Write $t = \sqrt{2\ln(1.25/\delta)}$, so that $\sigma = t\,\Delta_2 f/\varepsilon$ and $e^{-t^2/2} = \delta/1.25$. Steps 1 to 3 of the proof of Theorem 5.2 used only $\sigma = t\,\Delta_2 f/\varepsilon$, so they give $\Pr[L > \varepsilon] \le \Pr[W > u]$ with $u = t - \varepsilon/(2t)$.
2. Since $\varepsilon \le 1$, the value u is at least $t - 1/(2t)$, which increases in t. Also $u^2 = t^2 - \varepsilon + \varepsilon^2/(4t^2) \ge t^2 - 1$.
3. *Case* $\delta \le 0.74$. Then $t \ge \sqrt{2\ln(1.25/0.74)} \ge 1.0239$, and by step 2, $u \ge 0.5356$. By part 4 of Lemma 5.3 and step 2,

$$\Pr[W > u] \le \frac{e^{-u^2/2}}{u\sqrt{2\pi}} \le \frac{e^{1/2}\,e^{-t^2/2}}{0.5356\sqrt{2\pi}} = \frac{e^{1/2}}{1.25 \cdot 0.5356\sqrt{2\pi}}\,\delta \le 0.983\,\delta. \tag{5.22}$$

4. *Case* $\delta > 0.74$. Then $t > \sqrt{2\ln 1.25} \ge 0.6680$, and by step 2, $u > -0.0805$. The density of W is at most $1/\sqrt{2\pi}$, so $\Pr[W > u] \le \frac{1}{2} + 0.0805/\sqrt{2\pi} \le 0.533 < \delta$.
5. In both cases $\Pr[L > \varepsilon] \le \delta$, and part 2 of Theorem 5.1 gives (ε, δ)-DP.

□

This is the bound of Dwork and Roth (2014, Theorem A.1). On a grid of ε from 0.01 to 1 and δ from 0.89 down to 10^{-20}, the exact δ of Proposition 5.3 at this noise level is at most 0.233 times the target. Both constants are loose, because the tail bounds of Lemma 5.3 are loose. The exact δ comes from evaluating the identity of Lemma 5.1 for the Gaussian privacy loss, instead of bounding its tail.

Proposition 5.3. *Let Φ be the distribution function of $\mathcal{N}(0,1)$, let f have $0 < \Delta_2 f < \infty$, and let $s = \Delta_2 f/\sigma$. For every $\varepsilon \ge 0$, the Gaussian mechanism with noise level σ is $(\varepsilon, \delta(\varepsilon))$-DP for*

$$\delta(\varepsilon) = \Phi\Big(\frac{s}{2} - \frac{\varepsilon}{s}\Big) - e^{\varepsilon}\,\Phi\Big(-\frac{s}{2} - \frac{\varepsilon}{s}\Big), \tag{5.23}$$

and for no smaller δ when some pair of neighbours has $\|f(D) - f(D')\|_2 = \Delta_2 f$.

Proof.

1. Fix a pair of neighbours and let $r = \|v\|_2/\sigma \in (0, s]$, so that $\rho = r^2/2$ and $L = r^2/2 + rW$. By Lemma 5.1, the smallest δ for this pair is $\mathbb{E}[(1 - e^{\varepsilon - L})\mathbf{1}[L > \varepsilon]]$.
2. The event $L > \varepsilon$ is $W > c$ with $c = \varepsilon/r - r/2$, so it has probability $\Phi(r/2 - \varepsilon/r)$.
3. Write φ for the standard normal density. Completing the square gives $e^{-rz}\varphi(z) = e^{r^2/2}\varphi(z + r)$, so $\mathbb{E}[e^{-rW}\mathbf{1}[W > c]] = e^{r^2/2}\,\Phi(-c - r)$. Since $e^{\varepsilon - L} = e^{\varepsilon - r^2/2}e^{-rW}$ and $-c - r = -r/2 - \varepsilon/r$,

$$\mathbb{E}\big[e^{\varepsilon - L}\mathbf{1}[L > \varepsilon]\big] = e^{\varepsilon - r^2/2}\,e^{r^2/2}\,\Phi\Big(-\frac{r}{2} - \frac{\varepsilon}{r}\Big) = e^{\varepsilon}\,\Phi\Big(-\frac{r}{2} - \frac{\varepsilon}{r}\Big). \tag{5.24}$$

4. Subtracting step 3 from step 2 gives Equation 5.23 with r in place of s. Call that function $h(r)$.
5. Write $c_1 = r/2 - \varepsilon/r$ and $c_2 = -r/2 - \varepsilon/r$. Then $(c_1^2 - c_2^2)/2 = (c_1 - c_2)(c_1 + c_2)/2 = -\varepsilon$, so $e^{\varepsilon}\varphi(c_2) = \varphi(c_1)$. Differentiating, $h'(r) = \varphi(c_1)(\frac{1}{2} + \varepsilon/r^2) + e^{\varepsilon}\varphi(c_2)(\frac{1}{2} - \varepsilon/r^2) = \varphi(c_1) > 0$. 6. So h increases in r, the worst pair is the one with the largest r, and $h(r) \le h(s) = \delta(\varepsilon)$, with equality when some pair attains $\Delta_2 f$. □

The exact formula shows the graceful degradation promised in Section 5.3. At a single noise level, the Gaussian mechanism meets a whole family of guarantees.

Example 5.5. For $\varepsilon = 1$ and $\delta = 10^{-5}$, Equation 5.14 gives $\sigma = \sqrt{2\ln(10^5) + 1}\,\Delta_2 f = 4.9016\,\Delta_2 f$. At this noise level, Equation 5.23 gives the following guarantees, all at once:

- $(0.5,\ 6.1 \times 10^{-4})$-DP,
- $(1,\ 3.0 \times 10^{-8})$-DP,
- $(2,\ 3.0 \times 10^{-24})$-DP.

The guarantee at $\varepsilon = 1$ is far better than the $\delta = 10^{-5}$ that was asked for, and the smallest noise level that gives $(1, 10^{-5})$-DP is $\sigma = 3.7306\,\Delta_2 f$. As ε grows, δ falls faster than exponentially, so the privacy loss of this mechanism degrades gracefully.

5.5. Properties

The properties of pure DP proved in the Laplace mechanism chapter carry over to approximate DP. Post-processing and composition carry over unchanged, and group privacy carries over with a worse δ. Post-processing comes first. It says that whatever an analyst computes from a private output, without going back to the data, is as private as the output.

Theorem 5.3 (Post-processing)**.** *Let $M \colon \mathcal{X}^n \to \mathcal{Y}$ be (ε, δ)-DP, and let $F \colon \mathcal{Y} \to \mathcal{Z}$ be a randomized map whose coins are independent of those of M. Then $F \circ M$ is (ε, δ)-DP.*

Proof.

1. Write $F(y) = g_R(y)$, where R is the coins of F and each g_r is a deterministic function. So F is a random choice of a deterministic function.
2. For neighbours $D \sim D'$ and an event $S \subseteq \mathcal{Z}$, condition on R and apply Equation 5.2 to the event $g_R^{-1}(S)$:

$$\Pr[F(M(D)) \in S] = \mathbb{E}_R\Big[\Pr[M(D) \in g_R^{-1}(S)]\Big] \le \mathbb{E}_R\Big[e^{\varepsilon}\Pr[M(D') \in g_R^{-1}(S)] + \delta\Big] = e^{\varepsilon}\Pr[F(M(D')) \in S] + \delta. \tag{5.25}$$

□

For group privacy, recall that an ε-DP mechanism satisfies the bound with $e^{k\varepsilon}$ on datasets that differ in k records. With approximate DP, the δ is multiplied by the factors e^{ε} that it meets along the way.

Theorem 5.4 (Group privacy)**.** *Let M be (ε, δ)-DP, and let $D, D' \in \mathcal{X}^n$ differ in $k \ge 1$ records. For every event S,*

$$\Pr[M(D) \in S] \le e^{k\varepsilon}\Pr[M(D') \in S] + \delta\sum_{j=0}^{k-1} e^{j\varepsilon} \le e^{k\varepsilon}\Pr[M(D') \in S] + ke^{(k-1)\varepsilon}\delta. \tag{5.26}$$

Proof.

1. Change the differing records one at a time to get a chain $D = D^{(0)} \sim D^{(1)} \sim \cdots \sim D^{(k)} = D'$ of neighbours. Write $P_j = \Pr[M(D^{(j)}) \in S]$.
2. We show $P_0 \le e^{j\varepsilon} P_j + \delta \sum_{i=0}^{j-1} e^{i\varepsilon}$ by induction on j. The case $j = 1$ is Equation 5.2.
3. If the bound holds for j, apply Equation 5.2 to $D^{(j)} \sim D^{(j+1)}$, which gives $P_j \le e^{\varepsilon} P_{j+1} + \delta$. Substituting,

$$P_0 \le e^{j\varepsilon}(e^{\varepsilon} P_{j+1} + \delta) + \delta \sum_{i=0}^{j-1} e^{i\varepsilon} = e^{(j+1)\varepsilon} P_{j+1} + \delta \sum_{i=0}^{j} e^{i\varepsilon}. \tag{5.27}$$

4. At $j = k$ this is the first inequality in Equation 5.26. Each of the k terms of the sum is at most $e^{(k-1)\varepsilon}$, which gives the second.

□

The additive term grows exponentially in the group size. It exceeds 1, and the bound says nothing, once k is a little over $\ln(1/\delta)/\varepsilon$.

Composition needs one more tool. The definition bounds probabilities of events, and the proof also needs averages of functions with values in $[0, 1]$.

Lemma 5.5. *Let M be (ε, δ)-DP and let $h \colon \mathcal{Y} \to [0, 1]$. For all neighbours $D \sim D'$,*

$$\mathbb{E}[h(M(D))] \le e^{\varepsilon}\, \mathbb{E}[h(M(D'))] + \delta. \tag{5.28}$$

Proof.

1. A variable with values in $[0, 1]$ has mean equal to the integral of its tail, so $\mathbb{E}[h(M(D))] = \int_0^1 \Pr[h(M(D)) > t]\, \mathrm{d}t$.
2. Each tail is the probability of the event $\{y : h(y) > t\}$, which Equation 5.2 bounds by $e^{\varepsilon} \Pr[h(M(D')) > t] + \delta$.
3. Integrating over $t \in [0, 1]$ gives $e^{\varepsilon}\, \mathbb{E}[h(M(D'))] + \delta$.

□

In an adaptive composition, M_1 runs on the dataset, then M_2 is chosen in the light of the output y_1 of M_1 and runs on the same dataset, and so on. Formally, for each i and each sequence of earlier outputs $y_{<i} = (y_1, \ldots, y_{i-1})$, the mechanism $M_i(\cdot\,; y_{<i})$ is $(\varepsilon_i, \delta_i)$-DP, with fresh coins. The composition releases $(y_1, \ldots, y_k)$.

Theorem 5.5 (Basic composition)**.** *If each $M_i(\cdot\,; y_{<i})$ is $(\varepsilon_i, \delta_i)$-DP for every $y_{<i}$, where the parameters ε_i and δ_i are fixed in advance, then the adaptive composition of $M_1, \ldots, M_k$ is $(\sum_i \varepsilon_i, \sum_i \delta_i)$-DP.*

Proof.

1. The composition of the first $k-1$ mechanisms is itself a mechanism. So by induction on k it suffices to treat $k = 2$.
2. Fix neighbours $D \sim D'$ and an event S in $\mathcal{Y}_1 \times \mathcal{Y}_2$. For an output y_1, write $S_{y_1} = \{y_2 : (y_1, y_2) \in S\}$ and $g_D(y_1) = \Pr[M_2(D; y_1) \in S_{y_1}]$, and define $g_{D'}$ in the same way. Conditioning on the first output,

$$\Pr[M(D) \in S] = \mathbb{E}_{y_1 \sim M_1(D)}[g_D(y_1)]. \tag{5.29}$$

3. Since $M_2(\cdot\,; y_1)$ is $(\varepsilon_2, \delta_2)$-DP and $g_D \le 1$, we have $g_D(y_1) \le \min\{1, e^{\varepsilon_2} g_{D'}(y_1) + \delta_2\} \le h(y_1) + \delta_2$, where $h = \min\{1, e^{\varepsilon_2} g_{D'}\}$ takes values in $[0, 1]$.
4. By Lemma 5.5 applied to M_1 and h,

$$\Pr[M(D) \in S] \le \mathbb{E}_{M_1(D)}[h] + \delta_2 \le e^{\varepsilon_1}\, \mathbb{E}_{M_1(D')}[h] + \delta_1 + \delta_2. \tag{5.30}$$

5. Finally $h \le e^{\varepsilon_2} g_{D'}$, and $\mathbb{E}_{y_1 \sim M_1(D')}[g_{D'}(y_1)] = \Pr[M(D') \in S]$ by step 2 applied to D'. So $\Pr[M(D) \in S] \le e^{\varepsilon_1 + \varepsilon_2} \Pr[M(D') \in S] + \delta_1 + \delta_2$.

□

The parameters must be fixed in advance. If the analyst could choose ε_i in the light of earlier outputs, the sum would no longer be a valid bound without further work. With k mechanisms that are each (ε, δ)-DP, basic composition gives $(k\varepsilon, k\delta)$-DP. The next chapter proves that, for every $\delta' > 0$, the same composition is also $(\varepsilon', k\delta + \delta')$-DP with

$$\varepsilon' = \varepsilon \sqrt{2k \ln(1/\delta')} + k\varepsilon\,(e^{\varepsilon} - 1). \tag{5.31}$$

When ε is of order $1/\sqrt{k}$ or smaller, the second term is of the same order as the first or smaller, so ε' grows like $\sqrt{k}$ rather than k. That allows quadratically more analyses for the same privacy, and it requires $\delta' > 0$ even when every mechanism is pure.

5.6. Summary

Table 5.1.: The chapter's results.

Result	Statement
Definition 5.1	$\Pr[M(D) \in S] \le e^{\varepsilon} \Pr[M(D') \in S] + \delta$ for all neighbours and events
Lemma 5.1	the largest violation over events is $\mathbb{E}[(1 - e^{\varepsilon - L})_+]$
Theorem 5.1	pure DP is $L \le \varepsilon$ almost surely; $\Pr[L > \varepsilon] \le \delta$ implies (ε, δ)-DP; the converse holds with $\varepsilon + t$ and $\delta/(1 - e^{-t})$
Proposition 5.1	publishing each record with probability δ is $(0, \delta)$-DP, so δ must be well below $1/n$

Result	Statement
Example 5.3	for approximate DP, single outcomes do not suffice
Lemma 5.2	$\Delta_2 f \le \Delta_1 f \le \sqrt{d}\,\Delta_2 f$
Lemma 5.4	the Gaussian mechanism's privacy loss is $\mathcal{N}(\rho, 2\rho)$ with $\rho = \lVert f(D) - f(D')\rVert_2^2/(2\sigma^2)$
Theorem 5.2	$\sigma = \Delta_2 f\sqrt{2\ln(1/\delta) + \varepsilon}/\varepsilon$ gives (ε, δ)-DP
Corollary 5.1	for $\varepsilon \le 1$, $\sigma = \Delta_2 f\sqrt{2\ln(1.25/\delta)}/\varepsilon$ suffices
Proposition 5.3	exact $\delta(\varepsilon) = \Phi(s/2 - \varepsilon/s) - e^{\varepsilon}\Phi(-s/2 - \varepsilon/s)$ with $s = \Delta_2 f/\sigma$
Theorem 5.3	post-processing keeps (ε, δ)
Theorem 5.4	groups of k: $(k\varepsilon,\ ke^{(k-1)\varepsilon}\delta)$
Theorem 5.5	adaptive composition adds the ε_i and the δ_i

5.7. References

- Cynthia Dwork, Krishnaram Kenthapadi, Frank McSherry, Ilya Mironov and Moni Naor, "Our Data, Ourselves: Privacy via Distributed Noise Generation", 2006. The definition of (ε, δ)-DP and the first analysis of Gaussian noise.
- Cynthia Dwork and Aaron Roth, *The Algorithmic Foundations of Differential Privacy*, 2014. The privacy loss random variable (Section 3.5), the Gaussian mechanism with the constant $\sqrt{2\ln(1.25/\delta)}$ (Theorem A.1), and the composition theorems (Theorems 3.16 and 3.20).
- Salil Vadhan, "The Complexity of Differential Privacy", 2017. Section 1.6 explains why an additive bound alone fails, and the tutorial proves the properties of approximate DP.
- Shiva Prasad Kasiviswanathan and Adam Smith, "On the 'Semantics' of Differential Privacy: A Bayesian Formulation", 2014. The relation between (ε, δ)-DP and tail bounds on the privacy loss.
- Borja Balle and Yu-Xiang Wang, "Improving the Gaussian Mechanism for Differential Privacy: Analytical Calibration and Optimal Denoising", 2018. The exact δ of the Gaussian mechanism, Equation 5.23, and its calibration.
- Cynthia Dwork, Guy Rothblum and Salil Vadhan, "Boosting and Differential Privacy", 2010. Advanced composition.
- Ryan Rogers, Aaron Roth, Jonathan Ullman and Salil Vadhan, "Privacy Odometers and Filters: Pay-as-you-Go Composition", 2016. What goes wrong, and what can be saved, when the privacy parameters are chosen adaptively.
- Ilya Mironov, "Rényi Differential Privacy", 2017, Cynthia Dwork and Guy Rothblum, "Concentrated Differential Privacy", 2016, and Mark Bun and Thomas Steinke, "Concentrated Differential Privacy: Simplifications, Extensions, and Lower Bounds", 2016. Definitions that describe the whole distribution of the privacy loss rather than one threshold.

6. Advanced Composition

Running k private analyses on the same data costs privacy, and basic composition says the cost adds up: k mechanisms that are each (ε, δ)-differentially private are together $(k\varepsilon, k\delta)$-differentially private. This chapter shows that the privacy parameter grows much more slowly, about as $\varepsilon\sqrt{k}$, if you accept a small extra δ'. The proof has two steps. The first reduces any pair of private output distributions to one simple pair with four outcomes. The second analyses k copies of that pair with Hoeffding's inequality.

6.1. The theorem

Write $D \sim D'$ for neighbouring datasets, which differ in one record (the replace-one relation). The chapter on approximate differential privacy relaxed the definition of differential privacy (DP) by an additive term. A mechanism M is (ε, δ)**-DP** if, for all neighbours $D \sim D'$ and all events S,

$$\Pr[M(D) \in S] \leq e^{\varepsilon} \Pr[M(D') \in S] + \delta. \tag{6.1}$$

The proofs below compare two output distributions at a time, so it helps to name that condition for a pair of random variables.

Definition 6.1 (Indistinguishable pair)**.** Let $\varepsilon, \delta \geq 0$. Two random variables A and B with values in the same space are (ε, δ)**-indistinguishable** if, for every event S,

$$\Pr[A \in S] \leq e^{\varepsilon} \Pr[B \in S] + \delta \qquad \text{and} \qquad \Pr[B \in S] \leq e^{\varepsilon} \Pr[A \in S] + \delta. \tag{6.2}$$

So M is (ε, δ)-DP exactly when $M(D)$ and $M(D')$ are (ε, δ)-indistinguishable for every pair of neighbours.

The theorem is about analyses run one after another on the same data, where each analysis may be chosen in the light of the answers so far. The next definition makes this precise.

Definition 6.2 (Adaptive composition)**.** Let $M_1, \dots, M_k$ be mechanisms, where M_j takes the dataset D and the earlier outputs $y_{<j} = (y_1, \dots, y_{j-1})$ and returns y_j. Their **adaptive composition** M runs them in order on the same D, with y_j drawn from $M_j(D; y_{<j})$ with fresh coins, and outputs the whole transcript $(y_1, \dots, y_k)$.

Basic composition, proved in the two previous chapters, says that if each $M_j(\,\cdot\,; y_{<j})$ is (ε, δ)-DP for every fixed $y_{<j}$, then M is $(k\varepsilon, k\delta)$-DP. The theorem of this chapter improves the first parameter. Throughout, every output space is countable, so that probabilities are sums over outcomes.

Theorem 6.1 (Advanced composition)**.** *Let $\varepsilon \geq 0$, $\delta \in [0, 1)$ and $\delta' \in (0, 1)$. Let M be the adaptive composition of k mechanisms such that $M_j(\,\cdot\,; y_{<j})$ is (ε, δ)-DP for every j and every $y_{<j}$. Then M is $(\tilde{\varepsilon}, \tilde{\delta})$-DP with*

$$\tilde{\varepsilon} = \varepsilon\sqrt{2k\ln(1/\delta')} + k\varepsilon\,\frac{e^{\varepsilon}-1}{e^{\varepsilon}+1}, \qquad \tilde{\delta} = k\delta + \delta'. \tag{6.3}$$

The first term of $\tilde{\varepsilon}$ is the important one, and it grows as $\varepsilon\sqrt{k}$ instead of $k\varepsilon$. The second term is small when ε is small. Its factor is $(e^{\varepsilon}-1)/(e^{\varepsilon}+1) = \tanh(\varepsilon/2)$, and $\tanh x \leq x$ for $x \geq 0$, so the second term is at most $k\varepsilon^2/2$. That is also at most $k\varepsilon(e^{\varepsilon}-1)$, because $e^{\varepsilon}-1 \geq \varepsilon$, so the theorem implies the weaker form with $k\varepsilon(e^{\varepsilon}-1)$ that the previous chapter quoted. The extra slack δ' is a free parameter. A larger δ' makes $\ln(1/\delta')$ smaller and so lowers $\tilde{\varepsilon}$, which means the theorem trades a bigger δ for a smaller ε.

The proof of Theorem 6.1 is in Section 6.3. Before it, Example 6.1 and Figure 6.1 say how much the theorem gains, and Corollary 6.1 turns it into a rule for splitting a privacy budget.

Example 6.1. Take $\varepsilon = 0.1$ and $\delta' = 10^{-6}$, so that $\ln(1/\delta') = 13.8155$ and $\tanh(0.05) = 0.04996$. For $k = 100$ mechanisms, basic composition gives 10. Advanced composition gives $0.1\sqrt{200 \times 13.8155} + 100 \times 0.1 \times 0.04996 = 5.257 + 0.500 = 5.756$. For $k = 1{,}000$, basic composition gives 100 and advanced composition gives $16.623 + 4.996 = 21.618$. Advanced composition is the smaller of the two from $k = 31$ on.

The exact curve in Figure 6.1 shows that Theorem 6.1 is close to the truth, within 7% of it at $k = 10{,}000$. So growth as $\varepsilon\sqrt{k\ln(1/\delta')}$ is how privacy really degrades, not an artefact of the proof. Kairouz, Oh and Viswanath (2015) computed the exact optimal $\tilde{\varepsilon}$ for every k, ε, δ and δ'. Their formula is tighter but harder to use, while Equation 6.3 is a good general-purpose bound.

When you have a total budget $\varepsilon_{\text{total}}$ to spend on k analyses, you can invert Equation 6.3. The next corollary gives a simple choice that always works.

Corollary 6.1. *Let $\varepsilon_{\text{total}} \in (0, 1]$ and $\delta' \in (0, 1/2]$. If each of k adaptively composed mechanisms is (ε, δ)-DP with*

$$\varepsilon = \frac{\varepsilon_{\text{total}}}{2\sqrt{2k\ln(1/\delta')}}, \tag{6.4}$$

then the composition is $(\varepsilon_{\text{total}}, k\delta + \delta')$-DP.

Proof.

1. By Theorem 6.1, the composition is $(\tilde{\varepsilon}, k\delta + \delta')$-DP. Its first term is $\varepsilon\sqrt{2k\ln(1/\delta')} = \varepsilon_{\text{total}}/2$.
2. The second term is $k\varepsilon\tanh(\varepsilon/2) \leq k\varepsilon^2/2$, because $\tanh x \leq x$. Substituting ε gives $k\varepsilon^2/2 = \varepsilon_{\text{total}}^2/(16\ln(1/\delta'))$.
3. Since $\delta' \leq 1/2$, we have $\ln(1/\delta') \geq \ln 2 > 1/8$, and since $\varepsilon_{\text{total}} \leq 1$, the second term is at most $\varepsilon_{\text{total}}^2/2 \leq \varepsilon_{\text{total}}/2$.
4. Adding the two terms gives $\tilde{\varepsilon} \leq \varepsilon_{\text{total}}$. A mechanism that is $(\tilde{\varepsilon}, \tilde{\delta})$-DP is also $(\varepsilon_{\text{total}}, \tilde{\delta})$-DP for every $\varepsilon_{\text{total}} \geq \tilde{\varepsilon}$.

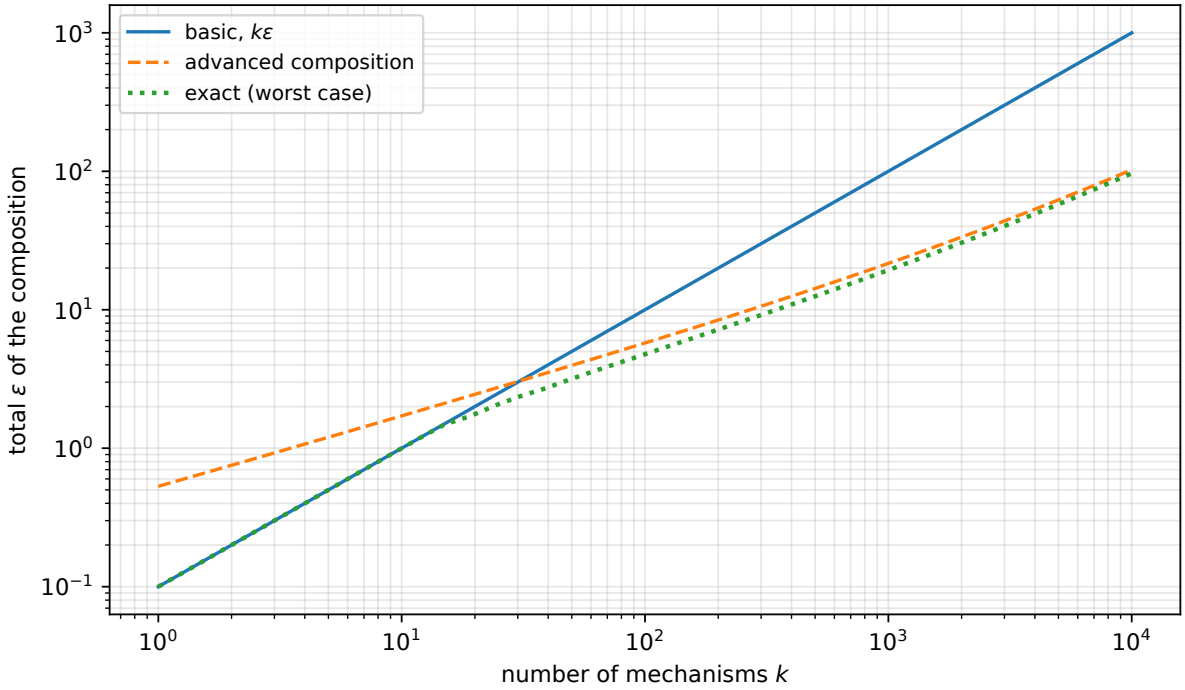


Figure 6.1.: The total privacy parameter of k mechanisms that are each 0.1-DP, at $\delta' = 10^{-6}$. The solid line is basic composition, the dashed line is Equation 6.3, and the dotted line is the exact smallest $\tilde{\varepsilon}$ for k independent randomized responses, which is the worst case for pure DP by Section 6.2. Advanced composition overtakes basic composition at $k = 31$ and is within 21% of the exact value for every k from 100 to 10,000.

□

The corollary is conservative by about a factor of 2. With $\varepsilon_{\text{total}} = 1$, $k = 1{,}000$ and $\delta' = 10^{-6}$, it allows $\varepsilon = 0.00301$ per mechanism, and Equation 6.3 then gives a total of only 0.5045. Solving Equation 6.3 numerically allows $\varepsilon = 0.00591$, almost twice as much. Basic composition with that same ε would give a total of 5.911.

The gain matters most when the number of analyses grows with the problem, as in the next example.

Example 6.2 (Mean estimation with the Laplace mechanism)**.** Each of n people holds a vector $x_i \in \{0,1\}^d$, and the analyst wants the mean $f(D) = \frac{1}{n}\sum_i x_i$. Recall the Laplace mechanism: if f has ℓ_1 sensitivity $\Delta_1 f$, meaning that $\|f(D) - f(D')\|_1 \le \Delta_1 f$ for all neighbours, then adding independent noise from the Laplace distribution with scale $\Delta_1 f/\varepsilon$ to each coordinate is ε-DP. A Laplace variable with scale b has mean 0 and variance $2b^2$, so the expected squared ℓ_2 error of the noise vector is $2db^2$.

All coordinates at once. Replacing one person changes each coordinate of the mean by at most $1/n$, so $\Delta_1 f = d/n$, and the noise has scale $d/(n\varepsilon)$. The root mean squared ℓ_2 error is $\sqrt{2d}\, d/(n\varepsilon) = \sqrt{2}\, d^{3/2}/(n\varepsilon)$.

One coordinate at a time. Each coordinate is a mechanism with sensitivity $1/n$, released with Laplace noise of scale $1/(n\varepsilon_1)$, which makes it ε_1-DP. With ε_1 from Corollary 6.1 for $k = d$ and $\delta = 0$, the d releases together are (ε, δ')-DP for $\varepsilon \le 1$. The root mean squared error is

$$\frac{\sqrt{2d}}{n\varepsilon_1} = \frac{\sqrt{2d} \cdot 2\sqrt{2d\ln(1/\delta')}}{n\varepsilon} = \frac{4d\sqrt{\ln(1/\delta')}}{n\varepsilon}. \tag{6.5}$$

The error now grows as d instead of $d^{3/2}$, which is the order the Gaussian mechanism achieves. Take $n = 100{,}000$, $\varepsilon = 1$ and $\delta' = 10^{-6}$. At $d = 1{,}000$, the two errors are 0.4472 and 0.1487. At $d = 10{,}000$, they are 14.14 and 1.487, nearly ten times apart. The price is the move from pure to approximate DP.

6.2. Reduction to four outcomes

A pair of output distributions $M_j(D)$ and $M_j(D')$ can be complicated. The first step of the proof replaces every such pair by one fixed pair, the simplest pair that is (ε, δ)-indistinguishable. It has two ordinary outcomes, 0 and 1, on which the probabilities differ by the factor e^{ε}. It also has two catastrophic outcomes, each with probability δ, which reveal which distribution the sample came from.

Definition 6.3 (The four-outcome pair)**.** For $\varepsilon \ge 0$ and $\delta \in [0, 1)$, let U and V take values in $\{0, 1, \text{I am } U, \text{I am } V\}$ with

$$\begin{aligned}
&\Pr[U=0]=\frac{e^{\varepsilon}(1-\delta)}{1+e^{\varepsilon}}, \quad \Pr[U=1]=\frac{1-\delta}{1+e^{\varepsilon}}, \qquad \Pr[U=\text{I am } U]=\delta, \quad \Pr[U=\text{I am } V]=0,\\
&\Pr[V=0]=\frac{1-\delta}{1+e^{\varepsilon}}, \qquad \Pr[V=1]=\frac{e^{\varepsilon}(1-\delta)}{1+e^{\varepsilon}}, \quad \Pr[V=\text{I am } U]=0, \quad \Pr[V=\text{I am } V]=\delta.
\end{aligned} \tag{6.6}$$

For $\delta = 0$ this is randomized response with truth probability $e^{\varepsilon}/(1+e^{\varepsilon})$, from the first chapter of this part, where U is the report of a person whose bit is 0 and V that of a person whose bit is 1. The first thing to check is that the pair meets the condition it stands for.

Lemma 6.1. *The pair U, V is (ε, δ)-indistinguishable.*

Proof.

1. On the outcomes 0 and 1, the ratio $\Pr[U = t]/\Pr[V = t]$ is e^{ε} or $e^{-\varepsilon}$, so $\Pr[U = t] \le e^{\varepsilon}\Pr[V = t]$.
2. Take an event S. Sum step 1 over the outcomes of S in $\{0, 1\}$. The outcome "I am U" adds at most δ to $\Pr[U \in S]$, and "I am V" adds nothing. So $\Pr[U \in S] \le e^{\varepsilon}\Pr[V \in S] + \delta$.
3. Swapping the labels $0 \leftrightarrow 1$ and "I am U" $\leftrightarrow$ "I am V" turns the distribution of U into that of V and back, so the same argument gives the other inequality.

□

Indistinguishability survives any further processing, as privacy does. The proof is that of post-processing in the approximate DP chapter, written for a pair of random variables. Its first step extends the definition from events to functions with values in $[0, 1]$, as that chapter does for mechanisms.

Lemma 6.2 (Post-processing)**.** *Let A, B be (ε, δ)-indistinguishable, and let F be a randomized map whose coins are independent of A and B. Then $F(A)$ and $F(B)$ are (ε, δ)-indistinguishable.*

Proof.

1. *Functions with values in* $[0, 1]$. Let g take values in $[0, 1]$. Write its average as an integral of tail probabilities, $\mathbb{E}[g(A)] = \int_0^1 \Pr[g(A) > s]\,\mathrm{d}s$. Each tail is the probability of the event $\{a : g(a) > s\}$, so it is at most $e^{\varepsilon}\Pr[g(B) > s] + \delta$. Integrating over $s \in [0, 1]$ gives $\mathbb{E}[g(A)] \le e^{\varepsilon}\,\mathbb{E}[g(B)] + \delta$.
2. *Post-processing.* Fix an event S and let $g(a) = \Pr[F(a) \in S]$, a probability over the coins of F. Because those coins are independent of A, $\Pr[F(A) \in S] = \mathbb{E}[g(A)]$, and likewise for B.
3. Step 1 applied to g gives $\Pr[F(A) \in S] \le e^{\varepsilon}\Pr[F(B) \in S] + \delta$. The other direction is the same with A and B swapped.

□

The four-outcome pair is the hardest pair to tell apart among all (ε, δ)-indistinguishable pairs, in the following strong sense. Any such pair A, B can be produced from U, V by one randomized map.

Theorem 6.2 (Simulation by the four-outcome pair)**.** *Let $\varepsilon \geq 0$ and $\delta \in [0, 1)$, and let A, B be (ε, δ)-indistinguishable random variables on a countable space. Then there is a randomized map F such that $F(U)$ has the distribution of A and $F(V)$ has the distribution of B.*

The map sends each of the four outcomes to a distribution on the space. The outcome "I am U" covers the excess of A on outcomes y that are more than e^{ε} times as likely under A as under B, and "I am V" covers the reverse excess. The outcome 0 puts more weight on outcomes where A is the more likely of the two, and 1 more weight where B is. For $\delta > 0$ these four distributions must be chosen with care. Kairouz, Oh and Viswanath (2015) prove the theorem by hypothesis testing. Indistinguishability says that every test between A and B has error rates that some test between U and V also has, and Blackwell's theorem on comparing statistical experiments turns that into a randomized map. We prove the case $\delta = 0$, which already gives the theorem for pure DP.

Proof. Proof for $\delta = 0$. Write $p(y) = \Pr[A = y]$ and $q(y) = \Pr[B = y]$.

1. Applying the definition to the single outcome $\{y\}$ gives $p(y) \leq e^{\varepsilon} q(y)$ and $q(y) \leq e^{\varepsilon} p(y)$. If $\varepsilon = 0$, then $p = q$, and F can ignore its input and output a sample of A. Assume $\varepsilon > 0$ from now on.

2. Define

$$P_0(y) = \frac{e^{\varepsilon} p(y) - q(y)}{e^{\varepsilon} - 1}, \qquad P_1(y) = \frac{e^{\varepsilon} q(y) - p(y)}{e^{\varepsilon} - 1}. \tag{6.7}$$

   Both are non-negative by step 1, and each sums to $(e^{\varepsilon} - 1)/(e^{\varepsilon} - 1) = 1$ over y, so they are probability distributions.

3. Let F map the input 0 to a sample from P_0 and the input 1 to a sample from P_1. With $\delta = 0$, U is 0 with probability $e^{\varepsilon}/(1 + e^{\varepsilon})$ and 1 otherwise. So $F(U)$ equals y with probability

$$\frac{e^{\varepsilon} P_0(y) + P_1(y)}{1 + e^{\varepsilon}} = \frac{e^{2\varepsilon} p(y) - e^{\varepsilon} q(y) + e^{\varepsilon} q(y) - p(y)}{(1 + e^{\varepsilon})(e^{\varepsilon} - 1)} = p(y). \tag{6.8}$$

4. In the same way, $F(V)$ equals y with probability $(P_0(y) + e^{\varepsilon} P_1(y))/(1 + e^{\varepsilon}) = q(y)$.

□

Applying Theorem 6.2 once per step of an adaptive composition reduces the whole transcript to a map of k independent copies of the four-outcome pair.

Theorem 6.3. *Let M be the adaptive composition of $M_1, \ldots, M_k$ as in Theorem 6.1, and fix neighbours $D \sim D'$. Let $U_1, \ldots, U_k$ be independent copies of U and $V_1, \ldots, V_k$ independent copies of V. Then there is a randomized map F^* such that $F^*(U_1, \ldots, U_k)$ has the distribution of $M(D)$ and $F^*(V_1, \ldots, V_k)$ has the distribution of $M(D')$.*

Proof.

1. Fix a step j and an earlier transcript $y_{<j}$. The mechanism $M_j(\,\cdot\,; y_{<j})$ is (ε, δ)-DP, so its outputs $M_j(D; y_{<j})$ and $M_j(D'; y_{<j})$ are (ε, δ)-indistinguishable. Theorem 6.2 gives a randomized map $F_j^{y_{<j}}$ that turns U into $M_j(D; y_{<j})$ and V into $M_j(D'; y_{<j})$.
2. Define F^* on an input $(z_1, \dots, z_k)$ step by step. Set $y_1 = F_1(z_1)$, then $y_2 = F_2^{y_1}(z_2)$, and in general $y_j = F_j^{y_{<j}}(z_j)$, with fresh coins at each step. Output $(y_1, \dots, y_k)$.
3. By induction on j, the prefix $(y_1, \dots, y_j)$ computed from $(U_1, \dots, U_k)$ has the distribution of the first j outputs of $M(D)$. For $j = 1$ this is step 1. Given the prefix $y_{<j}$, the input U_j is still a fresh copy of U, because it is independent of $U_1, \dots, U_{j-1}$ and of the coins used so far. So $y_j = F_j^{y_{<j}}(U_j)$ has the distribution of $M_j(D; y_{<j})$, which is how M draws its j-th output.
4. The same induction with $V_1, \dots, V_k$ gives the distribution of $M(D')$.

□

The map F^* depends on D and D', but that is harmless. Privacy is a statement about one pair of neighbours at a time, and Lemma 6.2 applies to any map. So it remains to show that $(U_1, \dots, U_k)$ and $(V_1, \dots, V_k)$ are $(\tilde{\varepsilon}, \tilde{\delta})$-indistinguishable.

6.3. Composing the four-outcome pair

Write $\mathbf{U} = (U_1, \dots, U_k)$ and $\mathbf{V} = (V_1, \dots, V_k)$. As in the approximate DP chapter, the privacy loss of an outcome compares how likely it is under $\mathbf{U}$ and under $\mathbf{V}$. An outcome that is possible under $\mathbf{U}$ but impossible under $\mathbf{V}$ has infinite loss. For an outcome z that is possible under $\mathbf{V}$, the **privacy loss** at z is the log of the ratio of the two probabilities,

$$L(z) = \ln \frac{\Pr[\mathbf{U} = z]}{\Pr[\mathbf{V} = z]}. \tag{6.9}$$

The plan is to show that, for z drawn from $\mathbf{U}$, the loss exceeds $\tilde{\varepsilon}$ only with small probability. Two kinds of outcome are bad. An outcome that contains "I am U" is impossible under $\mathbf{V}$, so its loss is infinite, and this happens with probability about $k\delta$. Among the other outcomes, the loss is a sum of k independent terms, and Hoeffding's inequality bounds the chance that it strays far above its mean by δ'.

The first step computes the loss on the outcomes with no catastrophe, those in which every coordinate is 0 or 1.

Lemma 6.3. *For every $z \in \{0,1\}^k$, the privacy loss is $L(z) = \varepsilon \sum_{j=1}^{k} (1 - 2z_j)$.*

Proof.

1. The coordinates are independent, so both probabilities factor and the log of the ratio is a sum: $L(z) = \sum_{j=1}^{k} \ln\big(\Pr[U = z_j] / \Pr[V = z_j]\big)$.
2. By Definition 6.3, the ratio is e^{ε} when $z_j = 0$ and $e^{-\varepsilon}$ when $z_j = 1$. In both cases its logarithm is $\varepsilon(1 - 2z_j)$.

□

Conditioned on containing no catastrophic outcome, the coordinates of $\mathbf{U}$ are independent bits. The loss is then a sum of independent terms, each equal to ε or $-\varepsilon$, and Hoeffding's inequality applies to it.

Take k independent random variables $W_1, \dots, W_k$, each with values in an interval $[l, u]$ of length $u - l$. **Hoeffding's inequality** bounds the chance that their sum exceeds its mean by a margin $\tau \geq 0$ or more,

$$\Pr\left[\sum_{j=1}^{k} W_j \geq \sum_{j=1}^{k} \mathbb{E}[W_j] + \tau\right] \leq \exp\left(-\frac{2\tau^2}{k(u-l)^2}\right). \tag{6.10}$$

With the loss in hand and Hoeffding's inequality to bound its tail, the main estimate follows. It says that k copies of the four-outcome pair are exactly as hard to tell apart as Theorem 6.1 claims.

Proposition 6.1. *Let $\varepsilon \geq 0$, $\delta \in [0, 1)$ and $\delta' \in (0, 1)$. Then $\mathbf{U}$ and $\mathbf{V}$ are $(\tilde{\varepsilon}, \tilde{\delta})$-indistinguishable, with $\tilde{\varepsilon}$ and $\tilde{\delta}$ as in Equation 6.3.*

Proof.

1. *The events.* Let E_1 be the event that some coordinate of $\mathbf{U}$ is "I am U". Let G be the set of good outcomes, the $z \in \{0,1\}^k$ with $L(z) \leq \tilde{\varepsilon}$. Every outcome z with $\Pr[\mathbf{U} = z] > 0$ that is not in G either lies in E_1 or lies in $\{0,1\}^k$ with $L(z) > \tilde{\varepsilon}$. Outcomes containing "I am V" have probability 0 under $\mathbf{U}$.

2. *The first bad event.* Each coordinate avoids "I am U" with probability $1 - \delta$, independently, so $\Pr[E_1] = 1 - (1 - \delta)^k$. Bernoulli's inequality $(1 - \delta)^k \geq 1 - k\delta$ gives $\Pr[E_1] \leq k\delta$.

3. *The coordinates given no catastrophe.* Given that U_j is not "I am U", it is 0 with probability $\frac{e^{\varepsilon}(1-\delta)/(1+e^{\varepsilon})}{1-\delta} = \frac{e^{\varepsilon}}{1+e^{\varepsilon}}$ and 1 otherwise. Conditioned on the complement $\bar{E}_1$, the coordinates stay independent, because $\bar{E}_1$ is an intersection of independent events, one per coordinate.

4. *The mean of the loss.* By Lemma 6.3, given $\bar{E}_1$ the loss is $\sum_j W_j$ with $W_j = \varepsilon(1 - 2U_j)$. Each W_j is ε with probability $e^{\varepsilon}/(1 + e^{\varepsilon})$ and $-\varepsilon$ otherwise, so

$$\mathbb{E}\Big[\sum_{j=1}^{k} W_j \,\Big|\, \bar{E}_1\Big] = k\varepsilon\,\frac{e^{\varepsilon} - 1}{e^{\varepsilon} + 1}, \tag{6.11}$$

   which is the second term of $\tilde{\varepsilon}$.

5. *The second bad event.* If $\varepsilon = 0$, the loss is $0 \leq \tilde{\varepsilon}$ and there is nothing to bound. Otherwise apply Equation 6.10, conditioned on $\bar{E}_1$, with $[l, u] = [-\varepsilon, \varepsilon]$ and $\tau = \varepsilon\sqrt{2k\ln(1/\delta')}$, the first term of $\tilde{\varepsilon}$. The exponent is $2\tau^2/(4k\varepsilon^2) = \ln(1/\delta')$, so

$$\Pr\left[L(\mathbf{U}) > \tilde{\varepsilon} \,\middle|\, \bar{E}_1\right] \leq \delta'. \tag{6.12}$$

   Multiplying by $\Pr[\bar{E}_1] \leq 1$ bounds the probability that $\mathbf{U} \in \{0,1\}^k$ and $L(\mathbf{U}) > \tilde{\varepsilon}$ by δ'.

6. *Assembly.* Take any event S. Split it into $S \cap G$ and the rest. On G, each outcome has $\Pr[\mathbf{U} = z] \le e^{\tilde{\varepsilon}} \Pr[\mathbf{V} = z]$ by the definition of L. By step 1, the rest has probability at most that of the two bad events. Steps 2 and 5 then give

$$\Pr[\mathbf{U} \in S] \le e^{\tilde{\varepsilon}} \Pr[\mathbf{V} \in S \cap G] + k\delta + \delta' \le e^{\tilde{\varepsilon}} \Pr[\mathbf{V} \in S] + \tilde{\delta}. \tag{6.13}$$

7. *The other direction.* Swapping the labels $0 \leftrightarrow 1$ and "I am U" $\leftrightarrow$ "I am V" in every coordinate maps $\mathbf{U}$ to $\mathbf{V}$ and back. The same argument then gives $\Pr[\mathbf{V} \in S] \le e^{\tilde{\varepsilon}} \Pr[\mathbf{U} \in S] + \tilde{\delta}$.

□

The two terms of $\tilde{\delta}$ now have a meaning. The term $k\delta$ pays for a catastrophic outcome in some step, and δ' pays for an unusually long run of outcomes that favour $\mathbf{U}$. The proof of the theorem only puts the pieces together.

Proof. Proof of Theorem 6.1. Fix neighbours $D \sim D'$.

1. By Theorem 6.3, $M(D)$ and $M(D')$ have the distributions of $F^*(\mathbf{U})$ and $F^*(\mathbf{V})$ for one randomized map F^*.
2. By Proposition 6.1, $\mathbf{U}$ and $\mathbf{V}$ are $(\tilde{\varepsilon}, \tilde{\delta})$-indistinguishable.
3. By Lemma 6.2, so are $F^*(\mathbf{U})$ and $F^*(\mathbf{V})$. The neighbours were arbitrary, so M is $(\tilde{\varepsilon}, \tilde{\delta})$-DP.

□

6.4. Summary

Table 6.1.: The chapter's results.

Result	Statement
Theorem 6.1	k adaptive (ε, δ)-DP steps are $(\varepsilon\sqrt{2k\ln(1/\delta')} + k\varepsilon\tanh(\varepsilon/2),\ k\delta + \delta')$-DP
Corollary 6.1	per-step $\varepsilon = \varepsilon_{\text{total}}/(2\sqrt{2k\ln(1/\delta')})$ gives a total of $\varepsilon_{\text{total}} \le 1$
Example 6.2	Laplace noise per coordinate estimates a mean in $\{0,1\}^d$ with error $4d\sqrt{\ln(1/\delta')}/(n\varepsilon)$
Lemma 6.1	the four-outcome pair is (ε, δ)-indistinguishable
Lemma 6.2	indistinguishability survives post-processing
Theorem 6.2	every (ε, δ)-indistinguishable pair is a randomized map of the four-outcome pair
Theorem 6.3	an adaptive composition is a map of k independent copies of the four-outcome pair
Lemma 6.3	on $\{0,1\}^k$, the privacy loss is $\varepsilon\sum_j(1-2z_j)$

Result	Statement
Proposition 6.1	k copies of the four-outcome pair are $(\tilde{\varepsilon}, \tilde{\delta})$-indistinguishable

6.5. References

- Cynthia Dwork, Guy N. Rothblum and Salil Vadhan, "Boosting and Differential Privacy", 2010. The advanced composition theorem, first proved there.
- Peter Kairouz, Sewoong Oh and Pramod Viswanath, "The Composition Theorem for Differential Privacy", 2015. The simulation of every private pair by the four-outcome pair, by hypothesis testing, and the exact optimal composition.
- Wassily Hoeffding, "Probability Inequalities for Sums of Bounded Random Variables", 1963. The tail bound used for the privacy loss.
- Cynthia Dwork and Aaron Roth, *The Algorithmic Foundations of Differential Privacy*, 2014. Section 3.5 proves advanced composition with the privacy loss random variable, and Theorem 3.20 states it with $k\varepsilon(e^{\varepsilon} - 1)$ as the second term.
- Salil Vadhan, "The Complexity of Differential Privacy", 2017. Section 2 presents basic and advanced composition.
- Jack Murtagh and Salil Vadhan, "The Complexity of Computing the Optimal Composition of Differential Privacy", 2016. Shows that the exact optimal composition of mechanisms with different parameters is hard to compute, which is one reason the simpler bound Equation 6.3 is used in practice.

7. The Exponential Mechanism

The Laplace and Gaussian mechanisms release numbers, and they work because a small change in the number does little harm. Many tasks instead ask for the best object from a set, such as a price, a classifier or a candidate, and there a small change in the output can ruin its quality. This chapter answers that need with the exponential mechanism. It samples each object with probability that grows exponentially with a score, so that good objects are likely while the choice reveals little about any one person, as differential privacy (DP) requires. The chapter proves that the mechanism is private and that its output is nearly optimal, and then applies it to three problems.

7.1. Selecting an object

A seller has a digital good, such as an e-book, that costs nothing to copy, and n buyers who each want one copy. Buyer i values the good at v_i and buys it if the price p is at most v_i. The seller's revenue at price p is the price times the number of buyers,

$$p \cdot \big|\{i : v_i \geq p\}\big|. \tag{7.1}$$

The seller could compute the price with the highest revenue from the valuations, but that price can reveal a single buyer. If one buyer values the good far above the others, the best price may be that buyer's valuation.

Adding noise to the best price does not help, because revenue is very sensitive to the price. Take four buyers, three who value the good at \$1 and one who values it at \$2.50. At a price of \$1, all four buy and the revenue is \$4. At \$1.01 only the last buyer buys, so the revenue falls to \$1.01. It climbs back to \$2.50 at a price of \$2.50, and falls to nothing at \$2.51. So a cent of noise can cost three quarters of the revenue. The prices are therefore better treated as objects with a quality, the revenue, rather than as numbers to perturb.

In general, the input to the selection problem has three parts:

- a dataset $D \in \mathcal{X}^n$,
- a set $\mathcal{R}$ of objects,
- a **score function** $u \colon \mathcal{X}^n \times \mathcal{R} \to \mathbb{R}$, where $u(D, r)$ says how good the object r is for the dataset D.

In the example, D is the list of valuations, $\mathcal{R}$ is the set of prices, and $u(D, p)$ is the revenue at price p. Only the dataset is private. The set of objects and the score function are public, so the sensitivity measures how much the score of a fixed object can move when one record of the dataset changes.

Definition 7.1 (Sensitivity of a score function)**.** The **sensitivity** of a score function u is

$$\Delta u = \sup_{r \in \mathcal{R}} \sup_{D \sim D'} \bigl|u(D, r) - u(D', r)\bigr|, \tag{7.2}$$

where the inner supremum is over neighbouring datasets, which differ in one record (the replace-one relation).

The score is a single real number, so no choice of norm is needed, unlike the ℓ_1 and ℓ_2 sensitivities of the Laplace and Gaussian mechanisms. The mechanism gives each object r a weight that grows exponentially with its score, $\exp(\varepsilon u(D, r)/(2\Delta))$. Here ε is the privacy parameter, and Δ is any upper bound on the sensitivity Δu. Dividing each weight by the sum of all of them, the normalizer, turns the weights into probabilities. For an infinite set of objects, such as an interval of prices, the weights become a density. The sum is then an integral against a fixed base measure μ on $\mathcal{R}$, such as length on the real line.

Definition 7.2 (Exponential mechanism)**.** Let $\varepsilon > 0$, and let $\Delta > 0$ satisfy $\Delta \geq \Delta u$. Let μ be a measure on $\mathcal{R}$, taken to be the counting measure, which gives each object weight 1, when $\mathcal{R}$ is finite. Assume that the normalizer

$$Z(D) = \int_{\mathcal{R}} \exp\left(\frac{\varepsilon\, u(D, r')}{2\Delta}\right) \mu(\mathrm{d}r') \tag{7.3}$$

is positive and finite for every D. The **exponential mechanism** M_E outputs an object $r \in \mathcal{R}$ drawn with density, with respect to μ,

$$p_D(r) = \frac{1}{Z(D)} \exp\left(\frac{\varepsilon\, u(D, r)}{2\Delta}\right). \tag{7.4}$$

For finite $\mathcal{R}$, Equation 7.4 says that $\Pr[M_E(D) = r]$ is proportional to $\exp(\varepsilon u(D, r)/(2\Delta))$. A higher score makes an object more likely. A smaller ε or a larger Δ flattens the distribution towards uniform, which is more private and less accurate.

The definition says nothing about how to sample. For a finite set, you can compute every weight, add them up and sample, in time proportional to $|\mathcal{R}|$. That is slow when $\mathcal{R}$ is exponentially large, which is common. For an infinite set, there is often no direct way to sample, and the usual remedy is to discretize, as Section 7.4.2 does. So the exponential mechanism always gives an algorithm, though not always an efficient one, and it is often the first step towards one.

7.2. Privacy

The score of r changes by at most Δ between neighbours, so the weight of r changes by at most a factor $e^{\varepsilon/2}$. The normalizer is a sum of such weights, so it also changes by at most $e^{\varepsilon/2}$. The two factors together give e^{ε}.

Theorem 7.1. *The exponential mechanism is ε-DP.*

Proof.

1. Fix neighbours $D \sim D'$ and an object r. By Equation 7.4, the ratio of the densities is

$$\frac{p_D(r)}{p_{D'}(r)} = \exp\left(\frac{\varepsilon\,(u(D,r) - u(D',r))}{2\Delta}\right) \cdot \frac{Z(D')}{Z(D)}. \tag{7.5}$$

2. By Definition 7.1 and $\Delta \geq \Delta u$, the difference $u(D,r) - u(D',r)$ is at most Δ, so the first factor is at most $e^{\varepsilon/2}$.

3. For every object r', the same bound gives $u(D',r') \leq u(D,r') + \Delta$, so $\exp(\varepsilon u(D',r')/(2\Delta)) \leq e^{\varepsilon/2}\exp(\varepsilon u(D,r')/(2\Delta))$. Integrating over r' gives $Z(D') \leq e^{\varepsilon/2} Z(D)$, so the second factor is at most $e^{\varepsilon/2}$.

4. So $p_D(r) \leq e^{\varepsilon} p_{D'}(r)$ for every r. For an event $S \subseteq \mathcal{R}$, integrating over S gives $\Pr[M_E(D) \in S] \leq e^{\varepsilon}\Pr[M_E(D') \in S]$.

□

Step 3 is where the factor 2 in 2Δ is spent, because it pays for the change in the normalizer. When the normalizer does not depend on the dataset, that half of the budget is not needed.

Corollary 7.1. *Suppose that $Z(D)$, computed with Δ in place of 2Δ, is the same for every dataset D. Then the mechanism with density proportional to $\exp(\varepsilon u(D,r)/\Delta)$ is ε-DP.*

Proof. In step 1 of the proof of Theorem 7.1, with Δ in place of 2Δ, the factor $Z(D')/Z(D)$ equals 1, and the first factor is at most $\exp(\varepsilon\Delta/\Delta) = e^{\varepsilon}$. Step 4 then applies unchanged. □

7.3. Utility

The mechanism should output an object whose score is close to the best one. Write the best score and the set of objects that attain it as

$$\mathrm{OPT}(D) = \max_{r \in \mathcal{R}} u(D,r), \qquad \mathcal{R}^*(D) = \{r \in \mathcal{R} : u(D,r) = \mathrm{OPT}(D)\}. \tag{7.6}$$

The theorem says that the output's score is rarely more than about $(2\Delta/\varepsilon)\ln|\mathcal{R}|$ below $\mathrm{OPT}(D)$. The loss grows linearly with the sensitivity and with $1/\varepsilon$, but only logarithmically with the number of objects.

Theorem 7.2. *Let $\mathcal{R}$ be finite. For every dataset D and every $t \geq 0$,*

$$\Pr\left[u(D, M_E(D)) \leq \mathrm{OPT}(D) - \frac{2\Delta}{\varepsilon}\left(\ln\frac{|\mathcal{R}|}{|\mathcal{R}^*(D)|} + t\right)\right] \leq e^{-t}. \tag{7.7}$$

Proof.

1. Fix a threshold τ. The output has score at most τ exactly when it is one of the objects with $u(D, r) \le \tau$, so

$$\Pr[u(D, M_E(D)) \le \tau] = \frac{\sum_{r : u(D,r) \le \tau} \exp(\varepsilon u(D, r)/(2\Delta))}{\sum_{r' \in \mathcal{R}} \exp(\varepsilon u(D, r')/(2\Delta))}. \tag{7.8}$$

2. *Numerator.* There are at most $|\mathcal{R}|$ terms, and each is at most $\exp(\varepsilon\tau/(2\Delta))$ because its score is at most τ.

3. *Denominator.* Every term is positive, so keeping only the $|\mathcal{R}^*(D)|$ terms with score $\mathrm{OPT}(D)$ gives a lower bound of $|\mathcal{R}^*(D)| \exp(\varepsilon \,\mathrm{OPT}(D)/(2\Delta))$.

4. Together, steps 2 and 3 give

$$\Pr[u(D, M_E(D)) \le \tau] \le \frac{|\mathcal{R}|}{|\mathcal{R}^*(D)|} \exp\left(\frac{\varepsilon\,(\tau - \mathrm{OPT}(D))}{2\Delta}\right). \tag{7.9}$$

5. Set $\tau = \mathrm{OPT}(D) - (2\Delta/\varepsilon)(\ln(|\mathcal{R}|/|\mathcal{R}^*(D)|) + t)$. The exponent becomes $-\ln(|\mathcal{R}|/|\mathcal{R}^*(D)|) - t$, and the right-hand side becomes e^{-t}.

□

At least one object attains the best score, so $|\mathcal{R}^*(D)| \ge 1$. Dropping it gives the form that is used most often.

Corollary 7.2. *Let $\mathcal{R}$ be finite. For every dataset D and every $t \ge 0$,*

$$\Pr\left[u(D, M_E(D)) \le \mathrm{OPT}(D) - \frac{2\Delta}{\varepsilon}(\ln|\mathcal{R}| + t)\right] \le e^{-t}. \tag{7.10}$$

Proof.

1. At least one object attains the best score, so $|\mathcal{R}^*(D)| \ge 1$ and $\ln(|\mathcal{R}|/|\mathcal{R}^*(D)|) \le \ln|\mathcal{R}|$. The threshold in Equation 7.10 is therefore at most the threshold of Theorem 7.2 for the same t, so the event here is contained in the event there, and its probability is at most e^{-t}.

□

7.4. Applications

Each application below follows the same recipe. It chooses the set of objects and the score function, bounds the sensitivity, and reads off the privacy and accuracy from Theorem 7.1 and Corollary 7.2.

7.4.1. The Laplace mechanism

Let $f\colon \mathcal{X}^n \to \mathbb{R}$ have sensitivity Δf, which means that $|f(D) - f(D')| \le \Delta f$ for all neighbours. Take the real line as the set of objects, with Lebesgue measure, and score each candidate answer by its distance from the true one,

$$u(D, r) = -|f(D) - r|. \tag{7.11}$$

By the triangle inequality, $\big||f(D) - r| - |f(D') - r|\big| \le |f(D) - f(D')| \le \Delta f$, so $\Delta u \le \Delta f$. With $\Delta = \Delta f$, the density Equation 7.4 is proportional to $\exp(-\varepsilon|f(D) - r|/(2\Delta f))$. That is the density of $f(D)$ plus Laplace noise of scale $2\Delta f/\varepsilon$, so the exponential mechanism contains a Laplace mechanism.

Its noise is twice what the Laplace mechanism needs, which adds noise of scale $\Delta f/\varepsilon$. Corollary 7.1 removes the extra factor.

Proposition 7.1. *Let $f\colon \mathcal{X}^n \to \mathbb{R}$ have sensitivity $\Delta f > 0$. The mechanism that outputs $f(D)$ plus noise from the Laplace distribution with scale $\Delta f/\varepsilon$ is ε-DP.*

Proof.

1. The output has density proportional to $\exp(-\varepsilon|f(D) - r|/(\Delta f)) = \exp(\varepsilon u(D, r)/(\Delta f))$, with the score above, whose sensitivity is at most Δf.
2. The normalizer is $\int_{\mathbb{R}} \exp(-\varepsilon|f(D) - r|/(\Delta f))\,\mathrm{d}r = 2\Delta f/\varepsilon$, after the substitution $u = r - f(D)$. It does not depend on D.
3. Corollary 7.1, with $\Delta = \Delta f$, applies and gives ε-DP.

□

7.4.2. Pricing a digital good

Return to the seller of Section 7.1, now with a dataset of valuations $D = (v_1, \dots, v_n) \in [0, 1]^n$ and prices in $[0, 1]$. Write γ for the best revenue over all prices,

$$\gamma = \max_{p \in [0,1]} p \cdot |\{i : v_i \ge p\}|. \tag{7.12}$$

The maximum exists, because the revenue increases linearly between consecutive valuations and drops just after each one, so it is attained at one of the v_i or at 0. A buyer also gains little from misreporting their valuation to a private mechanism, because their report barely changes the price. That makes such auctions approximately truthful, a link explored by McSherry and Talwar (2007).

The set of all prices is infinite, and the revenue has no structure that makes sampling easy, so the mechanism runs on a grid. Fix an integer $m \ge 1$ and the step $\alpha = 1/m$, and use the m prices

$$\mathcal{R} = \{\alpha, 2\alpha, \dots, 1\}. \tag{7.13}$$

Write $\mathrm{OPT}(D)$ for the best revenue over the grid. Rounding the best price down to the grid loses a little revenue, and the next lemma bounds that loss.

Lemma 7.1. *The best revenue on the grid satisfies* $\mathrm{OPT}(D) \geq \gamma - \alpha n$.

Proof.

1. Let p^* attain γ. If $p^* < \alpha$, then $\gamma \leq p^* n < \alpha n$, so $\gamma - \alpha n < 0$. Revenue is never negative, so $\mathrm{OPT}(D) \geq 0 > \gamma - \alpha n$.
2. Otherwise, round p^* down to the grid, $p' = \alpha \lfloor p^*/\alpha \rfloor \in \mathcal{R}$, so that $p^* - \alpha < p' \leq p^*$.
3. Every buyer with $v_i \geq p^*$ also has $v_i \geq p'$. So the revenue at p' is at least $p' \cdot |\{i : v_i \geq p^*\}| \geq (p^* - \alpha)|\{i : v_i \geq p^*\}| \geq \gamma - \alpha n$.

□

The private price therefore loses revenue twice, once to rounding and once to the random selection. The theorem adds the two losses, with a grid of about εn prices.

Theorem 7.3. *Let* $\varepsilon > 0$, $m = \lceil \varepsilon n \rceil$ *and* $\alpha = 1/m$. *The exponential mechanism with the grid* $\mathcal{R}$, *the revenue* $u(D, p) = p \cdot |\{i : v_i \geq p\}|$ *as score and* $\Delta = 1$ *is* ε*-DP. For every* $D \in [0,1]^n$ *and* $t \geq 0$, *with probability at least* $1 - e^{-t}$, *the revenue at its output price is at least*

$$\gamma - \frac{1 + 2\ln\lceil \varepsilon n \rceil + 2t}{\varepsilon}. \tag{7.14}$$

Proof.

1. *Sensitivity.* Changing one valuation changes the number of buyers at any price p by at most 1, and $p \leq 1$. So $\Delta u \leq 1 = \Delta$, and Theorem 7.1 gives ε-DP.
2. *Selection.* By Corollary 7.2 with $|\mathcal{R}| = m$, with probability at least $1 - e^{-t}$ the output has revenue at least $\mathrm{OPT}(D) - (2/\varepsilon)(\ln m + t)$.
3. *Discretization.* By Lemma 7.1, $\mathrm{OPT}(D) \geq \gamma - n/m$, and $n/m \leq n/(\varepsilon n) = 1/\varepsilon$ because $m \geq \varepsilon n$.
4. Adding the two losses gives Equation 7.14.

□

The loss is $O(\log(\varepsilon n)/\varepsilon)$. When a constant fraction of buyers values the good above a constant price, γ grows linearly in n, so the loss is a vanishing fraction of the best revenue. The grid size balances the two losses. A finer grid loses less to rounding, αn, but more to selection, $(2/\varepsilon)\ln(1/\alpha)$. An example shows how the guarantee compares with the revenue the mechanism actually earns.

Example 7.1. Draw $n = 1{,}000$ valuations uniformly from $[0, 1]$ and take $\varepsilon = 0.5$, so that the grid has $m = 500$ prices. The best revenue is $\gamma = 267.40$, and the best grid price earns 266.88. With $t = \ln 20$, Equation 7.14 guarantees a revenue of at least $267.40 - 38.84 = 228.56$ with probability 0.95. The exact output distribution does much better. Its mean revenue is 263.86, and its 5th percentile is 256.52, a loss of 10.88, so the guarantee is about 3.6 times too pessimistic. Figure 7.1 shows the output distribution on a smaller example.

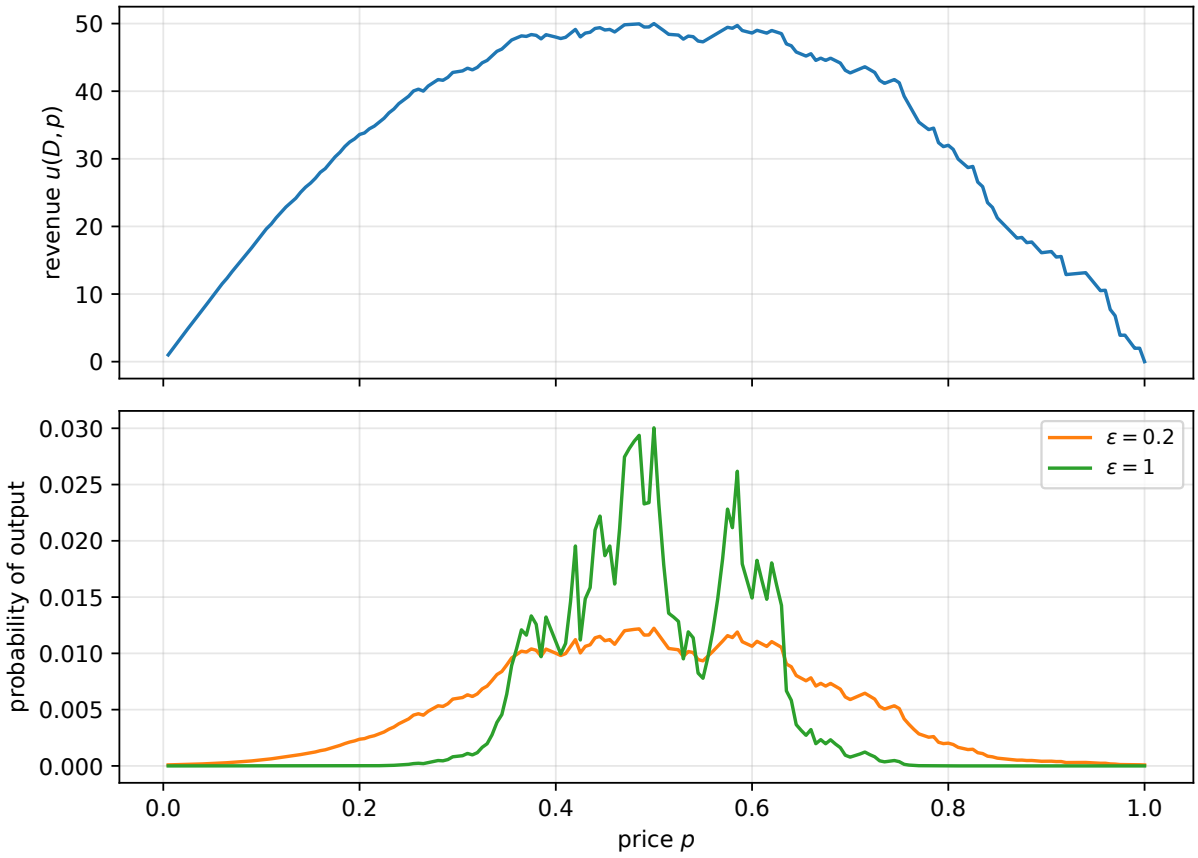


Figure 7.1.: The exponential mechanism for pricing. The top panel is the revenue of 200 buyers with valuations drawn uniformly from $[0, 1]$, on a grid of 200 prices. The bottom panel is the probability with which the mechanism outputs each price, for $\varepsilon = 0.2$ and $\varepsilon = 1$. Prices near the best revenue are the most likely, and a larger ε concentrates the distribution on them.

7.4.3. Private PAC learning

In the probably approximately correct (PAC) model of learning, a learner sees labelled records and must output a classifier that is accurate on new points from the same source. Here the records are points of the Boolean cube $\{0,1\}^d$, labelled 0 or 1. The definition measures accuracy by an error α and a failure probability β, and in this section α no longer denotes the grid step of the pricing example.

Definition 7.3 (PAC learning)**.** A **concept class** is a finite set $\mathcal{C}$ of functions $c\colon \{0,1\}^d \to \{0,1\}$. An unknown **target** $c^* \in \mathcal{C}$ and an unknown distribution $\mathcal{D}$ on $\{0,1\}^d$ generate the dataset. Each record is a point x_i, drawn independently from $\mathcal{D}$, together with its label $y_i = c^*(x_i)$. So the dataset is $D = (z_1, \dots, z_n)$ with labelled records $z_i = (x_i, y_i)$. The **error** of a function h is

$$\operatorname{err}(h) = \Pr_{x\sim\mathcal{D}}[h(x) \neq c^*(x)]. \tag{7.15}$$

A learner is (α,β)**-accurate** if, for every target and every $\mathcal{D}$, it outputs a function $\hat{h}$ with $\operatorname{err}(\hat{h}) \le \alpha$ with probability at least $1-\beta$.

For example, $\mathcal{C}$ might be a finite set of half-spaces, each labelling the points on one side of a hyperplane 1. The learner sees only the labelled points, and any half-space that separates them correctly is a reasonable guess.

Without privacy, the classic learner outputs any function in $\mathcal{C}$ that labels every record correctly. One exists, because the target does. It works because, for every function at once, the fraction of mistakes on the sample is close to the error. Write that fraction as

$$\widehat{\operatorname{err}}(h) = \frac{1}{n}\big|\{i : h(x_i) \neq y_i\}\big|. \tag{7.16}$$

The tool is again Hoeffding's inequality. Take n independent random variables with values in $[0,1]$. Hoeffding's inequality says that their average falls below its mean by a margin $s \ge 0$ or more with probability at most $\exp(-2ns^2)$. Applied to the mistakes of one function, and combined with a union bound over $\mathcal{C}$, it shows that the sample error is close to the true error for every function at once.

Lemma 7.2. *Let $\alpha,\beta \in (0,1)$ and $n \ge 2\ln(|\mathcal{C}|/\beta)/\alpha^2$. With probability at least $1-\beta$ over the sample, every $h \in \mathcal{C}$ satisfies $\operatorname{err}(h) \le \widehat{\operatorname{err}}(h) + \alpha/2$. In particular, every function in $\mathcal{C}$ that labels the whole sample correctly has error at most $\alpha/2$.*

Proof.

1. Fix $h \in \mathcal{C}$. The indicators $W_i = \mathbf{1}[h(x_i) \neq c^*(x_i)]$ are independent, lie in $[0,1]$, and have mean $\operatorname{err}(h)$, because the x_i are independent draws from $\mathcal{D}$. Their average is $\widehat{\operatorname{err}}(h)$.
2. Hoeffding's inequality with $s = \alpha/2$ gives $\Pr[\widehat{\operatorname{err}}(h) \le \operatorname{err}(h) - \alpha/2] \le \exp(-n\alpha^2/2) \le \beta/|\mathcal{C}|$, by the bound on n.
3. A union bound over the $|\mathcal{C}|$ functions shows that some h fails with probability at most β.
4. A function that labels the sample correctly has $\widehat{\operatorname{err}}(h) = 0$, so its error is at most $\alpha/2$.

□

So about $\ln|\mathcal{C}|/\alpha^2$ records suffice. For a function that labels the sample correctly, a multiplicative Chernoff bound improves this to about $\ln|\mathcal{C}|/\alpha$. For infinite classes, the uniform convergence theorem replaces $\ln|\mathcal{C}|$ by the Vapnik–Chervonenkis dimension.

For privacy, neighbouring datasets differ in one labelled record $z_i = (x_i, y_i)$, which may be replaced by any pair $z_i' = (x_i', y_i')$. The new point need not come from $\mathcal{D}$, and its label need not equal $c^*(x_i')$. Accuracy is proved for data drawn from the model, but privacy must hold for every dataset, so that a wrong model never costs anyone their privacy.

The private learner keeps the classic learner's aim of few mistakes on the sample, but it chooses the function with the exponential mechanism. Its score is minus the sample error, so a function with fewer mistakes is more likely to be chosen.

Theorem 7.4 (Private PAC learning)**.** *Let $\varepsilon > 0$ and $\alpha, \beta \in (0,1)$. Run the exponential mechanism with $\mathcal{R} = \mathcal{C}$, the score $u(D,h) = -\widehat{\mathrm{err}}(h)$ and $\Delta = 1/n$. It is ε-DP for every dataset. If*

$$n \geq \frac{2\ln(2|\mathcal{C}|/\beta)}{\alpha^2} \qquad \text{and} \qquad n \geq \frac{4\ln(2|\mathcal{C}|/\beta)}{\alpha\varepsilon}, \tag{7.17}$$

then it is an (α, β)-accurate learner.

Proof.

1. *Privacy.* Replacing one labelled record changes the number of mistakes of each h by at most 1, so $\Delta u \leq 1/n$. Theorem 7.1 gives ε-DP for all datasets, whatever their labels.
2. *Selection.* The target makes no mistakes, so $\mathrm{OPT} = 0$. By Corollary 7.2 with $t = \ln(2/\beta)$, with probability at least $1 - \beta/2$ over the mechanism's coins, the output $\hat{h}$ has

$$\widehat{\mathrm{err}}(\hat{h}) \leq \frac{2}{\varepsilon n}\bigl(\ln|\mathcal{C}| + \ln(2/\beta)\bigr) = \frac{2\ln(2|\mathcal{C}|/\beta)}{\varepsilon n} \leq \frac{\alpha}{2}, \tag{7.18}$$

   where the last step uses the second bound on n.
3. *Generalization.* The first bound on n is the condition of Lemma 7.2 with $\beta/2$ in place of β. So with probability at least $1 - \beta/2$ over the sample, every $h \in \mathcal{C}$, including $\hat{h}$, has $\mathrm{err}(h) \leq \widehat{\mathrm{err}}(h) + \alpha/2$.
4. By a union bound, both events hold with probability at least $1 - \beta$, and then $\mathrm{err}(\hat{h}) \leq \alpha/2 + \alpha/2 = \alpha$.

□

The first bound on n is what learning needs without privacy, and the second is the cost of privacy. Both are logarithmic in $|\mathcal{C}|$. For ε of constant size and small α, the cost of privacy is a lower-order term. For instance, with $|\mathcal{C}| = 2^{20}$, $\alpha = \beta = 0.05$ and $\varepsilon = 1$, we have $\ln(2|\mathcal{C}|/\beta) = 17.552$, and the two bounds are 14,041 and 1,404. The price is computation, since the mechanism scores every function in $\mathcal{C}$.

7.5. Summary

Table 7.1.: The chapter's results.

Result	Statement
Definition 7.2	output r with probability proportional to $\exp(\varepsilon u(D, r)/(2\Delta))$
Theorem 7.1	the exponential mechanism is ε-DP
Corollary 7.1	with a normalizer independent of D, $\exp(\varepsilon u/\Delta)$ is ε-DP
Theorem 7.2	the score falls below $\mathrm{OPT} - \frac{2\Delta}{\varepsilon}(\ln \frac{\lvert\mathcal{R}\rvert}{\lvert\mathcal{R}^*\rvert} + t)$ with probability at most e^{-t}
Corollary 7.2	the same with $\ln \lvert\mathcal{R}\rvert$
Proposition 7.1	the Laplace mechanism is the exponential mechanism with score $-\lvert f(D) - r\rvert$
Theorem 7.3	a private price earns $\gamma - O(\log(\varepsilon n)/\varepsilon)$
Lemma 7.2	$2\ln(\lvert\mathcal{C}\rvert/\beta)/\alpha^2$ records make every sample error close to the true error
Theorem 7.4	private PAC learning with $O(\ln(\lvert\mathcal{C}\rvert/\beta)(1/\alpha^2 + 1/(\alpha\varepsilon)))$ records

7.6. References

- Frank McSherry and Kunal Talwar, "Mechanism Design via Differential Privacy", 2007. The exponential mechanism, with the digital goods auction and its approximate truthfulness.
- Cynthia Dwork and Aaron Roth, *The Algorithmic Foundations of Differential Privacy*, 2014. Section 3.4 presents the exponential mechanism and its privacy and utility theorems, and Section 10.1 the link with truthful mechanism design.
- Salil Vadhan, "The Complexity of Differential Privacy", 2017. Section 8 treats private PAC learning, where the exponential mechanism learns every finite concept class.
- Andrew V. Goldberg, Jason D. Hartline and Andrew Wright, "Competitive Auctions and Digital Goods", 2001. The digital goods auction as a problem in mechanism design.
- Leslie G. Valiant, "A Theory of the Learnable", 1984. The PAC model of learning.
- Shiva Prasad Kasiviswanathan, Homin K. Lee, Kobbi Nissim, Sofya Raskhodnikova and Adam Smith, "What Can We Learn Privately?", 2008. Private PAC learning of every finite class with $O(\ln \lvert\mathcal{C}\rvert)$ examples, by the exponential mechanism.
- Wassily Hoeffding, "Probability Inequalities for Sums of Bounded Random Variables", 1963. The tail bound behind Lemma 7.2.

Part III.

Query Release and Its Limits

The second part built mechanisms that answer a single question, or a few, under differential privacy (DP). This part asks what happens when an analyst wants many answers from the same dataset, and how far the mechanisms of the second part can be pushed.

The first two chapters answer many queries while paying far less than one budget share per query. Private multiplicative weights answers a large set of linear queries with a number of records that grows only logarithmically in the number of queries and in the size of the universe. It borrows an algorithm from online learning and touches the data only through the exponential and Laplace mechanisms. Sparse vector handles queries that arrive one at a time. When only a few of them exceed a threshold, it pays for those few and only logarithmically for the rest. It also turns private multiplicative weights into an online algorithm.

The third chapter questions a choice made in every mechanism so far: noise calibrated to the global sensitivity, the worst case over all pairs of neighbouring datasets. Calibrating to the dataset at hand is tempting, and naively it is not private. The chapter shows four ways to do it safely: testing privately that the dataset is far from any dataset with high sensitivity, a stability argument built on that test for the mode and for histograms, bounding the local sensitivity privately before using it, and smoothing the sensitivity across nearby datasets.

The last chapter turns from algorithms to limits. Packing arguments prove that, under pure DP, no mechanism can answer one-way marginals or histograms with error much smaller than the simple mechanisms already achieve. These lower bounds hold for every algorithm, however clever.

Table 7.2.: The chapters of this part.

Chapter	Question it answers
Private multiplicative weights	How many records are needed to answer a large set of linear queries at once?
Sparse vector	How can a stream of queries be answered when the cost should depend only on the few large ones?
Beyond global sensitivity	When can the noise follow the sensitivity of the actual dataset instead of the worst case?
Packing lower bounds	How much error must any pure DP mechanism incur on marginals and histograms?

8. Private Multiplicative Weights

An analyst wants the answers to many statistical queries about a dataset, such as the fraction of records in each cell of a table or the fraction with a given pair of features. Adding independent noise to each answer makes the error grow polynomially with the number of queries. This chapter shows that, under approximate differential privacy (DP), the number of records needed grows only with the logarithm of the number of queries and the logarithm of the size of the universe. The algorithm that achieves it, private multiplicative weights, comes from online learning. It learns a distribution over the universe that agrees with the data on every query, and it touches the data only through a few private measurements.

8.1. Linear queries

Records take values in a set $\mathcal{X}$, the universe, and a dataset is a list of n records. A linear query asks for the average of a bounded function over the records.

Definition 8.1 (Linear query)**.** A **linear query** is a function $q\colon \mathcal{X} \to [0,1]$ on the universe. On a dataset $D \in \mathcal{X}^n$ with records $x_1, \dots, x_n$ it takes the value

$$q(D) = \frac{1}{n} \sum_{i=1}^{n} q(x_i). \tag{8.1}$$

A counting query is the special case where q takes only the values 0 and 1. Two families come up often. In a **histogram**, there is one query $q_s(x) = \mathbf{1}[x = s]$ for each element s of the universe, and the answers are the fractions of records equal to each s. In **one-way marginals**, the universe is $\{0,1\}^d$, there is one query $q_j(x) = x_j$ for each coordinate j, and the answers are the fractions of records with each feature set to 1. This chapter allows any finite set of linear queries, however irregular.

The analyst wants every query in the set answered at once, so accuracy is measured by the worst error over the set. A mechanism succeeds if, except with a small failure probability β, no answer is off by more than a target error α.

Definition 8.2 (Query release)**.** Let $\mathcal{Q}$ be a finite set of linear queries, $\alpha \in (0,1]$ a target error and $\beta \in (0,1)$ a failure probability. A mechanism $M\colon \mathcal{X}^n \to [0,1]^{\mathcal{Q}}$ outputs one answer $M(D)_q$ for each query $q \in \mathcal{Q}$. It answers $\mathcal{Q}$ **to accuracy** α with probability $1-\beta$ if, for every dataset D,

$$\Pr\left[\max_{q \in \mathcal{Q}} |M(D)_q - q(D)| \le \alpha\right] \ge 1 - \beta. \tag{8.2}$$

The question is how large n must be for a private mechanism to do this. The sample sizes below are sufficient conditions, and smaller is better.

The first answer uses only earlier chapters. Changing one record changes a linear query by at most $1/n$, so its **sensitivity**, the largest change over neighbouring datasets, is at most $1/n$. The Laplace mechanism adds noise from the Laplace distribution $\mathrm{Lap}(b)$ of scale b, whose density at z is $e^{-|z|/b}/(2b)$. It is ε-DP when b is the ℓ_1 sensitivity divided by ε. Its tails are $\Pr[|\mathrm{Lap}(b)| \geq tb] = e^{-t}$ for $t \geq 0$. Adding independent Laplace noise to every query in $\mathcal{Q}$, with the budget split evenly among them, gives the baseline that the rest of the chapter improves on.

Proposition 8.1 (Laplace baseline)**.** *Let $\varepsilon > 0$ and $\beta \in (0,1)$. Release $M(D)_q = q(D) + Z_q$ for every $q \in \mathcal{Q}$, with independent $Z_q \sim \mathrm{Lap}(|\mathcal{Q}|/(\varepsilon n))$. This mechanism is ε-DP, and it answers $\mathcal{Q}$ to accuracy α with probability $1-\beta$ whenever*

$$n \geq \frac{|\mathcal{Q}| \ln(|\mathcal{Q}|/\beta)}{\alpha\varepsilon}. \tag{8.3}$$

Proof.

1. The vector $(q(D))_{q\in\mathcal{Q}}$ changes by at most $1/n$ in each of its $|\mathcal{Q}|$ coordinates when one record changes, so its ℓ_1 sensitivity is at most $|\mathcal{Q}|/n$. The Laplace mechanism with scale $|\mathcal{Q}|/(\varepsilon n)$ is therefore ε-DP.
2. Write $b = |\mathcal{Q}|/(\varepsilon n)$ and $t = \ln(|\mathcal{Q}|/\beta)$. By the tail formula and a union bound over the $|\mathcal{Q}|$ coordinates, $\Pr[\max_q |Z_q| \geq tb] \leq |\mathcal{Q}|e^{-t} = \beta$.
3. The condition on n says exactly that $tb \leq \alpha$. So with probability at least $1-\beta$, every error is below α.

□

Approximate DP does better, through the advanced composition theorem. That theorem implies that running k mechanisms in sequence, each ε_0-DP and each possibly chosen after seeing the outputs of the earlier ones, is (ε', δ)-DP for every $\delta \in (0,1)$, with

$$\varepsilon' = \varepsilon_0\sqrt{2k\ln(1/\delta)} + k\varepsilon_0\left(e^{\varepsilon_0} - 1\right). \tag{8.4}$$

To use the theorem, one starts from a target budget ε and asks how much each of the k mechanisms may spend. The next lemma gives one answer, a per-mechanism budget ε_0 that keeps ε' at most ε. It is the budget-splitting corollary of that chapter, proved again here for the form Equation 8.4.

Lemma 8.1 (Splitting a budget)**.** *Let $\varepsilon \in (0,1]$, $\delta \in (0, 1/2]$ and $k \geq 1$. The adaptive composition of k mechanisms, each ε_0-DP with*

$$\varepsilon_0 = \frac{\varepsilon}{2\sqrt{2k\ln(1/\delta)}}, \tag{8.5}$$

is (ε, δ)-DP.

Proof.

1. The first term of Equation 8.4 is $\varepsilon_0\sqrt{2k\ln(1/\delta)} = \varepsilon/2$.
2. Since $\ln(1/\delta) \geq \ln 2$, $\varepsilon \leq 1$ and $k \geq 1$, we have $\varepsilon_0 \leq 1$. The function e^x is convex, so $e^x - 1$ lies below the chord $(e-1)x \leq 2x$ on $[0,1]$. Hence the second term is at most $2k\varepsilon_0^2 = \varepsilon^2/(4\ln(1/\delta))$.
3. That is at most $\varepsilon/2$ because $\varepsilon \leq 1 \leq 2\ln 2 \leq 2\ln(1/\delta)$. Adding the two terms gives $\varepsilon' \leq \varepsilon$.

□

Applying the lemma with one Laplace mechanism per query, so $k = |\mathcal{Q}|$, lets each query spend a budget of order $\varepsilon/\sqrt{|\mathcal{Q}|}$ instead of $\varepsilon/|\mathcal{Q}|$. The factor $|\mathcal{Q}|$ in the baseline then becomes $\sqrt{|\mathcal{Q}|}$.

Proposition 8.2 (Laplace with advanced composition)**.** *Let $\varepsilon \in (0,1]$, $\delta \in (0, 1/2]$ and $\beta \in (0,1)$. Release $M(D)_q = q(D) + Z_q$ with independent $Z_q \sim \mathrm{Lap}(1/(\varepsilon_0 n))$, where $\varepsilon_0 = \varepsilon/(2\sqrt{2|\mathcal{Q}|\ln(1/\delta)})$. This mechanism is (ε,δ)-DP, and it answers $\mathcal{Q}$ to accuracy α with probability $1-\beta$ whenever*

$$n \geq \frac{2\sqrt{2|\mathcal{Q}|\ln(1/\delta)}\,\ln(|\mathcal{Q}|/\beta)}{\alpha\varepsilon}. \tag{8.6}$$

Proof.

1. Each coordinate is the Laplace mechanism for a query of sensitivity $1/n$, so it is ε_0-DP. Lemma 8.1 with $k = |\mathcal{Q}|$ makes the whole release (ε,δ)-DP.
2. As in Proposition 8.1, $\Pr[\max_q |Z_q| \geq \ln(|\mathcal{Q}|/\beta)/(\varepsilon_0 n)] \leq \beta$, and the condition on n says that $\ln(|\mathcal{Q}|/\beta)/(\varepsilon_0 n) \leq \alpha$.

□

The Gaussian mechanism with advanced composition does a little better, $\sqrt{|\mathcal{Q}|\ln|\mathcal{Q}|\ln(1/\delta)}/(\alpha\varepsilon)$, because its lighter tails make the union bound cost $\sqrt{\ln|\mathcal{Q}|}$ instead of $\ln|\mathcal{Q}|$. These bounds are polynomial in $|\mathcal{Q}|$, and in the worst case over query sets nothing better is possible. A smaller universe helps, though. Table 8.1 collects what is known, up to constants.

Table 8.1.: Sufficient sample sizes for answering every query in $\mathcal{Q}$ to accuracy α.

Method	Privacy	Records needed, up to constants
Laplace, basic composition	ε-DP	$\lvert\mathcal{Q}\rvert\ln\lvert\mathcal{Q}\rvert/(\alpha\varepsilon)$, improvable to $\lvert\mathcal{Q}\rvert/(\alpha\varepsilon)$
Gaussian, advanced composition	(ε,δ)-DP	$\sqrt{\lvert\mathcal{Q}\rvert\ln\lvert\mathcal{Q}\rvert\ln(1/\delta)}/(\alpha\varepsilon)$
Noisy histogram	ε-DP	$\sqrt{\lvert\mathcal{X}\rvert\ln\lvert\mathcal{Q}\rvert}/(\alpha\varepsilon)$
Exponential mechanism over small datasets	ε-DP	$\ln\lvert\mathcal{Q}\rvert\ln\lvert\mathcal{X}\rvert/(\alpha^3\varepsilon)$, up to logarithmic factors
Private multiplicative weights	(ε,δ)-DP	$\ln\lvert\mathcal{Q}\rvert\sqrt{\ln\lvert\mathcal{X}\rvert\ln(1/\delta)}/(\alpha^2\varepsilon)$, up to logarithmic factors

The first two rows lose a logarithmic factor that can almost be removed: $|\mathcal{Q}|/(\alpha\varepsilon)$ suffices for pure DP, which is optimal, and the approximate bound improves to within a factor $\sqrt{\ln\ln\ln|\mathcal{Q}|}$ of $\sqrt{|\mathcal{Q}|\ln(1/\delta)}/(\alpha\varepsilon)$. The noisy histogram releases one noisy count per element of the universe and answers every query from it, so it pays polynomially in $|\mathcal{X}|$. That cost is large: with d binary features, $|\mathcal{X}| = 2^d$. The fourth row is logarithmic in both $|\mathcal{Q}|$ and $|\mathcal{X}|$, but its running time is $|\mathcal{X}|^{O(\ln|\mathcal{Q}|/\alpha^2)}$. Private multiplicative weights keeps both logarithms and runs in time polynomial in $|\mathcal{Q}|$, $|\mathcal{X}|$ and n. That is close to the best possible for an algorithm that outputs synthetic data, as this one does. Under standard cryptographic assumptions, some query sets of size polynomial in $\ln|\mathcal{X}|$ admit no private algorithm that outputs accurate synthetic data in time polynomial in $\ln|\mathcal{X}|$.

8.2. Learning from experts

Multiplicative weights is an algorithm for a game of repeated predictions, which has nothing to do with privacy. This section develops it from the simplest case.

A group of N experts makes predictions over T days. Each day, every expert predicts whether a stock goes up or down, and the algorithm sees these predictions before making its own. Then the true outcome is revealed, and a wrong prediction counts as a **mistake**. Suppose first that one expert, whose identity is unknown, is always right. Then the algorithm can afford to stop listening to any expert as soon as it errs.

Proposition 8.3 (Halving)**.** *Predict with the majority of the experts that have never erred, breaking ties towards "down", and discard every expert that errs. If some expert never errs, this algorithm makes at most* $\log_2 N$ *mistakes.*

Proof.

1. Let S^t be the set of experts still kept on day t, so $|S^1| = N$.
2. On a day with a mistake, the experts who predicted like the algorithm are wrong. They are at least half of S^t, because the algorithm followed the majority. All of them are discarded, so $|S^{t+1}| \le |S^t|/2$.
3. The perfect expert is never discarded, so $|S^t| \ge 1$ on every day. After M mistakes, $1 \le N/2^M$, so $M \le \log_2 N$.

□

The algorithm is measured against the best expert in hindsight. Its **regret** is its total loss minus the total loss of the best single expert. Here the best expert makes no mistakes, so the regret is at most $\log_2 N$.

Now suppose instead that the best expert makes OPT mistakes. Halving fails, because the best expert may err on the first day and be discarded, after which the rest may all be wrong forever. The repair is to lower trust instead of removing it.

Proposition 8.4 (Weighted majority)**.** *Give each expert weight* 1*. Each day, predict with the side whose experts have the larger total weight, breaking ties towards "down", and then halve the weight of every expert that erred. If the best expert makes* OPT *mistakes, the algorithm makes* M *mistakes with*

$$M \le \frac{\mathrm{OPT} + \log_2 N}{\log_2(4/3)} \le 2.41\left(\mathrm{OPT} + \log_2 N\right). \tag{8.7}$$

Proof.

1. Let W^t be the total weight at the start of day t, so $W^1 = N$. Weights never increase.
2. *Upper bound.* On a day with a mistake, the side the algorithm followed carries at least half of W^t, and all of its weight is halved. So $W^{t+1} \le W^t - W^t/4 = \frac{3}{4}W^t$. Over T days, $W^{T+1} \le N(3/4)^M$.
3. *Lower bound.* The best expert's weight is halved exactly OPT times, so $W^{T+1} \ge (1/2)^{\mathrm{OPT}}$.
4. Combining the bounds gives $(4/3)^M \le N\,2^{\mathrm{OPT}}$. Taking $\log_2$ of both sides gives the first inequality, and $1/\log_2(4/3) = 2.409\ldots$ gives the second.

□

The guarantee is competitive with the best expert, but its regret, $1.41\,\mathrm{OPT} + 2.41\log_2 N$, grows with OPT. A randomized algorithm with gentler updates removes the factor in front of OPT.

The general game has losses instead of mistakes. In each of T rounds, every expert i suffers a loss $\ell_i^t \in [-1, 1]$, chosen by an adversary, and a negative loss is a gain. A randomized algorithm picks one expert and suffers that expert's loss. It is cleaner to let the algorithm pick a distribution p^t over the experts and suffer the average loss $\ell^t \cdot p^t = \sum_i p_i^t \ell_i^t$. A distribution is the right object, because a mixture can do well where every single choice fails. In rock-paper-scissors, for example, no single move does well against every opponent, but the uniform mixture does.

Multiplicative weights keeps one weight per expert and plays the distribution proportional to the weights. After each round it multiplies each weight by a factor that falls as that expert's loss rises.

Definition 8.3 (Multiplicative weights)**.** Fix a learning rate $\eta \in (0, 1/2]$. Start with weights $w_i^1 = 1$ for $i = 1, \ldots, N$. In round $t = 1, 2, \ldots$:

1. Play the distribution $p_i^t = w_i^t / W^t$, where $W^t = \sum_{i=1}^N w_i^t$.
2. Observe the loss vector $\ell^t \in [-1, 1]^N$, which may depend on $p^1, \ldots, p^t$, and suffer $\ell^t \cdot p^t$.
3. Update $w_i^{t+1} = w_i^t\,(1 - \eta\ell_i^t)$ for every i.

The factor $1 - \eta\ell_i^t$ lies in $[1/2, 3/2]$, so the weights stay positive. An expert with a positive loss loses weight and one with a negative loss gains it. The analysis takes the logarithm of the total weight, which turns the products of factors $1 - \eta\ell_i^t$ into sums of $\ln(1 - \eta\ell_i^t)$. It needs two bounds on that logarithm, one from above and one from below.

Lemma 8.2. *For every* $x < 1$, $\ln(1 - x) \le -x$*. For every* $x \in [-1/2, 1/2]$, $\ln(1 - x) \ge -x - x^2$.

Proof.

1. The function $\ln(1-x)$ is concave, so it lies below its tangent line at 0, which is $-x$.
2. Let $h(x) = \ln(1-x) + x + x^2$. Then $h(0) = 0$ and $h'(x) = -\frac{1}{1-x} + 1 + 2x = \frac{x(1-2x)}{1-x}$.
3. On $[0, 1/2]$, $h' \geq 0$, so $h(x) \geq h(0) = 0$. On $[-1/2, 0]$, $h' \leq 0$, so h decreases to $h(0) = 0$ and is non-negative there too.

□

The main result compares the total loss of multiplicative weights with the total loss of any fixed distribution p chosen in hindsight. The difference is the regret, and the theorem bounds it by a term that grows with η and a term that shrinks with it.

Theorem 8.1 (Regret of multiplicative weights)**.** *Let $p^1, p^2, \ldots$ be the distributions that multiplicative weights plays. For every learning rate $\eta \in (0, 1/2]$, every number of rounds $T \geq 1$, every sequence of loss vectors $\ell^1, \ldots, \ell^T$ in $[-1, 1]^N$ and every fixed distribution p over the experts,*

$$\sum_{t=1}^{T} \ell^t \cdot p^t \leq \sum_{t=1}^{T} \ell^t \cdot p + \eta T + \frac{\ln N}{\eta}. \tag{8.8}$$

If $N \geq 2$, $T \geq 4 \ln N$ and $\eta = \sqrt{\ln N / T}$, the last two terms add up to $2\sqrt{T \ln N}$.

Proof.

1. *Upper bound on the total weight.* By the update rule and $p_i^t = w_i^t / W^t$,

$$W^{t+1} = \sum_{i=1}^{N} w_i^t (1 - \eta \ell_i^t) = W^t \left(1 - \eta\, \ell^t \cdot p^t\right). \tag{8.9}$$

Unrolling from $W^1 = N$ and applying the first bound of Lemma 8.2 to each factor,

$$\ln W^{T+1} = \ln N + \sum_{t=1}^{T} \ln\left(1 - \eta\, \ell^t \cdot p^t\right) \leq \ln N - \eta \sum_{t=1}^{T} \ell^t \cdot p^t. \tag{8.10}$$

2. *Lower bound through one expert.* The total weight is at least the weight of any expert i. Each $\eta \ell_i^t$ lies in $[-1/2, 1/2]$, so the second bound of Lemma 8.2 applies, and $(\ell_i^t)^2 \leq 1$:

$$\ln W^{T+1} \geq \ln w_i^{T+1} = \sum_{t=1}^{T} \ln(1 - \eta \ell_i^t) \geq -\eta \sum_{t=1}^{T} \ell_i^t - \eta^2 T. \tag{8.11}$$

3. *Lower bound through a distribution.* Multiply the inequality of step 2 by p_i and sum over i. The weights p_i sum to 1, so $\ln W^{T+1} \geq -\eta \sum_t \ell^t \cdot p - \eta^2 T$.
4. Combining steps 1 and 3 gives $-\eta \sum_t \ell^t \cdot p - \eta^2 T \leq \ln N - \eta \sum_t \ell^t \cdot p^t$. Dividing by η and rearranging gives Equation 8.8.
5. With $\eta = \sqrt{\ln N / T}$, both ηT and $\ln N / \eta$ equal $\sqrt{T \ln N}$. The condition $N \geq 2$ makes $\eta > 0$, and $T \geq 4 \ln N$ makes $\eta \leq 1/2$.

□

The regret is of order $\sqrt{T}$, while the losses themselves can add up to order T. So on average over rounds the algorithm does as well as the best fixed distribution, up to a term that vanishes as $\sqrt{\ln N/T}$. Taking p to be a point mass, the distribution that puts all its weight on one expert, recovers the randomized algorithm that picks one expert.

Corollary 8.1. *Let $N \geq 2$. Fix losses $\ell^1, \dots, \ell^T \in [-1,1]^N$ in advance, with $T \geq 4 \ln N$, and let $L_i = \sum_t \ell_i^t$. Pick expert i_t with probability $p_{i_t}^t$, where p^t comes from Definition 8.3 with $\eta = \sqrt{\ln N/T}$. Then*

$$\mathbb{E}\Big[\sum_{t=1}^{T} \ell_{i_t}^t\Big] \leq \min_i L_i + 2\sqrt{T \ln N}. \tag{8.12}$$

Proof. The losses are fixed, so each p^t is fixed too, and $\mathbb{E}[\ell_{i_t}^t] = \ell^t \cdot p^t$. Sum over t and apply Theorem 8.1 with p the point mass on the best expert. □

8.3. Multiplicative weights for queries

To use the game for query release, the experts become the elements of the universe, and the distribution p^t becomes a candidate dataset. To compare a candidate with the true dataset, the true dataset is also written as a distribution over the universe.

Definition 8.4 (Histogram of a dataset)**.** The **histogram** of $D \in \mathcal{X}^n$ is the distribution p_D on $\mathcal{X}$ with $p_D(x) = |\{i : x_i = x\}|/n$, the fraction of records equal to x. For a linear query q and a distribution p on $\mathcal{X}$, write $\langle q, p\rangle = \sum_{x \in \mathcal{X}} q(x)p(x)$.

Grouping the records by value gives $q(D) = \langle q, p_D\rangle$. So answering the queries means finding a distribution $\hat{p}$ with $\langle q, \hat{p}\rangle$ close to $\langle q, p_D\rangle$ for every $q \in \mathcal{Q}$.

Picture now a game in rounds. An adversary picks a query $q^t \in \mathcal{Q}$, and the algorithm, which has already committed to a distribution p^t, is charged its error $|\langle q^t, p^t\rangle - \langle q^t, p_D\rangle|$. The aim is a small total error. This looks pointless, since the algorithm could play $p^t = p_D$ from the start. But that move reads the whole dataset, and multiplicative weights reads it only through the choice of one query per round and the sign of its error, which are cheap to privatize. A bound on the total error also turns into a bound on the number of rounds in which the error can be large, as shown below.

The loss vector comes from the error itself. Write the error as a function of the candidate, $f(p') = |\langle q^t, p'\rangle - \langle q^t, p_D\rangle|$. This function is convex, and $f(p_D) = 0$. Where the error is not zero, its gradient at p^t is $s_t q^t$, where s_t is the sign of $\langle q^t, p^t - p_D\rangle$, taken as $+1$ at 0. Convexity gives $f(p^t) \leq s_t q^t \cdot (p^t - p_D)$, so the loss vector to feed the experts is $\ell^t = s_t q^t$. Here the inequality is an equality, and that is all the next theorem needs.

Lemma 8.3. *Let q be a linear query, p' and p distributions on $\mathcal{X}$, and s the sign of $\langle q, p' - p\rangle$, with $s = +1$ if it is 0. The vector $\ell = sq$ lies in $[-1,1]^{\mathcal{X}}$, and*

$$\ell \cdot (p' - p) = |\langle q, p'\rangle - \langle q, p\rangle|. \tag{8.13}$$

Proof. The entries of q lie in $[0,1]$, so those of sq lie in $[-1,1]$. By linearity, $\ell \cdot (p' - p) = s\langle q, p' - p\rangle$, which is the absolute value of $\langle q, p' - p\rangle$ by the choice of s. □

With these losses, the regret of multiplicative weights against the fixed distribution p_D is exactly the total error of the candidates. The regret bound therefore bounds that total.

Theorem 8.2. *Run multiplicative weights over the experts $\mathcal{X}$ with learning rate $\eta \in (0, 1/2]$, and let p^t be the distribution it plays in round t. Feed it the losses $\ell^t = s_t q^t$, where $q^1, q^2, \ldots$ are any queries in $\mathcal{Q}$ and s_t is the sign of $\langle q^t, p^t - p_D\rangle$. After any number U of rounds,*

$$\sum_{t=1}^{U} |\langle q^t, p^t\rangle - q^t(D)| \le \eta U + \frac{\ln|\mathcal{X}|}{\eta}. \tag{8.14}$$

With $|\mathcal{X}| \ge 2$, $U \ge 4\ln|\mathcal{X}|$ and $\eta = \sqrt{\ln|\mathcal{X}|/U}$, the right-hand side is $2\sqrt{U\ln|\mathcal{X}|}$.

Proof. By Lemma 8.3, the left-hand side equals $\sum_t \ell^t \cdot p^t - \sum_t \ell^t \cdot p_D$. Apply Theorem 8.1 with $N = |\mathcal{X}|$ and the fixed distribution $p = p_D$. □

This bounds a total, while the goal is one distribution that is accurate on every query. Call round t a **mistake** if the error there exceeds α. If the adversary always picks a query on which the current distribution makes a mistake, the total grows by more than α per round. The regret bound grows only by η per round, so the mistakes cannot go on for long. To exploit this, the algorithm plays the adversary against itself, always choosing the hardest query it can find.

Definition 8.5 (Multiplicative weights for queries)**.** Given D, $\mathcal{Q}$ and $\alpha \in (0,1]$, set $\eta = \alpha/2$ and let p^1 be uniform on $\mathcal{X}$. In round $t = 1, 2, \ldots$:

1. If $|\langle q, p^t\rangle - q(D)| \le \alpha$ for every $q \in \mathcal{Q}$, output p^t and stop.
2. Otherwise pick $q^t \in \mathcal{Q}$ with $|\langle q^t, p^t\rangle - q^t(D)| > \alpha$, and let s_t be the sign of $\langle q^t, p^t\rangle - q^t(D)$.
3. Update $p^{t+1}(x) \propto p^t(x)\,(1 - \eta s_t q^t(x))$.

In step 3, the symbol $\propto$ means that the right-hand side is divided by its sum over x, so that p^{t+1} is again a distribution. Every update adds more than α to the total error, while the bound on that total grows by only $\eta = \alpha/2$ per round. So the updates must stop soon.

Corollary 8.2 (Mistake bound)**.** *Let $|\mathcal{X}| \ge 2$. The algorithm of Definition 8.5 stops after fewer than $4\ln|\mathcal{X}|/\alpha^2$ updates, and its output $\hat{p}$ satisfies $|\langle q, \hat{p}\rangle - q(D)| \le \alpha$ for every $q \in \mathcal{Q}$.*

Proof.

1. The algorithm is multiplicative weights with $\eta = \alpha/2 \le 1/2$ and the losses of Theorem 8.2, since normalizing the weights gives the update in step 3.
2. Suppose it makes U updates. Each of them has error above α, so by Theorem 8.2

$$\alpha U < \frac{\alpha}{2} U + \frac{2 \ln |\mathcal{X}|}{\alpha}, \qquad \text{that is,} \qquad U < \frac{4 \ln |\mathcal{X}|}{\alpha^2}. \tag{8.15}$$

3. So the algorithm stops, which happens only at step 1, where every query is within α.

□

The bound is a worst case. Figure 8.1 runs the algorithm on a small example, where the number of updates stays far below the bound.

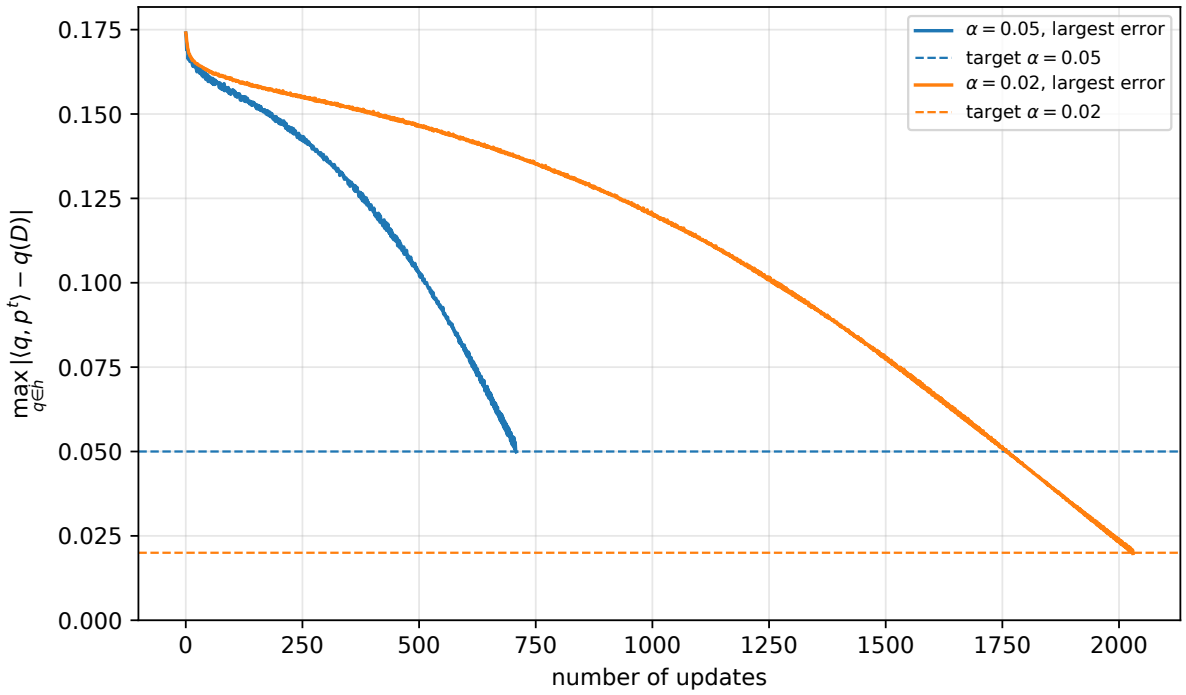


Figure 8.1.: Multiplicative weights for queries on the universe $\{0,1\}^8$, with the 128 one-way and two-way marginal queries $\mathbf{1}[x_j = a]$ and $\mathbf{1}[x_j = a, x_k = b]$. The dataset has 2,000 records whose eight features copy a hidden fair bit with probability 0.8. Each round picks the query with the largest error. The algorithm stops after about 700 updates for $\alpha = 0.05$ and about 2,000 for $\alpha = 0.02$, against the bounds $4 \ln 256/\alpha^2$ of 8,872 and 55,452.

8.4. Private multiplicative weights

The algorithm of Definition 8.5 touches the data in two places: it picks a query with a large error, and it reads the sign of that error. Both have private versions. Picking a query with nearly the

largest error is what the exponential mechanism does, and the error itself has sensitivity $1/n$, so the Laplace mechanism can measure it.

The exponential mechanism takes a finite set $\mathcal{R}$ of candidates and a score $u(D, r)$ whose sensitivity in D is at most Δ for each r. It outputs r with probability proportional to $\exp(\varepsilon u(D, r)/(2\Delta))$. It is ε-DP, and for every $t \geq 0$,

$$\Pr\left[u(D, \hat{r}) \leq \max_{r \in \mathcal{R}} u(D, r) - \frac{2\Delta}{\varepsilon}(\ln|\mathcal{R}| + t)\right] \leq e^{-t}, \tag{8.16}$$

where $\hat{r}$ is its output.

The private algorithm replaces the two steps of Definition 8.5 that read the data by these two mechanisms. It also fixes the number of rounds T in advance, so that the privacy budget can be split among them.

Definition 8.6 (Private multiplicative weights)**.** The inputs are a dataset $D \in \mathcal{X}^n$, a finite set $\mathcal{Q}$ of linear queries, an accuracy $\alpha \in (0, 1]$ and privacy parameters $\varepsilon \in (0, 1]$ and $\delta \in (0, 1/2]$. Set

$$T = \left\lceil \frac{64 \ln|\mathcal{X}|}{\alpha^2} \right\rceil, \qquad \eta = \frac{\alpha}{8}, \qquad \varepsilon_0 = \frac{\varepsilon}{4\sqrt{T \ln(1/\delta)}}, \tag{8.17}$$

and let p^1 be uniform on $\mathcal{X}$. In round $t = 1, \ldots, T$:

1. Choose $q^t \in \mathcal{Q}$ with the exponential mechanism, with privacy parameter ε_0 and score $u_t(D, q) = |\langle q, p^t\rangle - q(D)|$, whose sensitivity is $1/n$.
2. Measure $y_t = \langle q^t, p^t\rangle - q^t(D) + Z_t$, with $Z_t \sim \mathrm{Lap}(1/(\varepsilon_0 n))$.
3. If $|y_t| \leq \alpha/2$, output p^t and stop.
4. Otherwise let s_t be the sign of y_t and update $p^{t+1}(x) \propto p^t(x)\,(1 - \eta s_t q^t(x))$.

If no round stops, output p^{T+1}.

The score u_t has sensitivity $1/n$ because $\langle q, p^t\rangle$ does not depend on D, and $q(D)$ changes by at most $1/n$. The threshold $\alpha/2$ in step 3 sits between the errors that should trigger an update and those small enough to stop on, with room on both sides for the noise.

Every access to the data goes through one of the two mechanisms, so privacy follows from composing them over the T rounds.

Theorem 8.3. *Private multiplicative weights is (ε, δ)-DP.*

Proof.

1. In round t, the distribution p^t is a function of the earlier outputs $q^1, y_1, \ldots, q^{t-1}, y_{t-1}$ alone. So each round runs two mechanisms chosen from earlier outputs: the exponential mechanism, which is ε_0-DP, and the Laplace mechanism on a query of sensitivity $1/n$ with scale $1/(\varepsilon_0 n)$, which is also ε_0-DP.

2. A round after the algorithm stops can be treated as a mechanism with a fixed output, which is 0-DP. So the transcript of all T rounds is an adaptive composition of $k = 2T$ mechanisms, each ε_0-DP.
3. The value $\varepsilon_0 = \varepsilon/(4\sqrt{T\ln(1/\delta)})$ equals $\varepsilon/(2\sqrt{2k\ln(1/\delta)})$ for $k = 2T$. Lemma 8.1 makes the transcript (ε, δ)-DP.
4. The output is a function of the transcript, so it is (ε, δ)-DP by post-processing.

□

Basic composition would also work, and it would even give pure DP, but only with $\varepsilon_0 = \varepsilon/(2T)$. The condition on n in the next theorem would then grow like T instead of $\sqrt{T}$, which means $1/\alpha^3$ instead of $1/\alpha^2$.

Accuracy needs enough records for the errors of both mechanisms to stay below $\alpha/4$ in every round. Then each update still makes progress, and the algorithm stops on a distribution that is accurate on every query.

Theorem 8.4 (Accuracy of private multiplicative weights)**.** *Let $|\mathcal{X}| \geq 2$ and $\beta \in (0, 1)$, and let the other parameters be as in Definition 8.6. If*

$$n \geq \frac{32\sqrt{T\ln(1/\delta)}\,(\ln|\mathcal{Q}| + \ln(2T/\beta))}{\alpha\varepsilon}, \tag{8.18}$$

then with probability at least $1 - \beta$ the output $\hat{p}$ satisfies $|\langle q, \hat{p}\rangle - q(D)| \leq \alpha$ for every $q \in \mathcal{Q}$.

Proof.

1. *The good event.* Let G be the event that in every round t that runs, the chosen query has $u_t(D, q^t) \geq \max_q u_t(D, q) - \alpha/4$ and the noise has $|Z_t| < \alpha/4$. Condition Equation 8.18 is the same as $\frac{2}{n\varepsilon_0}(\ln|\mathcal{Q}| + \ln(2T/\beta)) \leq \frac{\alpha}{4}$, by the value of ε_0. So Equation 8.16 with $t = \ln(2T/\beta)$ bounds the first failure in each round by $\beta/(2T)$. The same condition gives $\ln(2T/\beta)/(\varepsilon_0 n) \leq \alpha/4$, so the Laplace tail bounds the second failure by $\beta/(2T)$. A union bound over the two failures and the T rounds gives $\Pr[G] \geq 1 - \beta$. The rest of the proof assumes G.
2. *Every update makes progress.* In a round that updates, $|y_t| > \alpha/2 > |Z_t|$. So the true error $\langle q^t, p^t\rangle - q^t(D) = y_t - Z_t$ has the sign s_t of y_t, and its absolute value exceeds $\alpha/2 - \alpha/4 = \alpha/4$.
3. *Few updates.* By step 2, the update is the one of Theorem 8.2, with $\eta = \alpha/8 \leq 1/2$. If U rounds update, each has error above $\alpha/4$, and the theorem gives

$$\frac{\alpha}{4}U < \frac{\alpha}{8}U + \frac{8\ln|\mathcal{X}|}{\alpha}, \qquad \text{that is,} \qquad U < \frac{64\ln|\mathcal{X}|}{\alpha^2} \leq T. \tag{8.19}$$

So the T rounds cannot all update, and the algorithm stops at step 3 in some round t. 4. *The output is accurate.* In that round, $|y_t| \leq \alpha/2$ and $|Z_t| < \alpha/4$, so $u_t(D, q^t) < 3\alpha/4$. On G, every query then has $u_t(D, q) \leq u_t(D, q^t) + \alpha/4 < \alpha$, and the output is $\hat{p} = p^t$. □

Since $\ln|\mathcal{X}|/\alpha^2 \geq \ln 2 > 1/2$, rounding up adds less than $2\ln|\mathcal{X}|/\alpha^2$ and $T \leq 66\ln|\mathcal{X}|/\alpha^2$. So condition Equation 8.18 asks for

$$n = O\left(\frac{\sqrt{\ln|\mathcal{X}|\,\ln(1/\delta)}\,(\ln|\mathcal{Q}| + \ln\frac{\ln|\mathcal{X}|}{\alpha\beta})}{\alpha^2\varepsilon}\right) \tag{8.20}$$

records, which is the last row of Table 8.1. It is logarithmic in both $|\mathcal{Q}|$ and $|\mathcal{X}|$. The price is a worse dependence on α, $1/\alpha^2$ instead of $1/\alpha$. Each round evaluates every query on p^t and on D, so the running time is $O(T|\mathcal{Q}|(|\mathcal{X}| + n))$, linear in $|\mathcal{X}|$.

The output is more than a list of answers. It is a distribution $\hat{p}$ over the universe, a **synthetic dataset** that an analyst can use like the original and that agrees with it on every query in $\mathcal{Q}$. That is useful in practice, because a dataset is easier to work with than a table of numbers.

The logarithms pay off only when the number of queries is large, as a comparison of the three sufficient sample sizes shows.

Example 8.1. Take records with $d = 20$ binary features, so $|\mathcal{X}| = 2^{20}$, and $\alpha = 0.05$, $\varepsilon = 1$, $\delta = 10^{-6}$, $\beta = 0.05$. Private multiplicative weights runs for at most $T = 354{,}892$ rounds. For $|\mathcal{Q}| = 10^6$ queries, the sufficient sample sizes of Proposition 8.1, Proposition 8.2 and Theorem 8.4 are about 3.4×10^8, 3.5×10^6 and 4.3×10^7 records. For $|\mathcal{Q}| = 10^9$ they become 4.7×10^{11}, 1.6×10^8 and 5.3×10^7. Private multiplicative weights wins only for very many queries, because the constants in Equation 8.18 are large and its accuracy costs $1/\alpha^2$.

8.5. Summary

Table 8.2.: The chapter's results.

Result	Statement
Proposition 8.1	Laplace noise on every query: $n \geq \lvert\mathcal{Q}\rvert\ln(\lvert\mathcal{Q}\rvert/\beta)/(\alpha\varepsilon)$ under ε-DP
Lemma 8.1	k mechanisms at $\varepsilon/(2\sqrt{2k\ln(1/\delta)})$ each compose to (ε,δ)-DP
Proposition 8.2	Laplace with advanced composition: $n \geq 2\sqrt{2\lvert\mathcal{Q}\rvert\ln(1/\delta)}\ln(\lvert\mathcal{Q}\rvert/\beta)/(\alpha\varepsilon)$
Proposition 8.3	with a perfect expert, halving makes at most $\log_2 N$ mistakes
Proposition 8.4	weighted majority makes at most $2.41(\mathrm{OPT} + \log_2 N)$ mistakes
Theorem 8.1	multiplicative weights has regret $\eta T + \ln N/\eta$, or $2\sqrt{T\ln N}$
Theorem 8.2	with losses $s_t q^t$, the total query error is at most $2\sqrt{U\ln\lvert\mathcal{X}\rvert}$
Corollary 8.2	fewer than $4\ln\lvert\mathcal{X}\rvert/\alpha^2$ updates, then every query is within α

Result	Statement
Theorem 8.3	private multiplicative weights is (ε, δ)-DP
Theorem 8.4	accuracy α from $n = O(\sqrt{\log\lvert\mathcal{X}\rvert \log(1/\delta)}\log\lvert\mathcal{Q}\rvert/(\alpha^2\varepsilon))$ records, up to lower-order logarithmic terms

8.6. References

- Moritz Hardt and Guy N. Rothblum, "A Multiplicative Weights Mechanism for Privacy-Preserving Data Analysis", 2010. Private multiplicative weights and its sample-size bound.
- Anupam Gupta, Aaron Roth and Jonathan Ullman, "Iterative Constructions and Private Data Release", 2012, and Moritz Hardt, Katrina Ligett and Frank McSherry, "A Simple and Practical Algorithm for Differentially Private Data Release", 2012. The offline version used here, which selects queries with the exponential mechanism, and a practical variant.
- Avrim Blum, Katrina Ligett and Aaron Roth, "A Learning Theory Approach to Noninteractive Database Privacy", 2008 (journal version 2013). The exponential mechanism over small datasets, the fourth row of Table 8.1.
- Nick Littlestone and Manfred K. Warmuth, "The Weighted Majority Algorithm", 1994. Halving and weighted majority.
- Sanjeev Arora, Elad Hazan and Satyen Kale, "The Multiplicative Weights Update Method: A Meta-Algorithm and Applications", 2012. A survey of the algorithm, the regret bound of Theorem 8.1, and its many uses.
- Thomas Steinke and Jonathan Ullman, "Between Pure and Approximate Differential Privacy", 2017. The bound $|\mathcal{Q}|/(\alpha\varepsilon)$ for pure DP and the approximate-DP bound within a factor $\sqrt{\ln\ln|\mathcal{Q}|}$ of $\sqrt{|\mathcal{Q}|\ln(1/\delta)}/(\alpha\varepsilon)$, both stated for one-way marginals.
- Arun Ganesh and Jiazheng Zhao, "Privately Answering Counting Queries with Generalized Gaussian Mechanisms", 2021. The approximate-DP bound within a factor $\sqrt{\ln\ln\ln|\mathcal{Q}|}$ of $\sqrt{|\mathcal{Q}|\ln(1/\delta)}/(\alpha\varepsilon)$.
- Salil Vadhan, "The Complexity of Differential Privacy", 2017. Theorem 2.9 is the noisy-histogram bound, stated there for counting queries.
- Jonathan Ullman and Salil Vadhan, "[Probabilistically Checkable Proofs] and the Hardness of Generating Private Synthetic Data" (the title abbreviates the bracketed words), 2011. Why no private algorithm that outputs synthetic data can be much faster in general.
- Cynthia Dwork and Aaron Roth, *The Algorithmic Foundations of Differential Privacy*, 2014. Section 3.5 states the advanced composition theorem in the form used here.

9. Sparse Vector

Queries often arrive one at a time, and each must be answered before the next is seen. Under differential privacy (DP), answering each of them with its own noise makes the error grow polynomially with the number of queries. Often the analyst only needs to know which queries exceed a threshold, and few of them do. This chapter shows that the cost then depends on the number of large queries and only logarithmically on the total. The tool is the sparse vector technique. Its privacy proof is short but delicate, because it rests on adding noise to the threshold as well as to the queries. At the end, the technique turns private multiplicative weights into an algorithm that answers queries as they arrive.

9.1. Answering queries online

The setting is a stream of queries $f_1, f_2, \dots, f_k \colon \mathcal{X}^n \to \mathbb{R}$, each with sensitivity at most 1. Here two datasets D and D' are neighbours, written $D \sim D'$, when they differ in one record, which is the replace-one relation. Sensitivity at most 1 means that $|f_i(D) - f_i(D')| \le 1$ for every such pair. Counting queries, which count the records with some property, are the typical example. Query f_i arrives at step i and must be answered before f_{i+1} is revealed.

The Laplace mechanism answers each query with noise from the Laplace distribution $\mathrm{Lap}(b)$ of scale b, whose density at z is $e^{-|z|/b}/(2b)$. For a sensitivity-1 query it is ε-DP when $b = 1/\varepsilon$. Its tails are $\Pr[|\mathrm{Lap}(b)| \ge tb] = e^{-t}$. Under pure DP with a total budget ε, basic composition forces the scale k/ε on each of k answers, so the error is of order k/ε. Under approximate DP, the advanced composition theorem brings this down to order $\sqrt{k \ln(1/\delta)}/\varepsilon$. Both are polynomial in k, and in the worst case no algorithm that answers every query can do much better.

A weaker question has a cheaper answer: which queries are large? If all k queries were known in advance, the exponential mechanism with score $f_i(D)$ would pick a query whose value is within $2(\ln k + t)/\varepsilon$ of the largest, except with probability e^{-t}. Repeating it c times with budget ε/c each finds c large queries with error of order $c \ln k/\varepsilon$, or $\sqrt{c \ln(1/\delta)} \ln k/\varepsilon$ with advanced composition. That is logarithmic in k, but it needs the whole list. Online, the first query cannot be compared with queries not yet seen, so the question becomes absolute.

> Given a public threshold τ, report the first c queries in the stream with $f_i(D) \ge \tau$, and stop.

The idea is to run the Laplace mechanism on each query and report only whether the noisy value exceeds the threshold. The essential detail is to compare it with a noisy threshold $\hat{\tau}$, not with τ. Getting the details exactly right is notoriously hard, and several published variants are not private.

9.2. AboveThreshold

The case $c = 1$ finds the first large query and halts. AboveThreshold draws one noisy threshold at the start, compares each noisy query with it, and answers with one of two symbols. The symbol $\top$ means "above the threshold" and $\bot$ means "below".

Definition 9.1 (AboveThreshold)**.** The inputs are a dataset D, a stream of sensitivity-1 queries $f_1, f_2, ...$, a threshold $\tau \in \mathbb{R}$ and $\varepsilon > 0$.

1. Draw $\hat{\tau} = \tau + \mathrm{Lap}(2/\varepsilon)$.
2. For each query f_i in turn, draw $\nu_i \sim \mathrm{Lap}(4/\varepsilon)$. If $f_i(D) + \nu_i \geq \hat{\tau}$, output $\top$ and halt. Otherwise output $\bot$.

All noise draws are independent.

Nothing in the algorithm depends on the number of queries, so the stream may even be infinite, and the privacy guarantee does not change. The accuracy guarantee does, because more noise draws mean a larger largest one.

The queries may be fixed in advance or chosen by an analyst who sees the earlier answers. Before the algorithm halts every answer is $\bot$, so in the second case each query is still a fixed function of the analyst's own coins, which do not depend on the dataset. The proof below applies to the stream fixed by each value of those coins, and averaging the inequality over the coins keeps it.

Theorem 9.1. *AboveThreshold is ε-DP.*

Proof.

1. Fix neighbours $D \sim D'$ and a stream, and write Y and Y' for the outputs on D and D'. The outputs are $y = \bot^{t-1}\top$, which halts at query t, and the output that never halts: $\bot^k$ for a stream of finite length k, or $\bot^\infty$ for an infinite one. These form a countable set, so by the pointwise lemma for discrete outputs it suffices to show $\Pr[Y = y] \leq e^\varepsilon \Pr[Y' = y]$ for each of them.
2. *Fix the noise of the first $t-1$ queries.* Condition on $\nu_i = v_i$ for $i < t$, the same values on both datasets, and let

$$g(D) = \max_{i<t} \left(f_i(D) + v_i\right), \tag{9.1}$$

with $g = -\infty$ if $t = 1$. With the v_i fixed, g is a number, not a random variable. Each f_i has sensitivity 1, so $|g(D) - g(D')| \leq 1$. 3. The output is y exactly when every earlier query falls below $\hat{\tau}$ and query t does not, that is, when $g(D) < \hat{\tau} \leq f_t(D) + \nu_t$. Write ϕ_2 and ϕ_4 for the densities of $\hat{\tau}$ and ν_t, whose subscripts recall the scales $2/\varepsilon$ and $4/\varepsilon$. Then

$$\Pr[Y = y] = \iint \phi_2(u)\, \phi_4(v)\, \mathbf{1}\big[g(D) < u \leq f_t(D) + v\big]\, \mathrm{d}u\, \mathrm{d}v. \tag{9.2}$$

4. *Change variables to move from D to D'.* Substitute

$$u = u' + g(D) - g(D'), \qquad v = v' + g(D) - g(D') + f_t(D') - f_t(D). \tag{9.3}$$

This is a shift, so $\mathrm{d}u\,\mathrm{d}v = \mathrm{d}u'\,\mathrm{d}v'$. The condition $g(D) < u$ becomes $g(D') < u'$, and $u \le f_t(D)+v$ becomes $u' \le f_t(D') + v'$, so the indicator becomes $\mathbf{1}[g(D') < u' \le f_t(D') + v']$. 5. *Pay for the shifts.* By step 2, $|u - u'| \le 1$. By step 2 and the sensitivity of f_t, $|v - v'| \le 2$. A Laplace density of scale b changes by at most a factor $e^{s/b}$ when its argument moves by s, so $\phi_2(u) \le e^{\varepsilon/2}\phi_2(u')$ and $\phi_4(v) \le e^{2\varepsilon/4}\phi_4(v') = e^{\varepsilon/2}\phi_4(v')$. Hence

$$\Pr[Y = y] \le e^{\varepsilon} \iint \phi_2(u')\,\phi_4(v')\,\mathbf{1}\big[g(D') < u' \le f_t(D') + v'\big]\,\mathrm{d}u'\,\mathrm{d}v' = e^{\varepsilon}\Pr[Y' = y], \tag{9.4}$$

where both probabilities are conditional on the fixed $v_1, \dots, v_{t-1}$. 6. The v_i have the same distribution on both datasets, so averaging step 5 over them gives the unconditional inequality. 7. *The output that never halts.* It occurs exactly when $\hat{\tau} > \sup_i (f_i(D) + \nu_i)$, the supremum over all queries of the stream. Fixing all ν_i makes this supremum a number $g(D)$, possibly $+\infty$, with $|g(D) - g(D')| \le 1$ when it is finite, and the output has probability 0 on both datasets when it is not. Shifting u by at most 1 as in steps 4 and 5 gives the factor $e^{\varepsilon/2} \le e^{\varepsilon}$. □

The two noises do different jobs. The noise on $\hat{\tau}$ pays, in one draw, for all the $\perp$ answers: the proof fixes their noise and absorbs the change in their maximum by shifting the threshold. The noise ν_t pays for the single $\top$. So the privacy cost does not grow with the number of $\perp$ answers.

The output carries no numbers, so accuracy means making no mistakes, with some slack around the threshold.

Definition 9.2 (Accuracy for threshold answers)**.** Let an algorithm answer a stream of queries $f_1, f_2, \dots$ with $\top$ or $\perp$, compared with a threshold τ, and write a_i for its answer to f_i. Let $\alpha \ge 0$ be a slack and $\beta \in (0,1)$ a failure probability. The algorithm is (α, β)**-accurate** on D if, with probability at least $1 - \beta$, every answered query satisfies

$$a_i = \top \implies f_i(D) \ge \tau - \alpha \qquad \text{and} \qquad a_i = \perp \implies f_i(D) \le \tau + \alpha. \tag{9.5}$$

In words, a $\top$ is never given to a query clearly below the threshold, and a $\perp$ is never given to a query clearly above it. AboveThreshold meets this definition with a slack α that grows only with the logarithm of the number of queries. The theorem also asks that every query before the last be clearly below the threshold, so that the algorithm does not halt early.

Theorem 9.2 (Accuracy of AboveThreshold)**.** *Let $\beta \in (0,1)$, and let $f_1, \dots, f_k$ be a stream with $f_i(D) \le \tau - \alpha$ for every $i < k$, where*

$$\alpha = \frac{8(\ln k + \ln(2/\beta))}{\varepsilon}. \tag{9.6}$$

Then AboveThreshold is (α, β)-accurate on D, and with probability at least $1 - \beta$ it does not halt before query k.

Proof.

1. *The good event.* Let G be the event that $|\hat{\tau} - \tau| < \alpha/2$ and $|\nu_i| < \alpha/2$ for every $i \le k$.
2. *Accuracy on G.* If $a_i = \top$, then $f_i(D) \ge \hat{\tau} - \nu_i > (\tau - \alpha/2) - \alpha/2 = \tau - \alpha$. If $a_i = \bot$, then $f_i(D) < \hat{\tau} - \nu_i < \tau + \alpha$.
3. *No early halt on G.* By step 2, a $\top$ at query i needs $f_i(D) > \tau - \alpha$, which the hypothesis rules out for $i < k$.
4. *G is likely.* The threshold noise has scale $2/\varepsilon$, so the tail formula gives $\Pr[|\hat{\tau} - \tau| \ge \alpha/2] = e^{-\varepsilon\alpha/4}$. Each ν_i has scale $4/\varepsilon$, so a union bound gives $\Pr[\max_{i \le k} |\nu_i| \ge \alpha/2] \le k e^{-\varepsilon\alpha/8}$.
5. With Equation 9.6, $k e^{-\varepsilon\alpha/8} = k \cdot \frac{1}{k} \cdot \frac{\beta}{2} = \beta/2$. Also $e^{-\varepsilon\alpha/4} = (\beta/(2k))^2 \le \beta/2$. So $\Pr[G] \ge 1 - \beta$.

□

The error grows only logarithmically in k, like the exponential mechanism's, but the queries can arrive online. A stream of a million queries shows how large the gain is.

Example 9.1. Take $k = 10^6$ counting queries, $\varepsilon = 1$ and $\beta = 0.05$. Then Equation 9.6 gives $\alpha = 8(\ln 10^6 + \ln 40) = 8(13.82 + 3.69) = 140.0$. On a dataset of $n = 10^5$ records, a query is placed correctly unless its count lies within 140 of the threshold. Answering every query with the Laplace mechanism and basic composition would instead add noise of scale 10^6, ten times the size of the dataset.

9.3. Sparse vector

To find c large queries, restart AboveThreshold after each $\top$, with a fresh threshold, until c of them have been found.

Definition 9.3 (Sparse vector)**.** The inputs are those of AboveThreshold, a count $c \ge 1$ of $\top$ answers to find, and a parameter δ, with either $\delta = 0$ or $\delta \in (0, 1/2]$. Set the Laplace noise scale b to

$$b = \frac{2c}{\varepsilon} \text{ if } \delta = 0, \qquad b = \frac{\sqrt{32c \ln(1/\delta)}}{\varepsilon} \text{ if } \delta > 0. \tag{9.7}$$

1. Draw $\hat{\tau} = \tau + \mathrm{Lap}(b)$ and set count $= 0$.
2. For each query f_i in turn, draw $\nu_i \sim \mathrm{Lap}(2b)$. If $f_i(D) + \nu_i < \hat{\tau}$, output $\bot$. Otherwise output $\top$ and add 1 to count. If count $= c$, halt. If not, draw a fresh $\hat{\tau} = \tau + \mathrm{Lap}(b)$ and continue.

Between two $\top$ answers the algorithm runs AboveThreshold with parameter $\varepsilon' = 2/b$, whose noise scales are $2/\varepsilon' = b$ and $4/\varepsilon' = 2b$. The fresh threshold matters, because reusing $\hat{\tau}$ after a $\top$ would break this equivalence, and with it the proof. With the equivalence, sparse vector is a composition of c runs of AboveThreshold, and its privacy follows from the composition theorems.

Theorem 9.3. *Sparse vector with $\delta = 0$ is ε-DP for every $\varepsilon > 0$. With $\varepsilon \in (0, 1]$ and $\delta \in (0, 1/2]$ it is (ε, δ)-DP.*

Proof.

1. The output is the concatenation of at most c runs of AboveThreshold, each with parameter $\varepsilon' = 2/b$ and fresh noise. Each run takes the rest of the stream, which may depend on the earlier outputs, so the whole is an adaptive composition of c mechanisms. A run that never starts, because the algorithm halted, has a fixed output. Each run is ε'-DP by Theorem 9.1.
2. If $\delta = 0$, then $\varepsilon' = \varepsilon/c$, and basic composition gives $c\varepsilon' = \varepsilon$.
3. If $\delta > 0$, then $\varepsilon' = 2\varepsilon/\sqrt{32c\ln(1/\delta)} = \varepsilon/(2\sqrt{2c\ln(1/\delta)})$. The budget-splitting lemma says that m adaptively composed mechanisms, each $\varepsilon/(2\sqrt{2m\ln(1/\delta)})$-DP, are together (ε, δ)-DP when $\varepsilon \le 1$ and $\delta \le 1/2$. Apply it with $m = c$.

□

Accuracy is proved as for AboveThreshold. The only change is that the union bound now runs over every threshold drawn, one per run, as well as over the noise on each query.

Theorem 9.4 (Accuracy of sparse vector)**.** *Let $\beta \in (0, 1)$, let $f_1, \dots, f_k$ be a stream, and let*

$$\alpha = 4b\ln(2k/\beta). \tag{9.8}$$

With probability at least $1 - \beta$, every answered query satisfies the two implications of Definition 9.2. If moreover fewer than c of the queries have $f_i(D) > \tau - \alpha$, the algorithm answers all k queries. In terms of the parameters, the accuracy is

$$\alpha = \frac{8c(\ln k + \ln(2/\beta))}{\varepsilon} \quad \text{if } \delta = 0, \qquad \alpha = \frac{\sqrt{512\,c\ln(1/\delta)}\,(\ln k + \ln(2/\beta))}{\varepsilon} \quad \text{if } \delta > 0. \tag{9.9}$$

Proof.

1. Let G be the event that every threshold drawn and every ν_i with $i \le k$ lies within $\alpha/2$ of its centre. On G, steps 2 and 3 of the proof of Theorem 9.2 apply to each query, with whichever threshold is current. So the implications hold, and a $\top$ needs $f_i(D) > \tau - \alpha$.
2. If fewer than c queries have $f_i(D) > \tau - \alpha$, then on G there are fewer than c answers $\top$, and the algorithm never halts.
3. The query noise has scale $2b$, and $\alpha/(4b) = \ln(2k/\beta)$. A union bound over the k draws bounds its failure by $ke^{-\alpha/(4b)} = \beta/2$.
4. At most $k+1$ thresholds are drawn, each with scale b. Their failure is at most $(k+1)e^{-\alpha/(2b)} = (k+1)\beta^2/(4k^2) \le \beta/2$, since $k + 1 \le 2k$ and $\beta \le 1$.
5. Substituting the two values of b into $4b\ln(2k/\beta)$ gives Equation 9.9, because $4\sqrt{32} = \sqrt{512}$.

□

With $c = 1$ and $\delta = 0$ this is Theorem 9.2 again. The error is polynomial in the number c of large queries, which every algorithm must pay in some form, and logarithmic in the total number k. So sparse vector can scan a very long stream, as long as few of its queries are large.

9.4. Numeric sparse vector

Sparse vector already pays in proportion to the number c of large queries. Releasing a noisy value for each of those c queries takes c more Laplace mechanisms, which costs the same order of budget. So sparse vector can also release the values of the large queries, at the cost of a constant factor.

Definition 9.4 (Numeric sparse vector)**.** Run sparse vector with $\delta = 0$ and budget $\varepsilon/2$, so $b = 4c/\varepsilon$. Whenever it would output $\top$ for query f_i, output instead $a_i = f_i(D) + \mathrm{Lap}(2c/\varepsilon)$, with fresh noise.

Each answer is now either $\perp$ or a number. The guarantee combines the two implications of Definition 9.2 with an error bound on every released number.

Theorem 9.5. *Numeric sparse vector is ε-DP. Let $f_1, \dots, f_k$ be a stream, $\beta \in (0,1)$ and $\alpha = 16c\ln(4k/\beta)/\varepsilon$. With probability at least $1-\beta$, every answered query satisfies the following, where $a_i \in \mathbb{R}$ means that the answer to f_i is a released number:*

$$a_i = \perp \implies f_i(D) \le \tau + \alpha, \qquad a_i \in \mathbb{R} \implies f_i(D) \ge \tau - \alpha \ \text{ and } \ |a_i - f_i(D)| \le \alpha, \tag{9.10}$$

and if fewer than c queries have $f_i(D) > \tau - \alpha$, all k are answered.

Proof.

1. *Privacy.* The algorithm is an adaptive composition of at most c runs of AboveThreshold, each with parameter $2/b = \varepsilon/(2c)$, and at most c Laplace mechanisms. Each Laplace mechanism measures one sensitivity-1 query, chosen from the earlier outputs, with scale $2c/\varepsilon$, so it is $\varepsilon/(2c)$-DP. Basic composition over the $2c$ mechanisms gives ε.
2. *Accuracy of the labels.* Theorem 9.4 with $b = 4c/\varepsilon$ and failure probability $\beta/2$ gives accuracy $4b\ln(4k/\beta) = \alpha$ for the labels, and no early halt, with probability at least $1 - \beta/2$.
3. *Accuracy of the values.* At most k values are released, one per query. By the tail formula and a union bound, some noise exceeds α with probability at most $k\,e^{-\alpha\varepsilon/(2c)} = k\,(\beta/(4k))^8 \le \beta/2$.

□

The fresh noise in the released value matters. Releasing $f_i(D) + \nu_i$, the value already compared with the threshold, is not private, and published variants that did so were wrong.

9.5. Online private multiplicative weights

The previous chapter built private multiplicative weights for a fixed set of linear queries $q\colon \mathcal{X} \to [0,1]$, with $q(D) = \frac{1}{n}\sum_i q(x_i)$. Its key steps were the following. For a distribution p on $\mathcal{X}$, write $\langle q, p\rangle = \sum_x q(x)p(x)$ for the average of q under p. The histogram p_D of the dataset, the fraction of records equal to each x, satisfies $q(D) = \langle q, p_D\rangle$. Multiplicative weights over the experts $\mathcal{X}$ with learning rate η and losses $s_t q^t$, where s_t is the sign of the error $\langle q^t, p^t\rangle - q^t(D)$, has total error at most $\eta U + \ln|\mathcal{X}|/\eta$ after U updates. And the offline algorithm used the exponential mechanism to find a query with a large error.

Online, the queries arrive one at a time, and the exponential mechanism cannot be used. Sparse vector replaces it. The algorithm keeps a distribution $\hat{p}$ and answers each query from it while it is accurate. When sparse vector reports a large error, the algorithm answers with a noisy true value and updates $\hat{p}$.

These queries have sensitivity $1/n$, not 1. Running sparse vector on f_i/Δ with threshold τ/Δ handles sensitivity Δ: it is the same algorithm with every noise scale multiplied by Δ, so it has the same privacy and Δ times the accuracy α.

The algorithm runs this rescaled sparse vector on the errors of $\hat{p}$, with threshold $\alpha/2$. Its count c of large errors is set to the bound on the number of updates from the previous chapter, and ε_0 is the budget of each of the $2c$ mechanisms it may run.

Definition 9.5 (Online private multiplicative weights)**.** The inputs are a dataset $D \in \mathcal{X}^n$, a stream of linear queries $f_1, f_2, \ldots$, an accuracy $\alpha \in (0,1]$, and $\varepsilon \in (0,1]$ and $\delta \in (0, 1/2]$. Set

$$c = \left\lceil \frac{64 \ln|\mathcal{X}|}{\alpha^2} \right\rceil, \qquad \eta = \frac{\alpha}{8}, \qquad \varepsilon_0 = \frac{\varepsilon}{4\sqrt{c\ln(1/\delta)}}, \tag{9.11}$$

let $\hat{p}$ be uniform on $\mathcal{X}$, set count $= 0$, and draw $\hat{\tau} = \alpha/2 + \mathrm{Lap}(2/(\varepsilon_0 n))$. For each query f_i in turn:

1. Let $e_i = |\langle f_i, \hat{p}\rangle - f_i(D)|$ and draw $\nu_i \sim \mathrm{Lap}(4/(\varepsilon_0 n))$.
2. If $e_i + \nu_i < \hat{\tau}$, answer $\langle f_i, \hat{p}\rangle$.
3. Otherwise answer $a_i = f_i(D) + Z_i$ with $Z_i \sim \mathrm{Lap}(1/(\varepsilon_0 n))$. Let s_i be the sign of $\langle f_i, \hat{p}\rangle - a_i$, update $\hat{p}(x) \propto \hat{p}(x)(1 - \eta s_i f_i(x))$, and add 1 to count. If count $= c$, halt. Otherwise draw a fresh $\hat{\tau} = \alpha/2 + \mathrm{Lap}(2/(\varepsilon_0 n))$.

Steps 1 to 3 follow numeric sparse vector, run on the errors e_i with threshold $\alpha/2$, except that the fresh measurement in step 3 is of $f_i(D)$ rather than of e_i. The answer in step 2 uses only $\hat{p}$, which depends on the earlier outputs alone, so it costs no privacy.

Theorem 9.6. *Online private multiplicative weights is (ε, δ)-DP.*

Proof.

1. Between two updates $\hat{p}$ is fixed, so the errors e_i form a stream of queries determined by the earlier outputs and the analyst. Each e_i has sensitivity $1/n$, since $\langle f_i, \hat{p}\rangle$ does not depend on D and $f_i(D)$ changes by at most $1/n$.

2. Each stretch between updates is therefore AboveThreshold with parameter ε_0, rescaled to sensitivity $1/n$, and so ε_0-DP by Theorem 9.1. Each measurement in step 3 is the Laplace mechanism on a query of sensitivity $1/n$ with scale $1/(\varepsilon_0 n)$, so it is also ε_0-DP. The answers of step 2 and the updates are post-processing.
3. There are at most c stretches and c measurements, so the algorithm is an adaptive composition of $2c$ mechanisms, each ε_0-DP. Since $\varepsilon_0 = \varepsilon/(2\sqrt{2 \cdot 2c \ln(1/\delta)})$, the budget-splitting lemma with $m = 2c$ mechanisms makes it (ε, δ)-DP.

□

Accuracy follows the proof for the offline algorithm. When all the noise is small, every answer from $\hat{p}$ is accurate, every measured answer is accurate, and every update makes progress, so the count of updates never reaches c.

Theorem 9.7 (Accuracy of online private multiplicative weights)**.** *Let $|\mathcal{X}| \geq 2$, $\beta \in (0, 1)$ and a stream of k queries, and let the other parameters be as in Definition 9.5. If*

$$n \geq \frac{128\sqrt{c \ln(1/\delta)}\, \ln\left((2c + k)/\beta\right)}{\alpha\varepsilon}, \tag{9.12}$$

then with probability at least $1 - \beta$ the algorithm answers all k queries, each within α of $f_i(D)$.

Proof.

1. *The good event.* Let G be the event that every threshold noise, every ν_i and every Z_i is less than $\alpha/8$ in absolute value. These are at most $2c + k$ draws, because there are at most c thresholds and c measurements. The widest scale is $4/(\varepsilon_0 n)$, so each draw fails with probability at most $e^{-\alpha\varepsilon_0 n/32}$. Condition Equation 9.12 is the same as $\alpha\varepsilon_0 n/32 \geq \ln((2c+k)/\beta)$, by the value of ε_0. A union bound gives $\Pr[G] \geq 1 - \beta$, and the rest assumes G.
2. *Answers from $\hat{p}$.* In step 2, $e_i < \hat{\tau} - \nu_i < \alpha/2 + \alpha/8 + \alpha/8 = 3\alpha/4$, so the answer is within α.
3. *Measured answers.* In step 3, $|a_i - f_i(D)| = |Z_i| < \alpha/8$. Also $e_i \geq \hat{\tau} - \nu_i > \alpha/2 - \alpha/4 = \alpha/4$. The difference $\langle f_i, \hat{p}\rangle - a_i$ is the true error minus Z_i, and $|Z_i| < \alpha/8 < e_i$, so s_i is the sign of the true error.
4. *Few updates.* By step 3, the updates are those of multiplicative weights for queries with $\eta = \alpha/8$, and each has error above $\alpha/4$. If there are U of them, the regret bound gives

$$\frac{\alpha}{4}U < \frac{\alpha}{8}U + \frac{8 \ln|\mathcal{X}|}{\alpha}, \qquad \text{that is,} \qquad U < \frac{64 \ln|\mathcal{X}|}{\alpha^2} \leq c. \tag{9.13}$$

So the count never reaches c, the algorithm never halts, and all k queries are answered. □

As in the previous chapter, $c \leq 66 \ln|\mathcal{X}|/\alpha^2$. So condition Equation 9.12 asks for

$$n = O\left(\frac{\sqrt{\ln|\mathcal{X}|\, \ln(1/\delta)}\, \ln\frac{k + \ln|\mathcal{X}|/\alpha^2}{\beta}}{\alpha^2\varepsilon}\right) \tag{9.14}$$

records. This matches the offline algorithm, with the length k of the stream in place of the number $|\mathcal{Q}|$ of queries. Sparse vector thus turns an algorithm that searches with the exponential mechanism into one that answers queries as they arrive.

9.6. Summary

Table 9.1.: The chapter's results.

Result	Statement
Theorem 9.1	AboveThreshold, with noise on the threshold and on each query, is ε-DP
Theorem 9.2	it finds the first large query with error $8(\ln k + \ln(2/\beta))/\varepsilon$
Theorem 9.3	sparse vector, c restarts with fresh thresholds, is ε-DP or (ε, δ)-DP
Theorem 9.4	error $8c(\ln k + \ln(2/\beta))/\varepsilon$, or $\sqrt{512c\ln(1/\delta)}(\ln k + \ln(2/\beta))/\varepsilon$
Theorem 9.5	releasing the large values with fresh noise costs a constant factor
Theorem 9.6	online private multiplicative weights is (ε, δ)-DP
Theorem 9.7	it answers k online queries to accuracy α from $n = O(\sqrt{\log\|\mathcal{X}\|\log(1/\delta)}\log k/(\alpha^2\varepsilon))$ records, up to lower-order logarithmic terms

9.7. References

- Cynthia Dwork, Moni Naor, Omer Reingold, Guy N. Rothblum and Salil Vadhan, "On the Complexity of Differentially Private Data Release: Efficient Algorithms and Hardness Results", 2009. The first use of the sparse vector idea.
- Aaron Roth and Tim Roughgarden, "Interactive Privacy via the Median Mechanism", 2010, and Moritz Hardt and Guy N. Rothblum, "A Multiplicative Weights Mechanism for Privacy-Preserving Data Analysis", 2010. Online query answering with sparse vector, including online private multiplicative weights.
- Cynthia Dwork and Aaron Roth, *The Algorithmic Foundations of Differential Privacy*, 2014. Section 3.6 presents AboveThreshold, sparse vector and its numeric variant in the form used here, and Section 4.2 the online multiplicative weights mechanism.
- Min Lyu, Dong Su and Ninghui Li, "Understanding the Sparse Vector Technique for Differential Privacy", 2017. A survey of published variants, several of them not private, including those that reuse the threshold or release the compared noisy values.

10. Beyond Global Sensitivity

The Laplace mechanism adds noise scaled to the global sensitivity of a function, the largest change one record can cause on any pair of neighbouring datasets. For many functions that worst case is far from the dataset at hand, or it is infinite. This chapter asks whether the noise can be scaled to the dataset actually held instead, while keeping differential privacy (DP).

Done naively, it cannot, because the amount of noise then depends on the data and leaks it. Four techniques make the idea work. Propose-test-release checks privately that the dataset is far from any dataset where the function is sensitive. A stability argument built on it finds the most frequent element, and estimates every count of a histogram, with no dependence on the size of the domain. The last two techniques estimate a bound on the sensitivity privately, or smooth it, before using it. All four can be accurate on well-behaved datasets even when no algorithm is accurate on the worst one.

10.1. Local sensitivity

Throughout the chapter, $f \colon \mathcal{X}^n \to \mathbb{R}$ is a real-valued function of the dataset. Its **global sensitivity** is the largest change that replacing one record can cause, over neighbouring datasets $D \sim D'$ (the replace-one relation),

$$\Delta f = \max_{D \sim D'} |f(D) - f(D')|, \tag{10.1}$$

and the Laplace mechanism adds to $f(D)$ a noise term drawn from $\mathrm{Lap}(\Delta f/\varepsilon)$, which makes the release ε-DP. The maximum ranges over all pairs of neighbours, including datasets that look nothing like the one held. Fixing the first dataset gives a local version.

Definition 10.1 (Local sensitivity)**.** The **local sensitivity** of f at a dataset $D \in \mathcal{X}^n$ is

$$\mathrm{LS}_f(D) = \max_{D' \sim D} |f(D) - f(D')|. \tag{10.2}$$

The global sensitivity is the largest local sensitivity, $\Delta f = \max_D \mathrm{LS}_f(D)$, and the local sensitivity can be much smaller. The natural mechanism releases $f(D) + \mathrm{Lap}(\mathrm{LS}_f(D)/\varepsilon)$. It fails, because the scale of the noise is itself a function of the data. The distance between the two closest records shows how the local sensitivity can jump from one dataset to a neighbour.

Example 10.1. Let $\mathcal{X} = \mathbb{R}$, $n = 3$, and let $f(D)$ be the distance between the two closest records. On the real line f has infinite global sensitivity. For example, $(0, c, 2c)$ and its neighbour $(0, 0, 2c)$ have f equal to c and 0, for every $c > 0$, so the Laplace mechanism does not apply. Take $D = (0, 0, 0)$. Replacing any one record leaves two records at 0, so f stays 0 and $\mathrm{LS}_f(D) = 0$. Now take the neighbour $D' = (0, 0, 5)$. It also has $f(D') = 0$, but replacing one of its zeros by 10 gives $(0, 10, 5)$ with closest distance 5, so $\mathrm{LS}_f(D') \geq 5$.

On these two neighbours, the naive mechanism adds no noise to one and some noise to the other, and that difference alone gives the dataset away.

Proposition 10.1. *The mechanism $M(D) = f(D) + \mathrm{Lap}(\mathrm{LS}_f(D)/\varepsilon)$, where $\mathrm{Lap}(0)$ means no noise, is not ε-DP for any finite ε when f is the closest-pair distance of Example 10.1.*

Proof.

1. On $D = (0, 0, 0)$ the local sensitivity is 0, so $M(D) = 0$ with probability 1.
2. On the neighbour $D' = (0, 0, 5)$ the local sensitivity is positive, so $M(D')$ has a continuous distribution and $\Pr[M(D') = 0] = 0$.
3. So the event $\{0\}$ has probability 1 on D and 0 on D', and no factor e^{ε} with finite ε can bound the first by the second.

□

The failure is that D has low local sensitivity while a neighbour has high local sensitivity, and the noise level tells them apart. The local sensitivity is itself a statistic of the data, and releasing it, even only through the scale of the noise, leaks information. The techniques below each account for the sensitivity of nearby datasets, not just of D.

10.2. Propose-test-release

The first technique proposes a bound b on the local sensitivity, tests privately that the bound holds at D and at every dataset near D, and releases $f(D)$ with noise scaled to b only if the test passes. The test measures how many records must change before the bound fails. The **Hamming distance** $d(D, D')$ between two datasets is the number of records in which they differ.

Definition 10.2 (Distance to instability)**.** For a proposed bound $b \geq 0$ on the local sensitivity LS_f, the **distance to instability** of D is

$$\gamma_b(D) = \min\{\, d(D, D') : D' \in \mathcal{X}^n,\ \mathrm{LS}_f(D') > b \,\}, \tag{10.3}$$

the fewest records that must change to reach a dataset whose local sensitivity exceeds b, where d is the Hamming distance. If no dataset has local sensitivity above b, set $\gamma_b(D) = \infty$ for every D.

So $\gamma_b(D) = 0$ exactly when $\mathrm{LS}_f(D) > b$. The distance is a count of records, and that gives it sensitivity 1 whatever f is.

Lemma 10.1. *If γ_b is finite, then $|\gamma_b(D) - \gamma_b(D')| \le 1$ for all neighbours $D \sim D'$.*

Proof.

1. If some dataset W has $\mathrm{LS}_f(W) > b$, every dataset is within distance n of W, so γ_b is finite everywhere.
2. Let W attain the minimum for D', so $d(D', W) = \gamma_b(D')$ and $\mathrm{LS}_f(W) > b$. The Hamming distance satisfies the triangle inequality, so $d(D, W) \le d(D, D') + d(D', W) = 1 + \gamma_b(D')$. Hence $\gamma_b(D) \le \gamma_b(D') + 1$.
3. Swapping the roles of D and D' gives the other inequality.

□

Because the distance has sensitivity 1, the Laplace mechanism can release it privately, and that noisy distance is the test. The threshold of the test is set so that a dataset whose local sensitivity exceeds b passes only with probability $\delta/2$.

Definition 10.3 (Propose-test-release)**.** The mechanism takes a proposed bound $b > 0$ and privacy parameters $\varepsilon > 0$ and $\delta \in (0, 1)$. On input D, it does three steps.

1. compute $\hat{\gamma} = \gamma_b(D) + Z_1$ with $Z_1 \sim \mathrm{Lap}(1/\varepsilon)$;
2. if $\hat{\gamma} \le \ln(1/\delta)/\varepsilon$, output $\perp$, meaning "no answer";
3. otherwise output $f(D) + Z_2$ with $Z_2 \sim \mathrm{Lap}(b/\varepsilon)$, independent of Z_1.

The test spends one budget of ε, and the release spends another. The test can wrongly pass on a sensitive dataset, and δ pays for that. The release step is the naive mechanism of Proposition 10.1 with the scale b in place of $\mathrm{LS}_f(D)$. It is safe now because b is fixed in advance, and the decision to use it depends on the data only through the private estimate $\hat{\gamma}$.

Theorem 10.1. *Propose-test-release is $(2\varepsilon, \delta)$-DP.*

Proof. Fix neighbours $D \sim D'$ and an event $S \subseteq \mathbb{R} \cup \{\perp\}$. Write $\tau = \ln(1/\delta)/\varepsilon$ for the threshold. If γ_b is infinite at every dataset, the test always passes and $b \ge \Delta f$, so the output is the Laplace mechanism and is ε-DP. Assume from now on that γ_b is finite.

1. *The test is ε-DP.* By Lemma 10.1, γ_b has sensitivity 1, so $\hat{\gamma}$ is the Laplace mechanism and is ε-DP. The events $\{\hat{\gamma} \le \tau\}$ and $\{\hat{\gamma} > \tau\}$ are post-processing of $\hat{\gamma}$, so each has probability on D at most e^{ε} times its probability on D'.
2. *Case* $\mathrm{LS}_f(D) > b$. Then $\gamma_b(D) = 0$, and the test passes only if $Z_1 > \tau$, which has probability $\frac{1}{2}e^{-\varepsilon\tau} = \delta/2$. Splitting on whether the output is $\perp$ and using step 1,

$$\Pr[M(D) \in S] \le \Pr[M(D) = \perp]\,\mathbf{1}[\perp \in S] + \Pr[M(D) \ne \perp] \le e^{\varepsilon}\Pr[M(D') \in S] + \tfrac{\delta}{2}. \quad (10.4)$$

3. *Case* $\mathrm{LS}_f(D) \le b$. Then $|f(D) - f(D')| \le \mathrm{LS}_f(D) \le b$, so the release step is the Laplace mechanism for a pair at distance at most b, with scale b/ε. Its density at any y changes by a factor at most e^{ε} between D and D'. Write $S_{\mathbb{R}} = S \cap \mathbb{R}$. The two steps use independent noise, so

$$\Pr[M(D) \in S] = \Pr[\hat{\gamma} \le \tau]\, \mathbf{1}[\perp \in S] + \Pr[\hat{\gamma} > \tau]\, \Pr[f(D) + Z_2 \in S_{\mathbb{R}}]. \tag{10.5}$$

Each probability on the right is at most e^{ε} times the same probability on D', by step 1 and by the Laplace mechanism. So $\Pr[M(D) \in S] \le e^{2\varepsilon} \Pr[M(D') \in S]$.

4. In both cases $\Pr[M(D) \in S] \le e^{2\varepsilon} \Pr[M(D') \in S] + \delta$. The pair was arbitrary, so this holds for every ordered pair of neighbours.

□

The test passes with high probability on datasets far from instability. If $\gamma_b(D) \ge 2\ln(1/\delta)/\varepsilon$, then the output is $\perp$ only if $Z_1 \le -\ln(1/\delta)/\varepsilon$, which has probability $\delta/2$. The price is computation. Evaluating LS_f at one dataset takes up to $n(|\mathcal{X}| - 1)$ evaluations of f, one per neighbour, and γ_b is a minimum over all $|\mathcal{X}|^n$ datasets. So the distance to instability has to be worked out by hand for each function. The median is the standard case where this can be done efficiently.

10.3. Stability: the mode and the histogram

Propose-test-release is most striking with the bound $b = 0$. Then the test checks that the answer does not change at all near D, and the release step adds no noise. The mode, the most frequent element, is the standard example. If all n records equal the same value v, the mode stays v until about $n/2$ records change, so the mode is far from instability.

Let $\mathcal{X}$ be any set with a fixed total order, used only to break ties. Write $c_x(D)$ for the number of records of D equal to x, and $c_{(1)}(D) \ge c_{(2)}(D)$ for the largest and second-largest of these counts, with $c_{(2)}(D) = 0$ if all records are equal. The mode $\mathrm{mode}(D)$ is the element with the largest count, the least one in the order among ties. To change the mode, records must move from the most frequent element to the runner-up until it overtakes. So half the gap between the two largest counts is, up to 1, the distance to instability, and it is computed in one pass over the data whatever the domain.

Lemma 10.2. *Let $g(D) = (c_{(1)}(D) - c_{(2)}(D))/2$. For all neighbours $D \sim D'$:*

1. $|g(D) - g(D')| \le 1$;
2. *if $g(D) > 1$, then* $\mathrm{mode}(D') = \mathrm{mode}(D)$.

Proof.

1. Replacing one record lowers one count by 1 and raises another by 1, so every count changes by at most 1.
2. The j-th largest count is the maximum, over sets S of j elements, of $\min_{x \in S} c_x$. If every count rises by at most 1, each such minimum rises by at most 1, and so does the maximum. The same holds for decreases, so $c_{(1)}$ and $c_{(2)}$ each change by at most 1, and g changes by at most $(1 + 1)/2 = 1$. This is claim 1.
3. For claim 2, let $a = \mathrm{mode}(D)$. In D', the count of a is at least $c_{(1)}(D) - 1$, and every other count is at most $c_{(2)}(D) + 1$. The condition $g(D) > 1$ means $c_{(1)}(D) - 1 > c_{(2)}(D) + 1$, so a has the strictly largest count in D'.

□

The stable mode is propose-test-release with $b = 0$, with the gap g in place of the distance to instability. Its threshold has an extra 1 because claim 2 protects the mode only when $g(D) > 1$.

Definition 10.4 (Stable mode)**.** Given $\varepsilon > 0$ and $\delta \in (0, 1)$, on input D, compute $\hat{g} = g(D) + Z$ with $Z \sim \mathrm{Lap}(1/\varepsilon)$, where $g(D)$ is half the gap between the two largest counts. If $\hat{g} \le 1 + \ln(1/\delta)/\varepsilon$, output $\perp$, and otherwise output $\mathrm{mode}(D)$ exactly.

The guarantee has two parts. The test costs δ in privacy, and it fails, by outputting $\perp$, only on datasets whose two largest counts are close.

Theorem 10.2. *The stable mode is (ε, δ)-DP. On every dataset whose two largest counts differ by at least $2 + 4\ln(1/\delta)/\varepsilon$, it outputs the mode with probability at least $1 - \delta/2$.*

Proof. Fix neighbours $D \sim D'$ and an event S, and write $\tau = 1 + \ln(1/\delta)/\varepsilon$.

1. *The gap is ε-DP.* By Lemma 10.2, claim 1, $\hat{g}$ is the Laplace mechanism for a function of sensitivity 1, so it is ε-DP.
2. *Case* $\mathrm{mode}(D') = \mathrm{mode}(D)$. The output on both datasets is the same function of $\hat{g}$, namely $\perp$ below the threshold and the common mode above it. Post-processing gives $\Pr[M(D) \in S] \le e^{\varepsilon} \Pr[M(D') \in S]$.
3. *Case* $\mathrm{mode}(D') \neq \mathrm{mode}(D)$. By Lemma 10.2, claim 2, $g(D) \le 1$. The mode is released only if $Z > \tau - g(D) \ge \ln(1/\delta)/\varepsilon$, which has probability at most $\delta/2$. As in the proof of Theorem 10.1, $\Pr[M(D) \in S] \le e^{\varepsilon} \Pr[M(D') \in S] + \delta/2$.
4. *Accuracy.* If $c_{(1)} - c_{(2)} \ge 2 + 4\ln(1/\delta)/\varepsilon$, then $g(D) \ge \tau + \ln(1/\delta)/\varepsilon$. The output is $\perp$ only if $Z \le -\ln(1/\delta)/\varepsilon$, which has probability $\delta/2$.

□

With typical parameters, the gap the theorem needs is a few dozen records.

Example 10.2. With $\varepsilon = 1$ and $\delta = 10^{-6}$, the threshold is $1 + \ln(10^6) = 14.82$. A gap of $2 + 4 \times 13.82 = 57.3$, so 58 records, between the most frequent element and the runner-up is enough. The domain plays no part. It may be the set of all strings, which is infinite.

The stable mode costs ε, not 2ε, because the release step adds no noise and spends no budget. Its accuracy depends on the gap, not on the domain. Under pure DP no such bound is possible. The Laplace histogram has error of order $\ln|\mathcal{X}|/\varepsilon$, which is useless for an infinite domain. The packing chapter shows that an ε-DP algorithm that finds the mode of a dataset over k elements needs at least $\ln(k/2)/\varepsilon$ records, even when all records are equal. The cost $\ln k$ of a large domain is replaced here by $\ln(1/\delta)$.

The same idea estimates every count at once. The algorithm adds noise only to the counts of elements that appear in the data, and it reports only those whose noisy count clears a threshold. An element that appears once in D and not at all in a neighbour is the only danger, and the threshold hides it except with probability $\delta/2$.

Definition 10.5 (Stability-based histogram)**.** Given $\varepsilon > 0$ and $\delta \in (0, 1)$, let $\tau = 1 + 2\ln(1/\delta)/\varepsilon$. On input D, for each element x whose count $c_x(D)$ of records is positive, draw an independent $Z_x \sim \mathrm{Lap}(2/\varepsilon)$ and set $\hat{c}_x = c_x(D) + Z_x$. Output the list of pairs $(x, \hat{c}_x)$ with $\hat{c}_x > \tau$. An element not on the list has estimate $\hat{c}_x = 0$.

The noise has scale $2/\varepsilon$ because replacing one record changes two counts. The threshold costs δ in privacy. In accuracy, it adds at most τ to the noise on each count, because a count is reported as 0 only when its noisy value is at most τ.

Theorem 10.3. *The stability-based histogram is (ε, δ)-DP. For every $\beta \in (0, 1)$, with probability at least $1 - \beta$, every element $x \in \mathcal{X}$ has*

$$|\hat{c}_x - c_x(D)| \le \tau + \frac{2}{\varepsilon} \ln \frac{n}{\beta}. \tag{10.6}$$

Proof. Privacy. Let D' be D with one record equal to a replaced by $a' \neq a$. Let C be the set of elements that appear in both datasets. Any element that appears only in D is a, with count 1, and any element that appears only in D' is a', with count 1. Write $O(D) = O_C(D) \cup O_A(D)$ for the output list on D, where O_C is the part of the list from C and O_A the part from elements only in D, which can only be a. Define $O_C(D')$ and $O_{A'}(D')$ in the same way, where $O_{A'}$ can only contain a'. Fix an event S and let $p_C = \Pr[O_C(D') \in S]$.

1. *The extra element of D.* An element with count 1 is listed only if its noise exceeds $\tau - 1 = 2\ln(1/\delta)/\varepsilon$, which has probability $\frac{1}{2}e^{-\ln(1/\delta)} = \delta/2$. When $O_A(D)$ is empty, $O(D) = O_C(D)$, so $\Pr[O(D) \in S] \le \Pr[O_C(D) \in S] + \delta/2$.
2. *The common elements.* On C, the counts of D and D' differ by at most 1 at a and at a', so in ℓ_1 by at most 2. Laplace noise of scale $2/\varepsilon$ makes the noisy counts on C ε-DP, and the list is post-processing of them. So $\Pr[O_C(D) \in S] \le \min(1, e^{\varepsilon} p_C)$.
3. *The extra element of D'.* The parts $O_C(D')$ and $O_{A'}(D')$ are independent, and $O_{A'}(D')$ is empty with probability at least $1 - \delta/2$ by the computation of step 1. When it is empty, $O(D') = O_C(D')$. So $\Pr[O(D') \in S] \ge (1 - \delta/2)\, p_C$.
4. *Combining.* If $e^{\varepsilon} p_C \le 1$, then $\min(1, e^{\varepsilon} p_C) - e^{\varepsilon}(1 - \delta/2) p_C = e^{\varepsilon} p_C\, \delta/2 \le \delta/2$. Otherwise $1 - e^{\varepsilon}(1 - \delta/2) p_C < 1 - (1 - \delta/2) = \delta/2$. In both cases, steps 1 to 3 give

$$\Pr[O(D) \in S] \le e^{\varepsilon}(1 - \delta/2)\, p_C + \tfrac{\delta}{2} + \tfrac{\delta}{2} \le e^{\varepsilon} \Pr[O(D') \in S] + \delta. \tag{10.7}$$

Accuracy.

5. At most n elements appear in D. A Laplace variable of scale $2/\varepsilon$ exceeds s in absolute value with probability $e^{-\varepsilon s/2}$, so by a union bound all of them satisfy $|Z_x| \le (2/\varepsilon)\ln(n/\beta)$ with probability at least $1 - \beta$.
6. On that event, take any element x. If it is listed, its error is $|Z_x|$. If it appears in D but is not listed, then $c_x + Z_x \le \tau$, so its error is $c_x \le \tau + |Z_x|$. If it does not appear, both its count and its estimate are 0. In every case the error is at most Equation 10.6.

□

In the usual regime $\delta \le 1/n$, the bound Equation 10.6 is at most $1 + (2/\varepsilon)(2\ln(1/\delta) + \ln(1/\beta))$. Every count is then accurate to within a number of order $\ln(1/\delta)/\varepsilon$, with no dependence on the size of the domain. The Laplace histogram under pure DP pays $\ln|\mathcal{X}|$ instead, and the packing chapter shows that pure DP cannot avoid that cost.

10.4. Privately bounding the local sensitivity

Propose-test-release guesses a bound and checks it. A second technique estimates the local sensitivity itself, privately, and uses the estimate as the noise scale. It applies when the local sensitivity does not change fast between neighbours, that is, when LS_f has bounded global sensitivity.

Definition 10.6 (Noisy local sensitivity)**.** Suppose that the local sensitivity changes by at most $\lambda > 0$ between neighbours, that is, $|\mathrm{LS}_f(D) - \mathrm{LS}_f(D')| \le \lambda$ for all $D \sim D'$. Given $\varepsilon > 0$ and $\delta \in (0,1)$, on input D:

1. compute $\hat{s} = \mathrm{LS}_f(D) + Z_1 + \lambda\ln(1/\delta)/\varepsilon$ with $Z_1 \sim \mathrm{Lap}(\lambda/\varepsilon)$;
2. if $\hat{s} \le 0$, output $\perp$, and otherwise output $f(D) + Z_2$ with $Z_2 \sim \mathrm{Lap}(\hat{s}/\varepsilon)$.

The shift $\lambda\ln(1/\delta)/\varepsilon$ makes the estimate $\hat{s}$ an overestimate of $\mathrm{LS}_f(D)$ except with probability $\delta/2$. On that good event, the noise is large enough for privacy. Paying both for the noisy estimate and for the release gives the factor 2ε in the next theorem.

Theorem 10.4. *Let $\lambda > 0$ satisfy $|\mathrm{LS}_f(D) - \mathrm{LS}_f(D')| \le \lambda$ for all neighbours $D \sim D'$, and let $\varepsilon > 0$ and $\delta \in (0,1)$. Then noisy local sensitivity with these parameters is $(2\varepsilon, \delta)$-DP.*

Proof. Fix neighbours $D \sim D'$. It suffices to prove the guarantee for the pair $(\hat{s}, \text{output})$, since the output alone is post-processing of it.

1. *The good event.* Let $G = \{\hat{s} \ge \mathrm{LS}_f(D)\}$, computed with the data D. It fails only if $Z_1 < -\lambda\ln(1/\delta)/\varepsilon$, which has probability $\delta/2$.
2. *Density of the scale.* The function LS_f has sensitivity λ, so the density of $\hat{s}$ on D is at most e^{ε} times its density on D', at every value s.
3. *Density of the release.* Fix a value s with $s \ge \mathrm{LS}_f(D)$ and $s > 0$. Then $|f(D) - f(D')| \le \mathrm{LS}_f(D) \le s$, so the Laplace density of scale s/ε at any y changes by a factor at most $e^{\varepsilon|f(D)-f(D')|/s} \le e^{\varepsilon}$ between D and D'.
4. *Combining.* On values of $\hat{s}$ in G, the joint density of $(\hat{s}, \text{output})$ on D is at most $e^{2\varepsilon}$ times that on D', by steps 2 and 3. For an event S, split on G:

$$\Pr[(\hat{s}, M(D)) \in S] \le e^{2\varepsilon}\Pr[(\hat{s}, M(D')) \in S] + \Pr[\text{not } G] \le e^{2\varepsilon}\Pr[(\hat{s}, M(D')) \in S] + \tfrac{\delta}{2}. \quad (10.8)$$

□

Two quantities must be worked out by hand for each function: the local sensitivity at D, and a bound λ on how fast it changes. The first takes up to $n(|\mathcal{X}| - 1)$ evaluations of f, one per neighbour, and can sometimes be computed even for an infinite domain. For some functions it is much faster, such as the mode, whose local sensitivity depends only on the two largest counts. The technique is simpler to run than propose-test-release, which needs the distance to instability, but it applies to fewer functions, because it needs the local sensitivity to change slowly everywhere.

10.5. Smooth sensitivity

The last technique also lets the local sensitivity grow away from D, but discounts each dataset by its distance. Take the closest-pair function of Example 10.1 on the bounded domain $\mathcal{X} = [0, 10]$, so that its global sensitivity is finite. The dataset $D = (0, 0, 0)$ still has local sensitivity 0, while its neighbour $(0, 0, 5)$ has local sensitivity at least 5. A safe noise scale at D must account for that neighbour, but less than for D itself.

Definition 10.7 (Smooth sensitivity)**.** For a smoothing parameter $s > 0$, the **smooth sensitivity** of f at D discounts the local sensitivity of every dataset D' by a factor e^{-s} for each record in which D' differs from D:

$$S_{f,s}(D) = \sup_{D' \in \mathcal{X}^n} \mathrm{LS}_f(D')\, e^{-s\, d(D,D')}, \tag{10.9}$$

where d is the Hamming distance.

The supremum ranges over all datasets, not only neighbours. Every dataset is within distance n of every other, so $e^{-sn}\Delta f \le S_{f,s}(D) \le \Delta f$. The smooth sensitivity is therefore finite exactly when the global sensitivity is, and it can be smaller by a factor up to e^{sn}. Two properties follow from the definition, and they are all that privacy needs.

Lemma 10.3. *For every D, $S_{f,s}(D) \ge \mathrm{LS}_f(D)$. For all neighbours $D \sim D'$, $S_{f,s}(D) \le e^s S_{f,s}(D')$.*

Proof.

1. The term $D' = D$ in the supremum is $\mathrm{LS}_f(D)$, so the supremum is at least that.
2. For any dataset W, the triangle inequality gives $d(D, W) \ge d(D', W) - 1$, so $\mathrm{LS}_f(W)e^{-s\, d(D,W)} \le e^s\, \mathrm{LS}_f(W)e^{-s\, d(D',W)} \le e^s S_{f,s}(D')$. Taking the supremum over W gives the second claim.

□

Laplace noise scaled to the smooth sensitivity is not enough for pure DP. The scale moves by a factor up to e^s between neighbours, and far in the tail a Laplace density changes by much more than that when its scale changes. Noise with polynomial tails changes little under rescaling, and it works. The standard Cauchy distribution is such a noise, with density at a point z equal to

$$h(z) = \frac{1}{\pi(1+z^2)}. \tag{10.10}$$

Theorem 10.5 (Smooth sensitivity with Cauchy noise)**.** *Let f be non-constant with finite global sensitivity Δf, let $\varepsilon > 0$, and write $S^* = S_{f,\varepsilon/2}$. The mechanism*

$$M(D) = f(D) + \frac{2\,S^*(D)}{\varepsilon}\,Z, \qquad Z \sim \textit{standard Cauchy}, \tag{10.11}$$

is ε-DP.

Proof.

1. *The scale is positive and finite.* Some dataset W_0 has $\mathrm{LS}_f(W_0) > 0$ because f is not constant, and every D is within distance n of it, so $S^*(D) \geq \mathrm{LS}_f(W_0)e^{-\varepsilon n/2} > 0$. Every term of the supremum is at most Δf, so $S^*(D) \leq \Delta f < \infty$.
2. *Two facts about h.* The derivative of $\ln(1+z^2)$ is $2z/(1+z^2)$, which lies in $[-1,1]$, so $h(z)/h(z+u) \leq e^{|u|}$ for all z and u. For every real κ and v, $e^{-\kappa}h(v)/h(e^{\kappa}v) = e^{-\kappa}(1+e^{2\kappa}v^2)/(1+v^2) \leq e^{-\kappa}\max(1, e^{2\kappa}) = e^{|\kappa|}$.
3. *The ratio of densities.* Fix neighbours $D \sim D'$ and write $a = \varepsilon/2$. The density of $M(D)$ at y is $(a/S^*(D))\,h(a(y-f(D))/S^*(D))$. Set $u = a(y-f(D))/S^*(D)$, $w = a(f(D)-f(D'))/S^*(D)$ and $e^{\kappa} = S^*(D)/S^*(D')$. Then $a(y-f(D'))/S^*(D') = e^{\kappa}(u+w)$, and the ratio of the densities on D and D' is

$$e^{-\kappa}\,\frac{h(u)}{h(e^{\kappa}(u+w))} = \frac{h(u)}{h(u+w)} \cdot \frac{e^{-\kappa}\,h(u+w)}{h(e^{\kappa}(u+w))} \leq e^{|w|}\,e^{|\kappa|}, \tag{10.12}$$

by the two facts of step 2.

4. *Bounding the exponents.* By Lemma 10.3, $|f(D)-f(D')| \leq \mathrm{LS}_f(D) \leq S^*(D)$, so $|w| \leq a = \varepsilon/2$. By the same lemma applied in both directions, $|\kappa| \leq \varepsilon/2$. The densities therefore differ by a factor at most e^{ε}, and integrating over any event gives ε-DP.

□

The Cauchy distribution has such heavy tails that it has no mean. Its tail is $\Pr[|Z| > z] = \frac{2}{\pi}\arctan(1/z) \leq 2/(\pi z)$, so the error exceeds $4S^*(D)/(\pi\varepsilon\beta)$ with probability at most β, which grows as $1/\beta$ rather than as $\ln(1/\beta)$. Laplace or Gaussian noise scaled to the smooth sensitivity, with a smaller smoothing parameter, gives approximate DP with lighter tails. Laplace and Gaussian tails decay exponentially, which is where the logarithms in earlier error bounds come from. Computing $S_{f,s}$ means a maximum over all datasets, which is hard in general. For the median it takes polynomial time, and near-linear time with more care.

10.6. Summary

Table 10.1.: The chapter's results.

Result	Statement
Proposition 10.1	Laplace noise scaled to the local sensitivity is not private
Lemma 10.1	the distance to instability has sensitivity 1
Theorem 10.1	propose-test-release is $(2\varepsilon, \delta)$-DP
Theorem 10.2	the stable mode is (ε, δ)-DP and exact when the top two counts differ by $2 + 4\ln(1/\delta)/\varepsilon$
Theorem 10.3	the stability-based histogram is (ε, δ)-DP with every count within $\tau + (2/\varepsilon)\ln(n/\beta)$
Theorem 10.4	a noisy overestimate of a Lipschitz local sensitivity gives $(2\varepsilon, \delta)$-DP
Lemma 10.3	smooth sensitivity bounds the local sensitivity and changes by at most e^s between neighbours
Theorem 10.5	Cauchy noise of scale $2S_{f,\varepsilon/2}(D)/\varepsilon$ gives ε-DP when $0 < \Delta f < \infty$

10.7. References

- Kobbi Nissim, Sofya Raskhodnikova and Adam Smith, "Smooth Sensitivity and Sampling in Private Data Analysis", 2007. Local and smooth sensitivity, the failure of naive local noise, and noise distributions for smooth sensitivity under pure and approximate DP.
- Cynthia Dwork and Jing Lei, "Differential Privacy and Robust Statistics", 2009. The propose-test-release framework.
- Abhradeep Thakurta and Adam Smith, "Differentially Private Feature Selection via Stability Arguments, and the Robustness of the Lasso", 2013. The distance-to-instability approach behind the stable mode.
- Aleksandra Korolova, Krishnaram Kenthapadi, Nina Mishra and Alexandros Ntoulas, "Releasing Search Queries and Clicks Privately", 2009. Thresholded noisy counts over an unbounded domain, the idea of the stability-based histogram.
- Cynthia Dwork and Aaron Roth, *The Algorithmic Foundations of Differential Privacy*, 2014. Chapter 7 covers propose-test-release and stability-based algorithms.
- Salil Vadhan, "The Complexity of Differential Privacy", 2017. Section 3 treats local sensitivity, propose-test-release, the stability-based histogram (Theorem 3.5) and privately bounding the local sensitivity, with an example on graphs.

11. Packing Lower Bounds

Every chapter so far has given an algorithm and proved how much data it needs. This chapter asks the reverse question: could a cleverer algorithm need less? For pure differential privacy (DP), a short argument called packing answers it. If an algorithm is accurate, it must put its output in a different region for each of many datasets. If it is private, its output distribution barely moves between those datasets. Both can hold only when the dataset is large enough.

The argument gives lower bounds that hold for every ε-DP algorithm, however it is built. Applied to one-way marginals and to histograms, it shows that the simple Laplace-noise algorithms are optimal up to a constant or a logarithmic factor, and that pure DP is strictly more expensive than approximate DP for these problems.

11.1. Two problems and their upper bounds

Both problems ask for many counting queries at once, answered as fractions of the dataset. The first has records that are d yes-or-no answers, and asks for the fraction of yes answers to each question.

Definition 11.1 (One-way marginals)**.** Let $\mathcal{X} = \{0,1\}^d$. For a record x_i, write x_{ij} for its coordinate j. The **one-way marginals** of a dataset $D \in \mathcal{X}^n$ are the d averages

$$f_j(D) = \frac{1}{n}\sum_{i=1}^{n} x_{ij}, \qquad j = 1, \dots, d. \tag{11.1}$$

The second problem has records that are symbols from a finite set, and asks for the fraction of records equal to each symbol.

Definition 11.2 (Histogram)**.** Let $\mathcal{X} = [k] = \{1, \dots, k\}$. The **normalized histogram** of a dataset $D \in [k]^n$ is the vector $h(D) \in [0,1]^k$ with coordinates

$$h_j(D) = \frac{1}{n}\sum_{i=1}^{n} \mathbf{1}[x_i = j], \qquad j = 1, \dots, k, \tag{11.2}$$

where $\mathbf{1}[x_i = j]$ is 1 if record i equals j and 0 otherwise.

For both problems, an answer y is judged by its worst coordinate, the ℓ_∞ error $\max_j |y_j - f_j(D)|$. So an answer with error at most α gets every query right to within α at once.

The upper bounds come from the Laplace mechanism. Suppose the vector of answers f has ℓ_1 sensitivity at most $\Delta_1 f$, which means that $\sum_j |f_j(D) - f_j(D')| \le \Delta_1 f$ for all neighbours $D \sim D'$, datasets that differ in one record (the replace-one relation). The Laplace mechanism releases $f(D) + (Z_1, \dots, Z_d)$, where the noise terms Z_j are independent Laplace variables with scale $\Delta_1 f/\varepsilon$, and it is ε-DP. The analysis needs only the tail of a Laplace variable Z with scale b, which for every $s \ge 0$ is

$$\Pr[|Z| \ge s] = e^{-s/b}. \tag{11.3}$$

Computing the ℓ_1 sensitivity of each problem and applying this tail to every coordinate gives the two upper bounds that the rest of the chapter tries to beat.

Proposition 11.1. *Let $\varepsilon > 0$, $\alpha > 0$ and $\beta \in (0,1)$. With probability at least $1 - \beta$:*

1. *the Laplace mechanism for the d one-way marginals, with scale $d/(n\varepsilon)$, has error at most $d \ln(d/\beta)/(n\varepsilon)$;*
2. *the Laplace mechanism for the histogram over $[k]$, with scale $2/(n\varepsilon)$, has error at most $2\ln(k/\beta)/(n\varepsilon)$.*

Both are ε-DP. So error at most α with probability at least $1/2$ needs only

$$n \ge \frac{d \ln(2d)}{\alpha\varepsilon} \quad \textit{for marginals}, \qquad n \ge \frac{2\ln(2k)}{\alpha\varepsilon} \quad \textit{for histograms}. \tag{11.4}$$

Proof.

1. *Sensitivity of the marginals.* Replacing one record changes each average f_j by at most $1/n$, so the ℓ_1 sensitivity of $(f_1, \dots, f_d)$ is at most d/n. The Laplace mechanism with scale $d/(n\varepsilon)$ is therefore ε-DP.
2. *Sensitivity of the histogram.* Replacing a record equal to a by a record equal to $a' \ne a$ lowers h_a by $1/n$ and raises $h_{a'}$ by $1/n$, so the ℓ_1 sensitivity of h is $2/n$. The scale $2/(n\varepsilon)$ gives ε-DP.
3. *Union bound.* For m Laplace variables with scale b, the chance that some $|Z_j|$ is at least s is at most $me^{-s/b}$. Setting this equal to β gives $s = b\ln(m/\beta)$. With $m = d$ and $b = d/(n\varepsilon)$ this is claim 1, and with $m = k$ and $b = 2/(n\varepsilon)$ it is claim 2.
4. *Sample size.* Each error bound is at most α exactly when n satisfies Equation 11.4 with $\beta = 1/2$.

□

Approximate DP does better on both problems. For the marginals, the Gaussian mechanism of the approximate DP chapter uses the ℓ_2 sensitivity $\sqrt{d}/n$ instead of the ℓ_1 sensitivity d/n, and needs n only of order $\sqrt{d\ln(1/\delta)\ln d}/(\alpha\varepsilon)$, a factor of about $\sqrt{d}$ less. For histograms, the stability-based histogram of the previous chapter needs n of order $\ln(1/\delta)/(\alpha\varepsilon)$ when $\delta \le 1/n$, with no dependence on k at all. The rest of the chapter shows that under pure DP neither gain is possible.

11.2. The packing argument

The name comes from packing rigid balls into a box. The balls cannot overlap, so only a limited number fit. Here the balls are regions of the output space, one for each of many datasets. Accuracy forces each dataset to send its output into its own region, and the regions do not overlap. Privacy forces every dataset to give every region some probability. The probabilities of disjoint regions add up to at most 1, so only a limited number of datasets can be packed.

The tool that connects datasets far apart is group privacy, from the Laplace mechanism chapter. The **Hamming distance** between two datasets in $\mathcal{X}^n$ is the number of records in which they differ. Group privacy says that an ε-DP mechanism M is $t\varepsilon$-DP for datasets at Hamming distance at most t. So for such datasets D and D' and every event S,

$$\Pr[M(D) \in S] \le e^{t\varepsilon} \Pr[M(D') \in S]. \tag{11.5}$$

11.2.1. A first example: the mode

The simplest instance of the argument is the mode of a dataset over $[k]$, the most frequent symbol. Take the k datasets $D_1, \dots, D_k$, where D_j consists of n copies of the symbol j. The mode of D_j is j, and any two of these datasets are at Hamming distance n. An accurate mechanism must answer j on D_j at least half the time, for every j. Group privacy then forces each answer j to keep some probability on every other dataset as well, and k such probabilities cannot all fit into a total of 1 unless n is large.

Proposition 11.2 (Lower bound for the mode)**.** *Let $k \ge 2$ and $\varepsilon > 0$, and for each $j \in [k]$ let D_j be the dataset of n copies of the symbol j. Let $M \colon [k]^n \to [k]$ be ε-DP, and suppose that $\Pr[M(D_j) = j] \ge 1/2$ for every $j \in [k]$. Then*

$$n \ge \frac{\ln(k/2)}{\varepsilon}. \tag{11.6}$$

Proof.

1. *Pigeonhole.* Fix any j. The k probabilities $\Pr[M(D_j) = \ell]$, for $\ell \in [k]$, sum to 1. So at least one symbol ℓ has $\Pr[M(D_j) = \ell] \le 1/k$, since otherwise the sum would exceed 1.
2. *Accuracy.* The mechanism is correct on D_ℓ with probability at least $1/2$, so $\Pr[M(D_\ell) = \ell] \ge 1/2$.
3. *Group privacy.* The datasets D_ℓ and D_j are at Hamming distance at most n. By Equation 11.5 with the event $\{\ell\}$,

$$\frac{1}{2} \le \Pr[M(D_\ell) = \ell] \le e^{n\varepsilon} \Pr[M(D_j) = \ell] \le \frac{e^{n\varepsilon}}{k}. \tag{11.7}$$

4. *Rearranging.* The chain gives $e^{n\varepsilon} \ge k/2$, and taking logarithms gives $n\varepsilon \ge \ln(k/2)$.

□

So a private mode needs a number of records of order $\ln k/\varepsilon$. With $k = 1{,}000$ symbols and $\varepsilon = 1$, the bound is $\ln 500 = 6.21$, so no 1-DP algorithm finds the mode reliably from 6 records.

11.2.2. The packing theorem

The proof used three things: many datasets, disjoint sets of correct answers, and group privacy across a bounded distance. The theorem below keeps exactly these. It also lets the datasets be close to one common centre D, at distance t that may be much smaller than n, which the histogram bound will need.

Theorem 11.1 (Packing)**.** *Let $D_1, \dots, D_m \in \mathcal{X}^n$ be datasets, each at Hamming distance at most t from a fixed dataset $D \in \mathcal{X}^n$. Let $Y_1, \dots, Y_m$ be pairwise disjoint subsets of the output space $\mathcal{Y}$. If $M \colon \mathcal{X}^n \to \mathcal{Y}$ is ε-DP and, for some $p > 0$,*

$$\Pr[M(D_\ell) \in S_\ell] \geq p \qquad \text{for every } \ell \in [m], \tag{11.8}$$

then

$$\frac{1}{m} \geq p\, e^{-t\varepsilon}. \tag{11.9}$$

Proof.

1. *Accuracy and group privacy.* Fix ℓ. The datasets D_ℓ and D are at Hamming distance at most t, so Equation 11.5 gives $p \leq \Pr[M(D_\ell) \in S_\ell] \leq e^{t\varepsilon} \Pr[M(D) \in S_\ell]$. Hence $\Pr[M(D) \in S_\ell] \geq pe^{-t\varepsilon}$.
2. *Disjointness.* The sets S_ℓ are disjoint, so the probability of their union is the sum of their probabilities:

$$1 \geq \Pr\left[M(D) \in \bigcup_{\ell=1}^{m} S_\ell\right] = \sum_{\ell=1}^{m} \Pr[M(D) \in S_\ell] \geq m\, p\, e^{-t\varepsilon}. \tag{11.10}$$

3. *Rearranging.* Dividing by m gives Equation 11.9.

□

Taking logarithms, the theorem says that $t\varepsilon \geq \ln(mp)$. With the trivial distance $t = n$ and $p = 1/2$, this reads $n \geq \ln(m/2)/\varepsilon$, which is Proposition 11.2 again with $m = k$. A lower bound on n therefore needs a packing with many datasets m, accuracy regions that do not overlap, and a small distance t in terms of n.

11.3. One-way marginals

To make the answers on different datasets as different as possible, take datasets made of one repeated record. For each $w \in \{0,1\}^d$, let D_w consist of n copies of w. Its marginals are $f(D_w) = w$, a corner of the cube $[0,1]^d$, and there are 2^d corners. Two distinct corners differ by 1 in some coordinate, so an answer with every error below $1/2$ points to a single corner. The packing theorem then applies with 2^d datasets.

Theorem 11.2 (Lower bound for one-way marginals)**.** *Let $\varepsilon > 0$. Let $M \colon (\{0,1\}^d)^n \to \mathbb{R}^d$ be ε-DP, and suppose that on every dataset D, with probability at least $1/2$, it answers every one-way marginal with error less than $1/2$:*

$$\Pr\Big[\max_j |M(D)_j - f_j(D)| < \tfrac{1}{2}\Big] \geq \tfrac{1}{2}. \tag{11.11}$$

Then

$$n \geq \frac{(d-1)\ln 2}{\varepsilon}. \tag{11.12}$$

Proof.

1. *The packing.* Use the $m = 2^d$ datasets D_w. For each w, let $S_w = \{y \in \mathbb{R}^d : |y_j - w_j| < 1/2 \text{ for all } j\}$ be the open ℓ_∞ ball of radius $1/2$ around the corner w. Since $f(D_w) = w$, the hypothesis says exactly that $\Pr[M(D_w) \in S_w] \geq 1/2$.
2. *Disjointness.* Let $w \neq w'$ differ in coordinate j, so $|w_j - w'_j| = 1$. For any y, the triangle inequality gives $|y_j - w_j| + |y_j - w'_j| \geq 1$, so the two distances cannot both be less than $1/2$. Hence S_w and $S_{w'}$ are disjoint.
3. *Distance.* Any two datasets in $\mathcal{X}^n$ differ in at most n records, so every D_w is within Hamming distance $t = n$ of $D = D_0$, the all-zero dataset.
4. *Packing.* Theorem 11.1 with $m = 2^d$, $p = 1/2$ and $t = n$ gives $2^{-d} \geq \frac{1}{2}e^{-n\varepsilon}$. Taking logarithms, $n\varepsilon \geq d\ln 2 - \ln 2$.

□

In two dimensions, the four sets S_w are open unit squares centred at the corners $(0,0)$, $(0,1)$, $(1,0)$ and $(1,1)$, and each meets $[0,1]^2$ in the open quarter square at its corner. An answer in the quarter at $(0,0)$ is acceptable for $D_{(0,0)}$ and for no other dataset in the packing.

So $n = \Omega(d/\varepsilon)$ records are necessary for error below $1/2$ under ε-DP, against the upper bound $d\ln(2d)/(\alpha\varepsilon)$ of Proposition 11.1. Under approximate DP, n of order $\sqrt{d}$ suffices. This gap of a factor $\sqrt{d}$ between pure and approximate DP is inherent. A matching lower bound of order $\sqrt{d}$ under approximate DP also holds, but packing cannot give it, because group privacy under (ε,δ)-DP degrades too fast with t. That bound needs fingerprinting codes, a more delicate technique.

The bound above does not depend on α, because the packing uses corners of the cube, which are far apart. Datasets with only a few records changed have marginals closer together, and they give the factor $1/\alpha$.

Corollary 11.1. *Let $\varepsilon > 0$ and $\alpha \in (0, 1/2)$. If an ε-DP mechanism $M \colon (\{0,1\}^d)^n \to \mathbb{R}^d$ answers every one-way marginal with error at most α, with probability at least $1/2$ on every dataset, then*

$$n \geq \frac{(d-1)\ln 2 - \varepsilon}{2\alpha\varepsilon}. \tag{11.13}$$

Proof.

1. *The packing.* Let $t = \lfloor 2\alpha n \rfloor + 1$, so that $t/n > 2\alpha$, and $t \leq n$ because $2\alpha < 1$. For each $w \in \{0,1\}^d$, let D_w consist of t copies of w and $n - t$ copies of the all-zero record. Its marginals are $f(D_w) = (t/n)\, w$.
2. *Disjointness.* Let S_w be the closed ℓ_∞ ball of radius α around $(t/n)\, w$, so that $\Pr[M(D_w) \in S_w] \geq 1/2$ by hypothesis. If $w \neq w'$ differ in coordinate j, the centres differ by $t/n > 2\alpha$ in that coordinate, and two closed intervals of half-width α whose centres are more than 2α apart do not meet. So the sets S_w are pairwise disjoint.
3. *Distance.* Each D_w differs from the all-zero dataset in at most t records.
4. *Packing.* Theorem 11.1 with $m = 2^d$ and $p = 1/2$ gives $\ln(2^{d-1}) \leq t\varepsilon \leq (2\alpha n + 1)\varepsilon$. Rearranging gives the claim.

□

For concrete parameters, the lower bound and the Laplace upper bound of Proposition 11.1 are within a modest factor of each other.

Example 11.1. With $d = 1{,}000$ questions, $\alpha = 0.1$ and $\varepsilon = 1$, Corollary 11.1 says that $n \geq 3{,}458$ is necessary, and Equation 11.4 says that $n \geq 76{,}010$ suffices. The ratio, 22.0, is the factor $\ln(2d) = 7.60$ times constants, and it grows only logarithmically with d.

11.4. Histograms

For histograms, the corner construction gives only the mode bound $\ln(k/2)/\varepsilon$, with no α. The factor $1/\alpha$ comes from datasets that are all close to one centre, as in Corollary 11.1. Take the centre D_1 to be n copies of the symbol 1, and move a small block of records to another symbol.

Theorem 11.3 (Lower bound for histograms)**.** *Let $k \geq 2$, $\varepsilon > 0$ and $\alpha \in (0, 1/2)$. If an ε-DP mechanism $M \colon [k]^n \to \mathbb{R}^k$ estimates the normalized histogram with error at most α,*

$$\Pr\Big[\max_j |M(D)_j - h_j(D)| \leq \alpha\Big] \geq \tfrac{1}{2} \qquad \text{for every } D \in [k]^n, \tag{11.14}$$

then

$$n \geq \frac{\ln(k/2) - \varepsilon}{2\alpha\varepsilon}. \tag{11.15}$$

Proof.

1. *The packing.* Let $t = \lfloor 2\alpha n \rfloor + 1$, so that $t/n > 2\alpha$, and $t \le n$ because $2\alpha < 1$. Let D_1 be n copies of the symbol 1, so its histogram puts all its mass on the symbol 1. For $\ell = 2, \dots, k$, let D_ℓ consist of t copies of ℓ and $n - t$ copies of 1. Its histogram $h(D_\ell)$ puts mass t/n on the symbol ℓ, mass $1 - t/n$ on the symbol 1, and nothing elsewhere.
2. *Accuracy regions.* Let S_ℓ be the closed ℓ_∞ ball of radius α around $h(D_\ell)$. The hypothesis applied to the dataset D_ℓ says that $\Pr[M(D_\ell) \in S_\ell] \ge 1/2$.
3. *Disjointness.* Take $\ell \ne \ell'$. At least one of them is not 1, say $\ell \ne 1$. In coordinate ℓ, the centre $h(D_\ell)$ has value t/n and the centre $h(D_{\ell'})$ has value 0, because $D_{\ell'}$ contains only the symbols 1 and ℓ'. These differ by more than 2α, so S_ℓ and $S_{\ell'}$ are disjoint.
4. *Distance.* Each D_ℓ differs from D_1 in at most t records, the ones not equal to 1.
5. *Packing.* Theorem 11.1 with $m = k$, $p = 1/2$ and centre D_1 gives $1/k \ge \frac{1}{2}e^{-t\varepsilon}$, that is, $\ln(k/2) \le t\varepsilon$. Since $t \le 2\alpha n + 1$, this gives $\ln(k/2) \le (2\alpha n + 1)\varepsilon$, which rearranges to Equation 11.15.

□

Figure 11.1 shows the construction for $k = 5$, $n = 20$ and $\alpha = 0.1$, where $t = 5$ and each moved block holds a quarter of the records. In coordinate ℓ, the band of acceptable answers for D_ℓ starts at 0.15, above the band of every other dataset, which ends at 0.1.

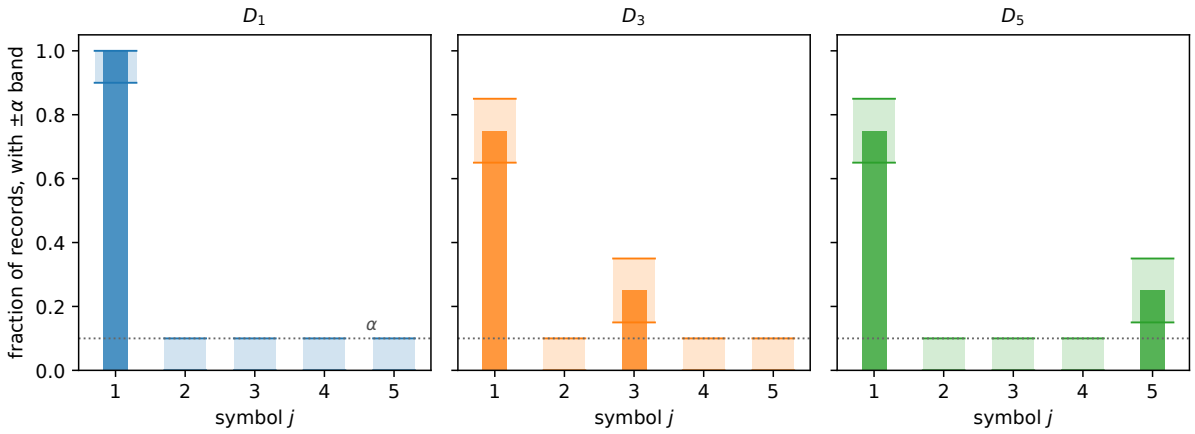


Figure 11.1.: The packing behind Theorem 11.3, for $k = 5$, $n = 20$ and $\alpha = 0.1$. Each panel shows the normalized histogram of one dataset as bars and the band of answers within α of each count as a shaded box. Three of the five datasets are shown. Any two datasets have non-overlapping bands in at least one coordinate, so no answer is acceptable for both.

The choice of t is what produces the factor $1/\alpha$. Moving the whole dataset, $t = n$, would also give disjoint regions, but the packing theorem would then bound only $n\varepsilon$, as in Proposition 11.2. A block of only about $2\alpha n$ records is just enough to separate the answers, and group privacy pays $e^{t\varepsilon}$ for those t records only. As for the marginals, numbers show how close this lower bound comes to the Laplace histogram.

Example 11.2. With $k = 10^6$ symbols, $\alpha = 0.01$ and $\varepsilon = 1$, Equation 11.15 says that $n \ge 607$ is necessary, and the Laplace histogram of Proposition 11.1 succeeds once $n \ge 2{,}902$. The ratio

is 4.8. As k grows, it falls towards 4, because the upper bound is $2\ln(2k)/(\alpha\varepsilon)$ and the lower bound is about $\ln k/(2\alpha\varepsilon)$. At $k = 10^{12}$ the ratio is 4.37. The union bound in Proposition 11.1 is nearly tight. In simulation with $k = 10^4$ and $n = 1{,}981$, the value Equation 11.4 gives for that k, the Laplace histogram met error α in 217 of 400 runs, 54%, against the 50% that the bound guarantees.

So under pure DP the Laplace histogram is optimal up to a constant factor, and the $\ln k$ in its sample size cannot be removed. The stability-based histogram, which replaces $\ln k$ by $\ln(1/\delta)$, shows that this cost belongs to pure DP and not to the problem.

11.5. Summary

Table 11.1.: The chapter's results. All lower bounds hold for every ε-DP mechanism that meets the accuracy requirement with probability at least 1/2.

Result	Statement
Proposition 11.1	Laplace noise answers d marginals with $n \geq d\ln(2d)/(\alpha\varepsilon)$ and a k-bin histogram with $n \geq 2\ln(2k)/(\alpha\varepsilon)$
Proposition 11.2	an ε-DP mode that is right half the time needs $n \geq \ln(k/2)/\varepsilon$
Theorem 11.1	m datasets within distance t of a centre, with disjoint accuracy regions hit with probability p, force $1/m \geq pe^{-t\varepsilon}$
Theorem 11.2	error below 1/2 on all one-way marginals needs $n \geq (d-1)\ln 2/\varepsilon$
Corollary 11.1	error at most α on all one-way marginals needs $n \geq ((d-1)\ln 2 - \varepsilon)/(2\alpha\varepsilon)$
Theorem 11.3	error at most α on a k-bin histogram needs $n \geq (\ln(k/2) - \varepsilon)/(2\alpha\varepsilon)$

11.6. References

- Moritz Hardt and Kunal Talwar, "On the Geometry of Differential Privacy", 2010. Packing lower bounds for linear queries under pure DP, stated geometrically.
- Amos Beimel, Hai Brenner, Shiva Prasad Kasiviswanathan and Kobbi Nissim, "Bounds on the Sample Complexity for Private Learning and Private Data Release", 2014. Packing arguments for the sample size of private learning and of releasing point functions, the histogram problem.
- Mark Bun, Jonathan Ullman and Salil Vadhan, "Fingerprinting Codes and the Price of Approximate Differential Privacy", 2014. The lower bound of order $\sqrt{d}$ for one-way marginals under approximate DP.

- Thomas Steinke and Jonathan Ullman, "Between Pure and Approximate Differential Privacy", 2016. Tight bounds for one-way marginals across pure and approximate DP.
- Salil Vadhan, "The Complexity of Differential Privacy", 2017. Section 5 presents packing lower bounds in the generality used here.

Part IV.

Private Learning and Statistics

The earlier parts released statistics of a fixed dataset under differential privacy (DP). This part turns to the goals of machine learning and statistics: fitting models, and learning about the population from which the data was drawn.

It opens by asking what privacy should protect. A trained model can predict a person's sensitive label, and the prediction can feel like a leak even when nothing about that person's record was used. The first chapter separates the inferences the definition is designed to allow from those it rules out. It proves that a private learner can still find a rule that holds in the population, and that it cannot memorize a secret held in one person's records.

The next two chapters train models privately. Private empirical risk minimization handles convex losses, and adds noise to the output of a solver, to its objective or to each gradient step. Noisy stochastic gradient descent comes with proofs of both privacy and accuracy. Private deep learning carries that algorithm over to neural networks by clipping gradients. It replaces composition with a tighter accounting based on the moments of the privacy loss, and it presents a second route, a private vote among models trained on disjoint parts of the data.

The last two chapters study the population. Private mean estimation asks how many samples a private algorithm needs to estimate the mean of a distribution, from bounded data to unbounded Gaussians and heavy-tailed data. Adaptive data analysis shows that reusing one dataset for questions chosen in sequence makes the answers overfit the sample. Answering through a private mechanism prevents this, because a private algorithm cannot find a question on which the sample and the population disagree.

Table 11.2.: The chapters of this part and the questions they answer.

Chapter	Question it answers
What differential privacy protects	Which inferences from a trained model are privacy violations, and which are learning at work?
Private empirical risk minimization	How can a convex model be fitted privately, and at what cost in accuracy?
Private deep learning	How are neural networks trained privately, and how tightly can their privacy be accounted?
Private mean estimation	How many samples does a private estimate of a population mean need?
Adaptive data analysis	How can one dataset answer adaptively chosen questions without overfitting?

12. What Differential Privacy Protects

Machine learning on sensitive data often produces results that feel invasive, and it is tempting to call each of them a privacy violation. This chapter separates the inferences that differential privacy (DP) is meant to prevent from those it is designed to allow. It gives a test for telling them apart, and proves the two facts behind it. A private learner can still discover a rule that predicts a person's sensitive label. A private learner cannot memorize a secret that appears only in one person's records.

12.1. What the definition promises

Recall the definition from the randomized response chapter. A mechanism M is ε-DP if, for all pairs of neighbouring datasets $D \sim D'$, which differ by replacing one record, and all sets of outputs S,

$$\Pr[M(D) \in S] \leq e^{\varepsilon} \Pr[M(D') \in S]. \tag{12.1}$$

The definition says that the output changes little when one record changes, and it says nothing else. In particular, it does not stop statistics or learning from working, nor anything that follows from their results.

Two properties from the Laplace mechanism chapter turn the definition into a statement about attackers. **Post-processing** says that a function of the output of an ε-DP mechanism, computed with coins independent of the data, is again ε-DP. **Group privacy** says that if D and D' differ in some number k of entries, then

$$\Pr[M(D) \in S] \leq e^{k\varepsilon} \Pr[M(D') \in S]. \tag{12.2}$$

Write $D^{(i \leftarrow z)}$ for the dataset D with entry i replaced by a record z. An attacker is any function $\mathcal{A}$ of the output, possibly randomized, that answers 1 when it believes it has learned something harmful. The attacker may know every other record and any side information, since all of that is fixed in $\mathcal{A}$.

The first fact behind the test says how much such an attacker gains from one person's record. It compares the attacker's chance of answering 1 on the real dataset with its chance when that person's record is replaced by any other.

Proposition 12.1 (What an attacker gains from one record)**.** *Let M be ε-DP, and let $\mathcal{A}$ be a randomized function of the output whose coins are independent of the coins of M. Then for every dataset D, every index i and every record z,*

$$\Pr\left[\mathcal{A}(M(D)) = 1\right] \leq e^{\varepsilon} \Pr\left[\mathcal{A}(M(D^{(i \leftarrow z)})) = 1\right]. \tag{12.3}$$

Proof.

1. The composed algorithm $\mathcal{A} \circ M$ is ε-DP by post-processing.
2. The datasets D and $D^{(i \leftarrow z)}$ differ in at most one entry, so they are neighbours or equal. Applying the definition to the event $\{1\}$ gives the inequality.

□

So any harm the output can cause a person is about as likely when their record is swapped for an arbitrary placeholder. An inference that would succeed just as well without the record is not something the definition tries to stop. This suggests three questions to ask of an alleged violation.

1. Would the result be substantially different if one person's record were changed or left out?
2. Is the harmful inference the same thing as the learned model being accurate?
3. If the inference still feels wrong, at what point, and by whom, was the information given away?

The second question deserves a word. If an inference about a person is exactly what an accurate model predicts, the only way to prevent it is to release an inaccurate model, for instance one that ignores the data. That is always possible and always useless.

12.2. Learning a rule is not a violation

A learning algorithm takes a training set of pairs (x_i, y_i), where x_i is a **feature vector** and y_i a **label**, and outputs a **classifier**, a function that maps a feature vector to a predicted label. Its quality is measured on test points, usually as the fraction of test labels it predicts correctly. In the private setting, the training set is the sensitive dataset and the classifier is the output that must be differentially private. Test points are not protected. Some work protects only the classifier's predictions instead of the classifier itself.

Consider a deliberately simple training set of real pairs,

$$(0, 1),\ (2, 3),\ (5, 6),\ (-6, -5),\ (0.5, 1.5),\ (\pi, \pi + 1),\ \ldots \tag{12.4}$$

in which every label is one more than its feature. Any reasonable learner finds the rule $y = x + 1$. Asked for the label of $x = 4$, which is not in the training set, it answers 5, and no one would call that a violation. Asked for the label of $x = \pi$, it answers $\pi + 1$, which is the true label of a training record. That is not a violation either. The learner found a fact about the population, and the answer at $x = \pi$ would be the same if the record $(\pi, \pi + 1)$ had never been collected.

The claim can be made exact. Each record determines an offset $y_i - x_i$, and a rule of the form $y = x + b$ is a choice of offset. The stability-based histogram of the chapter on beyond global sensitivity releases the offsets that many records share, even though offsets are real numbers and there are infinitely many possible bins. Its guarantee is approximate differential privacy, (ε, δ)-DP, from the approximate DP chapter, which allows the bound of the definition an additive slack δ:

$$\Pr[M(D) \in S] \le e^{\varepsilon} \Pr[M(D') \in S] + \delta \quad \text{for all } D \sim D' \text{ and all } S. \tag{12.5}$$

The algorithm counts how many records share each offset, adds Laplace noise to each count, and releases the offsets whose noisy count clears a threshold. A threshold that is high enough keeps an offset that only one record holds from being released.

Definition 12.1 (Stable histogram of offsets)**.** Let $\varepsilon > 0$ and $\delta \in (0, 1)$, and set $\tau = 1 + 2\ln(1/\delta)/\varepsilon$. On input $D = ((x_1, y_1), \dots, (x_n, y_n))$ with real entries, let $c_b = |\{i : y_i - x_i = b\}|$ be the number of records whose offset is b. For each b with $c_b > 0$, draw an independent Z_b from the Laplace distribution $\mathrm{Lap}(2/\varepsilon)$ with scale $2/\varepsilon$, and output the set

$$M(D) = \{\, b : c_b > 0 \text{ and } c_b + Z_b > \tau \,\}. \tag{12.6}$$

This is the stability-based histogram run on the offsets $y_i - x_i$, with the noisy counts dropped from its output. Bins with count zero are never released, so the output only ever names offsets that occur in the data. The threshold τ keeps a bin held by a single record from being released, except with small probability. The privacy guarantee is inherited from the histogram.

Theorem 12.1. *The stable histogram of offsets is (ε, δ)-DP.*

Proof.

1. Replacing one record changes at most one offset, so when $D \sim D'$ the two lists of offsets are neighbours or equal.
2. The stability-based histogram is (ε, δ)-DP on datasets of offsets, by the theorem of that chapter.
3. $M(D)$ is the set of elements on its list, a function of its output alone, so it is (ε, δ)-DP by post-processing, which holds for approximate DP as well.

□

The privacy cost does not depend on how many records follow the rule. Accuracy, on the other hand, needs only enough of them to clear the threshold.

Proposition 12.2. *Suppose every record satisfies $y_i = x_i + 1$, let τ be the threshold of Definition 12.1, and let $\beta \in (0, 1/2)$ be a failure probability. If*

$$n \ge \tau + \frac{2}{\varepsilon} \ln \frac{1}{2\beta}, \tag{12.7}$$

then with probability at least $1 - \beta$ the output is exactly $\{1\}$, and the released rule predicts $y = x + 1$ at every feature value.

Proof.

1. Only the offset 1 has a positive count, $c_1 = n$, so the output is either $\{1\}$ or empty.
2. It is empty when $n + Z_1 \le \tau$, which has probability $\Pr[Z_1 \le -(n-\tau)] = \frac{1}{2}e^{-\varepsilon(n-\tau)/2}$, because $n \ge \tau$ and a $\mathrm{Lap}(2/\varepsilon)$ variable satisfies $\Pr[Z \le -s] = \frac{1}{2}e^{-\varepsilon s/2}$ for $s \ge 0$.
3. The assumed bound on n makes this at most β.

□

Concrete values show how few records the rule needs.

Example 12.1. Take $\varepsilon = 1$, $\delta = 10^{-6}$ and $\beta = 0.01$. The threshold is $\tau = 1 + 2\ln(10^6) = 1 + 2 \times 13.82 = 28.63$. Adding $2\ln 50 = 7.82$ gives 36.45, so 37 records that follow the rule suffice for the private algorithm to output it with probability at least 0.99.

So the record $(\pi, \pi + 1)$ is protected in the sense of Proposition 12.1, and yet the released rule predicts its label exactly. The prediction is a consequence of the rule, and the rule is a fact about nature. If this still feels like a leak, ask the third question. The information was given away when someone revealed that a person's feature value is π. From then on, anyone who knows the rule can work out the label.

12.3. Inferences from what people reveal

The same pattern explains several examples that are often presented as privacy failures.

Smoking and cancer. A smoker decides whether to join a study of whether smoking causes cancer, and fears higher insurance premiums if the link is published. This is the rule $y = x + 1$ in another form. The feature is "smokes", the label is "likely to develop cancer", and the study would reach the same conclusion without this person. The only way to prevent the harm is not to run the study. Privacy was given up when the person's smoking became known.

A cast. Stripped of the study, the point is plainer. Someone who walks around with a cast on their foot probably has a broken foot. They may regard the injury as private, but they display the cast in public, and common knowledge links the two.

Social network likes. Kosinski, Stillwell and Graepel showed in 2013 that the pages a person likes on Facebook, which are public by default, predict sensitive traits such as sexual orientation and political views. The model was trained on public likes and on traits that many users state on their own profiles, so there was no sensitive dataset to protect, and the results would presumably hold under differential privacy too. Privacy was given up when the likes were made public. This example adds one lesson: whether a piece of information is revealing may become known only later, after it was shared in the belief that it was harmless.

Targeted advertising. A retailer is said to have mailed baby-product coupons to a teenager, having inferred her pregnancy from her purchases before her family knew. The story may be apocryphal. Even if it is true, the analysis was not the violation. The information was handed over when her purchases were linked to her identity through a loyalty programme, and paying anonymously would have prevented it.

12.4. What a violation looks like

A violation is an inference that depends on the person's own sensitive record being in the dataset.

The Netflix Prize. Netflix released movie ratings with dates under pseudonymous user identifiers. Narayanan and Shmatikov linked the records to public reviews on the Internet Movie Database and recovered, for the users they matched, films those users had rated but never reviewed in public. Had a user's record been withheld, that reidentification would have been impossible, so it fails the first question. Compare an analyst who sees only public reviews. If a user reviewed the first and third films of a trilogy, the analyst may guess that they watched the second. The guess may be right, but it uses only what the user chose to publish, so it is not a violation.

Memorized secrets. Carlini and coauthors trained language models on text that contained a secret, such as a nine-digit identification number, that one person had typed several times. The trained model assigned the secret an unusually high likelihood, and a search guided by those scores could extract it. This is a violation. A model that had learned only the language would have no reason to prefer one nine-digit number over another, so the preference comes from that person's records. Differential privacy rules it out, with group privacy accounting for the repetitions. The next proposition makes this exact. It bounds the chance that an attacker who sees the trained model guesses the secret, and compares it with blind guessing.

Proposition 12.3 (Secrets are not memorized)**.** *Let $\mathcal{S}$ be a finite set of N possible secrets. For each $s \in \mathcal{S}$, let $D(s)$ be a training set in which k entries depend on s and the other entries do not. Let M be ε-DP, and let G be any guessing rule, with coins independent of M, that maps the trained model to an element of $\mathcal{S}$. If the secret s is drawn uniformly from $\mathcal{S}$, then*

$$\Pr\left[G(M(D(s))) = s\right] \le \frac{e^{k\varepsilon}}{N}. \tag{12.8}$$

Proof.

1. Let D_0 be the training set with the k entries that depend on the secret replaced by fixed placeholder records. For every s, the datasets $D(s)$ and D_0 differ in at most k entries.
2. By post-processing and group privacy, $\Pr[G(M(D(s))) = s] \le e^{k\varepsilon} \Pr[G(M(D_0)) = s]$ for every s.
3. The guess on D_0 does not depend on s, and it equals at most one secret, so $\sum_{s \in \mathcal{S}} \Pr[G(M(D_0)) = s] \le 1$.
4. Averaging step 2 over the uniform secret and using step 3 gives

$$\Pr\left[G(M(D(s))) = s\right] = \frac{1}{N} \sum_{s \in \mathcal{S}} \Pr\left[G(M(D(s))) = s\right] \le \frac{e^{k\varepsilon}}{N}. \tag{12.9}$$

□

For a nine-digit number, $N = 10^9$. With $\varepsilon = 1$ and a secret typed $k = 4$ times, the attacker's chance is at most $e^4/10^9 = 5.5 \times 10^{-8}$, against 10^{-9} for blind guessing. The bound degrades exponentially in k, which is why a person who contributes many records needs a guarantee per person rather than per record.

The three questions settle most cases. The most useful is the second. If an alleged violation is at the same time the model doing what it is supposed to do, then it is the notion of violation, not the model, that needs another look.

12.5. Summary

Table 12.1.: The chapter's results.

Result	Statement
Proposition 12.1	an attacker succeeds about as often when a person's record is replaced by any other
Theorem 12.1	releasing offsets whose noisy count exceeds τ is (ε, δ)-DP
Proposition 12.2	with $n \geq \tau + \frac{2}{\varepsilon} \ln \frac{1}{2\beta}$ records, the private algorithm finds $y = x + 1$
Proposition 12.3	a secret in k records of one person is guessed with probability at most $e^{k\varepsilon}/N$

12.6. References

- Frank McSherry, "Statistical Inference Considered Harmful", 2016. A blog post arguing that inference from public facts, such as the cast, is not a privacy violation.
- Cynthia Dwork and Aaron Roth, *The Algorithmic Foundations of Differential Privacy*, 2014. Section 2.3 discusses what the definition promises and the smoking example.
- Michal Kosinski, David Stillwell and Thore Graepel, "Private Traits and Attributes Are Predictable from Digital Records of Human Behavior", 2013. Prediction of sensitive traits from public likes.
- Arvind Narayanan and Vitaly Shmatikov, "Robust De-anonymization of Large Sparse Datasets", 2008. The linkage attack on the Netflix Prize data.
- Nicholas Carlini, Chang Liu, Úlfar Erlingsson, Jernej Kos and Dawn Song, "The Secret Sharer: Evaluating and Testing Unintended Memorization in Neural Networks", 2019. Extraction of secrets that appear in few training records.
- Abhradeep Thakurta and Adam Smith, "Differentially Private Feature Selection via Stability Arguments, and the Robustness of the Lasso", 2013, and Salil Vadhan, "The Complexity of Differential Privacy", 2017, Section 3. Stability-based algorithms such as the histogram behind Theorem 12.1.

13. Private Empirical Risk Minimization

Most of machine learning fits a model by minimizing its average loss on a training set. This chapter asks how to do that under differential privacy (DP) when the loss is convex. Noise can enter at three places: in the output of a non-private solver, in the objective that the solver minimizes, or in each step of gradient descent. The chapter proves a privacy guarantee for each, and an accuracy guarantee for the third, which is the method that the next chapter scales to neural networks.

13.1. Empirical risk minimization

Each record $z = (x, y)$ is a pair of a feature vector x and a label y. The dataset $D = (z_1, \dots, z_n)$ holds n records.

A model is described by a parameter vector θ in $\mathbb{R}^d$. A **loss function** $\ell(\theta; z)$ measures how badly the model θ fits the record z. Fitting a model means choosing θ so that the total loss over the dataset is small, with θ restricted to a set of allowed values. The definition also says how to measure an algorithm that only comes close to the best choice.

Definition 13.1 (Empirical risk minimization)**.** Let $\Theta \subseteq \mathbb{R}^d$ be a nonempty, closed, convex set of allowed parameters. The empirical risk of θ on D is the total loss

$$\mathcal{L}(\theta; D) = \sum_{i=1}^{n} \ell(\theta; z_i), \tag{13.1}$$

and $\theta^* \in \Theta$ denotes a minimizer of $\mathcal{L}(\cdot\,; D)$ over Θ. An algorithm that outputs $\tilde{\theta} \in \Theta$ has **excess empirical risk** $\mathbb{E}[\mathcal{L}(\tilde{\theta}; D)] - \mathcal{L}(\theta^*; D)$, where the expectation is over the algorithm's coins only. The dataset is fixed.

The problem is called empirical risk minimization (ERM). In practice one cares about the loss on new data from the same distribution, which is **generalization**. Where a uniform convergence argument applies, a good empirical minimizer also generalizes, so this chapter studies the empirical problem only. A model θ comes with a predictor, such as $x \mapsto \langle x, \theta \rangle$ for linear regression, which the chapter does not need.

In this generality the problem is hopeless, so three assumptions are made throughout. The set of allowed parameters must be bounded, which its diameter measures. The loss must be convex, so that a local minimum is a global one, and Lipschitz, so that it cannot change quickly. The definitions below make these terms precise, together with the subgradient, which replaces the gradient where a convex function has a kink.

Definition 13.2 (Diameter, convexity, Lipschitz continuity)**.**

- The **diameter** of Θ is $\|\Theta\|_2 = \sup_{\theta,\theta' \in \Theta} \|\theta - \theta'\|_2$.
- A function f on a convex set K is **convex** if $f(t\theta + (1-t)\theta') \le t f(\theta) + (1-t) f(\theta')$ for all $\theta, \theta' \in K$ and $t \in [0,1]$.
- A function f on K is L**-Lipschitz** if $|f(\theta) - f(\theta')| \le L\|\theta - \theta'\|_2$ for all $\theta, \theta' \in K$.
- A vector g is a **subgradient** of a convex f at θ if $f(\theta') \ge f(\theta) + \langle g, \theta' - \theta \rangle$ for all θ'. Where f is differentiable, its gradient $\nabla f(\theta)$, the vector of partial derivatives, is the only subgradient.

The assumptions are that Θ has finite diameter, and that for every record z the loss $\theta \mapsto \ell(\theta; z)$ is convex and L-Lipschitz on $\mathbb{R}^d$. A convex function that is finite on all of $\mathbb{R}^d$ has a subgradient at every point, and we write $\nabla \ell(\theta; z)$ for a gradient, or a subgradient where ℓ is not differentiable. Lipschitz continuity matters to privacy because it bounds every gradient, and so the effect of one record on a gradient step.

Lemma 13.1. *If a convex function f is L-Lipschitz on $\mathbb{R}^d$, then every subgradient g of f satisfies $\|g\|_2 \le L$.*

Proof.

1. Suppose $g \neq 0$, take any point θ, and set $u = g/\|g\|_2$. The definition of a subgradient at $\theta' = \theta + u$ gives $f(\theta + u) - f(\theta) \ge \langle g, u \rangle = \|g\|_2$.
2. The Lipschitz condition gives $f(\theta + u) - f(\theta) \le L\|u\|_2 = L$. Together, $\|g\|_2 \le L$.

□

The losses used most often are convex, and most of them are Lipschitz when the features are bounded. The examples give the Lipschitz constant of each.

Example 13.1 (Common losses)**.** Take $x, \theta \in \mathbb{R}^d$.

- **Linear regression**, $y \in \mathbb{R}$: the squared loss $\ell(\theta; x, y) = (\langle x, \theta \rangle - y)^2$. It is convex, and Lipschitz only on bounded sets.
- **Logistic regression**, $y \in \{-1, 1\}$: $\ell(\theta; x, y) = \ln\left(1 + e^{-y\langle x, \theta \rangle}\right)$. Its gradient is $-y\,x/(1 + e^{y\langle x, \theta \rangle})$, of norm at most $\|x\|_2$, so it is $\|x\|_2$-Lipschitz.
- **Geometric median**, no label: $\ell(\theta; x) = \|\theta - x\|_2$. By the triangle inequality it is 1-Lipschitz. In one dimension it is $|\theta - x|$, and its minimizer is the median. The squared version $\|\theta - x\|_2^2$ is minimized by the mean.
- **Support vector machine (SVM)**, $y \in \{-1, 1\}$: the hinge loss $\ell(\theta; x, y) = \max(0, 1 - y\langle x, \theta \rangle)$. It is a maximum of two functions that are linear in θ, so it is convex and $\|x\|_2$-Lipschitz.

Other assumptions, such as strong convexity of the loss or smoothness, give faster rates, but they are not needed here. The non-private problem is a convex optimization problem, and until Section 13.5 we treat the solver as a black box.

13.2. Why the minimizer is not private

Releasing an exact minimizer can reveal records outright. The median is the simplest case.

Proposition 13.1. *Let $d = 1$, $\ell(\theta; x) = |\theta - x|$, records in $[-R, R]$ with $R > 0$, $\Theta = [-R, R]$, and n odd. Then the empirical risk has a unique minimizer, the median $x_{(m)}$ with $m = (n+1)/2$, where $x_{(1)} \le ... \le x_{(n)}$ are the sorted records. The algorithm that outputs it is not (ε, δ)-DP for any ε and any $\delta < 1$.*

Proof.

1. Away from the records, the function $\theta \mapsto \sum_i |\theta - x_i|$ has slope equal to the number of records below θ minus the number above. For $\theta < x_{(m)}$ at least m records lie above, so the slope is at most $(m-1) - m = -1$. For $\theta > x_{(m)}$ the slope is at least 1 in the same way. The function is continuous, so $x_{(m)}$ is its unique minimizer.
2. On the dataset with every record equal to $-R$ the output is $-R$, and on the dataset with every record equal to R it is R. A path of n replacements joins the two datasets, so some neighbours $D \sim D'$ have different outputs $v \neq v'$.
3. The algorithm is deterministic. For the event $S = \{v\}$ this gives $\Pr[M(D) \in S] = 1$ and $\Pr[M(D') \in S] = 0$, and $1 \le e^{\varepsilon} \cdot 0 + \delta$ fails for every $\delta < 1$.

□

The same happens for support vector machines. With separable data, the solution of maximum margin is a combination of the few records closest to the separating hyperplane, the **support vectors**. One added record close to that hyperplane becomes a support vector and can move the solution far.

One way out is **input perturbation**: privatize each record, then run any non-private algorithm, which is private by post-processing. Randomizing each record separately is local differential privacy, which needs far more data for the same accuracy, so we do not pursue it. The three methods below add noise after the data has been pooled.

13.3. Output perturbation

The obvious method computes the minimizer and adds noise scaled to its sensitivity. As Proposition 13.1 shows, the minimizer itself can have large sensitivity, and the fix is regularization.

Polynomial regression shows the problem. With records $(x_i, y_i) \in \mathbb{R}^2$ and the squared loss $\ell(\theta; x, y) = \left(\sum_{j=0}^{p} \theta_j x^j - y\right)^2$, the model is a polynomial of degree p. If the degree is not known in advance and is set as high as $p = n - 1$, then for distinct features the minimizer passes through every point. Its excess risk is zero, but it fits new data poorly and swings wildly when one point moves. These are two views of the same instability, as the left panel of Figure 13.1 shows.

Regularization adds a penalty that favours simple models. For a function N of the parameters alone, called the **regularizer**, and a weight $\lambda > 0$, the regularized objective is

$$F_D(\theta) = \mathcal{L}(\theta; D) + \lambda N(\theta). \tag{13.2}$$

At $\lambda = 0$ there is no regularization, and as $\lambda \to \infty$ the data is ignored. The usual choice is $N(\theta) = \frac{1}{2}\|\theta\|_2^2$. Its effect on sensitivity follows from strong convexity. This is a stronger form of convexity, which asks that the function curve upward by at least a fixed amount everywhere, and the weight λ will be that amount for the regularized objective.

Definition 13.3 (Strong convexity)**.** For $\lambda > 0$, a function f on a convex set K is **λ-strongly convex** if, for all $\theta, \theta' \in K$ and $t \in [0, 1]$,

$$f(t\theta + (1-t)\theta') \le tf(\theta) + (1-t)f(\theta') - \frac{\lambda}{2}\, t(1-t)\, \|\theta - \theta'\|_2^2. \tag{13.3}$$

The function $\frac{1}{2}\|\theta\|_2^2$ is 1-strongly convex, with equality, by the identity $t\|a\|^2 + (1-t)\|b\|^2 - \|ta + (1-t)b\|^2 = t(1-t)\|a-b\|^2$. A λ-strongly convex function plus a convex function is λ-strongly convex. A strongly convex function grows at least quadratically away from its minimizer.

Lemma 13.2. *Let f be λ-strongly convex on a convex set K, with a minimizer $\hat{\theta}$ over K. Then for every $\theta \in K$,*

$$f(\theta) - f(\hat{\theta}) \ge \frac{\lambda}{2}\|\theta - \hat{\theta}\|_2^2. \tag{13.4}$$

In particular, the minimizer is unique.

Proof.

1. For $t \in (0, 1)$, the point $\hat{\theta} + t(\theta - \hat{\theta})$ lies in K, so its value is at least $f(\hat{\theta})$.
2. Strong convexity, applied with weight t on θ, bounds that value by $tf(\theta) + (1-t)f(\hat{\theta}) - \frac{\lambda}{2}t(1-t)\|\theta - \hat{\theta}\|_2^2$.
3. Combining the two and dividing by t gives $f(\theta) - f(\hat{\theta}) \ge \frac{\lambda}{2}(1-t)\|\theta - \hat{\theta}\|_2^2$. Letting $t \to 0$ gives the claim.
4. A second minimizer θ would have $f(\theta) - f(\hat{\theta}) = 0$, which forces $\theta = \hat{\theta}$.

□

This growth is what keeps the regularized minimizer stable. When one record changes, the two regularized objectives differ by a Lipschitz function. Each objective grows quadratically around its own minimizer, so the two minimizers cannot be far apart.

Theorem 13.1 (Sensitivity of the regularized minimizer)**.** *Let Θ be nonempty, closed, convex and bounded. Suppose that for every record z the loss $\ell(\cdot\,; z)$ is convex and L-Lipschitz on Θ, and that N is continuous and 1-strongly convex. For each dataset D, let θ_D be the minimizer of $F_D = \mathcal{L}(\cdot\,; D) + \lambda N$ over Θ. Then for all neighbours $D \sim D'$, which differ in one replaced record,*

$$\|\theta_D - \theta_{D'}\|_2 \le \frac{2L}{\lambda}. \tag{13.5}$$

Proof.

1. F_D is continuous on the closed bounded set Θ, so it has a minimizer, and it is λ-strongly convex, so the minimizer is unique by Lemma 13.2. The same holds for $F_{D'}$.
2. Write $\theta = \theta_D$ and $\theta' = \theta_{D'}$. Lemma 13.2 applied to each function gives

$$F_D(\theta') - F_D(\theta) \geq \frac{\lambda}{2}\|\theta - \theta'\|_2^2, \qquad F_{D'}(\theta) - F_{D'}(\theta') \geq \frac{\lambda}{2}\|\theta - \theta'\|_2^2. \tag{13.6}$$

3. Add the two inequalities. The left side becomes $h(\theta') - h(\theta)$, where $h = F_D - F_{D'}$, and the right side becomes $\lambda\|\theta - \theta'\|_2^2$.
4. If D and D' differ in record j, then $h(\cdot) = \ell(\cdot\,; z_j) - \ell(\cdot\,; z'_j)$ is a difference of two L-Lipschitz functions, so $h(\theta') - h(\theta) \leq 2L\|\theta - \theta'\|_2$.
5. So $\lambda\|\theta - \theta'\|_2^2 \leq 2L\|\theta - \theta'\|_2$, and dividing by $\lambda\|\theta - \theta'\|_2$ gives the claim. If $\theta = \theta'$, there is nothing to prove.

□

The bound depends on neither n nor the dimension. With it, the Gaussian mechanism of the approximate DP chapter finishes the job. Take a function f whose ℓ_2 sensitivity, the largest change in Euclidean norm between neighbours, is $\Delta_2 f$. The mechanism adds Gaussian noise $\mathcal{N}(0, \sigma^2 I_d)$ with standard deviation σ in each coordinate. For every $\varepsilon > 0$ and $\delta \in (0, 1)$ it is (ε, δ)-DP, the relaxation that allows an additive slack δ, when

$$\sigma = \frac{\Delta_2 f}{\varepsilon}\sqrt{2\ln(1/\delta) + \varepsilon}. \tag{13.7}$$

Applying the mechanism to the regularized minimizer $f(D) = \theta_D$, with $\Delta_2 f = 2L/\lambda$, gives the first private algorithm.

Theorem 13.2 (Output perturbation)**.** *Under the assumptions of Theorem 13.1, let $\varepsilon > 0$ and $\delta \in (0, 1)$. The algorithm that outputs $\tilde{\theta} = \theta_D + b$, where θ_D is the regularized minimizer, with $b \sim \mathcal{N}(0, \sigma^2 I_d)$ and*

$$\sigma = \frac{2L}{\lambda} \cdot \frac{\sqrt{2\ln(1/\delta) + \varepsilon}}{\varepsilon}, \tag{13.8}$$

is (ε, δ)-DP. So is its projection onto Θ.

Proof. The map $D \mapsto \theta_D$ has ℓ_2 sensitivity at most $2L/\lambda$ by Theorem 13.1, so the Gaussian mechanism gives the guarantee. The projection is post-processing. □

Output perturbation is simple and works with any exact solver. Its accuracy is often weak, because the regularization strength that privacy needs can be far from the one that fits the data best.

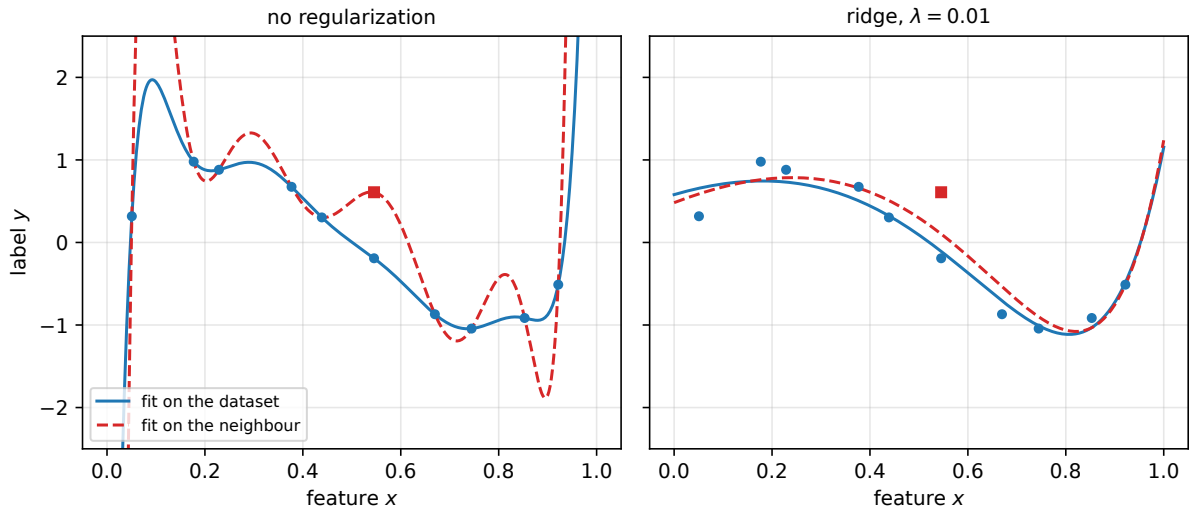


Figure 13.1.: Polynomials of degree 9 fitted by least squares to 10 points, on a dataset and on a neighbour in which one label is raised (the square). Without regularization, both fits pass through every point, and moving one point changes the curve over the whole interval. With the penalty $\frac{\lambda}{2}\|\theta\|_2^2$, the fit no longer interpolates, and the two fits stay close.

13.4. Objective perturbation

The second method adds noise to the objective instead of the output. It minimizes

$$J_b(\theta) = \mathcal{L}(\theta; D) + \frac{\lambda}{2}\|\theta\|_2^2 + \langle b, \theta \rangle \tag{13.9}$$

for a random vector b, and releases the exact minimizer. The linear term rewards or penalizes each direction of the parameter space by a random amount. With b Gaussian the result is approximately private. With b drawn from a density proportional to $e^{-c\|b\|_2}$ for a constant c, whose norm has a Gamma distribution, it is purely private.

The proof below is for losses of **generalized linear** form, which covers logistic regression and smoothed versions of the hinge loss. Such a loss is $\ell(\theta; x, y) = \phi(\langle x, \theta \rangle, y)$, where for each label y the scalar function $u \mapsto \phi(u, y)$ is convex and twice continuously differentiable with derivatives $|\phi'| \le L$ and $0 \le \phi'' \le \beta$. The features satisfy $\|x\|_2 \le 1$. Then $\nabla \ell = \phi'(\langle x, \theta \rangle, y)\, x$ has norm at most L. The Hessian, the matrix of second derivatives, is $\nabla^2 \ell = \phi''(\langle x, \theta \rangle, y)\, xx^\top$. It has rank at most one and eigenvalues in $[0, \beta]$.

The proof computes the density of the output on two neighbouring datasets and compares them. So it needs a criterion from the approximate DP chapter that is stated in terms of densities. Write p_D for the density of the output when the dataset is D. The **privacy loss** of an output θ is $\ln\left(p_D(\theta)/p_{D'}(\theta)\right)$, which is large when θ is much more likely under D than under D'. The criterion says that M is (ε, δ)-DP if the following holds for every ordered pair of neighbours. When the output is drawn from $M(D)$, its privacy loss exceeds ε with probability at most δ.

Theorem 13.3 (Objective perturbation)**.** *Let $\Theta = \mathbb{R}^d$, and let ℓ be a generalized linear loss as above, with $|\phi'| \le L$, $0 \le \phi'' \le \beta$ and features of norm at most 1. Let $\varepsilon \in (0, 4]$ and $\delta \in (0, 1)$, and set*

$$\lambda \ge \frac{\beta}{e^{\varepsilon/2} - 1}, \qquad \sigma = \frac{4L(1 + \sqrt{2\ln(1/\delta)})}{\varepsilon}. \tag{13.10}$$

Then the algorithm that draws $b \sim \mathcal{N}(0, \sigma^2 I_d)$ and outputs the exact minimizer $\tilde{\theta}$ of J_b is (ε, δ)-DP.

Proof.

1. *The output as a function of the noise.* For each b, the function J_b is continuous and λ-strongly convex on $\mathbb{R}^d$, so it has a unique minimizer, where its gradient vanishes. So $\tilde{\theta} = \theta$ exactly when $b = g_D(\theta)$, where $g_D(\theta) = -\nabla\mathcal{L}(\theta; D) - \lambda\theta$. Every θ is the minimizer for the noise $g_D(\theta)$, and each noise has one minimizer, so g_D is a bijection of $\mathbb{R}^d$ and $\tilde{\theta} = g_D^{-1}(b)$.
2. *The density of the output.* The map g_D is continuously differentiable with Jacobian $-(H_D(\theta) + \lambda I)$, where $H_D(\theta) = \sum_i \nabla^2 \ell(\theta; z_i)$ is positive semidefinite. By the change of variables formula, $\tilde{\theta}$ has density
$$p_D(\theta) = \varphi_\sigma(g_D(\theta)) \, \det\left(H_D(\theta) + \lambda I\right), \tag{13.11}$$
where φ_σ is the density of $\mathcal{N}(0, \sigma^2 I_d)$.
3. *The determinants.* Let D and D' differ in the last record, $z = (x, y)$ in D and $z' = (x', y')$ in D'. Fix θ and write $A = \sum_{i<n} \nabla^2 \ell(\theta; z_i) + \lambda I$, which satisfies $A \succeq \lambda I$. Then $H_D + \lambda I = A + a\, xx^\top$ with $a = \phi''(\langle x, \theta\rangle, y) \in [0, \beta]$. The matrix determinant lemma gives
$$\det(A + a\, xx^\top) = \det(A)\,(1 + a\, x^\top A^{-1} x) \in \left[\det A,\ (1 + \beta/\lambda)\det A\right], \tag{13.12}$$
because $0 \le x^\top A^{-1} x \le \|x\|_2^2/\lambda \le 1/\lambda$. The same holds for D'. So the ratio of the two determinants is at most $1 + \beta/\lambda \le e^{\varepsilon/2}$, by the choice of λ.
4. *The Gaussian factors.* The two noise vectors that lead to θ differ by $u = g_{D'}(\theta) - g_D(\theta) = \nabla\ell(\theta; z) - \nabla\ell(\theta; z')$. This vector has norm at most $2L$, and it lies in the subspace V spanned by x and x', which does not depend on θ. With $b = g_D(\theta)$ and P_V the orthogonal projection onto V,
$$\ln \frac{\varphi_\sigma(b)}{\varphi_\sigma(b + u)} = \frac{2\langle b, u\rangle + \|u\|_2^2}{2\sigma^2} \le \frac{2L\,\|P_V b\|_2}{\sigma^2} + \frac{2L^2}{\sigma^2}, \tag{13.13}$$
by the Cauchy–Schwarz inequality applied to $\langle b, u\rangle = \langle P_V b, u\rangle$.
5. *The privacy loss.* When the output is drawn from $M(D)$, the vector $g_D(\tilde{\theta})$ is the noise b itself. Steps 2 to 4 bound the privacy loss at the output by
$$\ln \frac{p_D(\tilde{\theta})}{p_{D'}(\tilde{\theta})} \le \frac{\varepsilon}{2} + \frac{2L\,\|P_V b\|_2}{\sigma^2} + \frac{2L^2}{\sigma^2}. \tag{13.14}$$
6. *The tail.* The subspace V has dimension $d_V \le 2$, so $\|P_V b\|_2^2/\sigma^2$ is chi-squared with d_V degrees of freedom. Its tail is at most that of two degrees of freedom, which is $e^{-r/2}$ beyond r. Write $s = \sqrt{2\ln(1/\delta)}$ for the cut-off. Then
$$\Pr\left[\|P_V b\|_2 > \sigma s\right] \le e^{-s^2/2} = \delta. \tag{13.15}$$

7. Outside that event, substituting $\sigma = 4L(1+s)/\varepsilon$ into step 5 gives

$$\ln \frac{p_D(\tilde{\theta})}{p_{D'}(\tilde{\theta})} \le \frac{\varepsilon}{2} + \frac{\varepsilon}{2}\left(\frac{s}{1+s} + \frac{\varepsilon}{4(1+s)^2}\right) \le \frac{\varepsilon}{2} + \frac{\varepsilon}{2}\cdot\frac{s+1}{1+s} = \varepsilon, \tag{13.16}$$

   where the second step uses $\varepsilon \le 4$ and $1/(1+s)^2 \le 1/(1+s)$.
8. The privacy loss criterion turns this bound into (ε, δ)-DP.

□

The constants come from this proof, and the original analyses obtain similar ones by the same change of variables. The theorem has a practical weakness: it holds only for the exact minimizer. A real solver never returns one. Floating-point arithmetic cannot represent most real numbers, and Mironov showed that even a naive floating-point implementation of the Laplace mechanism leaks the data through the low-order bits of its output. Iterative solvers such as gradient descent only approach the minimizer. Later versions of objective perturbation need only an approximate minimizer and perform well in practice for convex problems, and a version for non-convex losses keeps the privacy guarantee, though not the accuracy one.

13.5. Gradient perturbation

The third method, now the standard one, runs gradient descent and adds noise to each gradient. Its analysis starts from the non-private algorithm. Write $\Pi_\Theta(u) = \arg\min_{\theta\in\Theta} \|\theta - u\|_2$ for the **projection** onto Θ, the closest point of Θ, which exists and is unique because Θ is closed and convex.

The definition leaves the step direction G_t random. One analysis then covers exact gradients, gradients of one sampled record, and the noisy gradients of the private algorithm.

Definition 13.4 (Projected stochastic gradient descent)**.** Fix a start $\theta_0 \in \Theta$, a number of steps T and a learning rate $\eta > 0$, also called the step size. For $t = 1, \dots, T$, draw a random vector G_t and set

$$\theta_t = \Pi_\Theta(\theta_{t-1} - \eta\, G_t). \tag{13.17}$$

Output the average $\bar{\theta} = \frac{1}{T}\sum_{t=0}^{T-1} \theta_t$. With $G_t = \nabla\mathcal{L}(\theta_{t-1}; D)$ this is projected **gradient descent**. With $G_t = n\,\nabla\ell(\theta_{t-1}; z_{I_t})$, for an index I_t drawn uniformly from $\{1, \dots, n\}$, it is projected stochastic gradient descent (SGD).

A step moves against the gradient, the direction of steepest increase, and the projection brings a step that leaves Θ back to the nearest allowed point. Projection never moves a point farther from any point of Θ, and the convergence proof uses this to compare each iterate with the minimizer.

Lemma 13.3. *For every $u \in \mathbb{R}^d$ and every $\theta \in \Theta$, $\|\Pi_\Theta(u) - \theta\|_2 \le \|u - \theta\|_2$.*

Proof.

1. Let $v = \Pi_\Theta(u)$. For $t \in (0,1]$ the point $v + t(\theta - v)$ is in Θ, so it is no closer to u than v: $\|u - v - t(\theta - v)\|^2 \geq \|u - v\|^2$. Expanding, dividing by t and letting $t \to 0$ gives $\langle u - v, \theta - v \rangle \leq 0$.
2. Then $\|u - \theta\|^2 = \|u - v\|^2 + \|v - \theta\|^2 + 2\langle u - v, v - \theta \rangle \geq \|v - \theta\|^2$, because the last term is nonnegative by step 1.

□

The convergence theorem asks only that each G_t be right on average and not too large.

Theorem 13.4 (Convergence of projected SGD)**.** *Let F be convex on $\mathbb{R}^d$, and let θ^* minimize F over Θ. Suppose that, given $\theta_0, \dots, \theta_{t-1}$, the conditional mean $\mathbb{E}[G_t \mid \theta_0, \dots, \theta_{t-1}]$ is a subgradient of F at θ_{t-1}, and that $\mathbb{E}\|G_t\|_2^2 \leq G^2$ for every t. With learning rate $\eta = \|\Theta\|_2/(G\sqrt{T})$,*

$$\mathbb{E}[F(\bar{\theta})] - F(\theta^*) \leq \frac{\|\Theta\|_2 \, G}{\sqrt{T}}. \tag{13.18}$$

Proof.

1. By Lemma 13.3 with $\theta = \theta^*$,
$$\|\theta_t - \theta^*\|^2 \leq \|\theta_{t-1} - \eta G_t - \theta^*\|^2 = \|\theta_{t-1} - \theta^*\|^2 - 2\eta\langle G_t, \theta_{t-1} - \theta^* \rangle + \eta^2 \|G_t\|^2. \tag{13.19}$$
2. Condition on the past and let g be the conditional mean of G_t, a subgradient at θ_{t-1}. The definition of a subgradient gives $\langle g, \theta_{t-1} - \theta^* \rangle \geq F(\theta_{t-1}) - F(\theta^*)$. Taking expectations in step 1,
$$\mathbb{E}[F(\theta_{t-1})] - F(\theta^*) \leq \frac{\mathbb{E}\|\theta_{t-1} - \theta^*\|^2 - \mathbb{E}\|\theta_t - \theta^*\|^2}{2\eta} + \frac{\eta G^2}{2}. \tag{13.20}$$
3. Summing over $t = 1, \dots, T$, the distances telescope. The first is at most $\|\Theta\|_2^2$ and the last is nonnegative, so
$$\sum_{t=0}^{T-1} \left(\mathbb{E}[F(\theta_t)] - F(\theta^*)\right) \leq \frac{\|\Theta\|_2^2}{2\eta} + \frac{T\eta G^2}{2}. \tag{13.21}$$
4. The average $\bar{\theta}$ lies in Θ, and convexity gives $F(\bar{\theta}) \leq \frac{1}{T}\sum_t F(\theta_t)$. Dividing step 3 by T and substituting η makes both terms $\|\Theta\|_2 G/(2\sqrt{T})$.

□

The theorem is stated for the average of the iterates. Shamir and Zhang proved that the last iterate θ_T does almost as well, with learning rates $\eta_t = \|\Theta\|_2/(G\sqrt{t})$: its excess risk is at most $2\|\Theta\|_2 G(2 + \ln T)/\sqrt{T}$.

Applied to the empirical risk, the theorem says how many steps each method needs. The bound G on the step comes from the bound L on every gradient.

Corollary 13.1. *Suppose, as throughout the chapter, that Θ has finite diameter $\|\Theta\|_2$ and that every loss $\ell(\cdot\,; z)$ is convex and L-Lipschitz on $\mathbb{R}^d$. Then gradient descent and SGD both reach excess empirical risk at most α after $T = \lceil n^2 L^2 \|\Theta\|_2^2/\alpha^2 \rceil$ steps. Gradient descent evaluates nT gradients of ℓ, and SGD evaluates T.*

Proof.

1. For gradient descent, $G_t = \sum_i \nabla \ell(\theta_{t-1}; z_i)$ is a subgradient of $\mathcal{L}$, and by Lemma 13.1 its norm is at most nL.
2. For SGD, the conditional mean of $G_t = n\nabla \ell(\theta_{t-1}; z_{I_t})$ is $\frac{1}{n}\sum_i n\nabla \ell(\theta_{t-1}; z_i)$, the same subgradient. Its second moment is $\frac{1}{n}\sum_i n^2 \|\nabla \ell(\theta_{t-1}; z_i)\|^2 \le n^2 L^2$.
3. In both cases $G = nL$, and Theorem 13.4 gives excess risk at most $\|\Theta\|_2 nL/\sqrt{T} \le \alpha$.

□

For a constant α, gradient descent costs about n^3 gradient evaluations and SGD about n^2. SGD is also the easier one to make private, because each step touches one record, and the noise has only one natural place to go. Private SGD adds Gaussian noise to each sampled gradient. Its noise level σ is the one that the privacy proof below needs, and the learning rate accounts for the extra noise.

Definition 13.5 (Private SGD)**.** Let $\varepsilon \in (0, 1]$ and $\delta \in (0, 1)$. Run projected SGD for $T = n^2$ steps with

$$G_t = n\,\nabla \ell(\theta_{t-1}; z_{I_t}) + b_t, \qquad b_t \sim \mathcal{N}(0, \sigma^2 I_d), \qquad \sigma = \frac{16\, nL\sqrt{\ln(3.3n/\delta)\,\ln(2/\delta)}}{\varepsilon}, \tag{13.22}$$

with fresh, independent indices I_t and noise b_t, and learning rate $\eta = \|\Theta\|_2/(G\sqrt{T})$ where $G^2 = n^2L^2 + d\sigma^2$.

Accuracy is the easier half, because the noise has mean zero and only enlarges the second moment.

Theorem 13.5. *Under the assumptions of the chapter, with $\ell(\cdot\,; z)$ convex and L-Lipschitz on $\mathbb{R}^d$, private SGD satisfies*

$$\mathbb{E}[\mathcal{L}(\bar{\theta}; D)] - \mathcal{L}(\theta^*; D) \le \|\Theta\|_2 \sqrt{L^2 + \frac{d\sigma^2}{n^2}} \le \|\Theta\|_2\, L \left(1 + \frac{16\sqrt{d\,\ln(3.3n/\delta)\ln(2/\delta)}}{\varepsilon}\right). \tag{13.23}$$

Proof.

1. Given the past, b_t has mean zero and is independent of I_t, so the conditional mean of G_t is the subgradient of step 2 of Corollary 13.1.
2. Expanding the square,

$$\mathbb{E}\|G_t\|^2 = n^2\,\mathbb{E}\|\nabla \ell(\theta_{t-1}; z_{I_t})\|^2 + 2n\,\mathbb{E}\langle \nabla \ell(\theta_{t-1}; z_{I_t}), b_t\rangle + \mathbb{E}\|b_t\|^2 \le n^2L^2 + 0 + d\sigma^2. \tag{13.24}$$

   The first term is bounded by Lemma 13.1. The middle term vanishes because b_t is independent of the gradient and has mean zero. The last is a sum of d squared Gaussians of variance σ^2.

3. Theorem 13.4 with $T = n^2$ gives $\|\Theta\|_2 G/n$, which is the first bound. The second uses $\sqrt{a+b} \le \sqrt{a} + \sqrt{b}$ and the value of σ.

□

Up to logarithmic factors, the excess risk is $O(\|\Theta\|_2 L\sqrt{d\ln(1/\delta)}/\varepsilon)$, so the average excess risk per record falls as $1/n$. The price of privacy is the factor $\sqrt{d}$, and Bassily, Smith and Thakurta proved that some dependence of this kind is unavoidable. The constant 16, however, comes from a loose chain of bounds. For example, with $n = 10{,}000$, $\varepsilon = 1$ and $\delta = 10^{-6}$, the noise is $\sigma = 299.9\,nL$. With $d = 100$, the bound per record is $\|\Theta\|_2 L(1 + 2999)/10{,}000 = 0.30\,\|\Theta\|_2 L$. The next chapter replaces this chain with a much tighter one.

Privacy combines three tools. The first is the Gaussian mechanism, restated above. The second is **amplification by subsampling**: a mechanism run on one record picked at random is more private than the same mechanism run on a fixed record.

Lemma 13.4 (Amplification by sampling one record)**.** *Let M_0 be a mechanism whose input is a single record and which satisfies $\Pr[M_0(z) \in S] \le e^{\varepsilon_0} \Pr[M_0(z') \in S] + \delta_0$ for all records z, z' and all events S. Let $M(D) = M_0(z_I)$ with I uniform on $\{1, \dots, n\}$. Then M is $(\varepsilon_1, \delta_0/n)$-DP with*

$$\varepsilon_1 = \ln\left(1 + \frac{e^{\varepsilon_0} - 1}{n}\right) \le \frac{e^{\varepsilon_0} - 1}{n}. \tag{13.25}$$

Proof.

1. Let D and D' differ in record j, fix an event S, and write $p_i = \Pr[M_0(z_i) \in S]$ and $p'_j = \Pr[M_0(z'_j) \in S]$. Let r be the average of p_i over $i \ne j$, and write $q = 1/n$ for the sampling rate. Then

$$P = \Pr[M(D) \in S] = (1-q)r + qp_j, \qquad P' = \Pr[M(D') \in S] = (1-q)r + qp'_j. \tag{13.26}$$

2. The hypothesis, applied to z_j against z'_j and against each z_i and then averaged, gives two bounds:

$$p_j - p'_j \le (e^{\varepsilon_0} - 1)p'_j + \delta_0, \qquad p_j - p'_j \le e^{\varepsilon_0} r - p'_j + \delta_0. \tag{13.27}$$

3. Let $c = e^{\varepsilon_0} - 1$ and take the convex combination of the two bounds with weights w and $1 - w$, where $1 - w = c(1-q)/e^{\varepsilon_0} \in [0, 1)$. The coefficient of r becomes $c(1-q)$, and that of p'_j becomes $we^{\varepsilon_0} - 1 = cq$. So

$$p_j - p'_j \le c((1-q)r + qp'_j) + \delta_0 = cP' + \delta_0. \tag{13.28}$$

4. Then $P = P' + q(p_j - p'_j) \le (1 + qc)P' + q\delta_0$, which is the claim. The inequality $\ln(1+x) \le x$ gives the simpler bound on ε_1.

□

For $\varepsilon_0 \le 1$, convexity of the exponential gives $e^{\varepsilon_0} - 1 \le (e-1)\varepsilon_0 \le 2\varepsilon_0$, so sampling one record of n cuts the privacy parameter by a factor of about $n/2$.

The third tool is advanced composition, from the advanced composition chapter. It says that running k mechanisms in sequence, each $(\varepsilon_0, \delta_0)$-DP and each chosen adaptively from the outputs of the previous ones, is $(\varepsilon, k\delta_0 + \delta')$-DP for every $\delta' \in (0,1)$, where

$$\varepsilon = \varepsilon_0 \sqrt{2k \ln(1/\delta')} + k\varepsilon_0 \frac{e^{\varepsilon_0} - 1}{e^{\varepsilon_0} + 1} \le \varepsilon_0 \sqrt{2k \ln(1/\delta')} + \frac{k\varepsilon_0^2}{2}, \tag{13.29}$$

and the inequality uses $(e^x - 1)/(e^x + 1) = \tanh(x/2) \le x/2$ for $x \ge 0$.

The three tools now give the privacy of private SGD, one after another. The Gaussian noise protects the sampled record in one step, sampling one record of n makes the step more private, and advanced composition adds up the steps.

Theorem 13.6 (Privacy of private SGD)**.** *For $\varepsilon \in (0,1]$ and $\delta \in (0,1)$, private SGD is (ε, δ)-DP, even if it releases every iterate $\theta_1, \dots, \theta_T$.*

Proof.

1. *One step, fixed index.* Fix the previous iterate and the index i. The step releases a post-processing of $n\nabla\ell(\theta_{t-1}; z_i) + b_t$. As a function of the single record z_i, the vector $n\nabla\ell$ has ℓ_2 sensitivity at most $2nL$, by Lemma 13.1. Set $\varepsilon_0 = \varepsilon/(4\sqrt{2\ln(2/\delta)})$, which is below 1, and $\delta_0 = \delta/(2n)$. The Gaussian mechanism is $(\varepsilon_0, \delta_0)$-DP with respect to z_i at the noise level

$$\sigma_0 = \frac{2nL}{\varepsilon_0} \sqrt{2\ln(1/\delta_0) + \varepsilon_0} \le \frac{2nL}{\varepsilon_0} \sqrt{2 \ln \frac{2\sqrt{e}\, n}{\delta}} \le \frac{2nL}{\varepsilon_0} \sqrt{2 \ln \frac{3.3n}{\delta}} = \sigma, \tag{13.30}$$

   where the first inequality uses $\varepsilon_0 \le 1 = 2\ln\sqrt{e}$, and the last equality is the σ of Definition 13.5. Noise of variance $\sigma^2 \ge \sigma_0^2$ is noise of variance σ_0^2 plus an independent Gaussian of variance $\sigma^2 - \sigma_0^2$, a post-processing, so the step is $(\varepsilon_0, \delta_0)$-DP with respect to z_i.
2. *One step, random index.* By Lemma 13.4, each step is $(\varepsilon_1, \delta_0/n)$-DP with respect to the dataset, where $\varepsilon_1 \le (e^{\varepsilon_0} - 1)/n \le 2\varepsilon_0/n$.
3. *All steps.* Each step is chosen adaptively from the previous iterates. Advanced composition over $T = n^2$ steps with $\delta' = \delta/2$ gives total failure probability $n^2 \cdot \delta/(2n^2) + \delta/2 = \delta$, and privacy parameter

$$\varepsilon_1 \sqrt{2n^2 \ln(2/\delta)} + \frac{n^2 \varepsilon_1^2}{2} \le 2\varepsilon_0 \sqrt{2\ln(2/\delta)} + 2\varepsilon_0^2 = \frac{\varepsilon}{2} + 2\varepsilon_0^2. \tag{13.31}$$

4. Finally, the second term is small because $\varepsilon \le 1$ and $\ln(2/\delta) > \ln 2 > 1/8$:

$$2\varepsilon_0^2 = \frac{\varepsilon^2}{16 \ln(2/\delta)} \le \frac{\varepsilon}{2}. \tag{13.32}$$

   The average $\bar{\theta}$ is a post-processing of the iterates.

□

The privacy cost per step is about ε/n, and composition over n^2 steps multiplies it by only $\sqrt{n^2} = n$. Every piece of the argument is generic: noise in the right place, amplification, composition and post-processing. That modularity is why the same recipe extends to models for which no convergence theorem is available.

13.6. Summary

Table 13.1.: The chapter's results.

Result	Statement
Proposition 13.1	the exact median is not (ε, δ)-DP for any $\delta < 1$
Theorem 13.1	the minimizer of loss plus $\lambda\cdot$(1-strongly convex) has sensitivity $2L/\lambda$
Theorem 13.2	adding Gaussian noise to that minimizer is (ε, δ)-DP
Theorem 13.3	the exact minimizer of a Gaussian-perturbed objective is (ε, δ)-DP
Theorem 13.4	projected SGD has excess risk $\|\Theta\|_2 G/\sqrt{T}$
Lemma 13.4	sampling one record of n turns ε_0 into $\ln(1 + (e^{\varepsilon_0} - 1)/n)$
Theorem 13.6	private SGD with n^2 steps is (ε, δ)-DP
Theorem 13.5	its excess risk is $O(\|\Theta\|_2 L\sqrt{d \ln(1/\delta)}/\varepsilon)$ up to logarithmic factors

13.7. References

- Kamalika Chaudhuri, Claire Monteleoni and Anand Sarwate, "Differentially Private Empirical Risk Minimization", 2011. Output and objective perturbation for regularized convex losses, with the sensitivity bound of Theorem 13.1 for averaged losses.
- Benjamin Rubinstein, Peter Bartlett, Ling Huang and Nina Taft, "Learning in a Large Function Space: Privacy-Preserving Mechanisms for SVM Learning", 2012. Output perturbation for support vector machines.
- Daniel Kifer, Adam Smith and Abhradeep Thakurta, "Private Convex Empirical Risk Minimization and High-Dimensional Regression", 2012. Objective perturbation with Gaussian noise and the change-of-variables analysis used in Theorem 13.3.
- Ilya Mironov, "On Significance of the Least Significant Bits for Differential Privacy", 2012. The floating-point attack on the Laplace mechanism.
- Roger Iyengar, Joseph Near, Dawn Song, Om Thakkar, Abhradeep Thakurta and Lun Wang, "Towards Practical Differentially Private Convex Optimization", 2019. Objective perturbation with approximate minimizers.

- Seth Neel, Aaron Roth, Giuseppe Vietri and Zhiwei Steven Wu, "Oracle Efficient Private Non-Convex Optimization", 2020. Privacy of objective perturbation without convexity.
- Oliver Williams and Frank McSherry, "Probabilistic Inference and Differential Privacy", 2010, and Shuang Song, Kamalika Chaudhuri and Anand Sarwate, "Stochastic Gradient Descent with Differentially Private Updates", 2013. Early versions of noisy gradient descent.
- Raef Bassily, Adam Smith and Abhradeep Thakurta, "Private Empirical Risk Minimization: Efficient Algorithms and Tight Error Bounds", 2014. Private SGD with the privacy argument of this chapter (Gaussian noise, amplification and advanced composition), and matching lower bounds. Their constants differ from the ones derived here.
- Ohad Shamir and Tong Zhang, "Stochastic Gradient Descent for Non-Smooth Optimization: Convergence Results and Optimal Averaging Schemes", 2013. Theorem 2 is the bound on the last iterate.
- Borja Balle, Gilles Barthe and Marco Gaboardi, "Privacy Amplification by Subsampling: Tight Analyses via Couplings and Divergences", 2018. Tight amplification bounds for several sampling schemes, including Lemma 13.4.

14. Private Deep Learning

The private stochastic gradient descent (SGD) of the previous chapter needs a convex, Lipschitz loss, and its privacy analysis loses large constants. Neural networks have neither property, and in practice the constants decide whether a guarantee is meaningful. This chapter adapts noisy SGD to neural networks under differential privacy (DP) by clipping each gradient. It then replaces the composition argument with a tighter accounting built on the moments of the privacy loss. Finally, it presents a second approach that turns any non-private learner into a private one by a vote among models trained on disjoint data.

14.1. Neural networks

The setting is the one of the previous chapter. Each record is a pair of a feature vector and a label, and training minimizes the total loss of the records over a parameter vector θ,

$$\mathcal{L}(\theta; D) = \sum_{i=1}^{n} \ell(\theta; x_i, y_i). \tag{14.1}$$

Here a feature vector x lies in $\mathbb{R}^{m_0}$, and a label is one of K classes, $y \in \{1, \dots, K\}$. A neural network maps a feature vector to one score per class, and its parameters θ are what training adjusts. The simplest kind of network is a stack of layers, each a linear map followed by a fixed nonlinear function.

Definition 14.1 (Multilayer perceptron)**.** A network with J layers sets $v^{(0)} = x$ and, for $j = 1, \dots, J$,

$$v^{(j)} = \phi\big(W^{(j)} v^{(j-1)} + c^{(j)}\big), \tag{14.2}$$

where $W^{(j)}$ is an $m_j \times m_{j-1}$ matrix of weights, $c^{(j)} \in \mathbb{R}^{m_j}$ is a vector of offsets, and the **activation function** $\phi \colon \mathbb{R} \to \mathbb{R}$ is applied to each coordinate. The last layer has no activation and outputs $o \in \mathbb{R}^K$, one score per class. The parameter vector θ collects all weights and offsets.

Without the activation the network would be a composition of linear maps, and so a single linear map. Common activations are the rectified linear unit $\phi(u) = \max(0, u)$, the sigmoid $\phi(u) = 1/(1 + e^{-u})$ and $\tanh(u) = (e^u - e^{-u})/(e^u + e^{-u})$. The rectified linear unit is the usual choice without privacy.

The scores become a probability vector over the classes through the **softmax** function, $\hat{y}_k = e^{o_k} / \sum_{j=1}^{K} e^{o_j}$. This is a smooth stand-in for picking the largest score, and $\hat{y}$ reads as the

network's confidence in each class. The prediction is the class with the largest $\hat{y}_k$. The loss is the **cross-entropy**, the negative log of the probability given to the true class,

$$\ell(\theta; x, y) = -\ln \hat{y}_y = -\sum_{k=1}^{K} \mathbf{1}[k = y] \ln \hat{y}_k, \tag{14.3}$$

where $\hat{y}$ depends on θ through every layer. The analysis of the previous chapter does not apply to this loss. It is not convex in θ, so no convergence theorem holds, although SGD works well in practice. It is not Lipschitz either. In a two-layer network, the gradient with respect to the first layer's weights is multiplied by the second layer's weights, which can be arbitrarily large. Without Lipschitz continuity, a single gradient has unbounded sensitivity.

14.2. Differentially private SGD

Two changes make noisy SGD practical. Each gradient is **clipped** to a fixed norm, which bounds sensitivity directly. And each step uses a random batch of records, called a **lot**, chosen by Poisson sampling. Clipping comes first. It shrinks a vector that is too long and leaves a short one alone.

Definition 14.2 (Clipping)**.** For a clipping norm C, a positive number, the clipping of a vector g is

$$\text{clip}_C(g) = g \cdot \min\left(1, \frac{C}{\|g\|_2}\right). \tag{14.4}$$

A vector inside the ball of radius C is unchanged, and a longer one is scaled back onto the sphere. So every clipped vector has norm at most C.

DP-SGD combines the two changes. Each step clips the gradient of every record in the lot, adds Gaussian noise to the sum, and divides by the expected lot size qn. The noise is set by the **noise multiplier** σ, the ratio of its standard deviation to the clipping norm. Unlike the σ of the previous chapter, it is not the standard deviation itself.

Definition 14.3 (DP-SGD)**.** Let the parameter vector θ have d coordinates. Fix a sampling rate q in $(0, 1]$, a noise multiplier $\sigma > 0$, a clipping norm C, a number of steps T, learning rates η_t and a start θ_0. For $t = 1, \dots, T$:

1. Form the lot B_t by including each index $i \in \{1, \dots, n\}$ independently with probability q.
2. Compute the noisy average of clipped gradients,

$$\tilde{g}_t = \frac{1}{qn}\left(\sum_{i \in B_t} \text{clip}_C(\nabla \ell(\theta_{t-1}; z_i)) + b_t\right), \qquad b_t \sim \mathcal{N}(0, \sigma^2 C^2 I_d). \tag{14.5}$$

3. Step: $\theta_t = \theta_{t-1} - \eta_t \, \tilde{g}_t$.

Output θ_T.

Compared with the private SGD of the previous chapter, a step uses a lot of expected size qn instead of one record, and the gradient is clipped instead of bounded by assumption. The noise is calibrated to C, the most one record can contribute to the sum.

Lemma 14.1. *Adding or removing one record changes the clipped sum* $\sum_{i \in B} \operatorname{clip}_C(\nabla \ell(\theta; z_i))$ *by a vector of norm at most* C*. Replacing one record changes it by at most* $2C$*.*

Proof. Adding or removing a record adds or removes one clipped vector, of norm at most C. A replacement does both, and the triangle inequality gives $2C$. □

Clipping has a cost. The clipped average is a biased estimate of the gradient, so even the convergence theorem for convex losses would no longer apply. For neural networks no such theorem was available anyway, and clipped SGD works well in practice.

The privacy analysis depends on how the lot is drawn, and schemes that look alike differ for privacy. **Poisson sampling** includes each record independently with probability q, so the lot size is random. Even at $q = 1/n$ this is not the same as picking one record, because the lot may be empty or hold several records. **Sampling without replacement** draws a lot of fixed size m uniformly among subsets. Picking one record, as in the previous chapter, is the case $m = 1$. **Sampling with replacement** draws m records independently, so a record can appear more than once. The analyses below assume Poisson sampling, and they use the add/remove relation, under which neighbours $D \sim D'$ differ by one added or removed record and Lemma 14.1 gives sensitivity C. In practice, implementations often shuffle the data once and cut it into consecutive batches of fixed size. Every batch is then a uniform sample, but the batches are dependent, so the Poisson analysis does not apply to them. Reporting it for shuffled batches is a common error.

14.3. Privacy accounting

DP-SGD can be analysed exactly as in the previous chapter. Without sampling, a single step is the Gaussian mechanism with noise multiplier σ. Suppose it is $(\varepsilon_0, \delta_0)$-DP. Amplification by Poisson sampling then makes it about $(2q\varepsilon_0, q\delta_0)$-DP, and advanced composition over T steps gives

$$\Big(O(q\varepsilon_0 \sqrt{T \ln(1/\delta)}),\ Tq\delta_0 + \delta\Big)\text{-DP}. \tag{14.6}$$

Abadi and coauthors proved a better bound for the same algorithm, which needs less noise for the same guarantee.

Theorem 14.1 (Moments accountant, asymptotic form)**.** *There are constants* c_1 *and* c_2 *such that, for every sampling rate* q *and number of steps* T*, every* $\varepsilon < c_1 q^2 T$ *and every* $\delta > 0$*, DP-SGD is* (ε, δ)*-DP under add/remove neighbours if*

$$\sigma \geq c_2 \, \frac{q \sqrt{T \ln(1/\delta)}}{\varepsilon}. \tag{14.7}$$

The proof is Theorem 1 of Abadi and coauthors, 2016. In the composition bound, the Gaussian mechanism of the approximate DP chapter gives ε_0 with $\sigma = \sqrt{2\ln(1/\delta_0) + \varepsilon_0}/\varepsilon_0$, so ε_0 is about $\sqrt{2\ln(1/\delta_0)}/\sigma$, and the composition route needs σ larger by a factor of order $\sqrt{\ln(1/\delta_0)}$, and it pays $Tq\delta_0$ in failure probability. These look like small gains, but constant and logarithmic factors in ε separate a meaningful guarantee from an empty one. The rest of this section proves the ingredients of the accountant and shows the gain numerically.

The starting point is the privacy loss. For a mechanism M whose output has density p_D on dataset D, the **privacy loss** at an output y is $\ln\left(p_D(y)/p_{D'}(y)\right)$. When y is drawn from $M(D)$, it is a random variable. Pure ε-DP says that it never exceeds ε. By the privacy loss criterion of the approximate DP chapter, (ε, δ)-DP follows if, for every ordered pair of neighbours, it exceeds ε with probability at most δ. That criterion only counts the mass above ε, and says nothing about how far above it lies. Moments describe the whole distribution, and unlike tail bounds, they compose by simple addition. The accountant uses the moments of the likelihood ratio $p_D(y)/p_{D'}(y)$. For an order $\lambda > 0$, it takes the average of the ratio raised to the power λ, and then the logarithm of that average. The worst pair of neighbours gives the log-moment of M.

Definition 14.4 (Log-moments of the privacy loss)**.** Let p_D be the density of the output $M(D)$. For $\lambda > 0$, the λ-th log-moment of M is the logarithm of the average λ-th power of the density ratio, for the worst pair of neighbours $D \sim D'$,

$$\kappa_M(\lambda) = \sup_{D \sim D'} \ln \, \mathbb{E}_{y \sim M(D)} \left[\left(\frac{p_D(y)}{p_{D'}(y)} \right)^{\lambda} \right]. \tag{14.8}$$

Abadi and coauthors write it $\alpha_M(\lambda)$. Up to a factor, it is the Rényi divergence of order $\alpha = \lambda + 1$ between the two output distributions, and a bound on it for every pair of neighbours is Rényi differential privacy (RDP). A bound on one log-moment already implies approximate privacy, by Markov's inequality.

Theorem 14.2 (From moments to approximate DP)**.** *For every $\lambda > 0$ and $\varepsilon \geq 0$, the mechanism M is (ε, δ)-DP with*

$$\delta = \exp\left(\kappa_M(\lambda) - \lambda\varepsilon\right). \tag{14.9}$$

Proof.

1. Fix neighbours $D \sim D'$ and write $L = \ln(p_D(Y)/p_{D'}(Y))$ for the privacy loss random variable, with $Y \sim M(D)$. The map $u \mapsto e^{\lambda u}$ is increasing, so $\Pr[L > \varepsilon] \leq \Pr[e^{\lambda L} \geq e^{\lambda\varepsilon}]$.
2. Markov's inequality bounds the right side by $\mathbb{E}[e^{\lambda L}]\, e^{-\lambda\varepsilon}$, and $\mathbb{E}[e^{\lambda L}] \leq e^{\kappa_M(\lambda)}$ by definition.
3. So the privacy loss exceeds ε with probability at most δ, and the privacy loss criterion gives (ε, δ)-DP.

□

The conversion is useful because log-moments are easy to track over many steps. The next lemma shows that they add up when mechanisms run one after another, even when each mechanism is chosen from the outputs of the previous ones.

Lemma 14.2 (Composition of moments)**.** *Let M_1 have output y_1, and let $M_2(\cdot\,; y_1)$ be a mechanism chosen as a function of y_1 with $\kappa_{M_2(\cdot;y_1)}(\lambda) \le \kappa_2$ for every y_1. The mechanism M that outputs (y_1, y_2) satisfies $\kappa_M(\lambda) \le \kappa_{M_1}(\lambda) + \kappa_2$. By induction, T adaptive steps have log-moments at most the sum of theirs.*

Proof.

1. The density of (y_1, y_2) on D is $p_D(y_1)\, p_D(y_2 \mid y_1)$, so the ratio for neighbours factors as $r_1(y_1)\, r_2(y_2 \mid y_1)$. Here r_1 is the density ratio of y_1, and r_2 is the density ratio of y_2 given y_1.
2. Taking the expectation over y_2 first, given y_1, gives

$$\mathbb{E}[r_1^\lambda r_2^\lambda] = \mathbb{E}_{y_1}\Big[r_1(y_1)^\lambda\, \mathbb{E}_{y_2}[r_2(y_2 \mid y_1)^\lambda \,\big|\, y_1]\Big] \le e^{\kappa_2}\, \mathbb{E}_{y_1}[r_1(y_1)^\lambda] \le e^{\kappa_2 + \kappa_{M_1}(\lambda)}. \tag{14.10}$$

3. Taking logarithms and the supremum over neighbours gives the claim.

□

For the Gaussian mechanism the log-moments have a closed form.

Lemma 14.3. *Let $M(D) = f(D) + b$ with $b \sim \mathcal{N}(0, \sigma^2 I_d)$, where f has ℓ_2 sensitivity $\Delta_2 f$. Then for every $\lambda > 0$,*

$$\kappa_M(\lambda) = \frac{\lambda(\lambda+1)(\Delta_2 f)^2}{2\sigma^2}. \tag{14.11}$$

Proof.

1. Fix neighbours, write $\mu = f(D)$ and $\mu' = f(D')$, and let $r = \|\mu - \mu'\|_2 \le \Delta_2 f$. For $y = \mu + \sigma\xi$ with ξ standard normal, the privacy loss is

$$\frac{\|y-\mu'\|^2 - \|y-\mu\|^2}{2\sigma^2} = \frac{r^2}{2\sigma^2} + \frac{\langle \xi, \mu - \mu' \rangle}{\sigma}. \tag{14.12}$$

2. So the privacy loss is normal with mean $m = r^2/(2\sigma^2)$ and variance $r^2/\sigma^2 = 2m$.
3. A normal variable Z with mean m and variance s^2 has $\mathbb{E}[e^{\lambda Z}] = e^{\lambda m + \lambda^2 s^2/2}$. Here the exponent is $\lambda m + \lambda^2 m = \lambda(\lambda+1) r^2/(2\sigma^2)$.
4. This increases with r, and the supremum is reached at $r = \Delta_2 f$.

□

The three results combine into a guarantee for many Gaussian steps without sampling. The log-moments of the steps add up, and the conversion turns their sum into (ε, δ)-DP for the best order λ.

Corollary 14.1. *Let $\rho = (\Delta_2 f)^2/(2\sigma^2)$. Any adaptive sequence of T Gaussian mechanisms, each with sensitivity $\Delta_2 f$ and noise $\mathcal{N}(0, \sigma^2 I_d)$, is (ε, δ)-DP for every $\delta \in (0,1)$, with*

$$\varepsilon = T\rho + 2\sqrt{T\rho \ln(1/\delta)}. \tag{14.13}$$

Proof.

1. By Lemma 14.3 and Lemma 14.2, the sequence has $\kappa(\lambda) \le \lambda(\lambda+1)T\rho$.
2. By Theorem 14.2, it is (ε,δ)-DP whenever $\lambda(\lambda+1)T\rho - \lambda\varepsilon \le \ln\delta$, that is, whenever $\varepsilon \ge (\lambda+1)T\rho + \ln(1/\delta)/\lambda$.
3. The right side is smallest at $\lambda = \sqrt{\ln(1/\delta)/(T\rho)}$, where it equals $T\rho + 2\sqrt{T\rho\ln(1/\delta)}$.

□

For comparison, take the same T Gaussian steps with $\Delta_2 f = 1$. Advanced composition needs a per-step guarantee first, about $\varepsilon_0 = \sqrt{2\ln(1/\delta_0)}/\sigma$ by the Gaussian mechanism, and then gives $\varepsilon \approx \varepsilon_0\sqrt{2T\ln(1/\delta)}$, which is $2\sqrt{T\ln(1/\delta_0)\ln(1/\delta)}/\sigma$. The corollary gives $\sqrt{2T\ln(1/\delta)}/\sigma$ plus a lower-order term, smaller by a factor of about $\sqrt{2\ln(1/\delta_0)}$, and with no failure probability $T\delta_0$ to pay.

DP-SGD adds Poisson sampling to each step. Its log-moments come from a one-dimensional pair of distributions. By Lemma 14.1 and scaling by C, a step compares the noise distribution $\mathcal{N}(0,\sigma^2)$ with the mixture $(1-q)\,\mathcal{N}(0,\sigma^2) + q\,\mathcal{N}(1,\sigma^2)$, in which the extra record was sampled with probability q.

Theorem 14.3 (Log-moments of the sampled Gaussian)**.** *For every integer $\lambda \ge 1$, one step of DP-SGD with sampling rate q and noise multiplier σ has, under add/remove neighbours,*

$$\kappa(\lambda) \le \ln A_{\lambda+1}, \qquad A_a = \sum_{k=0}^{a} \binom{a}{k} (1-q)^{a-k} q^k \exp\left(\frac{k^2 - k}{2\sigma^2}\right). \tag{14.14}$$

The full computation is Section 3.3 of Mironov, Talwar and Zhang (2019), listed in the references; in outline, the quantity A_a is computed by expanding the a-th power of the density ratio between the mixture and $\mathcal{N}(0,\sigma^2)$ with the binomial theorem. Each term is then a Gaussian integral, $\mathbb{E}[e^{k\xi/\sigma^2}] = e^{k^2/(2\sigma^2)}$ for $\xi \sim \mathcal{N}(0,\sigma^2)$. The hard part is to show that this pair, in this order, is the worst case over all neighbours and over the reverse order. Mironov, Talwar and Zhang prove it in "Rényi Differential Privacy of the Sampled Gaussian Mechanism", 2019. Earlier, Abadi and coauthors proved, for $\sigma \ge 1$, small q and moderate λ, the asymptotic bound $\kappa(\lambda) \le q^2\lambda(\lambda+1)/((1-q)\sigma^2) + O(q^3\lambda^3/\sigma^3)$ that yields Theorem 14.1.

Putting the pieces together gives the **moments accountant**. Compute the log-moment of one step for every λ, multiply by T by Lemma 14.2, and convert with Theorem 14.2, choosing the best λ:

$$\varepsilon(\delta) = \min_{\lambda \ge 1} \frac{T \ln A_{\lambda+1} + \ln(1/\delta)}{\lambda}. \tag{14.15}$$

The computation depends only on q, σ, T and δ, never on the data. So the privacy cost of a training run can be fixed before the data is touched, and the noise or the number of steps adjusted to meet a target. Figure 14.1 and the example below compare the accountant with the composition route on a typical setting.

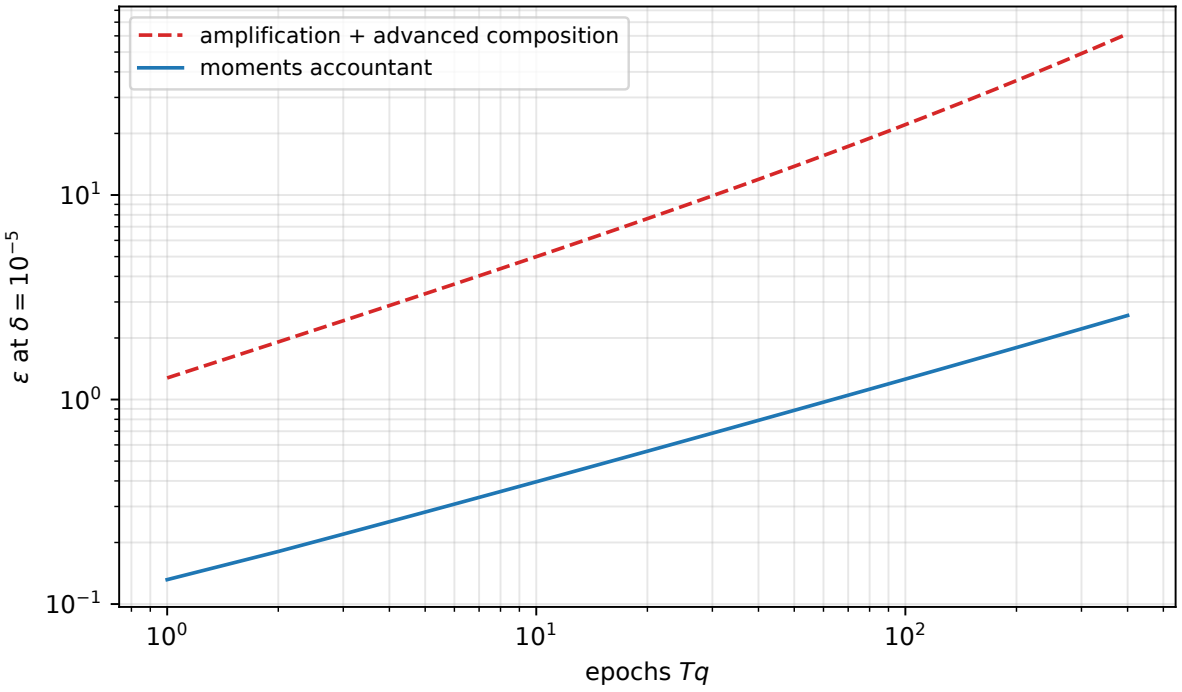


Figure 14.1.: Privacy parameter ε at $\delta = 10^{-5}$ of DP-SGD with the sampling rate q at 0.01 and noise multiplier $\sigma = 4$, against the number of epochs, where one epoch is $1/q$ steps. Dashed: each step is $(\varepsilon_0, \delta_0)$-DP by Corollary 14.1 with $T = 1$, Poisson sampling turns that into $(\ln(1 + q(e^{\varepsilon_0} - 1)), q\delta_0)$, and advanced composition combines the steps, with δ split evenly between the composition and the steps. Solid: the moments accountant Equation 14.15.

Example 14.1. With $q = 0.01$, $\sigma = 4$ and $\delta = 10^{-5}$, Figure 14.1 shows the two accountings of the same algorithm. After 100 epochs, or $T = 10{,}000$ steps, the composition route certifies $\varepsilon = 22.08$ and the moments accountant certifies $\varepsilon = 1.26$. After 400 epochs the values are 61.41 and 2.58. The first is no guarantee at all, and the second is a reasonable one.

The analyses certify upper bounds, and later work tightened them further. How much room is left can be tested from the other side. Jagielski, Ullman and Oprea trained models with DP-SGD on datasets with and without a few deliberately poisoned records, and measured how well an attacker could tell the two apart. Their attacks achieve privacy losses that come close to the certified bounds in some settings, so the accounting is not far from the truth there. Last, a caution about neighbours: these analyses use add/remove neighbours, while this book uses replace-one. A replacement is a removal followed by an addition, so a guarantee for one converts to the other at the cost of roughly doubling ε.

14.4. DP-SGD in practice

DP-SGD is designed as a drop-in replacement for SGD, but its behaviour differs in several ways.

Accuracy. On the handwritten-digit benchmark MNIST, a well-tuned DP-SGD reached 98.1% accuracy at $\varepsilon = 2.93$ and $\delta = 10^{-5}$, against about 99.8% without privacy. On the harder image benchmark CIFAR-10, it reached 66.2% at $\varepsilon = 7.53$, far below non-private accuracy. Both results are from Papernot and coauthors, 2020. The first is respectable, and the second is weak in both accuracy and privacy.

Speed. Standard frameworks compute the gradient of a whole batch at once on parallel hardware, and never form the gradient of each record. DP-SGD needs every per-example gradient in order to clip it, and computing them one at a time loses the parallelism. One workaround is **micro-batching**: average the gradients over small groups of s records, clip each group average, and add noise to their sum. One record changes one clipped group average by a vector of norm at most $2C$, so the noise cannot be smaller than with per-example clipping. The signal, on the other hand, is now a sum of averages, and it shrinks by a factor of about s. Faster training is paid for in accuracy. Other workarounds compute per-example gradients efficiently, using tricks with automatic differentiation or frameworks such as JAX with vectorized maps and just-in-time compilation.

Architecture. Bounded activations such as tanh work better than the rectified linear unit under DP-SGD, the opposite of the non-private experience. Size matters too. The noise vector b_t has expected squared norm $d\sigma^2C^2$, while the clipped sum has norm at most qnC. So the ratio of signal to noise is at most about $qn/(\sigma\sqrt{d})$, and it falls as the number of parameters d grows. Without privacy bigger networks are better, but under DP-SGD a network must be large enough to fit the task and no larger.

Hyperparameters. The learning rate, lot size, clipping norm and number of epochs must be tuned, and each training run on the private data spends privacy budget. Many reported results tune on the private data and do not count the cost, which overstates the privacy of the final model. A principled alternative is private selection, which picks the best of many private runs at a small constant factor in ε, as shown by Liu and Talwar in 2019. A practical one is public data. One can tune hyperparameters on a public dataset that resembles the private one. One can also pretrain a network on public data, freeze most of its layers, and train only the last layers privately. Fewer trained parameters mean a smaller d and less noise. Abadi and coauthors pretrained on CIFAR-100, treated as public, and trained privately on CIFAR-10 in this way.

14.5. Private aggregation of teacher ensembles

DP-SGD modifies the inside of one training algorithm. Private aggregation of teacher ensembles (PATE) treats the learner as a black box, so any improvement in non-private learning carries over. The price is a different model of the data. Besides a sensitive labelled training set, PATE needs a supply of **public unlabelled** feature vectors, for example from users who agreed to share their data.

PATE is an instance of **sample and aggregate**. The data is split into disjoint parts, a non-private algorithm runs on each part, and the answers are combined by a private aggregation. One record

affects only one part, so the aggregation has low sensitivity. Here the parts train m classifiers, the **teachers**, and the aggregation is a vote.

The simplest task is one private prediction for a public feature vector x. Each teacher votes for a label, and c_k is the number of votes for label k. The **mode** a is the label with the most votes, the smallest one in case of a tie, and the **gap** $\Gamma = c_a - \max_{k \neq a} c_k$ measures how far the vote is from changing its winner. The private mode releases the mode only when a noisy version of the gap clears a threshold. A large gap means that one changed record cannot change the winner.

Definition 14.5 (Private mode of the teachers)**.** Let $\varepsilon > 0$ and $\delta \in (0, 1/2)$. Split the dataset into m fixed disjoint parts by position, train one teacher on each, and collect their votes on x. Let a be the mode of the votes and Γ its gap, as above. Draw $Z \sim \mathrm{Lap}(2/\varepsilon)$, Laplace noise with scale $2/\varepsilon$. If

$$\Gamma + Z > \tau, \qquad \tau = 2 + \frac{2}{\varepsilon} \ln \frac{1}{2\delta}, \tag{14.16}$$

output the mode a. Otherwise output a symbol $\perp$ for "no answer", or a uniformly random label.

When the teachers agree, the mode is stable, and releasing it exactly costs little privacy. This is the propose-test-release idea of the chapter on beyond global sensitivity: test privately that the data is far from a dataset with another answer, and if so, release the answer without noise. The private mode is that chapter's stable mode applied to the votes, with the threshold lowered so that the failure probability is δ rather than $\delta/2$.

Theorem 14.4. *The private mode is (ε, δ)-DP. The teachers' training algorithm may be randomized, with coins independent of everything else.*

Proof.

1. Fix the teachers' coins; a guarantee for every fixed choice of coins implies one for random coins. Neighbours differ in one record, which lies in one part, so at most one teacher's vote changes. That moves one vote from one label to another, and changes each count by at most 1.
2. *The gap has sensitivity* 2. If the mode stays the same, both terms of Γ change by at most 1. If the mode changes from a on D to $a' \neq a$ on D', then $c_{a'} + 1 \geq c'_{a'} \geq c'_a \geq c_a - 1$, so $\Gamma \leq c_a - c_{a'} \leq 2$, and likewise $\Gamma' \leq 2$. In both cases $|\Gamma - \Gamma'| \leq 2$. So, by the Laplace mechanism, the indicator that the test passes is ε-DP.
3. *Same mode.* The output is a function of the test indicator and of the common mode, so it is ε-DP by post-processing.
4. *Different modes.* By step 2, $\Gamma \leq 2$. The test then passes on D with probability at most $\Pr[Z > \tau - 2] = \frac{1}{2} e^{-\varepsilon(\tau-2)/2} = \delta$. Outside that event the output is $\perp$, so for every set S of outputs,

$$\Pr[M(D) \in S] \leq \Pr[M(D) = \perp]\, \mathbf{1}[\perp \in S] + \delta \leq e^{\varepsilon} \Pr[M(D') = \perp]\, \mathbf{1}[\perp \in S] + \delta \leq e^{\varepsilon} \Pr[M(D') \in S] + \delta, \tag{14.17}$$

where the middle step is the ε-DP of the test from step 2. Replacing $\perp$ by a random label is post-processing.

□

The accuracy depends on how good the teachers are. Suppose the records are independent draws from a distribution. Then teachers trained on disjoint parts vote independently on a fixed x.

Proposition 14.1. *Let $\beta \in (0,1)$ be a failure probability, and let τ be the threshold of Definition 14.5. Suppose each teacher, independently, votes for the correct label y of x with probability at least $3/4$. If*

$$m \geq \max\left(4\tau + \frac{8}{\varepsilon}\ln\frac{1}{\beta},\ 32\ln\frac{2}{\beta}\right), \tag{14.18}$$

then the private mode outputs y with probability at least $1-\beta$.

Proof.

1. Let V be the number of correct votes. Its mean is at least $3m/4$, and Hoeffding's inequality for independent indicators gives $\Pr[V \leq 5m/8] \leq e^{-2m(1/8)^2} = e^{-m/32} \leq \beta/2$.
2. If $V > 5m/8$, then y has a majority, so it is the mode, and every other label has at most $m - V$ votes. So $\Gamma \geq 2V - m > m/4$.
3. The test then fails with probability $\Pr[Z \leq \tau - \Gamma] \leq \frac{1}{2}e^{-\varepsilon(m/4-\tau)/2}$, because $m/4 > \tau$. The first lower bound on m makes this at most $\beta/2$.
4. A union bound over the two failure events gives the claim.

□

Concrete values show how many teachers this takes.

Example 14.2. With $\varepsilon = 1$, $\delta = 10^{-5}$ and $\beta = 0.01$, the threshold is $\tau = 2 + 2\ln(50{,}000) = 23.64$. The two lower bounds are $94.56 + 8\ln 100 = 131.4$ and $32\ln 200 = 169.5$, so 170 teachers suffice. If a teacher needs n_0 records to be right three times in four, the private prediction needs $170\,n_0$.

The required number of teachers grows like $\ln(1/\delta)/\varepsilon$. This is the general price of sample and aggregate: about $1/\varepsilon$ times more data than the non-private algorithm, which is often more than the best private algorithm for a specific task needs.

Answering k queries by running the private mode k times costs, by advanced composition, about $\varepsilon\sqrt{2k\ln(1/\delta')}$. The sparse vector technique of the sparse vector chapter does better when most queries are stable. It answers a stream of threshold queries, and pays only for the queries whose answer is above the threshold. Here each query asks whether the gap on the next point is small, so the technique pays only for the unstable points, where it outputs a random label, and answers the stable ones with the exact mode. If the teachers are accurate, few points are unstable. Bassily, Thakkar and Thakurta prove this version, with a privacy cost that grows like the square root of the number of unstable queries rather than of all k.

The privately labelled points then serve a further purpose. The answers $(x_1, \hat{y}_1), \dots, (x_k, \hat{y}_k)$, with public features and private labels, form a labelled dataset that is safe to release. A non-private learner, the **student**, is trained on it. By post-processing, the student is as private as the labels:

if the labelling is (ε, δ)-DP and the public features do not depend on the private data, the student is (ε, δ)-DP too. The student can then answer any number of queries at no further cost, which is how PATE turns a bounded number of private predictions into a private model.

14.6. Summary

Table 14.1.: The chapter's results.

Result	Statement
Lemma 14.1	one record changes the clipped sum by at most C, or $2C$ under replacement
Theorem 14.1	DP-SGD is (ε, δ)-DP with $\sigma = O(q\sqrt{T\ln(1/\delta)}/\varepsilon)$
Theorem 14.2	a log-moment bound $\kappa(\lambda)$ gives $(\varepsilon, e^{\kappa(\lambda)-\lambda\varepsilon})$-DP
Lemma 14.2	log-moments add up under adaptive composition
Lemma 14.3	the Gaussian mechanism has $\kappa(\lambda) = \lambda(\lambda+1)(\Delta_2 f)^2/(2\sigma^2)$
	T Gaussian steps give $\varepsilon = T\rho + 2\sqrt{T\rho\ln(1/\delta)}$
Corollary 14.1	
Theorem 14.3	a sampled Gaussian step has $\kappa(\lambda) \le \ln A_{\lambda+1}$
Equation 14.15	the moments accountant for T steps of DP-SGD
Theorem 14.4	releasing the teachers' mode when its noisy gap exceeds τ is (ε, δ)-DP
Proposition 14.1	$O(\ln(1/\delta)/\varepsilon + \ln(1/\beta)/\varepsilon)$ accurate teachers suffice

14.7. References

- Martín Abadi, Andy Chu, Ian Goodfellow, H. Brendan McMahan, Ilya Mironov, Kunal Talwar and Li Zhang, "Deep Learning with Differential Privacy", 2016. DP-SGD with clipping and Poisson sampling, the moments accountant, and pretraining on public data.
- Ilya Mironov, "Rényi Differential Privacy", 2017. The privacy notion built on the moments of the privacy loss, and its composition and conversion rules.
- Ilya Mironov, Kunal Talwar and Li Zhang, "Rényi Differential Privacy of the Sampled Gaussian Mechanism", 2019. The exact log-moments of the sampled Gaussian, Theorem 14.3.
- Yu-Xiang Wang, Borja Balle and Shiva Kasiviswanathan, "Subsampled Rényi Differential Privacy and Analytical Moments Accountant", 2019. Log-moment bounds for subsampled mechanisms under sampling without replacement.

- Borja Balle, Gilles Barthe and Marco Gaboardi, "Privacy Amplification by Subsampling: Tight Analyses via Couplings and Divergences", 2018. Amplification by Poisson sampling, used in Figure 14.1.
- Matthew Jagielski, Jonathan Ullman and Alina Oprea, "Auditing Differentially Private Machine Learning: How Private is Private SGD?", 2020. Attacks that bound from below the privacy of DP-SGD.
- Nicolas Papernot, Abhradeep Thakurta, Shuang Song, Steve Chien and Úlfar Erlingsson, "Tempered Sigmoid Activations for Deep Learning with Differential Privacy", 2020. Bounded activations under DP-SGD, and the accuracy figures quoted above.
- Ian Goodfellow, "Efficient Per-Example Gradient Computations", 2015. A trick for per-example gradients in simple networks.
- Jingcheng Liu and Kunal Talwar, "Private Selection from Private Candidates", 2019. Hyperparameter selection at a constant factor in privacy.
- Kobbi Nissim, Sofya Raskhodnikova and Adam Smith, "Smooth Sensitivity and Sampling in Private Data Analysis", 2007. The sample and aggregate framework.
- Nicolas Papernot, Martín Abadi, Úlfar Erlingsson, Ian Goodfellow and Kunal Talwar, "Semi-Supervised Knowledge Transfer for Deep Learning from Private Training Data", 2017, and Nicolas Papernot, Shuang Song, Ilya Mironov, Ananth Raghunathan, Kunal Talwar and Úlfar Erlingsson, "Scalable Private Learning with PATE", 2018. The PATE framework and its refinements.
- Raef Bassily, Om Thakkar and Abhradeep Thakurta, "Model-Agnostic Private Learning", 2018. The analysis of PATE with the sparse vector technique.
- Abhradeep Thakurta and Adam Smith, "Differentially Private Feature Selection via Stability Arguments, and the Robustness of the Lasso", 2013. Stability-based release of a stable answer, as in Theorem 14.4.

15. Private Mean Estimation

Earlier chapters estimated properties of a fixed dataset. In statistics, the dataset is a sample, and the goal is a property of the distribution that generated it. This chapter asks how many samples an algorithm with differential privacy (DP) needs to estimate the mean of a distribution to within α, and how much more that is than without privacy.

For bounded data, privacy costs a lower-order term. For unbounded data it is impossible without prior knowledge, and a bound R on the mean is the least such knowledge. Clipping to $[-R, R]$ costs samples linear in R. Locating the data first with a histogram, or by shrinking a confidence interval step by step, brings the cost down to logarithmic in R. The same recipe extends to many dimensions. For heavy-tailed data, the clipping must be aggressive and introduce bias, and the price of privacy grows.

15.1. The mean of a dataset and of a distribution

Two facts from earlier chapters are used throughout. A Laplace variable with scale b, written $Z \sim \mathrm{Lap}(b)$, has density proportional to $e^{-|z|/b}$, so its tail is

$$\Pr[|Z| > t] = e^{-t/b} \qquad (t \geq 0). \tag{15.1}$$

The Laplace mechanism privatizes a function f of the dataset with values in $\mathbb{R}^d$. Its **sensitivity** $\Delta_1 f$ is the largest change that replacing one record can cause,

$$\Delta_1 f = \max_{D \sim D'} \|f(D) - f(D')\|_1. \tag{15.2}$$

If $\Delta_1 f$ is finite, releasing $f(D)$ plus independent $\mathrm{Lap}(\Delta_1 f/\varepsilon)$ noise in each coordinate is ε-DP.

Start with bits. Each of the n records is a bit $X_i \in \{0, 1\}$, and their empirical mean is $\tilde{p} = \frac{1}{n}\sum_i X_i$. Replacing one bit moves $\tilde{p}$ by at most $1/n$, so the Laplace mechanism releases

$$\hat{p} = \tilde{p} + \mathrm{Lap}\left(\frac{1}{\varepsilon n}\right). \tag{15.3}$$

For the dataset itself, the error $|\hat{p} - \tilde{p}|$ is the size of the noise, at most $\ln(1/\beta)/(\varepsilon n)$ with probability $1 - \beta$. Now suppose the bits are a sample, $X_1, \ldots, X_n$ independent draws from $\mathrm{Bern}(p)$, and the target is the unknown p.

Theorem 15.1 (Private Bernoulli mean)**.** *The estimator Equation 15.3 is ε-DP. If $X_1, \dots, X_n$ are independent* Bern(p) *variables, then for every $\beta \in (0,1)$, with probability at least $1-\beta$,*

$$|\hat{p} - p| \le \sqrt{\frac{\ln(4/\beta)}{2n}} + \frac{\ln(2/\beta)}{\varepsilon n}. \tag{15.4}$$

Proof.

1. Privacy is the Laplace mechanism with sensitivity $1/n$.
2. By the triangle inequality, $|\hat{p} - p| \le |\tilde{p} - p| + |\hat{p} - \tilde{p}|$.
3. The bits are independent, take values in $[0,1]$, and have mean p. Hoeffding's inequality gives $\Pr[|\tilde{p} - p| > t] \le 2e^{-2nt^2}$, which equals $\beta/2$ at $t = \sqrt{\ln(4/\beta)/(2n)}$.
4. The noise $\hat{p} - \tilde{p}$ is Lap$(1/(\varepsilon n))$, so by the Laplace tail it exceeds $\ln(2/\beta)/(\varepsilon n)$ with probability $\beta/2$.
5. A union bound over the two failures gives the claim.

□

Each term is at most $\alpha/2$ once

$$n \ge \frac{2\ln(4/\beta)}{\alpha^2} + \frac{2\ln(2/\beta)}{\alpha\varepsilon}, \tag{15.5}$$

so $n = O(1/\alpha^2 + 1/(\alpha\varepsilon))$ samples suffice for a fixed β. The first term is the sample complexity without privacy, and the second is the **cost of privacy**. For a constant ε and small α, the cost of privacy is smaller by a factor of α, so here privacy is nearly free, apart from constants that matter in practice.

The proof split the error into two parts, and that split recurs in the chapter. The **sampling error** $\tilde{p} - p$ is present without privacy, and the **noise error** $\hat{p} - \tilde{p}$ is what privacy adds.

15.2. Unbounded data

The Bernoulli case was easy because each record lies in $[0,1]$. Take instead real records, such as heights in centimetres. Replacing one height of 160 by a recorded 10^5 moves the empirical mean by about $10^5/n$, and nothing bounds the move. So the empirical mean has infinite sensitivity, and no private algorithm can fix that, as the next result shows. It uses **group privacy**. If a mechanism M is ε-DP and two datasets D, D' differ in k entries, then for every event S,

$$\Pr[M(D) \in S] \le e^{k\varepsilon} \Pr[M(D') \in S], \tag{15.6}$$

by applying the definition along a path of k neighbours. The next result uses this to show that every ε-DP algorithm is far from the empirical mean on some real dataset, however large the error α it is allowed.

Proposition 15.1. *Let $M : \mathbb{R}^n \to \mathbb{R}$ be ε-DP and $\alpha > 0$. Some dataset $D \in \mathbb{R}^n$ has $\Pr[|M(D) - \bar{X}| \le \alpha] < 1/2$, where $\bar{X}$ is its empirical mean.*

Proof.

1. Suppose instead that $\Pr[|M(D) - \bar{X}| \le \alpha] \ge 1/2$ for every D. Take an integer $m > 2e^{\varepsilon n}$ and, for $j = 0, \dots, m-1$, the dataset $D^{(j)}$ whose n entries all equal $3\alpha j$, so that its empirical mean is $3\alpha j$.
2. The intervals $I_j = [3\alpha j - \alpha, 3\alpha j + \alpha]$ are disjoint, and $\Pr[M(D^{(j)}) \in I_j] \ge 1/2$ by assumption.
3. The datasets $D^{(0)}$ and $D^{(j)}$ differ in n entries, so group privacy gives $\Pr[M(D^{(0)}) \in I_j] \ge e^{-\varepsilon n}/2$.
4. Summing over the disjoint intervals gives $1 \ge m e^{-\varepsilon n}/2 > 1$, a contradiction.

□

The way out is prior knowledge that does not come from the data. No human is taller than 300 cm, so a height can be forced into $[0, 300]$ before averaging.

Definition 15.1 (Clipping)**.** For $u \le v$, the **clipping** map onto $[u, v]$ is $g_{[u,v]}(x) = \min(\max(x, u), v)$. It leaves points of $[u, v]$ unchanged and moves every other point to the nearer endpoint.

Clipping to a range fixed in advance bounds how far one record can move the mean, so the Laplace mechanism applies to the clipped mean. When the data already lies in the range, clipping changes nothing and the only error is the noise.

Lemma 15.1. *Let $u < v$ be fixed before seeing the data, and let $g_{[u,v]}$ be the clipping map of Definition 15.1. The estimator*

$$M(D) = \frac{1}{n} \sum_{i=1}^{n} g_{[u,v]}(X_i) + \mathrm{Lap}\Big(\frac{v-u}{\varepsilon n}\Big) \tag{15.7}$$

is ε-DP. If every X_i lies in $[u, v]$, then $|M(D) - \bar{X}| \le (v-u)\ln(1/\beta)/(\varepsilon n)$ with probability at least $1 - \beta$.

Proof.

1. Each clipped value lies in $[u, v]$, so replacing one record moves the clipped mean by at most $(v-u)/n$. The Laplace mechanism gives privacy.
2. If every X_i lies in $[u, v]$, clipping changes nothing, and $M(D) - \bar{X}$ is the Laplace noise, whose tail gives the bound.

□

With heights in $[0, 300]$, $n = 1{,}000$ and $\varepsilon = 1$, the noise scale is $300/1{,}000 = 0.3$ cm. The range must be fixed before the data is seen. Choosing it after looking, for instance to exclude one absurd record, makes the range depend on that record and voids the guarantee.

15.3. Estimating a parameter of a distribution

From here on the data is a sample from an unknown distribution, and the target is the mean of that distribution. An estimator for this problem must meet two requirements of different kinds. It must be private on every dataset, and it must be accurate on samples from the distribution.

Definition 15.2 (Private estimation)**.** Let $\mathcal{P}$ be a class of distributions on $\mathbb{R}$, each with a mean $\mu(P)$. An algorithm $M \colon \mathbb{R}^n \to \mathbb{R}$ is an ε**-DP** (α, β)**-estimator** for $\mathcal{P}$ if

1. M is ε-DP, as a map on all datasets in $\mathbb{R}^n$, and
2. for every $P \in \mathcal{P}$, if $X_1, \dots, X_n$ are independent draws from P, then $\Pr[|M(D) - \mu(P)| \le \alpha] \ge 1 - \beta$.

The smallest such n is the **sample complexity** of the class.

Privacy is a worst-case guarantee. It holds even for a dataset that has probability zero under every $P \in \mathcal{P}$. If the analyst's model of the data is wrong, the people in the dataset must not pay for it. Average-case variants of privacy also lose composition and other properties that later chapters rely on.

Accuracy, in contrast, is an average over samples from P, and it needs assumptions on P even without privacy. The empirical mean of Gaussian samples converges to the mean, but the empirical mean of Cauchy samples does not converge at all, and a Cauchy distribution has no mean.

The main class is the Gaussians with unit variance and a bounded mean,

$$\mathcal{G}_R = \{\, \mathcal{N}(\mu, 1) : |\mu| \le R \,\}. \tag{15.8}$$

A bound on the mean is the analogue of the range of heights, and some bound is necessary, as the next result shows by the argument of Proposition 15.1.

Proposition 15.2. *Let $\alpha > 0$, $\beta \in (0, 1)$, and $m = \lfloor 2R/(3\alpha) \rfloor + 1$. Every ε-DP (α, β)-estimator for $\mathcal{G}_R$ uses*

$$n \ge \frac{\ln\big((1-\beta)m\big)}{\varepsilon} \tag{15.9}$$

samples. Without a bound on $|\mu|$ ($R = \infty$), no finite n suffices.

Proof.

1. The means $\mu_j = -R + 3\alpha j$, for $j = 0, \dots, m-1$, lie in $[-R, R]$, and the intervals $I_j = [\mu_j - \alpha, \mu_j + \alpha]$ are disjoint.
2. Let D be a sample from $\mathcal{N}(\mu_0, 1)$ and set $D^{(j)} = D + (\mu_j - \mu_0)$, shifting every entry. Then $D^{(j)}$ is a sample from $\mathcal{N}(\mu_j, 1)$, and it differs from D in all n entries.
3. For each fixed D, group privacy gives $\Pr_M[M(D) \in I_j] \ge e^{-\varepsilon n} \Pr_M[M(D^{(j)}) \in I_j]$. Averaging over D and using accuracy for $\mathcal{N}(\mu_j, 1)$ gives $\Pr[M(D) \in I_j] \ge e^{-\varepsilon n}(1 - \beta)$.

4. The intervals are disjoint, so summing over j gives $1 \geq m e^{-\varepsilon n}(1-\beta)$. Solving for n gives the bound, and $m = \infty$ when $R = \infty$.

□

So a term of order $\ln(R/\alpha)/\varepsilon$ is unavoidable. The rest of the chapter finds algorithms that pay little more than that.

15.4. Clipping the Gaussian

A bound on the mean is not a bound on the records. Combined with the Gaussian's light tails, though, it gives one with high probability.

Lemma 15.2 (Gaussian tail)**.** *If $Y \sim \mathcal{N}(0,1)$ and $t \geq 0$, then $\Pr[|Y| \geq t] \leq 2e^{-t^2/2}$.*

Proof.

1. For $\lambda > 0$, Markov's inequality applied to $e^{\lambda Y}$ gives $\Pr[Y \geq t] \leq e^{-\lambda t}\,\mathbb{E}[e^{\lambda Y}] = e^{-\lambda t + \lambda^2/2}$, using the Gaussian moment generating function.
2. Taking $\lambda = t$ gives $\Pr[Y \geq t] \leq e^{-t^2/2}$. The lower tail is the same by symmetry, and the two add.

□

Fix $\beta \in (0,1)$ and write

$$\tau = \sqrt{2\ln(2n/\beta)}. \tag{15.10}$$

For a sample from $\mathcal{N}(\mu,1)$, each point leaves $[\mu-\tau, \mu+\tau]$ with probability at most $2e^{-\tau^2/2} = \beta/n$, so by a union bound all n points stay inside with probability at least $1-\beta$. The width τ grows only like $\sqrt{\ln n}$, and the interval $\mu \pm 3$ already holds 99.7% of a Gaussian's mass.

The first estimator clips to an interval known to contain the mean, widened by τ. It is stated for any public centre c and radius r, because the later algorithms call it with a centre they computed privately.

Theorem 15.2 (Clipped Laplace estimator)**.** *Let $c \in \mathbb{R}$ and $r \geq 0$ be public, $\beta \in (0,1)$, and τ as in Equation 15.10. With $g_{[u,v]}$ the clipping map of Definition 15.1, the estimator*

$$M_{c,r}(D) = \frac{1}{n}\sum_{i=1}^{n} g_{[c-r-\tau,\; c+r+\tau]}(X_i) + Z, \qquad Z \sim \mathrm{Lap}\!\left(\frac{2(r+\tau)}{\varepsilon n}\right), \tag{15.11}$$

is ε-DP. Let $X_1, \dots, X_n$ be independent $\mathcal{N}(\mu,1)$ draws. With probability at least $1-2\beta$ over the sample, an event E holds on which every $M_{c,r}$ with $|\mu - c| \leq r$ satisfies

$$|M_{c,r}(D) - \mu| \le \sqrt{\frac{2\ln(2/\beta)}{n}} + |Z|. \tag{15.12}$$

In particular, if $|\mu - c| \le r$, then with probability at least $1 - 3\beta$,

$$|M_{c,r}(D) - \mu| \le s_\varepsilon(r) = \sqrt{\frac{2\ln(2/\beta)}{n}} + \frac{2(r+\tau)\ln(1/\beta)}{\varepsilon n}. \tag{15.13}$$

Proof.

1. *Privacy.* The clipping interval has length $2(r + \tau)$, so the clipped mean has sensitivity $2(r + \tau)/n$, and the Laplace mechanism applies.
2. *The event.* Let G be the event that every X_i lies in $[\mu - \tau, \mu + \tau]$, with probability at least $1 - \beta$ by the union bound above. Let S be the event that the empirical mean $\tilde{\mu} = \frac{1}{n}\sum_i X_i$ satisfies $|\tilde{\mu} - \mu| \le \sqrt{2\ln(2/\beta)/n}$. Since $\tilde{\mu} - \mu \sim \mathcal{N}(0, 1/n)$, Lemma 15.2 with $t = \sqrt{2\ln(2/\beta)}$ gives $\Pr[S] \ge 1 - \beta$. Set $E = G \cap S$.
3. *No clipping on E.* If $|\mu - c| \le r$, the interval $[\mu - \tau, \mu + \tau]$ lies inside $[c - r - \tau, c + r + \tau]$. So on G clipping changes no point, and $M_{c,r}(D) = \tilde{\mu} + Z$.
4. *Error on E.* By the triangle inequality, $|M_{c,r}(D) - \mu| \le |\tilde{\mu} - \mu| + |Z|$, and S bounds the first term.
5. *Noise.* The Laplace tail gives $|Z| \le 2(r + \tau)\ln(1/\beta)/(\varepsilon n)$ with probability $1 - \beta$. A union bound with E gives Equation 15.13.

□

The proof is an instance of a three-way split. Take any clipping map g, write $\bar{g} = \frac{1}{n}\sum_i g(X_i)$ for the clipped mean and $\tilde{\mu}$ for the empirical mean, and let the estimate be $\hat{\mu} = \bar{g} + Z$ with noise Z. Then

$$\hat{\mu} - \mu = \underbrace{(\tilde{\mu} - \mu)}_{\text{sampling}} + \underbrace{(\bar{g} - \tilde{\mu})}_{\text{clipping bias}} + \underbrace{Z}_{\text{noise}}. \tag{15.14}$$

Here the clipping interval was chosen wide enough that the bias term is zero with high probability. That choice holds until Section 15.8, where the bias must be traded against the noise.

With $c = 0$ and $r = R$, each term of Equation 15.13 is at most $\alpha/2$ when

$$n \ge \frac{8\ln(2/\beta)}{\alpha^2} + \frac{4(R+\tau)\ln(1/\beta)}{\alpha\varepsilon}, \tag{15.15}$$

that is, $n = \tilde{O}(1/\alpha^2 + R/(\alpha\varepsilon))$, where $\tilde{O}$ hides factors logarithmic in n and $1/\beta$. When R is a small constant, this matches the Bernoulli case. But R measures how little the analyst knows. For adult heights, a range of 140 to 200 cm gives $R = 30$ after centring, while for an unfamiliar quantity R may have to be 10^6, and then the noise swamps the signal. The goal of the next two sections is to replace R by $\ln R$.

15.5. A coarse estimate from a histogram

The first improvement locates the data before averaging it. Cut the range into bins of width 2τ. The whole sample fits in an interval of that width, so one or two bins hold all the points, and a private histogram can find one of them at a cost logarithmic in the number of bins.

Lemma 15.3 (Coarse estimate)**.** *Let τ be as in Equation 15.10, $m = \lfloor R/\tau \rfloor + 2$, and $a_j = -R - \tau + 2\tau j$ for $j = 0, \dots, m-1$. Let h_j count the records in the bin $B_j = [a_j, a_j + 2\tau)$. The algorithm outputs $c = a_{\hat{\jmath}} + \tau$, where $\hat{\jmath}$ maximizes $h_j + Z_j$ over j, with independent $Z_j \sim \mathrm{Lap}(2/\varepsilon)$. It is ε-DP. If $X_1, \dots, X_n$ are independent $\mathcal{N}(\mu, 1)$ draws with $|\mu| \le R$ and*

$$n > \frac{8}{\varepsilon} \ln \frac{m}{\beta}, \tag{15.16}$$

then $|\mu - c| \le 2\tau$ with probability at least $1 - 2\beta$.

Proof.

1. *Privacy.* Replacing one record moves it out of at most one bin and into at most one bin, so the vector of counts has sensitivity 2. Adding $\mathrm{Lap}(2/\varepsilon)$ to each count is ε-DP by the Laplace mechanism, and the argmax is post-processing.
2. *The bins catch the sample.* On the event G of Theorem 15.2, every record lies in $[\mu - \tau, \mu + \tau] \subseteq [-R - \tau, R + \tau]$. The bins cover $[-R - \tau, -R - \tau + 2\tau m)$, and $2\tau m > 2R + 2\tau$, so every record is counted.
3. *One or two bins hold everything.* The interval $[\mu - \tau, \mu + \tau]$ has length 2τ, so it meets at most two bins. Hence some bin has $h_j \ge n/2$, and every bin that misses the interval has $h_j = 0$.
4. *The noise is small.* Set $t = (2/\varepsilon)\ln(m/\beta)$. Each $|Z_j| > t$ with probability β/m, so with probability at least $1 - \beta$ all $|Z_j| \le t$. Then the fullest bin has noisy count at least $n/2 - t$, and an empty bin at most t. Since $n > 4t$, the winning bin $\hat{\jmath}$ is not empty.
5. *The winner is near the mean.* On G, the bin $B_{\hat{\jmath}}$ holds a record x with $|x - \mu| \le \tau$ and $a_{\hat{\jmath}} \le x < a_{\hat{\jmath}} + 2\tau$. So $\mu \in [a_{\hat{\jmath}} - \tau, a_{\hat{\jmath}} + 3\tau] = [c - 2\tau, c + 2\tau]$. The two failures have total probability at most 2β.

□

After the coarse step, the mean is known to within $2\tau = O(\sqrt{\ln(n/\beta)})$ instead of R. The fine step is Theorem 15.2 on that short interval. Basic composition, from the Laplace mechanism chapter, says that running an ε_1-DP algorithm and then an ε_2-DP algorithm that may depend on the first output is $(\varepsilon_1 + \varepsilon_2)$-DP. So each step gets half the budget.

Theorem 15.3 (Histogram, then clipping)**.** *Let m be the number of bins in Lemma 15.3. Run that algorithm with $\varepsilon/2$ (noise $\mathrm{Lap}(4/\varepsilon)$) to get c, then output $\hat{\mu} = M_{c,2\tau}(D)$ from Theorem 15.2 with $\varepsilon/2$. The result is ε-DP. If $X_1, \dots, X_n$ are independent $\mathcal{N}(\mu, 1)$ draws with $|\mu| \le R$ and $n > (16/\varepsilon)\ln(m/\beta)$, then with probability at least $1 - 4\beta$,*

$$|\hat{\mu} - \mu| \le \sqrt{\frac{2\ln(2/\beta)}{n}} + \frac{12\,\tau \ln(1/\beta)}{\varepsilon n}. \tag{15.17}$$

Proof.

1. *Privacy.* The two steps are $\varepsilon/2$-DP each, and basic composition adds them.
2. *Coarse step.* Lemma 15.3 with $\varepsilon/2$ needs $n > (16/\varepsilon)\ln(m/\beta)$. Its two failure events are G and the event that some $|Z_j|$ is large, of total probability at most 2β. Outside them, $|\mu - c| \le 2\tau$.
3. *Fine step.* On the event $E = G \cap S$ of Theorem 15.2, the fine estimate satisfies $|\hat{\mu} - \mu| \le \sqrt{2\ln(2/\beta)/n} + |Z|$, since $|\mu - c| \le 2\tau$. The event S adds probability β of failure.
4. The noise Z is $\mathrm{Lap}(2(2\tau + \tau)/((\varepsilon/2)n)) = \mathrm{Lap}(12\tau/(\varepsilon n))$, whatever c is, and it exceeds $12\tau\ln(1/\beta)/(\varepsilon n)$ with probability β. The union bound over G, the histogram noise, S and Z gives $1 - 4\beta$.

□

The error is at most α once $n \ge 8\ln(2/\beta)/\alpha^2 + 24\tau\ln(1/\beta)/(\alpha\varepsilon)$, and the coarse step needs $n > (16/\varepsilon)\ln(m/\beta)$ with $m \le R/\tau + 2$. So

$$n = \tilde{O}\Big(\frac{1}{\alpha^2} + \frac{1}{\alpha\varepsilon} + \frac{\ln R}{\varepsilon}\Big). \tag{15.18}$$

The dependence on R fell from linear to logarithmic. The first term is needed without privacy, Proposition 15.2 shows that the last is needed, and the middle one is needed too, up to logarithmic factors. The exponential mechanism over the bins can replace the noisy histogram in the coarse step with the same guarantee.

15.6. Shrinking confidence intervals

The histogram finds the data in one shot, but a histogram over a region of $\mathbb{R}^d$ needs a number of bins exponential in d. A second route uses only the clipped Laplace estimator, applied repeatedly. Each call returns a confidence interval instead of a point. The interval is a constant factor shorter than the one it started from, and with high probability it still contains μ.

One step takes an interval $[c - r, c + r]$ and a budget ε', computes the estimate $c' = M_{c,r}(D)$ with ε', and returns the interval $[c' - s, c' + s]$ with $s = s_{\varepsilon'}(r)$ from Equation 15.13. The radius s depends only on r and on public parameters, not on the data.

Lemma 15.4 (One step)**.** *Let τ be as in Equation 15.10, $n \ge 16$ and $n \ge 16\ln(1/\beta)/\varepsilon'$. Consider one step, as described above, from the interval $[c - r, c + r]$ to the interval $[c' - s, c' + s]$. Then*

$$s_{\varepsilon'}(r) \le \tfrac{1}{2}\max(r, \tau). \tag{15.19}$$

Moreover, on the event E of Theorem 15.2, if $|\mu - c| \le r$, then $\mu \in [c' - s, c' + s]$ with probability at least $1 - \beta$ over the step's noise.

Proof.

1. Since $2\ln(2/\beta) \le 2\ln(2n/\beta) = \tau^2$ and $n \ge 16$, the first term of $s_{\varepsilon'}(r)$ is at most $\tau/\sqrt{n} \le \tau/4$.

2. Since $n \geq 16\ln(1/\beta)/\varepsilon'$, the second term is at most $2(r+\tau)/16 = (r+\tau)/8 \leq \max(r,\tau)/4$.
3. Adding the two terms gives the first claim, because $\tau/4 \leq \max(r,\tau)/4$.
4. On E, Theorem 15.2 gives $|c' - \mu| \leq \sqrt{2\ln(2/\beta)/n} + |Z|$, where Z is the step's Laplace noise. That noise is at most $2(r+\tau)\ln(1/\beta)/(\varepsilon' n)$ with probability $1-\beta$, and then $|c' - \mu| \leq s$.

□

The first claim is deterministic, and the second needs the data to be Gaussian. Iterating halves the radius until it reaches the floor τ, where the tails of the data stop the clipping from doing better.

Theorem 15.4 (Shrinking intervals, then clipping)**.** *Let $T = \lceil \log_2(R/\tau) \rceil$ (or $T = 0$ if $R \leq \tau$), $\varepsilon' = \varepsilon/(2T)$, $r_0 = R$, $c_0 = 0$, and $r_t = s_{\varepsilon'}(r_{t-1})$. For $t = 1, \ldots, T$, let $c_t = M_{c_{t-1}, r_{t-1}}(D)$ computed with ε'. Output $\hat{\mu} = M_{c_T, r_T}(D)$ computed with $\varepsilon/2$. The result is ε-DP. Suppose $X_1, \ldots, X_n$ are independent $\mathcal{N}(\mu, 1)$ draws with $|\mu| \leq R$,*

$$n \geq 16 \qquad \text{and} \qquad n \geq \frac{32\,T\ln(1/\beta)}{\varepsilon}. \tag{15.20}$$

Then with probability at least $1 - (T+3)\beta$,

$$|\hat{\mu} - \mu| \leq \sqrt{\frac{2\ln(2/\beta)}{n}} + \frac{4(r_T + \tau)\ln(1/\beta)}{\varepsilon n} \leq \sqrt{\frac{2\ln(2/\beta)}{n}} + \frac{8\,\tau\ln(1/\beta)}{\varepsilon n}. \tag{15.21}$$

Proof.

1. *Privacy.* The radii r_t depend only on n, ε, β and R, so each step is an instance of Theorem 15.2 with public parameters. There are T steps of $\varepsilon/(2T)$ and one of $\varepsilon/2$, each using the outputs before it, and basic composition gives ε.
2. *The radii shrink.* The condition on n is the one of Lemma 15.4 with $\varepsilon' = \varepsilon/(2T)$, so $r_t \leq \max(r_{t-1}, \tau)/2$. By induction, $r_t \leq \max(R/2^t, \tau/2)$. Since $2^T \geq R/\tau$, this gives $r_T \leq \tau$.
3. *The intervals cover μ.* The event E holds with probability at least $1 - 2\beta$, and it concerns only the sample, so it serves every step. On E, by induction on t, Lemma 15.4 keeps $|\mu - c_t| \leq r_t$ unless the noise of step t fails, which has probability β.
4. *Final step.* On E and with $|\mu - c_T| \leq r_T$, Theorem 15.2 with $\varepsilon/2$ bounds the error by the first expression, unless its noise fails, with probability β. Step 2 gives $r_T \leq \tau$, hence the second expression.
5. The failures are E (2β), the T steps ($T\beta$) and the final noise (β).

□

The steps need $n = O(\ln(R)\ln(1/\beta)/\varepsilon)$, and the final step needs $n = \tilde{O}(1/\alpha^2 + 1/(\alpha\varepsilon))$ for error α. So shrinking intervals achieve the same rate Equation 15.18 as the histogram, with only the Laplace mechanism. The example below puts numbers on the three bounds when the prior bound R is large.

Example 15.1. Take $n = 4{,}000$, $\varepsilon = 1$, $R = 10^6$ and $\beta = 0.002$, so that $\tau = 5.51$. Then $T = \lceil \log_2(10^6/5.51) \rceil = 18$ steps, which need $n \geq 32 \cdot 18 \ln 500 = 3{,}580$, and all $T + 3 = 21$ events hold with probability at least $1 - 21 \cdot 0.002 = 0.958$. The radius reached by the recursion gives the error bound 0.098. The coarse step of Theorem 15.3 uses $m = 181{,}360$ bins, needs $n > 293$, and gives the bound 0.162. The clipped estimator on $[-R, R]$ alone has the bound 3,107.

The bounds are conservative. Figure 15.1 compares the simulated errors with each other.

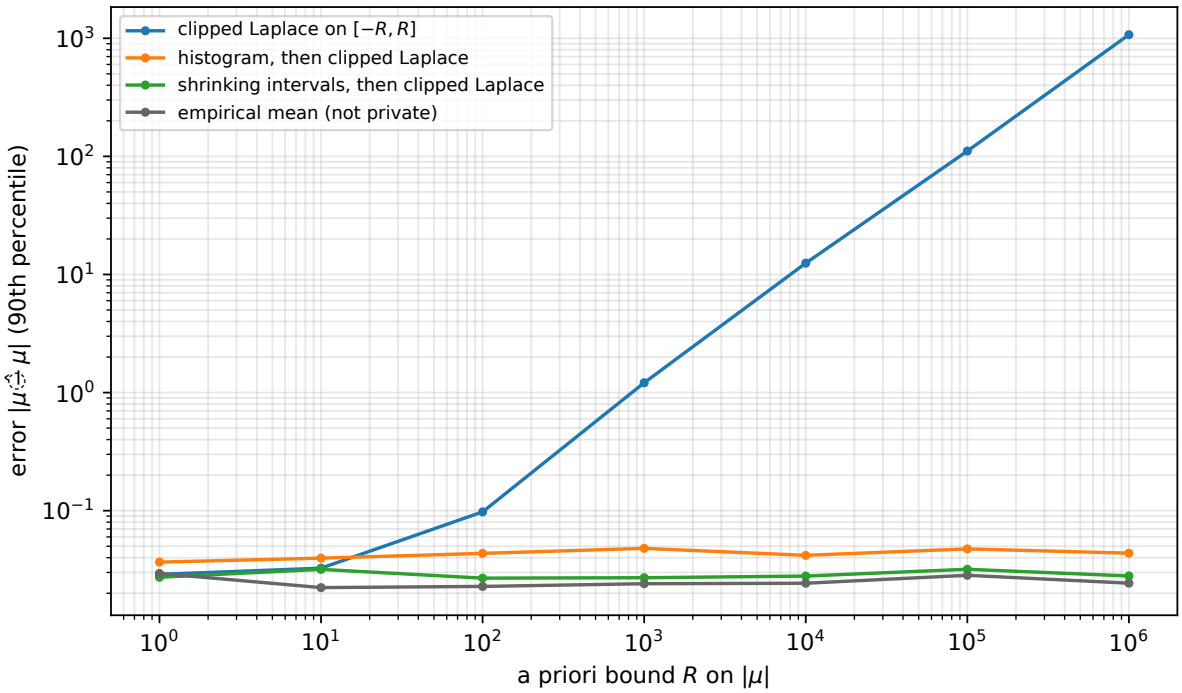


Figure 15.1.: Error of three ε-DP estimators of the mean of $\mathcal{N}(\mu, 1)$ with $\mu = R/3$, as the a priori bound R grows, for $n = 4{,}000$, $\varepsilon = 1$ and $\beta = 0.002$. Each point is the 90th percentile of $|\hat{\mu} - \mu|$ over 100 runs. The clipped estimator on $[-R, R]$ degrades linearly in R. The two estimators that first locate the data stay within a factor of about two of the non-private empirical mean for every R up to 10^6. At $R = 10^6$, over 300 runs, the 90th percentiles are 0.0250 without privacy, 0.0461 for the histogram and 0.0295 for shrinking intervals.

15.7. Many dimensions

The shrinking argument used three ingredients, and each has an analogue beyond one dimension. The recipe keeps a confidence set for the parameter, starting from the prior knowledge, and repeats three moves.

1. Clip the data to the confidence set, widened by a tail bound for the class of distributions.
2. Compute the empirical estimate of the parameter, and add noise scaled to its sensitivity, which the clipping bounds.

3. Replace the confidence set by one centred at the noisy estimate, whose size accounts for the sampling error and the noise.

For a Gaussian $\mathcal{N}(\mu, I_d)$ in $\mathbb{R}^d$ with $\|\mu\|_2 \le R$, the confidence set is a Euclidean ball, and the noise is Gaussian. Two facts replace those used above. The Gaussian mechanism adds $\mathcal{N}(0, \sigma^2 I_d)$ to a function f whose ℓ_2 sensitivity is $\Delta_2 f$. For every $\varepsilon > 0$ and $\delta \in (0,1)$ it is (ε, δ)-DP with $\sigma = \Delta_2 f \sqrt{2\ln(1/\delta) + \varepsilon}/\varepsilon$. Approximate DP also composes by adding the ε's and the δ's. Second, a standard Gaussian vector $Y \sim \mathcal{N}(0, I_d)$ has norm about $\sqrt{d}$, with

$$\Pr\left[\|Y\|_2 \ge \sqrt{d} + t\right] \le e^{-t^2/2} \qquad (t \ge 0). \tag{15.22}$$

This follows from the Gaussian concentration inequality for 1-Lipschitz functions (Boucheron, Lugosi and Massart, 2013, Theorem 5.6), applied to the norm, together with $\mathbb{E}\|Y\|_2 \le \sqrt{\mathbb{E}\|Y\|_2^2} = \sqrt{d}$. So a Gaussian sample sits on a shell at distance about $\sqrt{d}$ from its mean, even though the density is largest at the mean.

Clipping to the ball of centre c and radius ρ maps x to $c + (x - c)\min(1, \rho/\|x - c\|_2)$, which moves no point already inside the ball. Two radii play the roles that τ played in one dimension. They are

$$\tau_d = \sqrt{d} + \sqrt{2\ln(n/\beta)}, \qquad L = \sqrt{d} + \sqrt{2\ln(1/\beta)}. \tag{15.23}$$

The first is the amount by which each ball is widened before clipping. By Equation 15.22 and a union bound, every record lies within τ_d of μ with probability at least $1 - \beta$. The second is a radius that a single standard Gaussian vector stays within with probability at least $1 - \beta$, and it bounds both the sampling error and the noise.

The theorem runs the shrinking recipe with balls in place of intervals and the Gaussian mechanism in place of the Laplace mechanism.

Theorem 15.5 (Multivariate Gaussian mean)**.** *Let $\varepsilon > 0$ and $\delta \in (0,1)$. A step on the ball $B(c, r)$ with budget (ε', δ') clips every record to $B(c, r + \tau_d)$, averages, and adds $\mathcal{N}(0, \sigma^2 I_d)$ with $\sigma = 2(r + \tau_d)\sqrt{2\ln(1/\delta') + \varepsilon'}/(\varepsilon' n)$. Its new radius is*

$$s(r) = \frac{L}{\sqrt{n}} + \sigma L. \tag{15.24}$$

Run $T = \lceil \log_2(R/\tau_d) \rceil$ steps with $(\varepsilon/(2T), \delta/(2T))$ from $B(0, R)$, then one more with $(\varepsilon/2, \delta/2)$, and output its noisy mean $\hat{\mu}$. This is (ε, δ)-DP. If $X_1, \dots, X_n$ are independent $\mathcal{N}(\mu, I_d)$ draws with $\|\mu\|_2 \le R$, and

$$n \ge 16, \qquad n \ge \frac{32\, T L \sqrt{2\ln(2T/\delta) + \varepsilon}}{\varepsilon}, \tag{15.25}$$

then with probability at least $1 - (T+3)\beta$,

$$\|\hat{\mu} - \mu\|_2 \le \frac{L}{\sqrt{n}} + \frac{8\, \tau_d L \sqrt{2\ln(2/\delta) + \varepsilon}}{\varepsilon n}. \tag{15.26}$$

Proof. The proof follows that of Theorem 15.4, with these changes.

1. *Privacy.* The clipped mean has ℓ_2 sensitivity $2(r+\tau_d)/n$, so each step is the Gaussian mechanism. The radii are public, and composition adds the budgets to (ε, δ).
2. *No clipping.* By Equation 15.22 with $t = \sqrt{2\ln(n/\beta)}$ and a union bound, every $\|X_i - \mu\|_2 \le \tau_d$ with probability at least $1-\beta$. If $\|\mu - c\|_2 \le r$, the ball $B(c, r+\tau_d)$ then contains every record.
3. *Sampling.* The empirical mean satisfies $\tilde{\mu} - \mu = Y/\sqrt{n}$ with $Y \sim \mathcal{N}(0, I_d)$, so Equation 15.22 gives $\|\tilde{\mu} - \mu\|_2 \le L/\sqrt{n}$ with probability $1-\beta$.
4. *Noise.* The noise is $\sigma Y'$ with $Y' \sim \mathcal{N}(0, I_d)$, so its norm is at most σL with probability $1-\beta$. With steps 2 and 3, the new ball has radius $s(r)$ and contains μ.
5. *Halving.* Since $L \le \tau_d$ and $n \ge 16$, the first term of $s(r)$ is at most $\tau_d/4$. With $\varepsilon' = \varepsilon/(2T)$ and $\delta' = \delta/(2T)$, we have $2\ln(1/\delta') + \varepsilon' \le 2\ln(2T/\delta) + \varepsilon$, so $\sigma L \le 4T(r+\tau_d)L\sqrt{2\ln(2T/\delta)+\varepsilon}/(\varepsilon n)$, and the condition on n makes this at most $(r+\tau_d)/8$. So $s(r) \le \max(r, \tau_d)/2$, and after T steps the radius is at most τ_d.
6. *Final step.* With $r_T \le \tau_d$, the final noise, with budget $(\varepsilon/2, \delta/2)$, has $\sigma = 4(r_T + \tau_d)\sqrt{2\ln(2/\delta)+\varepsilon/2}/(\varepsilon n) \le 8\tau_d\sqrt{2\ln(2/\delta)+\varepsilon}/(\varepsilon n)$, and steps 3 and 4 give the bound. The failures add up as before.

□

Since L and τ_d are $\sqrt{d}$ up to logarithmic terms, error α with $\varepsilon \le 1$ needs

$$n = \tilde{O}\left(\frac{d}{\alpha^2} + \frac{d\sqrt{\ln(1/\delta)}}{\alpha\varepsilon} + \frac{\sqrt{d}\,\ln R\,\sqrt{\ln(1/\delta)}}{\varepsilon}\right). \tag{15.27}$$

The first two terms are d times the one-dimensional ones. The cost of locating the data grows only like $\sqrt{d}$, because each step's noise is compared with a radius that already includes the $\sqrt{d}$ spread of the data. The move to approximate DP is natural here, because clipping to a ball bounds the ℓ_2 sensitivity, which the Gaussian mechanism uses. The composition over the T steps can use advanced composition instead, which replaces the factor T in the condition on n by roughly $\sqrt{T\ln(1/\delta)}$. The same recipe also estimates the covariance matrix of a Gaussian, with confidence sets of matrices.

15.8. Heavy-tailed data

Everything so far used only that the data has light tails. The arguments therefore hold for **sub-Gaussian** distributions, whose tails decay at least as fast as a Gaussian's, with τ multiplied by the scale of the tails. Real data is often heavier tailed, and then only a few moments may be bounded. Take the class of distributions on $\mathbb{R}$ with

$$\mu = \mathbb{E}[X] \in [-1, 1], \qquad \mathrm{Var}(X) \le 1. \tag{15.28}$$

The narrow range for μ keeps the example simple, and the problem is to estimate μ to a precision α well below 1.

The previous strategy clips only where no point lands. Chebyshev's inequality gives $\Pr[|X-\mu| \ge \sqrt{n/\beta}] \le \beta/n$, so clipping to the interval $[-1-C, 1+C]$, with the clipping norm C set to $\sqrt{n/\beta}$, clips no point with probability $1-\beta$. But the noise scale is then $2(1+C)/(\varepsilon n)$, which is more than $2/(\varepsilon\sqrt{\beta n})$. The tail bound on the noise error is this scale times $\ln(1/\beta)$, and it is at most α only if

$$n \ge \frac{4\ln^2(1/\beta)}{\beta\,\alpha^2\varepsilon^2}, \tag{15.29}$$

which is worse than the Gaussian rate in both α and ε. The remedy is to clip more aggressively and accept a bias. The clipping bias of Equation 15.14 is no longer zero, but it can be bounded.

Lemma 15.5 (Clipping bias)**.** *Let X have mean μ and variance at most σ_X^2, and let $[u,v] \supseteq [\mu-C, \mu+C]$ with $C>0$. Then $g = g_{[u,v]}$ satisfies*

$$\big|\mathbb{E}[g(X)] - \mu\big| \le \frac{\sigma_X^2}{C}, \qquad \mathrm{Var}(g(X)) \le \sigma_X^2. \tag{15.30}$$

Proof.

1. Since $\mu = \mathbb{E}[X]$, $|\mathbb{E}[g(X)] - \mu| = |\mathbb{E}[g(X) - X]| \le \mathbb{E}|g(X) - X|$.
2. If $|X-\mu| \le C$, then $X \in [u,v]$ and $g(X) = X$. Otherwise X moves towards μ and does not pass it, because $\mu \in [u,v]$, so $|g(X)-X| \le |X-\mu|$.
3. So $|g(X)-X| \le |X-\mu|\,\mathbf{1}[|X-\mu| > C] \le (X-\mu)^2/C$, and taking expectations gives the first bound.
4. Clipping is 1-Lipschitz. With X' an independent copy of X, $\mathrm{Var}(g(X)) = \frac{1}{2}\mathbb{E}[(g(X) - g(X'))^2] \le \frac{1}{2}\mathbb{E}[(X-X')^2] = \mathrm{Var}(X)$.

□

With the bias bounded, the clipping interval no longer has to catch every point. The next estimator clips at a distance C from the range of possible means and lets the three terms of Equation 15.14 decide how large C should be.

Theorem 15.6 (Heavy-tailed mean)**.** *For $C>0$, the estimator*

$$\hat{\mu} = \frac{1}{n}\sum_{i=1}^{n} g_{[-1-C,\,1+C]}(X_i) + \mathrm{Lap}\!\left(\frac{2(1+C)}{\varepsilon n}\right) \tag{15.31}$$

is ε-DP. If $X_1, \dots, X_n$ are independent draws from a distribution with mean $\mu \in [-1,1]$ and variance at most 1, then with probability at least $1-2\beta$,

$$|\hat{\mu} - \mu| \le \frac{1}{C} + \frac{1}{\sqrt{\beta n}} + \frac{2(1+C)\ln(1/\beta)}{\varepsilon n}. \tag{15.32}$$

Proof.

1. Privacy is Lemma 15.1 with an interval of length $2(1 + C)$.
2. Since $|\mu| \le 1$, the interval $[-1 - C, 1 + C]$ contains $[\mu - C, \mu + C]$. Write $\nu = \mathbb{E}[g(X_1)]$ and $\bar{g}$ for the clipped mean. Then $\hat{\mu} - \mu = (\nu - \mu) + (\bar{g} - \nu) + Z$, with Z the noise.
3. *Bias.* Lemma 15.5 with $\sigma_X^2 = 1$ gives $|\nu - \mu| \le 1/C$.
4. *Sampling.* By Lemma 15.5, $\mathrm{Var}(\bar{g}) \le 1/n$, so Chebyshev's inequality gives $|\bar{g} - \nu| \le 1/\sqrt{\beta n}$ with probability at least $1 - \beta$.
5. *Noise.* The Laplace tail gives $|Z| \le 2(1 + C)\ln(1/\beta)/(\varepsilon n)$ with probability $1 - \beta$. A union bound finishes the proof.

□

The bias falls as C grows and the noise rises, so there is a best clipping norm. Setting each of the three terms of Equation 15.32 to at most $\alpha/3$ gives a clean sample complexity.

Corollary 15.1. *Let $\alpha \in (0, 1]$, and let $\hat{\mu}$ be the estimator of Theorem 15.6 with $C = 3/\alpha$. If*

$$n \ge \frac{9}{\beta\alpha^2} \qquad \text{and} \qquad n \ge \frac{24\ln(1/\beta)}{\alpha^2\varepsilon}, \tag{15.33}$$

then $|\hat{\mu} - \mu| \le \alpha$ with probability at least $1 - 2\beta$.

Proof.

1. The bias term is $1/C = \alpha/3$, and the first condition makes the sampling term at most $\alpha/3$.
2. Since $\alpha \le 1$, $1 + C \le 4/\alpha$, so the noise term is at most $8\ln(1/\beta)/(\alpha\varepsilon n)$, which the second condition makes at most $\alpha/3$.

□

So the clipping norm should be about $1/\alpha$, and $n = O(1/(\alpha^2\varepsilon))$ samples suffice for a fixed β and $\varepsilon \le 1$. This is more than the $O(1/\alpha^2 + 1/(\alpha\varepsilon))$ of the Gaussian case, because the two terms now multiply instead of adding. The increase is necessary. For distributions with only a bounded variance, no ε-DP estimator does better than error of order $1/\sqrt{\varepsilon n}$, by a lower bound of Barber and Duchi (2014). Knowing less about the distribution costs more samples. In terms of n, the clipping norm that minimizes $1/C + 2C\ln(1/\beta)/(\varepsilon n)$ is

$$C^* = \sqrt{\frac{\varepsilon n}{2\ln(1/\beta)}}, \tag{15.34}$$

and with $n = 4{,}000$, $\varepsilon = 1$ and $\beta = 0.05$ it is $C^* = 25.8$, where Equation 15.32 equals 0.150. Figure 15.2 shows the trade-off.

15.9. Summary

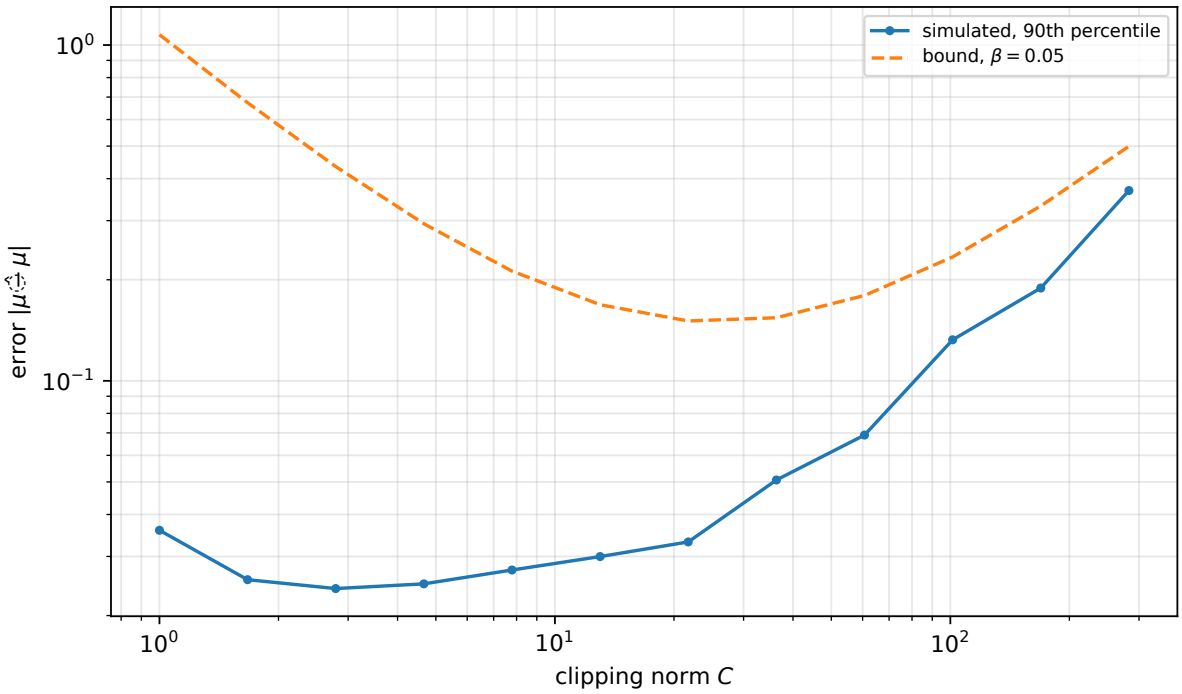


Figure 15.2.: Error of the clipped estimator of Theorem 15.6 on a Student t distribution with 3 degrees of freedom, scaled to variance 1, with $\mu = 0.5$, $n = 4{,}000$ and $\varepsilon = 1$, as the clipping norm C varies. The solid line is the 90th percentile of $|\hat{\mu} - \mu|$ over 300 runs, and the dashed line is the bound Equation 15.32 with $\beta = 0.05$. Both rise steeply at large C, where the noise dominates. The simulated error is smallest at a smaller C than the bound's, because this distribution's clipping bias is far below the worst case $1/C$ over all distributions with variance 1.

Table 15.1.: The chapter's results.

Result	Statement
Theorem 15.1	Bernoulli mean with error $\sqrt{\ln(4/\beta)/(2n)} + \ln(2/\beta)/(\varepsilon n)$
Proposition 15.1	no ε-DP algorithm estimates the mean of every real dataset to finite accuracy
Lemma 15.1	clipping to a public $[u, v]$ gives sensitivity $(v-u)/n$
Proposition 15.2	estimating a Gaussian mean with $\lvert\mu\rvert \le R$ needs $n \ge \ln((1-\beta)m)/\varepsilon$, with m about $2R/(3\alpha)$
Theorem 15.2	clipping to $[-R-\tau, R+\tau]$ needs $n = \tilde{O}(1/\alpha^2 + R/(\alpha\varepsilon))$
Lemma 15.3	a noisy histogram locates μ within 2τ once $n > (8/\varepsilon)\ln(m/\beta)$
Theorem 15.3	histogram, then clipping, needs $n = \tilde{O}(1/\alpha^2 + 1/(\alpha\varepsilon) + \ln R/\varepsilon)$
Lemma 15.4	one clipped step halves the confidence interval down to the floor τ
Theorem 15.4	$T = \lceil \log_2(R/\tau) \rceil$ shrinking steps, then clipping, reach the same rate
Theorem 15.5	in $\mathbb{R}^d$, $n = \tilde{O}(d/\alpha^2 + d/(\alpha\varepsilon) + \sqrt{d}\ln R/\varepsilon)$ under (ε, δ)-DP, up to factors in $\ln(1/\delta)$
Lemma 15.5	clipping at distance C from the mean biases it by at most σ_X^2/C
Corollary 15.1	with variance 1 and $C = 3/\alpha$, $n = O(1/(\alpha^2\varepsilon))$ samples suffice

15.10. References

- Cynthia Dwork, Frank McSherry, Kobbi Nissim and Adam Smith, "Calibrating Noise to Sensitivity in Private Data Analysis", 2006. The Laplace mechanism and the sensitivity of the mean of bounded data.
- Vishesh Karwa and Salil Vadhan, "Finite Sample Differentially Private Confidence Intervals", 2018. The histogram approach to the Gaussian mean with an unknown location, and lower bounds that show the rate Equation 15.18 is optimal up to logarithmic factors.
- Cynthia Dwork and Aaron Roth, *The Algorithmic Foundations of Differential Privacy*, 2014. The Laplace and Gaussian mechanisms and the composition theorems used in this chapter.
- Stéphane Boucheron, Gábor Lugosi and Pascal Massart, *Concentration Inequalities*, 2013. Theorem 5.6 is the Gaussian concentration inequality behind Equation 15.22.

- Rina Foygel Barber and John Duchi, "Privacy and Statistical Risk: Formalisms and Minimax Bounds", 2014. Minimax lower bounds for private mean estimation under moment assumptions. Their Proposition 2 bounds the squared error below by order $1/n+(\varepsilon n)^{-2(k-1)/k}$ for a bounded k-th moment, which for bounded variance ($k = 2$) gives the $1/\sqrt{\varepsilon n}$ rate in absolute error.

16. Adaptive Data Analysis

A lab spends years collecting a dataset and tests a hypothesis on it, which fails. It tests a second hypothesis, then a third, each suggested by what the earlier tests showed, until one comes out significant. Other labs then fail to reproduce the finding on fresh data. Every test may have been sound on its own, but the sequence used the same data both to choose the questions and to answer them, and the answers ended up describing the sample instead of the population.

This chapter makes that failure precise for averages of predicates, called statistical queries, and shows how differential privacy (DP) prevents it. Answers to queries fixed in advance generalize with a sample size logarithmic in the number k of queries. Adaptively chosen queries can overfit unless the sample size is of order k. A mechanism that answers through a DP algorithm needs only about $\sqrt{k}$, because a private algorithm cannot find a query on which the sample and the population disagree.

16.1. Statistical queries

The data is a sample from an unknown distribution $\mathcal{D}$ on a universe $\mathcal{X}$, and the analyst cares about the distribution, not the sample. The questions the analyst asks are the simplest kind, the fraction of records with a yes-or-no property, and each question has a value on the sample and a value on the distribution.

Definition 16.1 (Statistical query)**.** A **statistical query** is a function $q\colon \mathcal{X} \to \{0,1\}$. Its value on the distribution and its empirical value on a dataset $D = (X_1, \dots, X_n)$ are

$$q(\mathcal{D}) = \mathbb{E}_{x\sim\mathcal{D}}[q(x)], \qquad q(D) = \frac{1}{n}\sum_{j=1}^{n} q(X_j). \tag{16.1}$$

So $q(\mathcal{D})$ is the fraction of the population with the property q, such as having red hair, and $q(D)$ is the fraction in the sample. These are the counting queries of earlier chapters, normalized by n. If the records are drawn independently from $\mathcal{D}$, the empirical value is an unbiased estimate of the population value. A fixed list of queries can be answered this way with a sample whose size is only logarithmic in the length of the list.

Proposition 16.1. *Let $\alpha > 0$ be the allowed error and $\beta \in (0,1)$ the allowed failure probability. Let $q_1, \dots, q_k$ be statistical queries fixed before the sample is drawn, and let $X_1, \dots, X_n$ be independent draws from $\mathcal{D}$. If*

$$n \geq \frac{\ln(2k/\beta)}{2\alpha^2}, \tag{16.2}$$

then some query has error at least α with probability at most β,

$$\Pr\left[\exists i \colon |q_i(D) - q_i(\mathcal{D})| \geq \alpha\right] \leq \beta. \tag{16.3}$$

Proof.

1. For a fixed i, the values $q_i(X_1), \dots, q_i(X_n)$ are independent, lie in $[0,1]$, and have mean $q_i(\mathcal{D})$. Hoeffding's inequality gives $\Pr[|q_i(D) - q_i(\mathcal{D})| \geq \alpha] \leq 2e^{-2\alpha^2 n}$.
2. A union bound over the k queries gives failure probability at most $2ke^{-2\alpha^2 n}$, which the condition on n makes at most β.

□

Turned around, a sample of size n supports about $e^{2\alpha^2 n}$ queries. Step 1 needs the query to be independent of the sample, and that is where adaptivity breaks the argument.

16.2. Adaptivity breaks the bound

In an adaptive analysis, each query may depend on the answers to the earlier ones. The definition below sets up this interaction, and says what it means for the mechanism that answers to be private whatever the analyst does.

Definition 16.2 (Adaptive analysis)**.** An **analyst** $\mathcal{A}$ and a **mechanism** M interact for k rounds. In round i, the analyst chooses a statistical query q_i as a function of the earlier answers $a_1, \dots, a_{i-1}$ and its own random bits, called its coins. The mechanism returns an answer a_i, computed from the dataset D, its own coins and the history. The mechanism is (ε, δ)**-DP against adaptive analysts** if for every analyst $\mathcal{A}$ the map from D to the transcript $(q_1, a_1, \dots, q_k, a_k)$ is (ε, δ)-DP.

The empirical mechanism $a_i = q_i(D)$ answers every query exactly on the sample. Against an adaptive analyst it fails badly.

Proposition 16.2. *Let $\mathcal{X} = \{1, \dots, N\}$, let $\mathcal{D}$ be uniform on $\mathcal{X}$, and let the mechanism answer $a_i = q_i(D)$. Some analyst with $k = N + 1$ queries ends with a query q_k such that, for every sample,*

$$q_k(D) - q_k(\mathcal{D}) \geq 1 - \frac{n}{k-1}. \tag{16.4}$$

Proof.

1. For $i = 1, \dots, N$, the analyst asks $q_i(x) = \mathbf{1}[x = i]$. The answer a_i is the fraction of records equal to i, so after N rounds the analyst knows the set $S = \{X_1, \dots, X_n\}$ of values in the sample.
2. The last query is $q_k(x) = \mathbf{1}[x \in S]$. Every record lies in S, so $q_k(D) = 1$.
3. The set S has at most n elements, so $q_k(\mathcal{D}) = |S|/N \leq n/N = n/(k-1)$.

□

The first N queries are harmless one at a time, and each could have been fixed in advance. The last one is built from their answers, and on it the sample and the population disagree by almost 1 unless n is of order k. Compared with Proposition 16.1, the number of queries a sample supports drops from exponential in n to linear in n.

A simple defence is to give each query fresh data. The analyst can still choose adaptively, but each answer comes from records that no earlier answer has seen.

Proposition 16.3 (Sample splitting)**.** *Let k divide n. Split D into k disjoint parts $D^{(1)}, \dots, D^{(k)}$ of n/k records each, and answer $a_i = q_i(D^{(i)})$. If the records are independent draws from $\mathcal{D}$ and*

$$n \geq \frac{k \ln(2k/\beta)}{2\alpha^2}, \tag{16.5}$$

then for every analyst $\Pr\left[\exists i : |a_i - q_i(\mathcal{D})| \geq \alpha\right] \leq \beta$.

Proof.

1. The query q_i depends only on the analyst's coins and on $a_1, \dots, a_{i-1}$, which depend only on the parts $D^{(1)}, \dots, D^{(i-1)}$. So q_i is independent of $D^{(i)}$.
2. Conditioned on q_i, the part $D^{(i)}$ is still n/k independent draws from $\mathcal{D}$. Step 1 of the proof of Proposition 16.1 gives $\Pr[|a_i - q_i(\mathcal{D})| \geq \alpha] \leq 2e^{-2\alpha^2 n/k}$.
3. A union bound over i gives at most $2ke^{-2\alpha^2 n/k}$, which the condition makes at most β.

□

Sample splitting is valid but expensive, because it needs k times the data of the non-adaptive case. Answering through a DP mechanism does much better. With Gaussian noise, a sample of size

$$n = O\left(\frac{\sqrt{k}\,\ln^{3/2}(k/(\alpha\beta))}{\alpha^2}\right) \tag{16.6}$$

suffices for error α (Corollary 16.1). Under a standard cryptographic assumption, no computationally efficient mechanism can do better than $\sqrt{k}$ in general. For a finite universe and a very large k, a different private algorithm does better still.

16.3. Privacy implies generalization

Overfitting needs the analyst to learn too much about individual records, as the analyst of Proposition 16.2 did. If every answer passes through a DP mechanism, no answer reveals much about one record, and the analyst can learn only about the population. A **transfer theorem** makes this precise. It turns accuracy on the sample and privacy into accuracy on the population.

Theorem 16.1 (Transfer theorem)**.** *Let $\alpha, \beta \in (0, 0.1)$, and let M answer k adaptively chosen statistical queries. Suppose that*

1. *the mechanism is* $(\alpha/64, \alpha\beta/32)$*-DP against adaptive analysts, and*
2. *the mechanism is accurate on the sample, that is, for every analyst and every dataset* D*,* $\Pr\left[\exists i \colon |a_i - q_i(D)| > \alpha/8\right] \le \alpha\beta/16$*, with probability over the coins of* M *and of the analyst.*

Then for every analyst and every distribution $\mathcal{D}$*, if the* n *records are independent draws from* $\mathcal{D}$*,*

$$\Pr\left[\exists i \colon |a_i - q_i(\mathcal{D})| > \alpha\right] \le \beta. \tag{16.7}$$

This is Theorem 3.4 of Bassily, Nissim, Smith, Steinke, Stemmer and Ullman (2016), for queries of sensitivity $1/n$. The failure probability on the sample must be of order $\alpha\beta$, not β, for the conclusion to hold with failure probability β. The proof is long, so only its plan is given here. It runs about $1/\beta$ independent copies of the interaction, each on its own sample. It then imagines a monitor, which knows $\mathcal{D}$ and exists only in the proof, and which picks the copy and the query with the largest error on the population. The monitor only post-processes the outputs of DP interactions, so its choice is itself DP. A version of Lemma 16.1 below for several samples shows that such a choice cannot find a large population error on average. Hence each copy errs with probability $O(\beta)$. A shorter proof of a stronger theorem is in Jung et al. (2020).

Two ingredients from earlier chapters instantiate the theorem. The Gaussian mechanism answers a query f with Gaussian noise, and its sensitivity Δf sets the standard deviation. With $\sigma = \Delta f \sqrt{2\ln(1/\delta_0) + \varepsilon_0}/\varepsilon_0$, the mechanism is $(\varepsilon_0, \delta_0)$-DP for every $\varepsilon_0 > 0$ and $\delta_0 \in (0, 1)$. A corollary of advanced composition says that, for $\varepsilon \in (0, 1]$ and $\delta' \in (0, 1/2]$, k adaptively chosen mechanisms that are each $(\varepsilon_0, \delta_0)$-DP with

$$\varepsilon_0 = \frac{\varepsilon}{2\sqrt{2k\ln(1/\delta')}} \tag{16.8}$$

are, together, $(\varepsilon, k\delta_0 + \delta')$-DP. A statistical query has sensitivity $1/n$, because replacing one record changes one term of the average by at most 1.

The corollary below combines these ingredients. It adds Gaussian noise to each empirical answer, with a scale chosen so that the whole interaction meets both hypotheses of Theorem 16.1, and reads off the sample size that this needs.

Corollary 16.1 (Gaussian answers generalize)**.** *Let* $\alpha, \beta \in (0, 0.1)$*. Answer each query by* $a_i = q_i(D) + Z_i$ *with independent* $Z_i \sim \mathcal{N}(0, \sigma^2)$*, where*

$$\varepsilon_0 = \frac{\alpha}{128\sqrt{2k\ln(64/(\alpha\beta))}}, \qquad \sigma = \frac{\sqrt{2\ln(64k/(\alpha\beta)) + \varepsilon_0}}{\varepsilon_0\, n}. \tag{16.9}$$

If

$$n \ge \frac{2048\sqrt{3}}{\alpha^2}\sqrt{k\ln\frac{64}{\alpha\beta}}\ \ln\frac{64k}{\alpha\beta}, \tag{16.10}$$

then for every analyst, $\Pr\left[\exists i \colon |a_i - q_i(\mathcal{D})| > \alpha\right] \le \beta$*.*

Proof.

1. *Each answer is private.* With $\delta_0 = \alpha\beta/(64k)$, the noise is the Gaussian mechanism's for sensitivity $1/n$ and budget $(\varepsilon_0, \delta_0)$. For a fixed analyst, the i-th answer is this mechanism applied to a query fixed by the earlier answers and the analyst's coins.
2. *All answers are* $(\alpha/64, \alpha\beta/32)$*-DP.* Apply the corollary of advanced composition with $\varepsilon = \alpha/64$ and $\delta' = \alpha\beta/64$. The total δ is $k\delta_0 + \delta' = \alpha\beta/32$. The queries in the transcript are a post-processing of the answers and the analyst's coins, so the transcript has the same guarantee.
3. *Accuracy on the sample.* The error $a_i - q_i(D)$ is the noise Z_i, and a Gaussian variable exceeds t standard deviations in absolute value with probability at most $2e^{-t^2/2}$. With $t = \sqrt{2\ln(32k/(\alpha\beta))}$, each $|Z_i|$ exceeds $t\sigma$ with probability at most $\alpha\beta/(16k)$, so by a union bound all k answers are within $t\sigma$ of the sample values with probability at least $1 - \alpha\beta/16$.
4. *The threshold.* Write $\ell = \ln(64k/(\alpha\beta))$. Since $\varepsilon_0 < 1 < \ell$, we have $2\ell + \varepsilon_0 \le 3\ell$, and $t \le \sqrt{2\ell}$. So

$$t\sigma \le \frac{\sqrt{6}\,\ell}{\varepsilon_0 n} = \frac{256\sqrt{3}}{\alpha n}\sqrt{k \ln \frac{64}{\alpha\beta}}\;\ell, \tag{16.11}$$

which Equation 16.10 makes at most $\alpha/8$. Steps 2 to 4 are the hypotheses of Theorem 16.1, which gives the claim. □

The $\sqrt{k}$ in Equation 16.10 is the $\sqrt{k}$ of advanced composition. The constant is large because the transfer theorem's constants are, but the number of queries a sample supports grows from linear in n, for sample splitting, to quadratic in n. Other private mechanisms give other trade-offs. For a finite universe, answering with private multiplicative weights gives, by an analogous transfer theorem for the average error over the queries (Bassily et al., 2016, Corollary 6.3), a sample size of order

$$n = \frac{\sqrt{\ln|\mathcal{X}|}\,\ln k\,\ln^{3/2}(1/(\alpha\beta))}{\alpha^3}, \tag{16.12}$$

which is only logarithmic in k, and better when k is very large compared with $\ln|\mathcal{X}|$.

16.4. An easier transfer theorem

A weaker transfer theorem, with accuracy measured in expectation instead of with high probability, has a short complete proof. Its hypothesis bounds the expected worst error on the sample, and its conclusion bounds the expected worst error on the population.

Theorem 16.2 (Transfer in expectation)**.** *Let M be (ε, δ)-DP against adaptive analysts. Suppose that for every analyst, if the records are independent draws from $\mathcal{D}$, then*

$$\mathbb{E}\Big[\max_{i \le k} |a_i - q_i(D)|\Big] \le \alpha. \tag{16.13}$$

Then for every analyst,

$$\mathbb{E}\Big[\max_{i\le k}|a_i - q_i(\mathcal{D})|\Big] \le \alpha + e^{\varepsilon} - 1 + \delta. \tag{16.14}$$

For small ε, the excess is about $\varepsilon + \delta$. Most of the work is in a lemma about one query. It says that a query chosen by a DP algorithm takes about the same value on the sample as on the population, on average. The last query of Proposition 16.2 is exactly a query on which the two disagree, so a DP algorithm cannot produce it.

Lemma 16.1 (A private query generalizes)**.** *Let W be an (ε, δ)-DP algorithm that maps a dataset in $\mathcal{X}^n$ to a statistical query. Let $X_1, \dots, X_n$ be independent draws from $\mathcal{D}$ and $q = W(D)$. Then*

$$\big|\mathbb{E}[q(D)] - \mathbb{E}[q(\mathcal{D})]\big| \le e^{\varepsilon} - 1 + \delta, \tag{16.15}$$

with expectations over the sample and the coins of W.

Proof.

1. *Linearity.* Since q takes values in $\{0, 1\}$,

$$\mathbb{E}[q(D)] = \frac{1}{n}\sum_{i=1}^{n}\mathbb{E}[q(X_i)] = \frac{1}{n}\sum_{i=1}^{n}\Pr[q(X_i) = 1]. \tag{16.16}$$

2. *A fresh record.* Fix i, draw X_i' from $\mathcal{D}$ independently of everything else, and let $D^{(i)}$ be D with entry i replaced by X_i'. The datasets D and $D^{(i)}$ are neighbours, $D \sim D^{(i)}$, since they differ in one replaced record.
3. *Privacy.* Write $W(D)(X_i)$ for the value of the query $W(D)$ on the record X_i. For fixed values of D and X_i', the queries q' with $q'(X_i) = 1$ form an event about the output of W. DP for the neighbours D and $D^{(i)}$ gives

$$\Pr_W\big[W(D)(X_i) = 1\big] \le e^{\varepsilon}\Pr_W\big[W(D^{(i)})(X_i) = 1\big] + \delta. \tag{16.17}$$

Averaging over D and X_i' keeps the inequality. 4. *Swapping.* The pair $(D^{(i)}, X_i)$ has the same distribution as the pair (D, X_i'), since each is n independent draws from $\mathcal{D}$ plus one more independent draw. Hence $\Pr[W(D^{(i)})(X_i) = 1] = \Pr[q(X_i') = 1]$. The record X_i' is independent of $q = W(D)$, so this probability is $\mathbb{E}[q(\mathcal{D})]$. 5. *One direction.* Steps 1, 3 and 4 give $\mathbb{E}[q(D)] \le e^{\varepsilon}\,\mathbb{E}[q(\mathcal{D})] + \delta$. Since $\mathbb{E}[q(\mathcal{D})] \le 1$,

$$\mathbb{E}[q(D)] - \mathbb{E}[q(\mathcal{D})] \le (e^{\varepsilon} - 1)\,\mathbb{E}[q(\mathcal{D})] + \delta \le e^{\varepsilon} - 1 + \delta. \tag{16.18}$$

6. *Other direction.* DP applied from $D^{(i)}$ to D gives $\Pr[W(D^{(i)})(X_i) = 1] \le e^{\varepsilon}\Pr[W(D)(X_i) = 1] + \delta$. With steps 1 and 4, this is $\mathbb{E}[q(\mathcal{D})] \le e^{\varepsilon}\,\mathbb{E}[q(D)] + \delta$, and the same argument as in step 5 bounds $\mathbb{E}[q(\mathcal{D})] - \mathbb{E}[q(D)]$ by $e^{\varepsilon} - 1 + \delta$.

□

The proof of Theorem 16.2 lets a private algorithm pick out the query with the worst error. Two tricks make that query fit the lemma.

Proof.

1. *Removing the absolute value.* Whenever the analyst asks q_i and receives a_i, add the query $1 - q_i$ with the answer $1 - a_i$. The extended transcript is a post-processing of the original, so it is still (ε, δ)-DP. The errors on the sample are unchanged, because $|(1 - a_i) - (1 - q_i)(D)| = |a_i - q_i(D)|$. Since $(1 - q_i)(\mathcal{D}) - (1 - a_i) = -(q_i(\mathcal{D}) - a_i)$, the $2k$ queries q_j satisfy

$$\max_{i \le k} |q_i(\mathcal{D}) - a_i| = \max_{j \le 2k} (q_j(\mathcal{D}) - a_j). \tag{16.19}$$

2. *The selector.* Let W run the interaction on D and output the query $q^* = q_{j^*}$, with answer $a^* = a_{j^*}$, where j^* maximizes $q_j(\mathcal{D}) - a_j$. The selector uses $\mathcal{D}$, which is fixed and is not the data, so W is a post-processing of the transcript and is (ε, δ)-DP. It is a device of the proof and is never run.
3. *Splitting the error.* By linearity of expectation,

$$\mathbb{E}\Big[\max_i |q_i(\mathcal{D}) - a_i|\Big] = \mathbb{E}\big[q^*(\mathcal{D}) - q^*(D)\big] + \mathbb{E}\big[q^*(D) - a^*\big]. \tag{16.20}$$

4. *The two terms.* Lemma 16.1, applied to W, bounds the first term by $e^{\varepsilon} - 1 + \delta$. The second is at most $\mathbb{E}[\max_j |q_j(D) - a_j|] \le \alpha$, by the accuracy hypothesis and step 1.

□

The Gaussian mechanism instantiates this theorem too, with a proof that needs nothing cited. The only change from Corollary 16.1 is that the error on the sample is bounded in expectation.

Corollary 16.2. *Let $\alpha \in (0, 1]$. Answer each query by $a_i = q_i(D) + Z_i$ with independent $Z_i \sim \mathcal{N}(0, \sigma^2)$, where*

$$\varepsilon_0 = \frac{\alpha}{2\sqrt{2k \ln(2/\alpha)}}, \qquad \sigma = \frac{\sqrt{2 \ln(2k/\alpha) + \varepsilon_0}}{\varepsilon_0 \, n}. \tag{16.21}$$

If

$$n \ge \frac{4\sqrt{3}}{\alpha^2} \sqrt{k \ln \frac{2}{\alpha}} \; \ln \frac{2k}{\alpha}, \tag{16.22}$$

then for every analyst, $\mathbb{E}[\max_i |a_i - q_i(\mathcal{D})|] \le 4\alpha$.

Proof.

1. *Privacy.* As in steps 1 and 2 of the proof of Corollary 16.1, each answer is the Gaussian mechanism with budget $(\varepsilon_0, \alpha/(2k))$, and the corollary of advanced composition with $\varepsilon = \alpha$ and $\delta' = \alpha/2$ makes the mechanism (α, α)-DP against adaptive analysts.
2. *Expected worst noise.* For $\lambda > 0$, Jensen's inequality and $e^{\lambda \max_i |Z_i|} \le \sum_i e^{\lambda |Z_i|}$ give

$$e^{\lambda\, \mathbb{E}[\max_i |Z_i|]} \le \sum_{i=1}^{k} \mathbb{E}\big[e^{\lambda |Z_i|}\big] \le 2k\, e^{\lambda^2 \sigma^2/2}, \tag{16.23}$$

because $\mathbb{E}[e^{\lambda|Z|}] \le \mathbb{E}[e^{\lambda Z}] + \mathbb{E}[e^{-\lambda Z}] = 2e^{\lambda^2\sigma^2/2}$. Taking logarithms with $\lambda = \sqrt{2\ln(2k)}/\sigma$ gives $\mathbb{E}[\max_i |Z_i|] \le \sigma\sqrt{2\ln(2k)}$. 3. *The threshold.* Write $\ell = \ln(2k/\alpha) \ge \ln 2$. Since $\varepsilon_0 \le 1/(2\sqrt{2\ln 2}) < \ln 2 \le \ell$, we have $2\ell + \varepsilon_0 \le 3\ell$, and $\ln(2k) \le \ell$. So

$$\sigma\sqrt{2\ln(2k)} \le \frac{\sqrt{6}\,\ell}{\varepsilon_0 n} = \frac{4\sqrt{3}}{\alpha n}\sqrt{k \ln\frac{2}{\alpha}}\;\ell, \tag{16.24}$$

which the condition on n makes at most α. Hence the expected worst error on the sample is at most α. 4. Theorem 16.2 with $\varepsilon = \delta = \alpha$ gives at most $\alpha + (e^\alpha - 1) + \alpha \le 4\alpha$, since $e^\alpha - 1 \le 2\alpha$ for $\alpha \le 1$. □

16.5. Summary

Table 16.1.: The chapter's results.

Result	Statement
Proposition 16.1	k fixed queries generalize once $n \ge \ln(2k/\beta)/(2\alpha^2)$
Proposition 16.2	exact answers to k adaptive queries can be off by $1 - n/(k-1)$ on the population
Proposition 16.3	sample splitting generalizes once $n \ge k\ln(2k/\beta)/(2\alpha^2)$
Theorem 16.1	an $(\alpha/64, \alpha\beta/32)$-DP mechanism with error $\alpha/8$ on the sample, except with probability $\alpha\beta/16$, has error α on the population, except with probability β
Corollary 16.1	Gaussian answers generalize with n of order $\sqrt{k}$ up to logarithmic factors
Lemma 16.1	a query chosen by an (ε, δ)-DP algorithm has $\lvert\mathbb{E}[q(D)] - \mathbb{E}[q(\mathcal{D})]\rvert \le e^\varepsilon - 1 + \delta$
Theorem 16.2	expected worst error on the population is at most $\alpha + e^\varepsilon - 1 + \delta$
Corollary 16.2	Gaussian answers give expected worst error 4α with n of order $\sqrt{k}$

16.6. References

- John P. A. Ioannidis, "Why Most Published Research Findings Are False", 2005. The reproducibility problem that motivates the chapter.
- Cynthia Dwork, Vitaly Feldman, Moritz Hardt, Toniann Pitassi, Omer Reingold and Aaron Roth, "Preserving Statistical Validity in Adaptive Data Analysis", 2015. The first use of differential privacy to answer adaptively chosen queries with guarantees on the population.
- Raef Bassily, Kobbi Nissim, Adam Smith, Thomas Steinke, Uri Stemmer and Jonathan Ullman, "Algorithmic Stability for Adaptive Data Analysis", 2016. The transfer theorem with high probability (Theorem 3.4), the bound in expectation of Lemma 16.1 (Lemma 3.1), and instantiations with the Gaussian mechanism and private multiplicative weights.
- Christopher Jung, Katrina Ligett, Seth Neel, Aaron Roth, Saeed Sharifi-Malvajerdi and Moshe Shenfeld, "A New Analysis of Differential Privacy's Generalization Guarantees", 2020. A simpler and stronger transfer theorem.
- Moritz Hardt and Jonathan Ullman, "Preventing False Discovery in Interactive Data Analysis Is Hard", 2014, and Thomas Steinke and Jonathan Ullman, "Interactive Fingerprinting Codes and the Hardness of Preventing False Discovery", 2015. Assuming one-way functions with exponential hardness exist, computationally efficient mechanisms cannot answer many more than n^2 adaptive queries in general, so the $\sqrt{k}$ dependence is optimal for them.
- Cynthia Dwork, Guy N. Rothblum and Salil Vadhan, "Boosting and Differential Privacy", 2010. Advanced composition.
- Cynthia Dwork and Aaron Roth, *The Algorithmic Foundations of Differential Privacy*, 2014. The Gaussian mechanism and advanced composition (Theorem 3.20).

Part V.

Deployments

The earlier parts built differential privacy (DP) as mathematics, with a trusted curator who holds the data and releases noisy answers. This part asks what happens when the definition meets real systems, with millions of users, daily data collection and analysts outside the company. The deployments studied here are not just applications of earlier results. Each forced a new idea, and most of those ideas turned into theory of their own.

The first chapter removes the trusted curator. In the local model each person randomizes their own data before sending it, so no one ever sees a raw record. The price is an error about $\sqrt{n}$ times larger than in the central model for n people, and a lower bound shows that no noninteractive protocol can avoid it. Google, Apple and Microsoft accepted this price because they have hundreds of millions of users, and each solved a practical problem along the way. They had to handle questions repeated every day, find values that no one knew to ask about, and send numbers in a single bit.

The second chapter looks for a middle ground and then returns to the central model. A shuffler that only mixes messages lets each person add far less noise, and a sum of bits can then be estimated with an error that does not grow with n. A query service for advertisers brings a sharper analysis of the exponential mechanism and a fast way to select the top items of a histogram. Two large public data releases show how deployments choose what a neighbouring dataset is, and how they account for privacy with zero-concentrated differential privacy.

Table 16.2.: The chapters of this part.

Chapter	Question it answers
Local differential privacy	What does it cost to trust no one, and how do large deployments live with that cost?
The shuffle model and central deployments	How much does an anonymizing shuffler recover, and how do central deployments select, compose and budget?

17. Local Differential Privacy

Almost every mechanism so far ran on a trusted curator who saw the raw data. This chapter drops that trust. Each person randomizes their own record before it leaves their device, so that no one, the curator included, ever sees it. We define this local model and show that, with n people, it costs a factor of about $\sqrt{n}$ in accuracy on the simplest task, and that no noninteractive protocol can avoid the cost.

The rest of the chapter studies three large deployments of the local model, at Google, Apple and Microsoft. Each solves a problem the textbook protocol ignores. Google's protects people who answer the same question every day, Apple's finds values that no one knew to ask about, and Microsoft's sends numbers in a single bit and copes with values that drift.

17.1. The local model

In the central model of the first chapter on the definition, people send their data to a curator, and the trust barrier sits at the curator's output. In the local model the barrier moves to each person. Anything that leaves a person's device may be seen by an adversary, and the curator is now just an analyst who receives messages.

Definition 17.1 (Local randomizer)**.** Let $\varepsilon \geq 0$. A randomized algorithm $R\colon \mathcal{X} \to \mathcal{Y}$ is an **ε-local randomizer** if for all inputs $x, x' \in \mathcal{X}$ and all events $S \subseteq \mathcal{Y}$,

$$\Pr[R(x) \in S] \leq e^{\varepsilon} \Pr[R(x') \in S]. \tag{17.1}$$

This is the definition of differential privacy (DP) applied to a dataset with one record. In the simplest protocol, everyone speaks once and at the same time.

Definition 17.2 (Noninteractive local protocol)**.** A **noninteractive ε-local protocol** gives each person i an ε-local randomizer R_i. Person i sends $Y_i = R_i(X_i)$, using randomness independent of everyone else's, and the analyst outputs a function of the messages $(Y_1, \dots, Y_n)$.

Local privacy implies central privacy, because the analyst's view is a vector of messages and only one of them depends on any one record.

Proposition 17.1. *The messages $(Y_1, \dots, Y_n)$ of a noninteractive ε-local protocol, viewed as the output of a mechanism on D, are ε-DP. So is any function of them that the analyst computes.*

Proof.

1. Let $D \sim D'$ be neighbours, which differ by replacing entry j, and take an event S. Write y_{-j} for a value of all the messages except the j-th, and $S_{y_{-j}} = \{y_j : (y_j, y_{-j}) \in S\}$ for the slice of S through it.
2. The other messages have the same distribution under D and D', because they depend only on the other records and on independent randomness. Conditioning on them gives

$$\Pr_{D}[Y \in S] = \mathbb{E}_{Y_{-j}}\left[\Pr[R_j(X_j) \in S_{Y_{-j}}]\right] \le e^{\varepsilon}\, \mathbb{E}_{Y_{-j}}\left[\Pr[R_j(X'_j) \in S_{Y_{-j}}]\right] = e^{\varepsilon} \Pr_{D'}[Y \in S], \qquad (17.2)$$

where the inequality applies Definition 17.1 to each slice. 3. The analyst's output is a post-processing of the messages, so it is ε-DP by the post-processing property of the Laplace mechanism chapter. □

Randomized response, with each person reporting their true bit with probability $e^{\varepsilon}/(1+e^{\varepsilon})$, is an ε-local randomizer. That was the local protocol hiding inside the first private algorithm of this book.

Protocols can also interact. In a **sequentially interactive** protocol the analyst addresses people one at a time. The randomizer given to person i may depend on the messages of the people before them, but each person speaks only once. In a **fully interactive** protocol the analyst may return to a person many times. The requirement is then that all the messages a person ever sends, taken together, form an ε-local randomizer of their record.

The power of interaction is only partly understood. For some problems a fully interactive protocol needs exponentially fewer people than any noninteractive one, while for others interaction does not help. Pure and approximate privacy are another puzzle, because no separation between them is known in the local model. For noninteractive protocols, an approximate protocol with a small enough δ can even be turned into a pure one at a constant-factor cost in ε.

17.2. Counting bits, central and local

The cost of moving the barrier shows on the simplest task. Each person holds a bit $X_i \in \{0,1\}$, and the analyst wants the mean $\mu = \frac{1}{n}\sum_{i=1}^{n} X_i$. Write $\mathrm{Lap}(b)$ for the Laplace distribution with density $\frac{1}{2b}e^{-|z|/b}$. The Laplace mechanism adds $\mathrm{Lap}(\Delta f/\varepsilon)$ noise to a query f whose value changes by at most Δf between neighbours, and it is ε-DP. A central curator, who sees the bits, can therefore add Laplace noise to their mean.

Proposition 17.2. *For every $\varepsilon > 0$ and $\beta \in (0,1)$, the central estimate $\hat{\mu} = \mu + \mathrm{Lap}(1/(\varepsilon n))$ is ε-DP, and with probability at least $1-\beta$,*

$$|\hat{\mu} - \mu| \le \frac{\ln(1/\beta)}{\varepsilon n}. \qquad (17.3)$$

Proof.

1. Changing one bit changes μ by at most $1/n$, so the Laplace mechanism with $\Delta f = 1/n$ is ε-DP.
2. The Laplace tail is $\Pr[|\mathrm{Lap}(b)| > t] = 2\int_t^\infty \frac{1}{2b} e^{-z/b}\,dz = e^{-t/b}$. Setting $b = 1/(\varepsilon n)$ and $t = b\ln(1/\beta)$ makes the tail equal to β.

□

In the local model each person runs randomized response and sends Y_i, equal to X_i with probability $e^\varepsilon/(1+e^\varepsilon)$ and to $1 - X_i$ otherwise. A report equals 1 with probability $(1 + (e^\varepsilon - 1)X_i)/(e^\varepsilon + 1)$, which is linear in X_i. Solving for X_i gives the analyst's unbiased estimate

$$\hat{\mu} = \frac{1}{n}\sum_{i=1}^{n} \frac{(e^\varepsilon + 1)\,Y_i - 1}{e^\varepsilon - 1}. \tag{17.4}$$

This is the debiased estimator of the randomized response chapter with bias $2\kappa = (e^\varepsilon - 1)/(e^\varepsilon + 1)$, written in terms of ε. That chapter proved that, with probability at least $1 - \beta$,

$$|\hat{\mu} - \mu| \le \coth\!\Big(\frac{\varepsilon}{2}\Big)\sqrt{\frac{\ln(2/\beta)}{2n}}, \qquad \coth\!\Big(\frac{\varepsilon}{2}\Big) \ge \frac{2}{\varepsilon}. \tag{17.5}$$

The two bounds reach an error of α at different sizes. The central bound does so once $n \ge \ln(1/\beta)/(\varepsilon\alpha)$. The local bound does so once $n \ge \coth^2(\varepsilon/2)\ln(2/\beta)/(2\alpha^2)$, a threshold that is at least $2\ln(2/\beta)/(\varepsilon^2\alpha^2)$. So the local estimate needs quadratically more data in $1/\alpha$ and in $1/\varepsilon$, and Figure 17.1 shows the two error curves.

Example 17.1. Take $\varepsilon = 1$, $\alpha = 0.01$ and $\beta = 0.05$. The central estimate needs $n \ge \ln 20/0.01$, that is, 300 people. The local one needs $n \ge \coth^2(0.5)\ln 40/(2\times 10^{-4})$, that is, 86,370 people, about 288 times as many.

This gap is why local privacy has been deployed only by companies with hundreds of millions of users. With a few thousand users, a local protocol returns mostly noise.

17.3. The price is unavoidable

The local estimator is not merely a poor choice. No noninteractive local protocol can do better than error of order $1/(\varepsilon\sqrt{n})$, and the proof shows why. A local randomizer shrinks the difference between two input distributions by a factor that depends only on ε, and n independent messages cannot recover more than n such shrunken differences.

Two distances between distributions P and Q on a countable set measure this. The **total variation distance** is $\mathrm{TV}(P,Q) = \frac{1}{2}\sum_y |P(y) - Q(y)|$, the largest gap $P(S) - Q(S)$ over events S. The **chi-squared divergence** is

$$\chi^2(P\,\|\,Q) = \sum_y \frac{(P(y) - Q(y))^2}{Q(y)}, \tag{17.6}$$

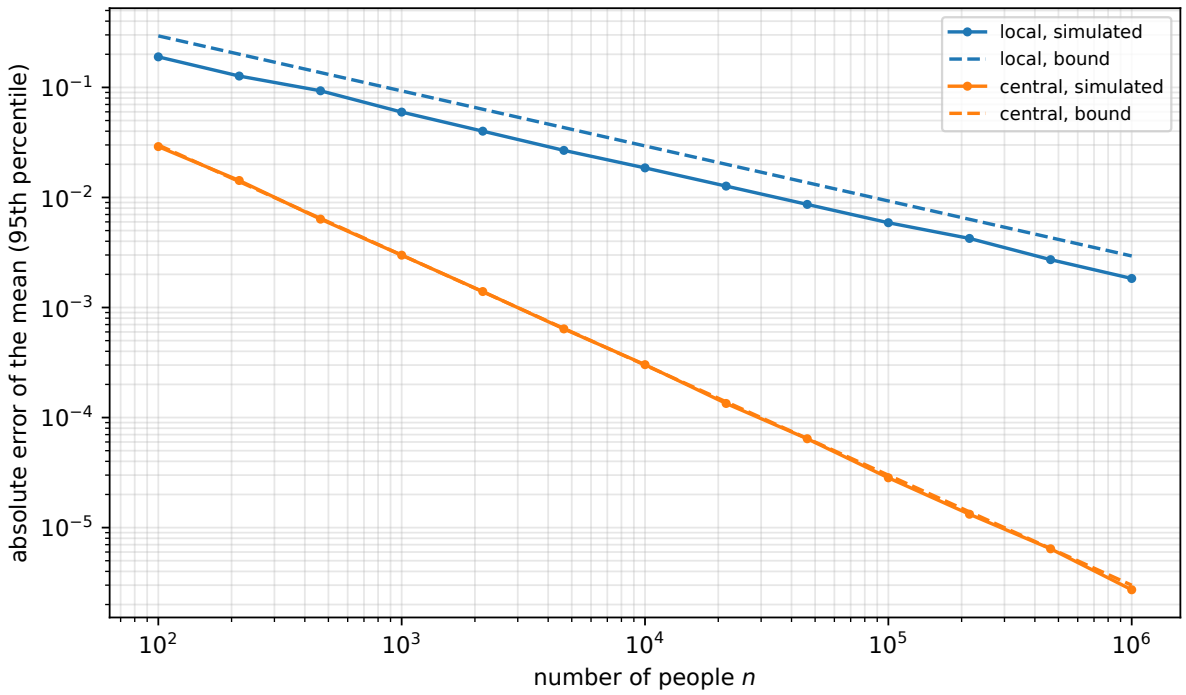


Figure 17.1.: Error of the mean of n bits at $\varepsilon = 1$. Solid lines are the 95th percentile of the error over 2,000 simulated runs with $\mu = 0.3$, dashed lines the bounds of Proposition 17.2 and Equation 17.5 with $\beta = 0.05$. The central error falls as $1/n$ and the local one as $1/\sqrt{n}$, so the gap widens with n.

where a term in which both $P(y)$ and $Q(y)$ are zero is dropped, and a term in which $Q(y)$ is zero but $P(y)$ is not makes the divergence infinite. The lower bound needs two facts about these distances. Total variation is at most a constant times the square root of the chi-squared divergence, and the chi-squared divergence behaves simply on products, which is where the n independent messages enter.

Lemma 17.1. *For distributions on a countable set,*

1. $\mathrm{TV}(P, Q) \le \frac{1}{2}\sqrt{\chi^2(P \,\|\, Q)}$;
2. *for product distributions,* $1 + \chi^2(P_1 \times P_2 \,\|\, Q_1 \times Q_2) = (1 + \chi^2(P_1 \,\|\, Q_1))(1 + \chi^2(P_2 \,\|\, Q_2))$.

Proof.

1. Write $|P(y) - Q(y)| = \frac{|P(y)-Q(y)|}{\sqrt{Q(y)}}\sqrt{Q(y)}$ and apply the Cauchy–Schwarz inequality to the sum. Since $\sum_y Q(y) = 1$, this gives $\sum_y |P(y) - Q(y)| \le \sqrt{\chi^2(P \,\|\, Q)}$.
2. Expanding the square shows $1 + \chi^2(P \,\|\, Q) = \sum_y P(y)^2/Q(y)$. For a product, this sum factorizes as $\sum_{y_1} P_1(y_1)^2/Q_1(y_1) \cdot \sum_{y_2} P_2(y_2)^2/Q_2(y_2)$.

□

The key step is a contraction. Feed a local randomizer an input drawn from P_0 or from P_1, and write W_0 and W_1 for the two resulting distributions of its message.

Lemma 17.2. *Let R be an ε-local randomizer with a countable message set, and let P_0, P_1 be distributions on $\mathcal{X}$. Write W_0 and W_1 for the distributions of the message $R(x)$ when x is drawn from P_0 and from P_1. Then*

$$\chi^2(W_0 \,\|\, W_1) \le (e^\varepsilon - 1)^2 \,\mathrm{TV}(P_0, P_1)^2. \tag{17.7}$$

Proof.

1. Fix a message y and write $Q(y \mid x) = \Pr[R(x) = y]$ and $c_y = \inf_x Q(y \mid x)$. By Definition 17.1, every $Q(y \mid x)$ lies in $[c_y, e^\varepsilon c_y]$. If $c_y = 0$, then $Q(y \mid x) = 0$ for every x, and y contributes nothing.
2. Since $\sum_x (P_0(x) - P_1(x)) = 0$, subtracting the midpoint $m_y = (1 + e^\varepsilon)c_y/2$ changes nothing:

$$W_0(y) - W_1(y) = \sum_x (Q(y \mid x) - m_y)(P_0(x) - P_1(x)). \tag{17.8}$$

3. Each factor $|Q(y \mid x) - m_y|$ is at most $(e^\varepsilon - 1)c_y/2$, and $\sum_x |P_0(x) - P_1(x)| = 2\,\mathrm{TV}(P_0, P_1)$. So $|W_0(y) - W_1(y)| \le (e^\varepsilon - 1)c_y\,\mathrm{TV}(P_0, P_1)$.
4. The denominator satisfies $W_1(y) \ge c_y$. Hence each term of χ^2 is at most $(e^\varepsilon - 1)^2 c_y\,\mathrm{TV}(P_0, P_1)^2$.
5. Summing over y and using $\sum_y c_y \le \sum_y Q(y \mid x) = 1$ for any fixed x gives the bound.

□

Now let the data be random. If each X_i is drawn independently from $\mathrm{Bern}(p)$, the Bernoulli distribution with mean p, an accurate estimate of the mean must tell $p = \frac{1}{2} - \alpha$ from $p = \frac{1}{2} + \alpha$.

Theorem 17.1 (Lower bound for local counting)**.** *Let $\alpha \in (0, \frac{1}{2}]$, and let $X_1, \dots, X_n$ be independent draws from $\mathrm{Bern}(\frac{1}{2} + \alpha)$ or from $\mathrm{Bern}(\frac{1}{2} - \alpha)$. Write P_+ and P_- for the two distributions of the messages of a noninteractive ε-local protocol with countable message sets. Then*

$$\mathrm{TV}(P_+, P_-) \le \frac{1}{2}\sqrt{\exp(4n(e^\varepsilon - 1)^2\alpha^2) - 1}. \tag{17.9}$$

Consequently, if an estimate $\tilde{\mu}$ computed from the messages satisfies $|\tilde{\mu} - p| < \alpha$ with probability at least $3/4$ in both cases, then

$$n \ge \frac{\ln 2}{4(e^\varepsilon - 1)^2\alpha^2}, \qquad \text{and for } \varepsilon \le 1, \quad n \ge \frac{1}{18\,\varepsilon^2\alpha^2}. \tag{17.10}$$

Proof.

1. The two input distributions have $\mathrm{TV}(\mathrm{Bern}(\frac{1}{2} + \alpha), \mathrm{Bern}(\frac{1}{2} - \alpha)) = 2\alpha$. By Lemma 17.2, each person's message satisfies $\chi^2 \le 4(e^\varepsilon - 1)^2\alpha^2$ between the two cases.
2. The messages are independent, because the records and the coins are. By the product rule of Lemma 17.1 and $1 + z \le e^z$, we get $1 + \chi^2(P_+ \,\|\, P_-) \le \exp(4n(e^\varepsilon - 1)^2\alpha^2)$. The first part of the same lemma gives the bound on TV.
3. Take the test that answers "$p = \frac{1}{2} + \alpha$" when $\tilde{\mu} \ge \frac{1}{2}$. In each case it errs only if $|\tilde{\mu} - p| \ge \alpha$, so each error probability is at most $\frac{1}{4}$ and the two add up to at most $\frac{1}{2}$. For any test, the error probabilities add up to at least $1 - \mathrm{TV}(P_+, P_-)$, because the chance of answering "+" differs between the cases by at most TV. So $\mathrm{TV}(P_+, P_-) \ge \frac{1}{2}$.
4. Combining steps 2 and 3 gives $\exp(4n(e^\varepsilon - 1)^2\alpha^2) \ge 2$, which is the first bound on n.
5. For $\varepsilon \le 1$, convexity of e^ε gives $e^\varepsilon - 1 \le (e - 1)\varepsilon$, and $\ln 2/(4(e-1)^2) = 0.0587 > 1/18$.

□

The theorem is about data drawn at random, but it also constrains estimates of the empirical mean on fixed data. Suppose a protocol is within $\alpha/2$ of the empirical mean with probability at least $7/8$ on every dataset. When $n \ge 2\ln(16)/\alpha^2$, Hoeffding's inequality puts the empirical mean of the sample strictly within $\alpha/2$ of p with probability at least $7/8$. So the protocol is strictly within α of p with probability at least $3/4$, and the theorem applies to it. This reduction says nothing about smaller n. A direct argument for fixed data, showing that a local protocol for constant ε has error of order $\sqrt{n}$ on the sum of n bits, is due to Chan, Shi and Song (2012).

For $\varepsilon \le 1$, comparing with Equation 17.5 shows that randomized response is optimal up to a constant factor. The central estimate, which needs only about $1/(\varepsilon\alpha)$ people, is out of reach of every noninteractive local protocol.

17.4. Repeated questions: RAPPOR

Google's first large deployment, RAPPOR, collected statistics from the Chrome browser. It met a problem that a one-off survey never has. The same question is asked of the same person every day, and fresh randomness each day leaks the answer.

Example 17.2 (Averaging attack)**.** Take $\varepsilon = \ln 4$, so that randomized response tells the truth with probability $4/5$. A single "yes" from a person whose answer is "no" happens with probability $1/5$, which leaves good deniability. Now ask the same question on 50 days with fresh coins, and suppose 40 reports say "yes". If the answer is "no", this happens with probability $\binom{50}{40}(1/5)^{40}(4/5)^{10} = 1.21 \times 10^{-19}$. If it is "yes", it happens with probability 0.140. The likelihood ratio is $4^{30} \approx 1.15 \times 10^{18}$, and the deniability is gone.

The attack is composition at work. Basic composition, from the Laplace mechanism chapter, says that k answers with fresh coins are $k\varepsilon$-DP, which here is $50 \ln 4 = 69.3$. This bound is met exactly, because the all-"yes" outcome has likelihood ratio 4^{50}. So the fix is to randomize once and repeat.

Definition 17.3 (Memoization)**.** A person who is asked a question for the first time runs their randomizer and stores the result. When asked the same question again, they send the stored result.

Memoization keeps the privacy cost of a single answer, however often the question is asked.

Proposition 17.3. *If the stored result comes from an ε-local randomizer, then any number of memoized answers to the same question is ε-locally private.*

Proof. The sequence of answers is a function of the single stored result, so the claim follows from post-processing. □

Memoization protects honest repetition, not an adversarial analyst. An analyst can ask for "your PIN" and then for "your PIN reversed", two questions with the same information that each get a fresh randomization.

RAPPOR handles values from any finite domain, not just bits. It first encodes a value as a bit vector.

Definition 17.4 (Bloom filter)**.** Fix a length K and h hash functions $H_1, \dots, H_h \colon \mathcal{X} \to \{1, \dots, K\}$. The Bloom filter of a value v is the vector $B \in \{0,1\}^K$ with $B_j = 1$ exactly when $H_l(v) = j$ for some l.

With $h = 1$ and no collisions the Bloom filter is a one-hot encoding of v. In general, B has at most h ones. RAPPOR then applies two layers of randomized response, bit by bit.

Definition 17.5 (RAPPOR)**.** Fix $\nu \in (0,1)$, the noise level of the permanent layer, and probabilities $p_0 < p_1$ in $(0,1)$, the chances of reporting a one in the instantaneous layer. Erlingsson and coauthors write these three parameters f, p and q. A person with value v and Bloom filter B of length K, from Definition 17.4, does the following.

1. **Permanent randomized response.** Once, they draw $B' \in \{0,1\}^K$ with independent coordinates, where B'_j is 1 with probability $\nu/2$, 0 with probability $\nu/2$, and B_j otherwise. They memoize B'.
2. **Instantaneous randomized response.** Each time they report, they send a fresh $Y \in \{0,1\}^K$ with independent coordinates, where $Y_j = 1$ with probability p_1 if $B'_j = 1$ and with probability p_0 if $B'_j = 0$.

The permanent layer carries the long-term guarantee. It holds against an analyst who sees every report the person ever sends, because all of them are computed from the one memoized vector.

Theorem 17.2. *Run RAPPOR with h hash functions and permanent noise level ν. As a function of the value v, the memoized vector B' is an ε_∞-local randomizer with*

$$\varepsilon_\infty = 2h \ln \frac{1-\nu/2}{\nu/2}. \tag{17.11}$$

So is the whole sequence of reports Y, however long.

Proof.

1. Let v, v' have Bloom filters $B, \tilde{B}$. Each has at most h ones, so they differ in at most $2h$ coordinates.
2. Fix an outcome $b \in \{0,1\}^K$. Coordinates are independent, so $\Pr[B' = b]$ is a product over j. The factors at coordinates where $B_j = \tilde{B}_j$ are equal for v and v' and cancel in the ratio.
3. Each remaining factor is $1-\nu/2$ or $\nu/2$, so each ratio is at most $(1-\nu/2)/(\nu/2)$, and there are at most $2h$ of them. By the pointwise lemma of the randomized response chapter, this bound on single outcomes gives ε_∞-privacy.
4. The reports depend on v only through B', so they are a post-processing of it.

□

The instantaneous layer gives a stronger guarantee for an analyst who sees one report. Combining the two layers, a reported bit is 1 with probability $p_1^* = (1-\frac{\nu}{2})p_1 + \frac{\nu}{2}p_0$ when $B_j = 1$ and $p_0^* = \frac{\nu}{2}p_1 + (1-\frac{\nu}{2})p_0$ when $B_j = 0$. Then $p_1^* > p_0^*$, because $p_1^* - p_0^* = (1-\nu)(p_1 - p_0)$.

Theorem 17.3. *Run RAPPOR with h hash functions, and let p_1^* and p_0^* be the combined probabilities above. A single report Y, as a function of v, is an ε_1-local randomizer with*

$$\varepsilon_1 = h \ln \frac{p_1^*(1-p_0^*)}{p_0^*(1-p_1^*)}. \tag{17.12}$$

Proof.

1. Given B, the coordinates of Y are independent, with $\Pr[Y_j = 1] = p_1^*$ or p_0^* according to B_j. As before, coordinates where $B_j = \tilde{B}_j$ cancel.
2. Let A be the set of coordinates with $B_j = 1$ and $\tilde{B}_j = 0$, and A' the set with the reverse. Both have at most h elements, since each lies inside the ones of one filter.

3. On A the ratio is $p_1^*/p_0^* > 1$ if $y_j = 1$ and $(1-p_1^*)/(1-p_0^*) < 1$ if $y_j = 0$, so it is at most p_1^*/p_0^*. On A' it is at most $(1-p_0^*)/(1-p_1^*)$ by the same reasoning.
4. Multiplying, the ratio is at most $\left(p_1^*(1-p_0^*)/(p_0^*(1-p_1^*))\right)^h$, and the pointwise lemma finishes the proof.

□

Typical parameters show how far apart the two guarantees can be.

Example 17.3. With $h = 2$, $\nu = 1/2$, $p_0 = 1/2$ and $p_1 = 3/4$, the long-term guarantee is $\varepsilon_\infty = 4 \ln 3 = 4.39$. The combined probabilities are $p_1^* = 0.6875$ and $p_0^* = 0.5625$, so a single report has $\varepsilon_1 = 1.07$.

The second layer has a further purpose that local privacy does not capture. A memoized B' of a few hundred bits is almost surely unique, so seeing it twice links two reports to the same person, as a cookie would. Fresh noise on each report breaks that link for an analyst who sees few reports. An analyst who sees many reports can average the noise away and recover B', and then still faces the ε_∞ guarantee of Theorem 17.2.

The analyst decodes the reports by debiasing the bit counts and fitting candidate values against them, in the spirit of Equation 17.4. Memoization fails when the value drifts. If the question is a person's age in days, the answer changes daily, each new value needs a new B', and k distinct values cost $k\,\varepsilon_\infty$ by composition. Chrome used RAPPOR on an opt-in basis, for example to measure the distribution of browser home pages, where software that hijacks a browser shows up as an unexpected popular value, and to survey the processes running on users' machines.

17.5. Values no one knew: Apple

Apple deployed local privacy on its phones and computers. Around the randomizers, the system adds measures outside the model of differential privacy. It discards IP addresses, sends reports after a delay and only a random sample of them, encrypts them in transit with Transport Layer Security (TLS), and limits how many reports each use sends per day. Apple did not publish all its parameters at first. Independent researchers then measured them from the software of macOS and found that each report had ε of 1 or 2, but that the budget allowed a privacy loss of up to 16 per day, renewed every day.

Most of Apple's protocols build on a private count mean sketch, a hashed table for estimating frequencies in the spirit of the count sketch. A value is hashed into a short table by one of several hash functions, and the table is randomized before it is sent. A variant based on the Hadamard transform lets each device send a single bit. Communication matters on a phone, where a report should cost a few bytes at most, and the last section returns to it.

The new problem Apple solved is **discovery**. Hashing maps a value to its encoding easily, but the encoding cannot be inverted. To estimate how often a word is typed, the analyst must already have the word in its dictionary and test its encoding against the reports. Slang and new words are exactly the values missing from any dictionary, and searching all strings is hopeless: there are $26^8 \approx 2.1 \times 10^{11}$ strings of eight letters.

The **sequence fragment puzzle** breaks a word into pieces short enough to enumerate. Fix a hash function H from words to $\{0, \dots, 255\}$, an eight-bit hash. Take, for illustration, eight-letter words and fragments of two letters. A person with an eight-letter word w picks one of the starting positions $j \in \{1, 3, 5, 7\}$ uniformly at random. They send j, which needs no protection, together with a private report of the pair $(w_j w_{j+1}, H(w))$: the two letters at position j and the hash of the whole word, the *puzzle piece*. The pairs come from a set of $26^2 \times 256 = 173{,}056$ candidates, small enough for the analyst to estimate the frequency of every one, separately for each position. Sending the position in the clear costs no privacy, because the person chooses it at random, independently of the word.

Proposition 17.4. *If the private report of the pair is an ε-local randomizer, then the whole message (j, report) is ε-locally private.*

Proof. The position j is drawn independently of w. Conditional on j, the message is an ε-local randomizer of w. For an event S, write S_j for the set of reports r with $(j, r) \in S$. Then $\Pr[(j, \text{report}) \in S \mid w] = \sum_j \frac{1}{4} \Pr[\text{report} \in S_j \mid w, j]$ is at most e^{ε} times the same sum for any other word. □

An example shows how the analyst puts the pieces back together.

Example 17.4. Suppose many people type “snowfall”, and $H(\text{snowfall}) = 117$. Their messages are spread over $(1, \text{sn}, 117)$, $(3, \text{ow}, 117)$, $(5, \text{fa}, 117)$ and $(7, \text{ll}, 117)$. Each of the four pairs is frequent at its position. Pieces that share the hash 117 are joined in order of position, giving the candidate “snowfall”. A second round of ordinary frequency estimation over the discovered words then counts it, and it removes candidates that a hash collision assembled from different words.

The method found abbreviations and slang that no dictionary held. It also found “lov”, “th” and “kno”, which are “love”, “the” and “know” cut short. People had pressed by accident the leftmost suggestion above the keyboard, which holds the letters typed so far, instead of the final letter. Other uses were the frequency of each emoji, which differed between English and French keyboards, the websites on which users wanted media to play automatically, which websites used the most energy or memory, and which data types were most used in the health app.

17.6. Numbers in one bit: Microsoft

Microsoft collects usage telemetry from Windows, such as how long people spend in each application. The values are numbers in a range $[0, m]$, and the analyst wants their mean μ. The Laplace mechanism works locally: person i sends $Y_i = X_i + \text{Lap}(m/\varepsilon)$, which is ε-local because the identity map on $[0, m]$ changes by at most m. The mean of the reports has error $\frac{1}{n} \sum_i \text{Lap}(m/\varepsilon)$, with standard deviation exactly $\sqrt{2}\, m/(\varepsilon\sqrt{n})$. This rate is optimal for local protocols, by Theorem 17.1 applied to values in $\{0, m\}$, but each report is a real number of size about m/ε, so it costs many bits. A single bit suffices, sent with a probability that grows linearly with the value.

Definition 17.6 (One-bit mean protocol)**.** Person i, with $X_i \in [0, m]$, sends a bit Y_i with

$$\Pr[Y_i = 1] = \frac{1}{e^\varepsilon + 1} + \frac{X_i}{m} \cdot \frac{e^\varepsilon - 1}{e^\varepsilon + 1}. \tag{17.13}$$

The analyst outputs

$$\hat{\mu} = \frac{m}{n} \sum_{i=1}^{n} \frac{(e^\varepsilon + 1)\, Y_i - 1}{e^\varepsilon - 1}. \tag{17.14}$$

At $X_i = 0$ and $X_i = m$ this is randomized response, and between them it interpolates linearly. The analyst's estimate inverts that linear relation, as Equation 17.4 did for bits, and the next theorem shows that it keeps the privacy and the accuracy of randomized response.

Theorem 17.4. *The one-bit protocol is ε-locally private, and $\hat{\mu}$ is unbiased. With probability at least $1 - \beta$,*

$$|\hat{\mu} - \mu| \le m \coth\Big(\frac{\varepsilon}{2}\Big) \sqrt{\frac{\ln(2/\beta)}{2n}}, \tag{17.15}$$

and the standard deviation of $\hat{\mu}$ is at most $m \coth(\varepsilon/2)/(2\sqrt{n})$.

Proof.

1. *Privacy.* For every $x \in [0, m]$, $\Pr[Y = 1]$ lies between $1/(e^\varepsilon + 1)$ and $e^\varepsilon/(e^\varepsilon + 1)$, and so does $\Pr[Y = 0] = 1 - \Pr[Y = 1]$. The ratio of two numbers in that interval is at most e^ε, and the pointwise lemma applies.
2. *Unbiasedness.* Substituting $\mathbb{E}[Y_i] = \Pr[Y_i = 1]$ gives $(e^\varepsilon + 1)\mathbb{E}[Y_i] - 1 = (X_i/m)(e^\varepsilon - 1)$, so each term of the sum has mean X_i/m, and $\mathbb{E}[\hat{\mu}] = \mu$.
3. *Error.* Each term $Z_i = m((e^\varepsilon + 1)Y_i - 1)/(e^\varepsilon - 1)$ takes the two values $-m/(e^\varepsilon - 1)$ and $me^\varepsilon/(e^\varepsilon - 1)$, an interval of length $m(e^\varepsilon + 1)/(e^\varepsilon - 1) = m \coth(\varepsilon/2)$. The Z_i are independent, and Hoeffding's inequality for their average gives the bound, as in the randomized response chapter.
4. *Variance.* A bit has variance at most $1/4$, so $\mathrm{Var}(Z_i) \le m^2 \coth^2(\varepsilon/2)/4$, and the average of n independent terms has variance at most $m^2 \coth^2(\varepsilon/2)/(4n)$.

□

Since $\coth(\varepsilon/2)/2 \approx 1/\varepsilon$ for small ε, one bit loses nothing against the Laplace protocol's $\sqrt{2}/\varepsilon$. At $\varepsilon = 1$ the one-bit standard deviation is at most 0.765 times the Laplace one.

17.6.1. Memoization with drifting values

Daily usage is never exactly the same two days running, so memoization as in RAPPOR does not apply. Rounding to bins before memoizing poses a dilemma. Narrow bins change often and force new randomizations, while wide bins make each report wrong by up to half a bin. Randomized rounding solves the second problem. It rounds each value to 0 or to m at random, with probabilities chosen so that the rounded value is right on average, so two bins suffice without bias.

Definition 17.7 (Memoized one-bit protocol)**.** Before any data arrive, person i does three things, independently of their data and of each other.

1. They run the one-bit protocol on the input 0 and store the bit $A_i(0)$.
2. They run it on the input m and store the bit $A_i(m)$.
3. They draw an offset u_i uniformly from $[0, m)$.

On a day when their value is $x \in [0, m]$, they round it to $r_i(x) = m$ if $x + u_i \geq m$ and to $r_i(x) = 0$ otherwise, and they send $A_i(r_i(x))$.

The random offset is what makes the rounding right on average.

Lemma 17.3. *For every $x \in [0, m]$, $\Pr[r_i(x) = m] = x/m$, so $\mathbb{E}[r_i(x)] = x$.*

Proof. The offset is uniform on $[0, m)$, so $\Pr[u_i \geq m - x] = (m - (m - x))/m = x/m$. □

For integer values, an integer offset drawn uniformly from $\{0, \dots, m-1\}$ gives the same probability x/m for every integer $x \in [0, m]$, since exactly x offsets reach $m - x$ or more.

Unbiased rounding, together with stored bits that are themselves outputs of the one-bit protocol, makes each day's bit look like a fresh run of that protocol.

Theorem 17.5. *On each day, the bit sent by the memoized protocol has exactly the distribution of the one-bit protocol run on that day's value x. So each day's bit is ε-locally private, and the estimator of Definition 17.6, applied to one day's bits, is unbiased with the error bound of Theorem 17.4.*

Proof. The offset and the two stored bits are independent. By Lemma 17.3,

$$\Pr[\text{bit} = 1] = \frac{x}{m} \cdot \frac{e^{\varepsilon}}{e^{\varepsilon} + 1} + \left(1 - \frac{x}{m}\right) \frac{1}{e^{\varepsilon} + 1} = \frac{1}{e^{\varepsilon} + 1} + \frac{x}{m} \cdot \frac{e^{\varepsilon} - 1}{e^{\varepsilon} + 1}, \tag{17.16}$$

which is the probability in Definition 17.6. Different people use independent coins, so the bits of one day are independent, as the error bound requires. □

Over many days, the bits depend on the data only through the sequence of rounded values. As long as a person's value stays on one side of their threshold $m - u_i$, they repeat one stored bit, and the analyst learns nothing new. A crossing shows in the reports only if the two stored bits differ, which happens with probability $(1 + e^{2\varepsilon})/(1 + e^{\varepsilon})^2$, or 0.607 at $\varepsilon = 1$. Then the analyst learns the day on which the value crossed a threshold drawn uniformly at random. For the age-in-days question, this is all an observer ever sees, a run of $A_i(0)$ followed by a run of $A_i(m)$. This leak is not covered by an ε-local guarantee for arbitrary sequences of values, and the deployed protocol adds a further layer of randomness to the sent bits to blur the crossing points.

17.7. Summary

Table 17.1.: The chapter's results.

Result	Statement
Proposition 17.1	a noninteractive ε-local protocol is ε-DP in the central sense
Proposition 17.2	central mean of bits: error $\ln(1/\beta)/(\varepsilon n)$
Equation 17.5	local mean of bits by randomized response: error $\coth(\varepsilon/2)\sqrt{\ln(2/\beta)/(2n)}$
Lemma 17.2	an ε-local randomizer shrinks χ^2 to at most $(e^{\varepsilon} - 1)^2\,\mathrm{TV}^2$
Theorem 17.1	every noninteractive local protocol needs $n \geq \ln 2/(4(e^{\varepsilon} - 1)^2\alpha^2)$
Proposition 17.3	memoized answers to one question cost ε once
Theorem 17.2	RAPPOR's memoized vector is $2h\ln((1 - \nu/2)/(\nu/2))$-locally private
Theorem 17.3	one RAPPOR report is $h\ln\left(p_1^*(1 - p_0^*)/(p_0^*(1 - p_1^*))\right)$-locally private
Proposition 17.4	a fragment report with a public random position keeps its ε
Theorem 17.4	one bit per person: unbiased, ε-local, error $m\coth(\varepsilon/2)\sqrt{\ln(2/\beta)/(2n)}$
Theorem 17.5	memoizing two bits with a random offset reproduces the one-bit protocol each day

17.8. References

- Stanley L. Warner, "Randomized Response: A Survey Technique for Eliminating Evasive Answer Bias", 1965. The first local protocol.
- Alexandre Evfimievski, Johannes Gehrke and Ramakrishnan Srikant, "Limiting Privacy Breaches in Privacy Preserving Data Mining", 2003. A ratio condition on a person's randomized data that anticipates the local definition.

- Shiva Prasad Kasiviswanathan, Homin K. Lee, Kobbi Nissim, Sofya Raskhodnikova and Adam Smith, "What Can We Learn Privately?", 2008. The local model in the language of differential privacy, and a separation between interactive and noninteractive protocols.
- T.-H. Hubert Chan, Elaine Shi and Dawn Song, "Optimal Lower Bound for Differentially Private Multi-party Aggregation", 2012. An $\Omega(\sqrt{n})$ lower bound, for constant ε, on the error of a sum of bits in the local model, whatever the number of rounds.
- John C. Duchi, Michael I. Jordan and Martin J. Wainwright, "Local Privacy and Statistical Minimax Rates", 2013. Contraction of divergences under local randomizers, the idea behind Lemma 17.2, and optimal rates for local estimation, including means in $[0, m]$.
- Mark Bun, Jelani Nelson and Uri Stemmer, "Heavy Hitters and the Structure of Local Privacy", 2018. Noninteractive approximate local protocols can be made pure.
- Matthew Joseph, Jieming Mao, Seth Neel and Aaron Roth, "The Role of Interactivity in Local Differential Privacy", 2019. Separations and equivalences between the noninteractive, sequentially interactive and fully interactive models.
- Burton H. Bloom, "Space/Time Trade-offs in Hash Coding with Allowable Errors", 1970. The Bloom filter.
- Úlfar Erlingsson, Vasyl Pihur and Aleksandra Korolova, "RAPPOR: Randomized Aggregatable Privacy-Preserving Ordinal Response", 2014. Memoization, the two layers of randomized response, Theorem 17.2 and Theorem 17.3, and the Chrome measurements.
- Moses Charikar, Kevin Chen and Martin Farach-Colton, "Finding Frequent Items in Data Streams", 2002. The count sketch, the frequency sketch that Apple's private sketches follow.
- Differential Privacy Team, Apple, "Learning with Privacy at Scale", 2017. Apple's architecture, its count mean sketch protocols, the sequence fragment puzzle and its applications.
- Jun Tang, Aleksandra Korolova, Xiaolong Bai, Xueqiang Wang and XiaoFeng Wang, "Privacy Loss in Apple's Implementation of Differential Privacy on macOS 10.12", 2017. An independent measurement of Apple's privacy parameters, including the daily loss of up to 16.
- Bolin Ding, Janardhan Kulkarni and Sergey Yekhanin, "Collecting Telemetry Data Privately", 2017. The one-bit mean protocol, memoization with randomized rounding, and the output perturbation added on top, deployed in Windows.

18. The Shuffle Model and Central Deployments

The local model of the previous chapter asks no one to be trusted, and pays with an error $\sqrt{n}$ times larger than the central model's. This chapter first studies a middle ground for differential privacy (DP), the shuffle model, in which a party that only mixes messages lets each person add far less noise. It then turns to deployments in the central model, where a company holds the data and releases statistics or answers queries. These are an analytics service that answers queries from outside analysts, and two large public data releases.

Each deployment brings a piece of theory with it. The shuffle model comes with amplification by shuffling and a counting protocol whose error does not grow with n. The query service comes with a sharper way to account for the exponential mechanism and a fast way to run it. The data releases come with zero-concentrated differential privacy and with a choice between protecting each action and protecting each person.

18.1. Encode, shuffle, analyze

Google's Prochlo system started from a failure of local privacy. Its designers wanted to learn which applications call which system interfaces, so that rarely used interfaces could be retired and old ones tracked. With a thousand applications and a hundred interfaces there are 10^5 pairs to count. By their estimate, a local protocol such as RAPPOR would need reports from about a trillion people to count them usefully, more people than exist.

Their architecture has three stages, **encode, shuffle, analyze**. Each person *encodes* their data on their own device, for instance with a local randomizer. A *shuffler* collects the encoded messages, strips their metadata such as arrival times and network addresses, and forwards them in a uniformly random order. The *analyst* sees only the shuffled batch and decodes statistics from it. The shuffler's only job is to break the link between a message and its sender.

Someone must run the shuffler, and the users must trust it not to reveal the order. Prochlo names three options: a trusted third party, a split of the work between several parties, or trusted hardware. It implements the last one, running the shuffle inside a trusted execution environment (TEE), an Intel SGX enclave, which is hardware designed so that even the machine's owner cannot see the computation inside. Trust then rests on the hardware vendor, and side-channel attacks such as Spectre have shown that such hardware can leak. Splitting the work is the idea of a mix network, in which several servers each shuffle and re-encrypt the batch, so that the order stays hidden unless all of them collude. In Prochlo's experiments on word frequencies, RAPPOR recovered less than 5% as many distinct words as the shuffled system.

18.2. The shuffle model

Prochlo motivated a formal model that sits between the local and central ones. The trust barrier passes after the shuffler: the adversary sees the shuffled messages but not who sent which.

Definition 18.1 (Shuffle model)**.** A protocol in the **shuffle model** gives each person the same randomizer $R\colon \mathcal{X} \to \mathcal{Y}$. Person i sends $Y_i = R(X_i)$, using randomness independent of everyone else's, and the **shuffler** outputs $(Y_{\pi(1)}, \dots, Y_{\pi(n)})$ for a permutation π drawn uniformly at random. The protocol is (ε, δ)**-DP in the shuffle model** if this shuffled output, as a mechanism on the dataset D, is (ε, δ)-DP in the central sense.

The shuffled output carries the same information as the multiset of messages, because the uniform order adds nothing. A local protocol is also a shuffle protocol with the same guarantee, since the shuffled output is a post-processing of the messages. The interesting protocols are those whose messages are *not* private one by one but become private once mixed.

The benchmark task is again a sum of bits. In the central model, a count plus Laplace noise of scale $1/\varepsilon$ has error of order $1/\varepsilon$. In the local model each person adds noise of that size, and the n independent noises add up to a standard deviation of order $\sqrt{n}/\varepsilon$, which no local protocol can beat by the previous chapter's lower bound. The question is how close the shuffle model can get to $1/\varepsilon$.

18.3. Amplification by shuffling

The first answer is general. Shuffling makes any local randomizer more private, just as subsampling did in an earlier chapter. Here is the result in the form in which it was first proved. If each person's message is ε_0-private on its own, the shuffled batch is private with a parameter smaller by a factor of order $\sqrt{n}$.

Theorem 18.1 (Amplification by shuffling)**.** *Let $n \geq 1000$, $\delta \in (0, \frac{1}{100})$, and let R be an ε_0-local randomizer with $0 < \varepsilon_0 < \frac{1}{2}$. Then the shuffled output of $R(X_1), \dots, R(X_n)$ is (ε, δ)-DP with*

$$\varepsilon = 12\,\varepsilon_0 \sqrt{\frac{\ln(1/\delta)}{n}}. \tag{18.1}$$

The proof is in the 2019 paper by Erlingsson and coauthors listed in the references (Theorem 7 and Corollary 9 there). It shuffles the records before randomizing them, which gives the same output distribution when everyone uses the same randomizer, and which also allows each randomizer to depend on the messages before it. The proof moves the one differing record to the front and swaps it into a uniformly random position, and then bounds how much each successive message reveals about that position.

The gain is a factor of order $\sqrt{n}$, but it applies only to randomizers that are already very private. Few deployments want $\varepsilon_0 \leq \frac{1}{2}$ locally, because that is the regime where local accuracy is worst. The useful regime is ε_0 of order $\ln n$, which the theorem does not cover. A later analysis by Feldman, McMillan and Talwar (2021) does, and still yields a constant ε there. It covers

every $\varepsilon_0 \le \ln\left(n/(16\ln(2/\delta))\right)$, with $\varepsilon = O((1-e^{-\varepsilon_0})\sqrt{e^{\varepsilon_0}\ln(1/\delta)/n})$. Its proof splits each other person's randomizer into a mixture, in which with probability $e^{-\varepsilon_0}$ the message looks like a "clone" of what the differing person might have sent, and the count of clones hides that person. The next section carries out this idea in full for one simple randomizer.

18.4. Counting bits in the shuffle model

For a sum of bits, a direct protocol reaches error of order $\sqrt{\ln(1/\delta)}/\varepsilon$, independent of n. Its randomizer is randomized response with a twist: almost everyone tells the truth.

Definition 18.2 (Shuffled randomized response)**.** Fix a coin probability $\nu \in (0, 1]$. Person i, with bit X_i, sends $Y_i = X_i$ with probability $1-\nu$, and otherwise sends a fresh fair coin. The analyst sees the shuffled reports, counts their ones as $Z = \sum_{i=1}^{n} Y_i$, and estimates the count $c = \sum_{i=1}^{n} X_i$ by

$$\tilde{c} = \frac{Z - n\nu/2}{1-\nu}. \tag{18.2}$$

In the local model, ν would be a constant. Here it will be of order $\ln(1/\delta)/(\varepsilon^2 n)$, so only about $\ln(1/\delta)/\varepsilon^2$ people in total send coins. The coins of the others hide any one person, much as the clones above. Two lemmas prepare the proof, and both use $\tau = \tanh(\varepsilon/2)$, the number for which $(1+\tau)/(1-\tau) = e^{\varepsilon}$. The first says that the number of ones among s fair coins hides whether one more bit, of any distribution, is added to it, as long as s is large.

Lemma 18.1. *Let $\varepsilon > 0$, let $\tau = \tanh(\varepsilon/2)$, and let $s \ge 2/\tau$ be an integer. Let N be binomial with s trials and success probability $\frac{1}{2}$, and let B and B' be Bernoulli variables independent of N, with any parameters. Then for every set S of integers,*

$$\Pr[N + B \in S] \le e^{\varepsilon} \Pr[N + B' \in S] + 2e^{-s\tau^2/8}. \tag{18.3}$$

Proof.

1. Write $P(k) = \Pr[N = k] = \binom{s}{k} 2^{-s}$. For $1 \le k \le s$, the ratio of consecutive probabilities is $P(k-1)/P(k) = k/(s+1-k)$.
2. Let K be the set of integers k with $|k - \frac{s+1}{2}| \le \frac{s+1}{2}\tau$. Writing $k = \frac{s+1}{2}(1+v)$ with $|v| \le \tau$, the ratio becomes $(1+v)/(1-v)$, which lies in $[e^{-\varepsilon}, e^{\varepsilon}]$. Since $\tau < 1$, every $k \in K$ satisfies $1 \le k \le s$.
3. For $k \in K$ and any Bernoulli parameter b, $\Pr[N + B = k] = (1-b)P(k) + bP(k-1)$ is at most $\max(P(k), P(k-1))$. By step 2 this is at most $e^{\varepsilon}\min(P(k), P(k-1))$, which is at most $e^{\varepsilon}\Pr[N + B' = k]$.
4. Outside K, use $|N + B - \frac{s+1}{2}| \le |N - \frac{s}{2}| + \frac{1}{2}$. So $N + B \notin K$ requires $|N - \frac{s}{2}| > \frac{(s+1)\tau - 1}{2}$, and this threshold is at least $s\tau/4$ because $s\tau \ge 2$. Hoeffding's inequality for s fair coins gives $\Pr[N + B \notin K] \le 2\exp(-2(s\tau/4)^2/s) = 2e^{-s\tau^2/8}$.
5. Split S into $S \cap K$, bounded by step 3, and the rest, bounded by step 4.

□

The second lemma ensures that enough people send coins. It is a lower tail for the binomial distribution, sharper than Hoeffding's inequality when the mean is small.

Lemma 18.2. *Let N be binomial with mean μ. Then*

$$\Pr\Big[N \le \frac{\mu}{2}\Big] \le \exp\!\left(-\frac{(1-\ln 2)\,\mu}{2}\right). \tag{18.4}$$

Proof.

1. Let N have m trials with success probability p, so $\mu = mp$. For $\lambda > 0$, Markov's inequality applied to $e^{-\lambda N}$ gives $\Pr[N \le a] \le e^{\lambda a}\,\mathbb{E}[e^{-\lambda N}]$.
2. The trials are independent, so $\mathbb{E}[e^{-\lambda N}] = (1 - p + pe^{-\lambda})^m$. Using $1 + z \le e^z$, this is at most $\exp\left(\mu(e^{-\lambda} - 1)\right)$.
3. Take $a = \mu/2$ and $\lambda = \ln 2$. The exponent becomes $\mu \ln 2/2 - \mu/2$.

□

The privacy theorem chooses the coin probability ν so that, by Lemma 18.2, enough people send coins with high probability. Those coins then hide the one person in question, by Lemma 18.1.

Theorem 18.2 (Privacy of shuffled randomized response)**.** *Let $\varepsilon > 0$, $\delta \in (0,1)$, $\tau = \tanh(\varepsilon/2)$ and $n \ge 1 + 16\ln(4/\delta)/\tau^2$, and set*

$$\nu = \frac{16\ln(4/\delta)}{\tau^2(n-1)}. \tag{18.5}$$

Then shuffled randomized response is (ε, δ)-DP in the shuffle model.

Proof.

1. *Reduce to the count.* Given the count Z, every arrangement of the shuffled bits with Z ones is equally likely. So the shuffled output and Z are functions of each other, up to independent randomness, and it suffices to prove that Z is (ε, δ)-DP.
2. *Isolate one person.* The count does not depend on the order of the people, so we may take neighbours $D \sim D'$ that differ by replacing the record of person 1. Let J be the set of persons $i \ge 2$ who send coins. Its law does not depend on the data, and given J,

$$Z = C_J + N + Y_1, \qquad C_J = \sum_{i \ge 2,\ i \notin J} X_i, \tag{18.6}$$

where N counts the ones among the $|J|$ fair coins. Both C_J and the law of N are the same under D and D', and only the Bernoulli variable Y_1 changes. 3. *Enough coins.* Let $s_0 = 8\ln(4/\delta)/\tau^2$. The size $|J|$ is binomial with mean $(n-1)\nu = 2s_0$. By Lemma 18.2, $\Pr[|J| < s_0] \le \exp(-(1 - \ln 2)\,8\ln(4/\delta)/\tau^2)$. Since $8(1-\ln 2) > 2$ and $\tau < 1$, this is at most $(\delta/4)^2 \le \delta/4$. 4. *Each large J is private.* If $|J| \ge s_0$, then $|J|\tau \ge 8\ln 4/\tau > 2$, so Lemma 18.1 applies. Its error term is $2e^{-|J|\tau^2/8} \le 2e^{-\ln(4/\delta)} = \delta/2$. Shifting by the constant C_J preserves the inequality, so for every

event S, $\Pr_D[Z \in S \mid J] \le e^{\varepsilon} \Pr_{D'}[Z \in S \mid J] + \delta/2$. 5. *Average over J.* Summing over sets with $|J| \ge s_0$ and adding the probability of the others,

$$\Pr_D[Z \in S] \le e^{\varepsilon} \Pr_{D'}[Z \in S] + \frac{\delta}{2} + \frac{\delta}{4} \le e^{\varepsilon} \Pr_{D'}[Z \in S] + \delta. \tag{18.7}$$

□

The coins also add error to the count. But only about $n\nu$ people send them, so the error depends on $n\nu$, which Equation 18.5 keeps independent of n.

Theorem 18.3 (Accuracy of shuffled randomized response)**.** *For every dataset, $\tilde{c}$ is an unbiased estimate of c with variance at most $n\nu/(2(1-\nu)^2)$. With ν as in Equation 18.5 and $\beta \in (0,1)$, with probability at least $1-\beta$,*

$$|\tilde{c} - c| \le \frac{1}{1-\nu}\sqrt{\frac{n\nu}{2\beta}} \le \frac{4}{(1-\nu)\,\tau}\sqrt{\frac{\ln(4/\delta)}{\beta}}. \tag{18.8}$$

For $\varepsilon \le 1$ and $\nu \le \frac{1}{2}$, the right-hand side is at most $18\sqrt{\ln(4/\delta)/\beta}\,/\varepsilon$.

Proof.

1. Each report is 1 with probability $(1-\nu)X_i + \nu/2$, so $\mathbb{E}[Z] = (1-\nu)c + n\nu/2$, and $\tilde{c}$ is unbiased.
2. Each report is a Bernoulli variable with parameter $\nu/2$ or $1-\nu/2$, so its variance is $\frac{\nu}{2}(1-\frac{\nu}{2}) \le \frac{\nu}{2}$. The reports are independent, so $\mathrm{Var}(Z) \le n\nu/2$ and $\mathrm{Var}(\tilde{c}) \le n\nu/(2(1-\nu)^2)$.
3. Chebyshev's inequality gives the first bound. For the second, $n\nu = 16\ln(4/\delta)\,n/(\tau^2(n-1)) \le 32\ln(4/\delta)/\tau^2$ when $n \ge 2$.
4. The function tanh is concave on $[0,\infty)$, so $\tau \ge \tanh(\frac{1}{2})\,\varepsilon = 0.462\,\varepsilon$ for $\varepsilon \le 1$. With $1/(1-\nu) \le 2$, the constant is $8/0.462 < 18$.

□

So the shuffled count has error $O(\sqrt{\ln(1/\delta)}/\varepsilon)$, which does not grow with n. It is within a factor $\sqrt{\ln(1/\delta)}$ of the central model and far from the local model's $\sqrt{n}/\varepsilon$, as Figure 18.1 shows. The example below compares the three models at one setting.

Example 18.1. Take $\varepsilon = 1$, $\delta = 10^{-6}$ and $n = 10^6$. Then $\tau = 0.4621$, and Equation 18.5 gives $\nu = 0.001139$, so about 1,139 people send coins. The standard deviation of $\tilde{c}$ is at most 23.9. Local randomized response at the same ε has standard deviation $\sqrt{n}\,e^{1/2}/(e-1) = 959.5$, and central Laplace noise has $\sqrt{2} = 1.414$.

The price is trust in the shuffler. Each message on its own is randomized response that keeps the true bit with probability $1-\nu/2$, so it is only ε_0-locally private with $\varepsilon_0 = \ln((2-\nu)/\nu) = 7.47$ in the example. A shuffler that leaked the order would expose almost every bit. That ε_0 is below $\ln\left(n/(16\ln(2/\delta))\right) = 8.37$, so this protocol lies in the range that the clone analysis of the previous section covers. Prochlo's engineering, with its anonymization and mixing, is a practice already common in industry, and the shuffle model shows that, combined with a little local noise, it yields a formal guarantee close to the central one.

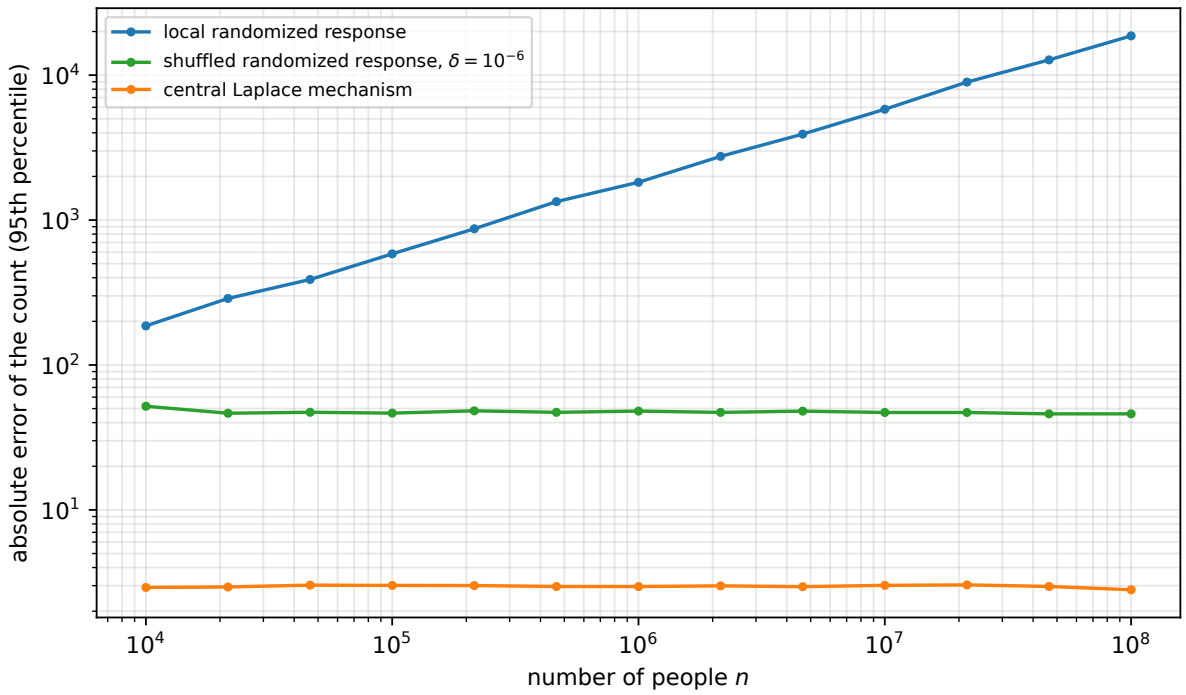


Figure 18.1.: Error of three private counts of n bits at $\varepsilon = 1$, as the 95th percentile of the absolute error over 2,000 simulated runs on data with half the bits equal to one. The local error grows as $\sqrt{n}$. The shuffled error, with ν from Equation 18.5 and $\delta = 10^{-6}$, stays near 47, and the central Laplace error stays near 3.

18.5. A query service: LinkedIn

LinkedIn's audience engagements service lets advertisers ask how members who match a target audience engage with content, such as which articles or topics they click on most. It runs in the central model: LinkedIn holds the data and answers each query with a private mechanism. Its design rests on research done for it, of which this section covers two pieces, a sharper analysis of the exponential mechanism and a fast way to run it.

18.5.1. Bounded range

Recall the exponential mechanism. Given a score $u(D, r)$ for each candidate r in a countable set, whose value changes by at most Δu between neighbours, it outputs r with probability proportional to $\exp(\varepsilon u(D, r)/(2\Delta u))$, and it is ε-DP. Write

$$L(y) = \ln \frac{\Pr[M(D) = y]}{\Pr[M(D') = y]} \tag{18.9}$$

for the privacy loss of output y between neighbours D and D'. Pure ε-DP says $|L(y)| \le \varepsilon$, an interval of width 2ε. The exponential mechanism satisfies a stronger condition.

Definition 18.3 (Bounded range)**.** A mechanism M with a countable output set is **ε-bounded range** if for all neighbours $D \sim D'$ and all outputs y, y',

$$\Pr[M(D) = y]\ \Pr[M(D') = y'] \le e^{\varepsilon}\ \Pr[M(D') = y]\ \Pr[M(D) = y']. \tag{18.10}$$

When the outputs have positive probability, this says $L(y) - L(y') \le \varepsilon$. So the privacy losses lie in an interval of width ε, wherever that interval sits. The next result shows that the two notions are within a factor of two of each other.

Proposition 18.1. *An ε-bounded range mechanism is ε-DP. An ε-DP mechanism with a countable output set is 2ε-bounded range.*

Proof.

1. *Same support.* Let M be ε-bounded range. Suppose some output y has $\Pr[M(D) = y] > 0$ but $\Pr[M(D') = y] = 0$. With this y, the right-hand side of the definition is zero for every y', and the left-hand side is $\Pr[M(D) = y]$ times $\Pr[M(D') = y']$. So the definition forces $\Pr[M(D') = y'] = 0$ for every y', which is impossible. So $M(D)$ and $M(D')$ give positive probability to the same outputs, and L is finite on them.
2. *Both signs occur.* Both distributions sum to one over the common support, so some output y_+ has $L(y_+) \ge 0$ and some output y_- has $L(y_-) \le 0$.
3. *Width ε around zero.* For every y, bounded range gives $L(y) \le L(y_-) + \varepsilon \le \varepsilon$ and $L(y) \ge L(y_+) - \varepsilon \ge -\varepsilon$. The pointwise lemma of the randomized response chapter turns $|L| \le \varepsilon$ into ε-DP.
4. *Converse.* If $|L| \le \varepsilon$ everywhere, then $L(y) - L(y') \le 2\varepsilon$.

□

The general conversion loses that factor of two, but the exponential mechanism does not. Its privacy losses already lie in an interval of width ε.

Proposition 18.2. *The exponential mechanism with parameter ε is ε-bounded range.*

Proof.

1. Write $Z_D = \sum_r \exp(\varepsilon u(D,r)/(2\Delta u))$ for the normalizing constant. Then $L(r) = \frac{\varepsilon}{2\Delta u}(u(D,r) - u(D',r)) + \ln(Z_{D'}/Z_D)$.
2. The constant cancels in a difference: $L(r) - L(r') = \frac{\varepsilon}{2\Delta u}\big[(u(D,r) - u(D',r)) - (u(D,r') - u(D',r'))\big]$.
3. Each bracketed difference is at most Δu in absolute value, so $L(r) - L(r') \le \varepsilon$.

□

The factor of two matters when many selections are composed. Two lemmas make the composition precise. The first converts a tail bound on the privacy loss into a privacy guarantee. It holds for any output space, with L the logarithm of the ratio of probability mass functions or of densities.

Lemma 18.3. *Let $\varepsilon' \ge 0$ and $\delta \in [0,1]$. If for all neighbours $D \sim D'$ the loss satisfies $\Pr_{y\sim M(D)}[L(y) > \varepsilon'] \le \delta$, then M is (ε', δ)-DP.*

Proof. Fix an event S and let $G = \{y : L(y) \le \varepsilon'\}$. On $S \cap G$ the probability under D is at most $e^{\varepsilon'}$ times that under D', so $\Pr[M(D) \in S \cap G] \le e^{\varepsilon'} \Pr[M(D') \in S]$. The rest of S has probability at most $\Pr[M(D) \notin G] \le \delta$ under D. □

The second lemma bounds the mean loss, using Hoeffding's lemma: a random variable W in an interval of length ℓ satisfies $\mathbb{E}[e^{\lambda(W - \mathbb{E}W)}] \le e^{\lambda^2\ell^2/8}$ for every real λ. This is the step at the heart of the proof of Hoeffding's inequality.

Lemma 18.4. *If M is ε-bounded range and $D \sim D'$, then $\mathbb{E}_{y\sim M(D)}[L(y)] \le \varepsilon^2/8$.*

Proof.

1. Under $y \sim M(D)$, the variable $W = -L(y)$ lies in an interval of length ε, and $\mathbb{E}[e^W] = \sum_y \Pr[M(D) = y] \frac{\Pr[M(D')=y]}{\Pr[M(D)=y]} = 1$.
2. Hoeffding's lemma with $\lambda = 1$ gives $1 = \mathbb{E}[e^W] \le e^{\mathbb{E}W + \varepsilon^2/8}$, so $\mathbb{E}[L] = -\mathbb{E}[W] \le \varepsilon^2/8$.

□

The two lemmas combine into a composition theorem. The total privacy loss of k steps is a sum of k losses, each with a small mean and a range of width ε. So the sum concentrates around $k\varepsilon^2/8$, and Lemma 18.3 turns that concentration into a guarantee.

Theorem 18.4 (Composition of bounded-range mechanisms)**.** *Run k mechanisms in sequence on the same dataset, where the j-th may be chosen as a function of the outputs of the first $j-1$, and each is ε-bounded range whatever those outputs are. For every $\delta \in (0,1)$, the sequence of outputs is (ε', δ)-DP with*

$$\varepsilon' = \frac{k\varepsilon^2}{8} + \varepsilon\sqrt{\frac{k\ln(1/\delta)}{2}}. \tag{18.11}$$

Proof.

1. The probability of an output sequence is the product of the conditional probabilities of its terms, so the loss is a sum $L = \sum_{j=1}^{k} L_j$, where L_j is the loss of the j-th output given the earlier ones.
2. Given the earlier outputs, L_j lies in an interval of length ε, by bounded range, and its conditional mean is at most $\varepsilon^2/8$, by Lemma 18.4. Let V_j be L_j minus its conditional mean, so $L \le \sum_j V_j + k\varepsilon^2/8$.
3. For $\lambda > 0$, Hoeffding's lemma applied conditionally gives $\mathbb{E}[e^{\lambda V_j} \mid \text{earlier outputs}] \le e^{\lambda^2\varepsilon^2/8}$. Taking conditional expectations from the last term to the first gives $\mathbb{E}[e^{\lambda \sum_j V_j}] \le e^{k\lambda^2\varepsilon^2/8}$.
4. By Markov's inequality, $\Pr[\sum_j V_j \ge t] \le e^{-\lambda t + k\lambda^2\varepsilon^2/8}$. Choosing $\lambda = 4t/(k\varepsilon^2)$ makes this $e^{-2t^2/(k\varepsilon^2)}$, which equals δ at $t = \varepsilon\sqrt{k\ln(1/\delta)/2}$.
5. So $\Pr[L > \varepsilon'] \le \delta$ for every pair of neighbours, and Lemma 18.3 finishes the proof.

□

Advanced composition gives $\varepsilon\sqrt{2k\ln(1/\delta)} + k\varepsilon\tanh(\varepsilon/2)$ for k mechanisms that are only known to be ε-DP. The square-root term of Theorem 18.4 is half as large. Its linear term $k\varepsilon^2/8$ is also smaller, about a quarter of $k\varepsilon\tanh(\varepsilon/2) \approx k\varepsilon^2/2$ for small ε. A numerical example shows the size of the gain.

Example 18.2. Compose $k = 100$ exponential mechanisms, each with $\varepsilon = 0.1$, at $\delta = 10^{-6}$. Basic composition gives 10.00, advanced composition 5.76, and Theorem 18.4 gives 2.75.

18.5.2. Top-k selection with Gumbel noise

An advertiser usually wants the few most popular items in a histogram, such as the top topics among an audience, not every count. LinkedIn's algorithms split along two lines. The **domain** of items may be known in advance, or not. A person's **contribution** may be m**-restricted**, meaning each person adds one to at most m counts, or unrestricted. Hires are 1-restricted, since a hire adds one to one employer, but skills are not, since one person may list many.

With m-restricted contributions, replacing one person changes the vector of counts by at most $2m$ in ℓ_1 norm. Laplace noise of scale $2m/\varepsilon$ on every count, followed by taking the top k, is then ε-DP. Its noise grows with m. For unrestricted contributions, a better route is to select with the exponential mechanism, whose score for an item is its count. A count changes by at most one between neighbours, however many counts a person touches. **Peeling** selects the top item, removes it, and repeats k times. Each round is the exponential mechanism, so by Proposition 18.2 and Theorem 18.4 its cost depends on k and not on m.

Sampling from the exponential mechanism needs its normalizing constant, and peeling needs it k times. Adding noise of the right shape and taking a maximum avoids both. The right shape is the Gumbel distribution.

Definition 18.4 (Gumbel distribution)**.** The **Gumbel distribution** with scale $b > 0$ has cumulative distribution function $\Pr[G \le g] = \exp(-e^{-g/b})$ for real g.

The next theorem is the reason for this choice. Adding independent Gumbel noise to each score and taking the largest gives exactly the exponential mechanism's distribution, and taking the k largest gives exactly peeling.

Theorem 18.5 (Gumbel-max trick)**.** *Let $u_1, \dots, u_d$ be scores, let $G_1, \dots, G_d$ be independent Gumbel variables with scale b, and write $w_i = e^{u_i/b}$. Then*

1. *the index of the largest noisy score $u_i + G_i$ is i with probability $w_i / \sum_j w_j$;*
2. *the indices of the k largest noisy scores, in decreasing order, have the distribution of k rounds of peeling, where each round picks a remaining item i with probability proportional to w_i.*

Proof.

1. *Exponential race.* Let $E_i = e^{-(u_i+G_i)/b}$. Then $\Pr[E_i > t] = \Pr[G_i < -b \ln t - u_i] = \exp(-w_i t)$ for $t > 0$. So the E_i are independent exponential variables with rates w_i, and the noisy scores in decreasing order are the E_i in increasing order.
2. *The first finisher.* With $W = \sum_j w_j$, the chance that E_i is the smallest is

$$\int_0^\infty w_i e^{-w_i t} \prod_{j \ne i} e^{-w_j t} \, dt = \frac{w_i}{W}. \tag{18.12}$$

This proves the first claim. 3. *Memorylessness.* Condition on E_i being smallest, with value t. This is the event that $E_j > t$ for every $j \ne i$, and given it, the excesses $E_j - t$ are independent exponential variables with the original rates w_j, because $\Pr[E_j > t + s \mid E_j > t] = e^{-w_j s}$. 4. *Induction.* By step 3, the order of the remaining E_j is a fresh exponential race among the remaining items. Step 2 applied to it shows that the next item is chosen with probability proportional to its w_j, and repeating gives the second claim. □

Counts change by at most one between neighbours, so take $b = 2/\varepsilon_0$. Then $w_i = e^{\varepsilon_0 u_i/2}$, and the first claim is the exponential mechanism with parameter ε_0. By the second claim, adding Gumbel noise once and reporting the k largest noisy counts *is* peeling, so it inherits peeling's guarantee from Theorem 18.4. It draws d noise values and selects the top k once, in time of order d plus the time to select the top k. Peeling would take time of order kd.

18.5.3. Queries from outside

Unlike the deployments of the previous chapter, here outside analysts choose the queries, and they may be careless or malicious. Each advertiser has a privacy budget for its own queries. If two advertisers with budgets ε each combined their answers, basic composition would only give 2ε about the members, and with many advertisers the guarantee would fade. The mechanisms alone cannot rule this out when each advertiser has its own budget. LinkedIn relies instead on a vetting process and on agreements in which advertisers undertake not to share results. The service also sets its parameters by modelling the attacks an analyst might mount, deliberately or by accident, and choosing ε and the number of queries allowed per month so that those attacks fail. By that accounting, its monthly privacy loss is lower than the reported values for the local deployments of the previous chapter, though those systems protect very different data.

LinkedIn used the same algorithms for its public labour market insights, which show which employers are hiring, which job titles are growing, and which skills new hires have, by region and industry. Hires by employer and by title are 1-restricted counts, so Laplace noise suits them. Skills are unrestricted, so their top lists come from Gumbel noise and peeling.

18.6. Data releases

The last two deployments do not answer queries. They publish a fixed dataset once, computed in the central model.

18.6.1. Mobility reports

During the Covid-19 pandemic, Google published community mobility reports, built from the location history of users who had opted in. For each region and day, they gave the change in visits to six kinds of places, such as retail and recreation, grocery stores and pharmacies, parks, transit stations, workplaces and residences, relative to a baseline from before the pandemic. Weekly patterns are visible. Visits to workplaces fell on weekdays but hardly changed at weekends, when most people stay away from work anyway, and time at home rose. The counts behind these percentages are noised with Laplace noise to satisfy differential privacy. To bound each person's influence, a person counts at most once per category and place, and at most four such pairs per day. A person with more has four pairs kept at random and the rest dropped, which is the same idea as m-restricted contribution. Regions smaller than 3 square kilometres, and metrics whose noisy count of people is below 100, are discarded, and so are changes that the noise makes too uncertain.

18.6.2. A URL dataset with zero-concentrated privacy

Facebook and the academic organization Social Science One released a dataset on how users interact with web links shared on Facebook. It covers more than 38 million URLs, and researchers obtain access through a proposal process. The raw data are events, such as "a woman aged 35 to 44 in Germany shared this URL on this date". The release aggregates them by URL, month,

country, age and gender, and publishes noisy counts of views, clicks, shares, likes and other reactions. A noisy count can be negative, a visible sign of the noise.

The noise is Gaussian, and the guarantee is zero-concentrated differential privacy, a variant of differential privacy that suits Gaussian noise. It is stated with the Rényi divergence. For an order $\alpha > 1$ and two distributions P and Q, write p and q for their densities, or their mass functions on a countable set. The **Rényi divergence** of order α is

$$D_\alpha(P \,\|\, Q) = \frac{1}{\alpha - 1} \ln \int p(y)^\alpha q(y)^{1-\alpha} \, dy, \tag{18.13}$$

with a sum in place of the integral on a countable set. It is zero when $P = Q$ and grows as they move apart. Zero-concentrated privacy asks that the divergence between the outputs on neighbours grow at most linearly in the order α.

Definition 18.5 (Zero-concentrated differential privacy)**.** Let $\rho \geq 0$. A mechanism M is **ρ-zero-concentrated differentially private** if for all neighbours $D \sim D'$ and all $\alpha > 1$,

$$D_\alpha(M(D) \,\|\, M(D')) \leq \rho \, \alpha. \tag{18.14}$$

It has one parameter, like pure privacy, but composes like advanced composition, as the next three results show. It lies between the two: an ε-DP mechanism is $(\varepsilon^2/2)$-zero-concentrated, and the converse direction holds with a δ.

Lemma 18.5. *For product distributions, $D_\alpha(P_1 \times P_2 \,\|\, Q_1 \times Q_2) = D_\alpha(P_1 \,\|\, Q_1) + D_\alpha(P_2 \,\|\, Q_2)$. So releasing the outputs of a ρ_1- and a ρ_2-zero-concentrated mechanism, run with independent coins, is $(\rho_1 + \rho_2)$-zero-concentrated.*

Proof. The integral of $p^\alpha q^{1-\alpha}$ over a product space factorizes into the two integrals, and the logarithm turns the product into a sum. □

The same additivity holds when the second mechanism is chosen after seeing the first output, as Bun and Steinke (2016) show. The second result gives the guarantee of the Gaussian mechanism, which is why the release uses it.

Proposition 18.3 (Gaussian mechanism)**.** *Let $f \colon \mathcal{X}^n \to \mathbb{R}^d$ change by at most $\Delta_2 f$ in Euclidean norm between neighbours. Then $M(D) = f(D) + Z$, with Z a vector of d independent normal variables of mean 0 and variance σ^2, is $(\Delta_2 f)^2/(2\sigma^2)$-zero-concentrated.*

Proof.

1. In one dimension, let p and q be normal densities with means a and b and variance σ^2. Completing the square,

$$\alpha(y-a)^2 + (1-\alpha)(y-b)^2 = \big(y - \alpha a - (1-\alpha)b\big)^2 + \alpha(1-\alpha)(a-b)^2. \tag{18.15}$$

2. The first term integrates against the normal density to one, so $\int p^{\alpha} q^{1-\alpha} = \exp\left(\alpha(\alpha-1)(a-b)^2/(2\sigma^2)\right)$, and $D_{\alpha} = \alpha(a-b)^2/(2\sigma^2)$.
3. In d dimensions the coordinates are independent, so by Lemma 18.5 the divergences add to $\alpha\|f(D) - f(D')\|_2^2/(2\sigma^2) \le \alpha(\Delta_2 f)^2/(2\sigma^2)$.

□

The third result converts a zero-concentrated guarantee into an (ε, δ) guarantee, the form in which a release usually states its budget.

Proposition 18.4. *A ρ-zero-concentrated mechanism is $(\rho + 2\sqrt{\rho \ln(1/\delta)}, \delta)$-DP for every $\delta \in (0, 1)$.*

Proof.

1. If $\rho = 0$, every Rényi divergence is 0, so $M(D)$ and $M(D')$ have the same distribution and the claim holds. Otherwise let $y \sim M(D)$, with density p, and let $L = \ln(p(y)/q(y))$, where q is the density of $M(D')$. For $\alpha > 1$, $\mathbb{E}[e^{(\alpha-1)L}] = \int p^{\alpha} q^{1-\alpha} = e^{(\alpha-1)D_{\alpha}} \le e^{(\alpha-1)\alpha\rho}$.
2. Markov's inequality gives $\Pr[L > t] \le e^{(\alpha-1)(\alpha\rho - t)}$.
3. Take $t = \rho + 2\sqrt{\rho \ln(1/\delta)}$ and $\alpha = 1 + \sqrt{\ln(1/\delta)/\rho}$. Writing $\ell = \ln(1/\delta)$, the exponent is $\sqrt{\ell/\rho}\,(\rho + \sqrt{\rho\ell} - \rho - 2\sqrt{\rho\ell}) = -\ell$, so $\Pr[L > t] \le \delta$.
4. Lemma 18.3 finishes the proof.

□

The release's budget raises a question that recurs in deployments: what is a neighbour? **User-level** privacy hides everything one person did, and **action-level** privacy hides one action, such as one click. User-level privacy is the goal, but a few users act thousands of times, and hiding them would drown every count in noise. The release instead protects every action strongly enough that almost all users get a user-level guarantee too. A user whose actions number at most a high percentile of the activity distribution enjoys a stated user-level guarantee, and heavier users still have the action-level one.

Example 18.3. Suppose each action adds one to a single cell, and the actions of one user fall in different cells. Replacing one action changes two counts by one each, a Euclidean change of $\sqrt{2}$. Replacing all k actions of a user changes at most $2k$ counts, a change of at most $\sqrt{2k}$. Then by Proposition 18.3, k actions cost k times the action-level ρ. To give users with at most 20 actions $(0.5, 10^{-6})$-DP, solve $\rho + 2\sqrt{\rho \ln 10^6} = 0.5$ for $\rho = 0.00444$ at the user level. The action level is then $\rho = 0.000222$, which Gaussian noise of standard deviation $\sigma = \sqrt{2}/\sqrt{2\rho} = 67.1$ achieves. A single action alone is protected at $(0.111, 10^{-6})$-DP.

Many other deployments exist. The largest is the US Census Bureau's use of differential privacy for the 2020 census, a system complex enough to need a separate treatment.

18.7. Summary

Table 18.1.: The chapter's results.

Result	Statement
Definition 18.1	the adversary sees the messages in a uniformly random order, not who sent which
Theorem 18.1	shuffling turns ε_0-local into $(12\varepsilon_0\sqrt{\ln(1/\delta)/n}, \delta)$-DP, for $\varepsilon_0 < \frac{1}{2}$ and $n \geq 1000$
Theorem 18.2	randomized response with $\nu = 16\ln(4/\delta)/(\tau^2(n-1))$ is (ε, δ)-DP once shuffled
Theorem 18.3	its count has error $O(\sqrt{\ln(1/\delta)/\beta}/\varepsilon)$, independent of n
Proposition 18.1	ε-bounded range implies ε-DP, which implies 2ε-bounded range
Proposition 18.2	the exponential mechanism is ε-bounded range
Theorem 18.4	k bounded-range steps cost $k\varepsilon^2/8 + \varepsilon\sqrt{k\ln(1/\delta)/2}$
Theorem 18.5	Gumbel noise and a maximum sample the exponential mechanism, and the top k is peeling
Proposition 18.3	Gaussian noise with sensitivity $\Delta_2 f$ is $(\Delta_2 f)^2/(2\sigma^2)$-zero-concentrated
Proposition 18.4	ρ-zero-concentrated implies $(\rho + 2\sqrt{\rho\ln(1/\delta)}, \delta)$-DP

18.8. References

- Andrea Bittau, Úlfar Erlingsson, Petros Maniatis, Ilya Mironov, Ananth Raghunathan, David Lie, Mitch Rudominer, Ushasree Kode, Julien Tinnes and Bernhard Seefeld, "Prochlo: Strong Privacy for Analytics in the Crowd", 2017. The encode, shuffle, analyse architecture and its implementation with hardware enclaves.
- Albert Cheu, Adam Smith, Jonathan Ullman, David Zeber and Maxim Zhilyaev, "Distributed Differential Privacy via Shuffling", 2019. The shuffle model as defined here, a one-message protocol for sums of bits with error $O(\sqrt{\ln(1/\delta)}/\varepsilon)$, and a protocol with many one-bit messages for sums of numbers in $[0, 1]$ with error $O(\ln(n/\delta)/\varepsilon)$. The constants in Theorem 18.2 are this chapter's own.
- Úlfar Erlingsson, Vitaly Feldman, Ilya Mironov, Ananth Raghunathan, Kunal Talwar and Abhradeep Thakurta, "Amplification by Shuffling: From Local to Central Differential Privacy via Anonymity", 2019. Theorem 18.1.

- Vitaly Feldman, Audra McMillan and Kunal Talwar, "Hiding Among the Clones: A Simple and Nearly Optimal Analysis of Privacy Amplification by Shuffling", 2021. Amplification for ε_0 up to about $\ln n$, by the clone decomposition.
- Ryan Rogers, Subbu Subramaniam, Sean Peng, David Durfee, Seunghyun Lee, Santosh Kumar Kancha, Shraddha Sahay and Parvez Ahammad, "LinkedIn's Audience Engagements API: A Privacy Preserving Data Analytics System at Scale", 2020. The query service, its algorithms for known and unknown domains, and its choice of parameters.
- David Durfee and Ryan Rogers, "Practical Differentially Private Top-k Selection with Pay-what-you-get Composition", 2019. Bounded range, Gumbel-noise top-k selection and its equivalence with peeling.
- Jinshuo Dong, David Durfee and Ryan Rogers, "Optimal Differential Privacy Composition for Exponential Mechanisms", 2020. The optimal composition of bounded-range mechanisms, and efficiently computable bounds that improve slightly on Theorem 18.4.
- Emil Julius Gumbel, *Statistical Theory of Extreme Values and Some Practical Applications*, 1954. The Gumbel distribution and its use for maxima.
- Ryan Rogers, Adrian Rivera Cardoso, Koray Mancuhan, Akash Kaura, Nikhil Gahlawat, Neha Jain, Paul Ko and Parvez Ahammad, "A Members First Approach to Enabling LinkedIn's Labor Market Insights at Scale", 2020. The labour market release and its choice of algorithms.
- Ahmet Aktay and coauthors, "Google Covid-19 Community Mobility Reports: Anonymization Process Description", 2020. How the mobility reports bound contributions and add noise.
- Solomon Messing, Christina DeGregorio, Bennett Hillenbrand, Gary King, Saurav Mahanti, Zagreb Mukerjee, Chaya Nayak, Nate Persily, Bogdan State and Arjun Wilkins, "Facebook Privacy-Protected Full URLs Data Set", 2020. The URL release, its Gaussian noise and its action-level and user-level budget, set so that the 99th percentile of users by activity gets the user-level guarantee.
- Cynthia Dwork and Guy N. Rothblum, "Concentrated Differential Privacy", 2016, and Mark Bun and Thomas Steinke, "Concentrated Differential Privacy: Simplifications, Extensions, and Lower Bounds", 2016. Concentrated privacy, and its zero-concentrated form with Proposition 18.3, Proposition 18.4 and adaptive composition.
- John M. Abowd and coauthors, "The 2020 Census Disclosure Avoidance System TopDown Algorithm", 2022. The census deployment.

A. Glossary

Acronyms used in the book and the pages that use them.

API Application programming interface. Used in The Shuffle Model and Central Deployments.
DNA Deoxyribonucleic acid. Used in Attempts at Data Privacy.
DP Differential privacy. Used in The Basics of Differential Privacy, Randomized Response and the Definition, The Laplace Mechanism, Approximate Differential Privacy, Advanced Composition, The Exponential Mechanism, Query Release and Its Limits, Private Multiplicative Weights, Sparse Vector, Beyond Global Sensitivity, Packing Lower Bounds, Private Learning and Statistics, What Differential Privacy Protects, Private Empirical Risk Minimization, Private Deep Learning, Private Mean Estimation, Adaptive Data Analysis, Deployments, Local Differential Privacy, The Shuffle Model and Central Deployments.
ERM Empirical risk minimization. Used in Private Empirical Risk Minimization.
IP Internet Protocol, as in IP address. Used in Local Differential Privacy.
LP Linear programming. Used in Reconstruction Attacks.
PAC Probably approximately correct. Used in The Exponential Mechanism.
PATE Private aggregation of teacher ensembles. Used in Private Deep Learning.
PIN Personal identification number. Used in Local Differential Privacy.
RDP Rényi differential privacy. Used in Private Deep Learning.
SGD Stochastic gradient descent. Used in Private Empirical Risk Minimization, Private Deep Learning.
SVM Support vector machine. Used in Private Empirical Risk Minimization.
TEE Trusted execution environment. Used in The Shuffle Model and Central Deployments.
TLS Transport Layer Security. Used in Local Differential Privacy.
ZIP Zone Improvement Plan code, the US postal code. Used in Attempts at Data Privacy.

B. Index

AboveThreshold, 94–98, 100, 101
adaptive composition, 40, 42, 55, 57–59, 63, 66, 81, 90, 97, 98, 100, 153, 211
advanced composition, 24, 42, 57, 59, 60, 67, 82, 91, 92, 140, 142, 149, 176, 179, 180, 205
amplification by shuffling, 198, 210, 211
amplification by subsampling, 142, 154
approximate differential privacy, 24, 43, 57, 61, 92, 111, 119, 120, 125, 146
auxiliary knowledge, 3

basic composition, 40, 42, 55, 59, 60, 82, 96–98, 162, 163, 205
Bernoulli distribution, 26, 65, 156, 170, 199–201
binomial distribution, 199, 200
blatantly non-private, 14, 17, 22

canary, 6, 10
Chernoff bound, 200
clipping, 122, 144, 153, 157, 159–161, 163, 164, 166, 167, 169, 170
composition, 11, 24, 32, 40–42, 55–60, 63, 66, 67, 81, 82, 90–92, 96–98, 100, 122, 140, 142, 147, 149, 153, 162, 163, 166, 170, 176, 179, 180, 205, 210, 211
composition attack, 11

differential privacy, 3, 22, 24, 27, 32, 42, 43, 57, 67, 77, 79, 92, 101, 111, 119, 120, 122, 128, 141, 142, 153, 154, 170, 180, 182, 196, 208, 210, 211
distance to instability, 103, 111
distributed differential privacy, 210
DP-SGD, 144, 145, 148, 149, 153, 154

empirical risk minimization, 122, 129, 133, 137–142
equivalence class, 8
exponential mechanism, 24, 69, 73, 74, 76, 77, 82, 89, 92, 182, 204, 205, 210, 211

fingerprinting code, 119, 180
frequency oracle, 192

Gaussian differential privacy, 32
Gaussian mechanism, 24, 45, 49–53, 57, 61, 92, 133, 140, 148, 153, 166, 170, 176, 179, 180, 208
global sensitivity, 79, 103, 110
group privacy, 39, 42, 54, 114, 115, 127, 157, 158
Gumbel-max trick, 206

k-anonymity, 3, 8, 10

Laplace distribution, 35, 61, 72, 125
Laplace mechanism, 24, 35–38, 42, 45, 50, 61, 77, 79, 81, 82, 89, 96, 98, 100, 103–106, 113, 141, 151, 156, 157, 160, 161, 170, 184, 185
linkage attack, 5, 128
Lipschitz, 111, 130, 132, 133, 137, 138, 167
local differential privacy, 182, 196
local randomizer, 183, 187, 189, 190, 192, 195, 196, 198
local sensitivity, 79, 102, 103, 108, 109, 111

moment generating function, 51, 159
moments accountant, 145, 149, 153

Narayanan–Shmatikov attack, 10, 128
neighbouring datasets, 27, 69, 79, 182

objective perturbation, 135, 141, 142
one-way marginal, 79, 92, 112, 113, 116, 117, 119, 120
optimal composition, 32, 67, 211
output perturbation, 133, 141, 196

packing lower bound, 79, 119, 120
Paley–Zygmund inequality, 16
Poisson distribution, 149, 153, 154
Poisson sampling, 149, 153, 154
post-processing, 24, 39, 41, 42, 54, 57, 62, 66, 90, 100, 104, 106–108, 124, 125, 127, 133, 140, 151, 161, 176, 178, 184, 189, 190
privacy loss, 44–47, 51, 54, 57, 64, 66, 67, 122, 135, 136, 146, 147, 153, 196
privacy loss random variable, 44–46, 51, 57, 67, 146
privacy odometer, 57
propose-test-release, 104, 111
pure differential privacy, 46, 56, 60, 79, 92, 119

quasi-identifier, 8, 10

randomized response, 24, 25, 30–33, 36, 37, 40, 60, 189, 190, 193, 195, 196, 199–201, 203, 210
re-identification, 22
reconstruction attack, 3, 22, 42
Rényi differential privacy, 57, 153
Rényi divergence, 209

sample complexity, 119, 158
sampling without replacement, 153
sensitivity, 32, 34, 36–38, 41, 42, 45, 49, 61, 69, 70, 72, 73, 79, 81, 82, 89, 94, 95, 98–100, 102–104, 106, 108–111, 113, 132, 133, 140, 141, 147, 148, 151, 154, 156, 160, 161, 164, 166, 170, 176, 210
shuffle model, 182, 198, 200, 210
smooth sensitivity, 109–111, 154
sparse vector technique, 79, 96–98, 101, 154
stochastic gradient descent, 122, 136–142, 144, 145, 148, 149, 153, 154
strong convexity, 132, 133, 135, 141

total variation distance, 32

union bound, 17, 18, 20, 38, 48, 75, 76, 81, 90, 96–98, 100, 107, 113, 118, 152, 156, 160, 162, 166, 168, 173, 174, 176

variance, 26, 33, 41, 45, 50, 61, 138, 140, 147, 167, 169–171, 193, 201, 208

zero-concentrated differential privacy, 182, 208, 210, 211